



CIS Benchmarks

Security Configuration Assessment Report for gn-VirtualBox

Target IP Address: 192.168.1.14

CIS Ubuntu Linux 20.04 LTS Benchmark v3.0.0

Level 1 - Workstation

Sunday, July 20 2025 12:00:58

Assessment Duration: 1 minute, 36 seconds

Report generated by the Center for Internet Security's Configuration Assessment Tool (CIS-CAT Pro Assessor) v4.55.0.

For further information, please visit [The Center for Internet Security](#) or our [Product Support](#) page.

Copyright ©2025, The Center for Internet Security

Content generated on 07/20/2025 12:02 PM. Content last obtained on 06/25/2025 19:40 PM.

Summary

Description	Tests						Scoring		
	Pass	Fail	Error	Unkn.	Man.	Exc.	Score	Max	Percent
1 Initial Setup	27	22	0	0	4	0	27.0	49.0	55%
1.1 Filesystem	19	7	0	0	1	0	19.0	26.0	73%
1.1.1 Configure Filesystem Kernel Modules	0	5	0	0	1	0	0.0	5.0	0%
1.1.2 Configure Filesystem Partitions	19	2	0	0	0	0	19.0	21.0	90%
1.1.2.1 Configure /tmp	3	1	0	0	0	0	3.0	4.0	75%
1.1.2.2 Configure /dev/shm	3	1	0	0	0	0	3.0	4.0	75%
1.1.2.3 Configure /home	2	0	0	0	0	0	2.0	2.0	100%
1.1.2.4 Configure /var	2	0	0	0	0	0	2.0	2.0	100%
1.1.2.5 Configure /var/tmp	3	0	0	0	0	0	3.0	3.0	100%
1.1.2.6 Configure /var/log	3	0	0	0	0	0	3.0	3.0	100%
1.1.2.7 Configure /var/log/audit	3	0	0	0	0	0	3.0	3.0	100%
1.2 Package Management	0	0	0	0	3	0	0.0	0.0	0%
1.2.1 Configure Package Repositories	0	0	0	0	2	0	0.0	0.0	0%
1.2.2 Configure Package Updates	0	0	0	0	1	0	0.0	0.0	0%
1.3 Mandatory Access Control	1	2	0	0	0	0	1.0	3.0	33%
1.3.1 Configure AppArmor	1	2	0	0	0	0	1.0	3.0	33%
1.4 Configure Bootloader	0	2	0	0	0	0	0.0	2.0	0%
1.5 Configure Additional Process Hardening	2	3	0	0	0	0	2.0	5.0	40%
1.6 Configure Command Line Warning Banners	4	2	0	0	0	0	4.0	6.0	67%
1.7 Configure GNOME Display Manager	1	6	0	0	0	0	1.0	7.0	14%
2 Services	28	11	0	0	1	0	28.0	39.0	72%
2.1 Configure Server Services	16	1	0	0	1	0	16.0	17.0	94%
2.2 Configure Client Services	4	2	0	0	0	0	4.0	6.0	67%
2.3 Configure Time Synchronization	5	1	0	0	0	0	5.0	6.0	83%
2.3.1 Ensure time synchronization is in use	1	0	0	0	0	0	1.0	1.0	100%
2.3.2 Configure systemd-timesyncd	1	1	0	0	0	0	1.0	2.0	50%
2.3.3 Configure chrony	3	0	0	0	0	0	3.0	3.0	100%
2.4 Job Schedulers	3	7	0	0	0	0	3.0	10.0	30%
2.4.1 Configure cron	2	7	0	0	0	0	2.0	9.0	22%
2.4.2 Configure at	1	0	0	0	0	0	1.0	1.0	100%
3 Network	2	9	0	0	1	0	2.0	11.0	18%
3.1 Configure Network Devices	0	0	0	0	1	0	0.0	0.0	0%
3.2 Configure Network Kernel Modules	0	0	0	0	0	0	0.0	0.0	0%
3.3 Configure Network Kernel Parameters	2	9	0	0	0	0	2.0	11.0	18%
4 Host Based Firewall	5	20	0	0	5	0	5.0	25.0	20%
4.1 Configure a single firewall utility	1	0	0	0	0	0	1.0	1.0	100%
4.2 Configure UncomplicatedFirewall	3	4	0	0	1	0	3.0	7.0	43%
4.3 Configure nftables	0	8	0	0	2	0	0.0	8.0	0%
4.4 Configure iptables	1	8	0	0	2	0	1.0	9.0	11%
4.4.1 Configure iptables software	1	2	0	0	0	0	1.0	3.0	33%
4.4.2 Configure IPv4 iptables	0	3	0	0	1	0	0.0	3.0	0%
4.4.3 Configure IPv6 ip6tables	0	3	0	0	1	0	0.0	3.0	0%
5 Access Control	42	24	0	0	1	0	42.0	66.0	64%
5.1 Configure SSH Server	22	0	0	0	0	0	22.0	22.0	100%
5.2 Configure privilege escalation	3	3	0	0	0	0	3.0	6.0	50%
5.3 Pluggable Authentication Modules	7	16	0	0	1	0	7.0	23.0	30%
5.3.1 Configure PAM software packages	2	1	0	0	0	0	2.0	3.0	67%
5.3.2 Configure pam-auth-update profiles	1	3	0	0	0	0	1.0	4.0	25%
5.3.3 Configure PAM Arguments	4	12	0	0	1	0	4.0	16.0	25%
5.3.3.1 Configure pam_faillock module	0	2	0	0	0	0	0.0	2.0	0%
5.3.3.2 Configure pam_pwquality module	2	5	0	0	1	0	2.0	7.0	29%
5.3.3.3 Configure pam_pwhistory module	0	3	0	0	0	0	0.0	3.0	0%
5.3.3.4 Configure pam_unix module	2	2	0	0	0	0	2.0	4.0	50%
5.4 User Accounts and Environment	10	5	0	0	0	0	10.0	15.0	67%
5.4.1 Configure shadow password suite parameters	3	2	0	0	0	0	3.0	5.0	60%
5.4.2 Configure root and system accounts and environment	7	1	0	0	0	0	7.0	8.0	88%
5.4.3 Configure user default environment	0	2	0	0	0	0	0.0	2.0	0%
6 Logging and Auditing	11	5	0	0	6	0	11.0	16.0	69%
6.1 Configure Integrity Checking	0	3	0	0	0	0	0.0	3.0	0%
6.2 System Logging	11	2	0	0	6	0	11.0	13.0	85%
6.2.1 Configure systemd-journald service	1	0	0	0	2	0	1.0	1.0	100%

Description	Tests						Scoring		
	Pass	Fail	Error	Unkn.	Man.	Exc.	Score	Max	Percent
6.2.2 Configure journald	6	0	0	0	1	0	6.0	6.0	100%
6.2.2.1 Configure systemd-journal-remote	3	0	0	0	1	0	3.0	3.0	100%
6.2.3 Configure rsyslog	4	1	0	0	3	0	4.0	5.0	80%
6.2.4 Configure Logfiles	0	1	0	0	0	0	0.0	1.0	0%
6.3 System Auditing	0	0	0	0	0	0	0.0	0.0	0%
6.3.1 Configure auditd Service	0	0	0	0	0	0	0.0	0.0	0%
6.3.2 Configure Data Retention	0	0	0	0	0	0	0.0	0.0	0%
6.3.3 Configure auditd Rules	0	0	0	0	0	0	0.0	0.0	0%
6.3.4 Configure auditd File Access	0	0	0	0	0	0	0.0	0.0	0%
7 System Maintenance	20	2	0	0	1	0	20.0	22.0	91%
7.1 Configure system file and directory access	12	0	0	0	1	0	12.0	12.0	100%
7.2 Local User and Group Settings	8	2	0	0	0	0	8.0	10.0	80%
Total	135	93	0	0	19	0	135.0	228.0	59%

Note: Actual scores are subject to rounding errors. The sum of these values may not result in the exact overall score.

The 'Exc' column only applies to Exceptions that are generated using CIS-CAT Pro Dashboard and is not utilized by CIS-CAT Pro Assessor.

Profiles

This benchmark contains 4 profiles. The **Level 1 - Workstation** profile was used for this assessment.

Title	Description
Level 1 - Server	<p>Items in this profile intend to:</p> <ul style="list-style-type: none">• be practical and prudent;• provide a clear security benefit; and• not inhibit the utility of the technology beyond acceptable means. <p>This profile is intended for servers.</p> <p>Show Profile XML</p>
Level 2 - Server	<p>This profile extends the "Level 1 - Server" profile. Items in this profile exhibit one or more of the following characteristics:</p> <ul style="list-style-type: none">• are intended for environments or use cases where security is paramount.• acts as defense in depth measure.• may negatively inhibit the utility or performance of the technology. <p>This profile is intended for servers.</p> <p>Show Profile XML</p>
Level 1 - Workstation	<p>Items in this profile intend to:</p> <ul style="list-style-type: none">• be practical and prudent;• provide a clear security benefit; and• not inhibit the utility of the technology beyond acceptable means. <p>This profile is intended for workstations.</p> <p>Show Profile XML</p>
Level 2 - Workstation	<p>This profile extends the "Level 1 - Workstation" profile. Items in this profile exhibit one or more of the following characteristics:</p> <ul style="list-style-type: none">• are intended for environments or use cases where security is paramount.• acts as defense in depth measure.• may negatively inhibit the utility or performance of the technology. <p>This profile is intended for workstations.</p> <p>Show Profile XML</p>

[↑](#)

Assessment Results

☐ Display Only Essential Hygiene (CIS Critical Security Controls V8- IG-1)

☐ Display Only Failures

[More](#)

w/	Benchmark Item	Result
1 Initial Setup		
1.1 Filesystem		
1.1.1 Configure Filesystem Kernel Modules		

W	Benchmark Item	Result
1.0	1.1.1.1 Ensure cramfs kernel module is not available	Fail
1.0	1.1.1.2 Ensure freevxfs kernel module is not available	Fail
1.0	1.1.1.3 Ensure hfs kernel module is not available	Fail
1.0	1.1.1.4 Ensure hfsplus kernel module is not available	Fail
1.0	1.1.1.5 Ensure jffs2 kernel module is not available	Fail
	1.1.1.10 Ensure unused filesystems kernel modules are not available	Manual
1.1.2 Configure Filesystem Partitions		
1.1.2.1 Configure /tmp		
1.0	1.1.2.1.1 Ensure /tmp is a separate partition	Fail
1.0	1.1.2.1.2 Ensure nodev option set on /tmp partition	Pass
1.0	1.1.2.1.3 Ensure nosuid option set on /tmp partition	Pass
1.0	1.1.2.1.4 Ensure noexec option set on /tmp partition	Pass
1.1.2.2 Configure /dev/shm		
1.0	1.1.2.2.1 Ensure /dev/shm is a separate partition	Pass
1.0	1.1.2.2.2 Ensure nodev option set on /dev/shm partition	Pass
1.0	1.1.2.2.3 Ensure nosuid option set on /dev/shm partition	Pass
1.0	1.1.2.2.4 Ensure noexec option set on /dev/shm partition	Fail
1.1.2.3 Configure /home		
1.0	1.1.2.3.2 Ensure nodev option set on /home partition	Pass
1.0	1.1.2.3.3 Ensure nosuid option set on /home partition	Pass
1.1.2.4 Configure /var		
1.0	1.1.2.4.2 Ensure nodev option set on /var partition	Pass
1.0	1.1.2.4.3 Ensure nosuid option set on /var partition	Pass
1.1.2.5 Configure /var/tmp		
1.0	1.1.2.5.2 Ensure nodev option set on /var/tmp partition	Pass
1.0	1.1.2.5.3 Ensure nosuid option set on /var/tmp partition	Pass
1.0	1.1.2.5.4 Ensure noexec option set on /var/tmp partition	Pass
1.1.2.6 Configure /var/log		
1.0	1.1.2.6.2 Ensure nodev option set on /var/log partition	Pass
1.0	1.1.2.6.3 Ensure nosuid option set on /var/log partition	Pass
1.0	1.1.2.6.4 Ensure noexec option set on /var/log partition	Pass
1.1.2.7 Configure /var/log/audit		
1.0	1.1.2.7.2 Ensure nodev option set on /var/log/audit partition	Pass
1.0	1.1.2.7.3 Ensure nosuid option set on /var/log/audit partition	Pass
1.0	1.1.2.7.4 Ensure noexec option set on /var/log/audit partition	Pass
1.2 Package Management		
1.2.1 Configure Package Repositories		
	1.2.1.1 Ensure GPG keys are configured	Manual
	1.2.1.2 Ensure package manager repositories are configured	Manual
1.2.2 Configure Package Updates		
	1.2.2.1 Ensure updates, patches, and additional security software are installed	Manual
1.3 Mandatory Access Control		
1.3.1 Configure AppArmor		
1.0	1.3.1.1 Ensure latest versions of the apparmor packages are installed	Fail
1.0	1.3.1.2 Ensure AppArmor is enabled in the bootloader configuration	Fail
1.0	1.3.1.3 Ensure all AppArmor Profiles are not disabled	Pass
1.4 Configure Bootloader		
1.0	1.4.1 Ensure bootloader password is set	Fail
1.0	1.4.2 Ensure access to bootloader config is configured	Fail
1.5 Configure Additional Process Hardening		
1.0	1.5.1 Ensure address space layout randomization is enabled	Fail
1.0	1.5.2 Ensure ptrace_scope is restricted	Pass
1.0	1.5.3 Ensure core dumps are restricted	Fail
1.0	1.5.4 Ensure prelink is not installed	Pass
1.0	1.5.5 Ensure Automatic Error Reporting is not enabled	Fail
1.6 Configure Command Line Warning Banners		
1.0	1.6.1 Ensure /etc/motd is configured	Pass

W	Benchmark Item	Result
1.0	1.6.2 Ensure /etc/issue is configured	Fail
1.0	1.6.3 Ensure /etc/issue.net is configured	Fail
1.0	1.6.4 Ensure access to /etc/motd is configured	Pass
1.0	1.6.5 Ensure access to /etc/issue is configured	Pass
1.0	1.6.6 Ensure access to /etc/issue.net is configured	Pass
1.7 Configure GNOME Display Manager		
1.0	1.7.2 Ensure GDM login banner is configured	Fail
1.0	1.7.3 Ensure GDM disable-user-list option is enabled	Fail
1.0	1.7.4 Ensure GDM screen locks when the user is idle	Fail
1.0	1.7.5 Ensure GDM screen locks cannot be overridden	Fail
1.0	1.7.8 Ensure GDM autorun-never is enabled	Fail
1.0	1.7.9 Ensure GDM autorun-never is not overridden	Fail
1.0	1.7.10 Ensure XDMCP is not enabled	Pass
2 Services		
2.1 Configure Server Services		
1.0	2.1.3 Ensure dhcp server services are not in use	Pass
1.0	2.1.4 Ensure dns server services are not in use	Pass
1.0	2.1.5 Ensure dnsmasq services are not in use	Pass
1.0	2.1.6 Ensure ftp server services are not in use	Pass
1.0	2.1.7 Ensure ldap server services are not in use	Pass
1.0	2.1.8 Ensure message access server services are not in use	Pass
1.0	2.1.9 Ensure network file system services are not in use	Pass
1.0	2.1.10 Ensure nis server services are not in use	Pass
1.0	2.1.12 Ensure rpcbind services are not in use	Pass
1.0	2.1.13 Ensure rsync services are not in use	Fail
1.0	2.1.14 Ensure samba file server services are not in use	Pass
1.0	2.1.15 Ensure snmp services are not in use	Pass
1.0	2.1.16 Ensure tftp server services are not in use	Pass
1.0	2.1.17 Ensure web proxy server services are not in use	Pass
1.0	2.1.18 Ensure web server services are not in use	Pass
1.0	2.1.19 Ensure xinetd services are not in use	Pass
1.0	2.1.21 Ensure mail transfer agents are configured for local-only mode	Pass
	2.1.22 Ensure only approved services are listening on a network interface	Manual
2.2 Configure Client Services		
1.0	2.2.1 Ensure nis client is not installed	Pass
1.0	2.2.2 Ensure rsh client is not installed	Pass
1.0	2.2.3 Ensure talk client is not installed	Pass
1.0	2.2.4 Ensure telnet client is not installed	Fail
1.0	2.2.5 Ensure ldap client is not installed	Pass
1.0	2.2.6 Ensure ftp client is not installed	Fail
2.3 Configure Time Synchronization		
2.3.1 Ensure time synchronization is in use		
1.0	2.3.1.1 Ensure a single time synchronization daemon is in use	Pass
2.3.2 Configure systemd-timesyncd		
1.0	2.3.2.1 Ensure systemd-timesyncd configured with authorized timeserver	Fail
1.0	2.3.2.2 Ensure systemd-timesyncd is enabled and running	Pass
2.3.3 Configure chrony		
1.0	2.3.3.1 Ensure chrony is configured with authorized timeserver	Pass
1.0	2.3.3.2 Ensure chrony is running as user _chrony	Pass
1.0	2.3.3.3 Ensure chrony is enabled and running	Pass
2.4 Job Schedulers		
2.4.1 Configure cron		
1.0	2.4.1.1 Ensure cron daemon is enabled and active	Pass
1.0	2.4.1.2 Ensure access to /etc/crontab is configured	Fail
1.0	2.4.1.3 Ensure access to /etc/cron.hourly is configured	Fail
1.0	2.4.1.4 Ensure access to /etc/cron.daily is configured	Fail
1.0	2.4.1.5 Ensure access to /etc/cron.weekly is configured	Fail
1.0	2.4.1.6 Ensure access to /etc/cron.monthly is configured	Fail

W	Benchmark Item	Result
1.0	2.4.1.7 Ensure access to /etc/cron.yearly is configured	Pass
1.0	2.4.1.8 Ensure access to /etc/cron.d is configured	Fail
1.0	2.4.1.9 Ensure access to crontab is configured	Fail
2.4.2 Configure at		
1.0	2.4.2.1 Ensure access to at is configured	Pass
3 Network		
3.1 Configure Network Devices		
	3.1.1 Ensure IPv6 status is identified	Manual
3.2 Configure Network Kernel Modules		
3.3 Configure Network Kernel Parameters		
1.0	3.3.1 Ensure ip forwarding is disabled	Fail
1.0	3.3.2 Ensure packet redirect sending is disabled	Fail
1.0	3.3.3 Ensure bogus icmp responses are ignored	Pass
1.0	3.3.4 Ensure broadcast icmp requests are ignored	Pass
1.0	3.3.5 Ensure icmp redirects are not accepted	Fail
1.0	3.3.6 Ensure secure icmp redirects are not accepted	Fail
1.0	3.3.7 Ensure reverse path filtering is enabled	Fail
1.0	3.3.8 Ensure source routed packets are not accepted	Fail
1.0	3.3.9 Ensure suspicious packets are logged	Fail
1.0	3.3.10 Ensure tcp syn cookies is enabled	Fail
1.0	3.3.11 Ensure ipv6 router advertisements are not accepted	Fail
4 Host Based Firewall		
4.1 Configure a single firewall utility		
1.0	4.1.1 Ensure a single firewall configuration utility is in use	Pass
4.2 Configure UncomplicatedFirewall		
1.0	4.2.1 Ensure ufw is installed	Pass
1.0	4.2.2 Ensure nftables is not in use with ufw	Pass
1.0	4.2.3 Ensure iptables-persistent is not installed with ufw	Pass
1.0	4.2.4 Ensure ufw service is enabled	Fail
1.0	4.2.5 Ensure ufw loopback traffic is configured	Fail
	4.2.6 Ensure ufw outbound connections are configured	Manual
1.0	4.2.7 Ensure ufw firewall rules exist for all open ports	Fail
1.0	4.2.8 Ensure ufw default deny firewall policy	Fail
4.3 Configure nftables		
1.0	4.3.1 Ensure nftables is installed	Fail
1.0	4.3.2 Ensure ufw is uninstalled or disabled with nftables	Fail
	4.3.3 Ensure iptables are flushed with nftables	Manual
1.0	4.3.4 Ensure a nftables table exists	Fail
1.0	4.3.5 Ensure nftables base chains exist	Fail
1.0	4.3.6 Ensure nftables loopback traffic is configured	Fail
	4.3.7 Ensure nftables outbound and established connections are configured	Manual
1.0	4.3.8 Ensure nftables default deny firewall policy	Fail
1.0	4.3.9 Ensure nftables service is enabled	Fail
1.0	4.3.10 Ensure nftables rules are permanent	Fail
4.4 Configure iptables		
4.4.1 Configure iptables software		
1.0	4.4.1.1 Ensure iptables packages are installed	Fail
1.0	4.4.1.2 Ensure nftables is not in use with iptables	Pass
1.0	4.4.1.3 Ensure ufw is not in use with iptables	Fail
4.4.2 Configure IPv4 iptables		
1.0	4.4.2.1 Ensure iptables default deny firewall policy	Fail
1.0	4.4.2.2 Ensure iptables loopback traffic is configured	Fail
	4.4.2.3 Ensure iptables outbound and established connections are configured	Manual
1.0	4.4.2.4 Ensure iptables firewall rules exist for all open ports	Fail
4.4.3 Configure IPv6 ip6tables		
1.0	4.4.3.1 Ensure ip6tables default deny firewall policy	Fail

W	Benchmark Item	Result
1.0	4.4.3.2 Ensure iptables loopback traffic is configured	Fail
	4.4.3.3 Ensure iptables outbound and established connections are configured	Manual
1.0	4.4.3.4 Ensure iptables firewall rules exist for all open ports	Fail
5 Access Control		
5.1 Configure SSH Server		
1.0	5.1.1 Ensure access to /etc/ssh/sshd_config is configured	Pass
1.0	5.1.2 Ensure access to SSH private host key files is configured	Pass
1.0	5.1.3 Ensure access to SSH public host key files is configured	Pass
1.0	5.1.4 Ensure sshd access is configured	Pass
1.0	5.1.5 Ensure sshd Banner is configured	Pass
1.0	5.1.6 Ensure sshd Ciphers are configured	Pass
1.0	5.1.7 Ensure sshd ClientAliveInterval and ClientAliveCountMax are configured	Pass
1.0	5.1.8 Ensure sshd DisableForwarding is enabled	Pass
1.0	5.1.9 Ensure sshd GSSAPIAuthentication is disabled	Pass
1.0	5.1.10 Ensure sshd HostbasedAuthentication is disabled	Pass
1.0	5.1.11 Ensure sshd IgnoreRhosts is enabled	Pass
1.0	5.1.12 Ensure sshd KexAlgorithms is configured	Pass
1.0	5.1.13 Ensure sshd LoginGraceTime is configured	Pass
1.0	5.1.14 Ensure sshd LogLevel is configured	Pass
1.0	5.1.15 Ensure sshd MACs are configured	Pass
1.0	5.1.16 Ensure sshd MaxAuthTries is configured	Pass
1.0	5.1.17 Ensure sshd MaxSessions is configured	Pass
1.0	5.1.18 Ensure sshd MaxStartups is configured	Pass
1.0	5.1.19 Ensure sshd PermitEmptyPasswords is disabled	Pass
1.0	5.1.20 Ensure sshd PermitRootLogin is disabled	Pass
1.0	5.1.21 Ensure sshd PermitUserEnvironment is disabled	Pass
1.0	5.1.22 Ensure sshd UsePAM is enabled	Pass
5.2 Configure privilege escalation		
1.0	5.2.1 Ensure sudo is installed	Pass
1.0	5.2.2 Ensure sudo commands use pty	Fail
1.0	5.2.3 Ensure sudo log file exists	Fail
1.0	5.2.5 Ensure re-authentication for privilege escalation is not disabled globally	Pass
1.0	5.2.6 Ensure sudo authentication timeout is configured	Pass
1.0	5.2.7 Ensure access to the su command is restricted	Fail
5.3 Pluggable Authentication Modules		
5.3.1 Configure PAM software packages		
1.0	5.3.1.1 Ensure latest version of pam is installed	Pass
1.0	5.3.1.2 Ensure latest version of libpam-modules is installed	Pass
1.0	5.3.1.3 Ensure latest version of libpam-pwquality is installed	Fail
5.3.2 Configure pam-auth-update profiles		
1.0	5.3.2.1 Ensure pam_unix module is enabled	Pass
1.0	5.3.2.2 Ensure pam_faillock module is enabled	Fail
1.0	5.3.2.3 Ensure pam_pwquality module is enabled	Fail
1.0	5.3.2.4 Ensure pam_pwhistory module is enabled	Fail
5.3.3 Configure PAM Arguments		
5.3.3.1 Configure pam_faillock module		
1.0	5.3.3.1.1 Ensure password failed attempts lockout is configured	Fail
1.0	5.3.3.1.2 Ensure password unlock time is configured	Fail
5.3.3.2 Configure pam_pwquality module		
1.0	5.3.3.2.1 Ensure password number of changed characters is configured	Fail
1.0	5.3.3.2.2 Ensure minimum password length is configured	Fail
	5.3.3.2.3 Ensure password complexity is configured	Manual
1.0	5.3.3.2.4 Ensure password same consecutive characters is configured	Fail
1.0	5.3.3.2.5 Ensure password maximum sequential characters is configured	Fail
1.0	5.3.3.2.6 Ensure password dictionary check is enabled	Pass
1.0	5.3.3.2.7 Ensure password quality checking is enforced	Pass
1.0	5.3.3.2.8 Ensure password quality is enforced for the root user	Fail

W	Benchmark Item	Result
5.3.3.3 Configure pam_pwhistory module		
1.0	5.3.3.3.1 Ensure password history remember is configured	Fail
1.0	5.3.3.3.2 Ensure password history is enforced for the root user	Fail
1.0	5.3.3.3.3 Ensure pam_pwhistory includes use_authok	Fail
5.3.3.4 Configure pam_unix module		
1.0	5.3.3.4.1 Ensure pam_unix does not include nullok	Fail
1.0	5.3.3.4.2 Ensure pam_unix does not include remember	Pass
1.0	5.3.3.4.3 Ensure pam_unix includes a strong password hashing algorithm	Pass
1.0	5.3.3.4.4 Ensure pam_unix includes use_authok	Fail
5.4 User Accounts and Environment		
5.4.1 Configure shadow password suite parameters		
1.0	5.4.1.1 Ensure password expiration is configured	Fail
1.0	5.4.1.3 Ensure password expiration warning days is configured	Pass
1.0	5.4.1.4 Ensure strong password hashing algorithm is configured	Pass
1.0	5.4.1.5 Ensure inactive password lock is configured	Fail
1.0	5.4.1.6 Ensure all users last password change date is in the past	Pass
5.4.2 Configure root and system accounts and environment		
1.0	5.4.2.1 Ensure root is the only UID 0 account	Pass
1.0	5.4.2.2 Ensure root is the only GID 0 account	Pass
1.0	5.4.2.3 Ensure group root is the only GID 0 group	Pass
1.0	5.4.2.4 Ensure root account access is controlled	Pass
1.0	5.4.2.5 Ensure root path integrity	Pass
1.0	5.4.2.6 Ensure root user umask is configured	Fail
1.0	5.4.2.7 Ensure system accounts do not have a valid login shell	Pass
1.0	5.4.2.8 Ensure accounts without a valid login shell are locked	Pass
5.4.3 Configure user default environment		
1.0	5.4.3.2 Ensure default user shell timeout is configured	Fail
1.0	5.4.3.3 Ensure default user umask is configured	Fail
6 Logging and Auditing		
6.1 Configure Integrity Checking		
1.0	6.1.1 Ensure AIDE is installed	Fail
1.0	6.1.2 Ensure filesystem integrity is regularly checked	Fail
1.0	6.1.3 Ensure cryptographic mechanisms are used to protect the integrity of audit tools	Fail
6.2 System Logging		
6.2.1 Configure systemd-journald service		
1.0	6.2.1.1 Ensure journald service is enabled and active	Pass
	6.2.1.2 Ensure journald log file access is configured	Manual
	6.2.1.3 Ensure journald log file rotation is configured	Manual
6.2.2 Configure journald		
6.2.2.1 Configure systemd-journal-remote		
1.0	6.2.2.1.1 Ensure systemd-journal-remote is installed	Pass
	6.2.2.1.2 Ensure systemd-journal-upload authentication is configured	Manual
1.0	6.2.2.1.3 Ensure systemd-journal-upload is enabled and active	Pass
1.0	6.2.2.1.4 Ensure systemd-journal-remote service is not in use	Pass
1.0	6.2.2.2 Ensure journald ForwardToSyslog is disabled	Pass
1.0	6.2.2.3 Ensure journald Compress is configured	Pass
1.0	6.2.2.4 Ensure journald Storage is configured	Pass
6.2.3 Configure rsyslog		
1.0	6.2.3.1 Ensure rsyslog service is enabled and active	Pass
1.0	6.2.3.2 Ensure rsyslog is installed	Pass
1.0	6.2.3.3 Ensure journald is configured to send logs to rsyslog	Fail
1.0	6.2.3.4 Ensure rsyslog log file creation mode is configured	Pass
	6.2.3.5 Ensure rsyslog logging is configured	Manual
	6.2.3.6 Ensure rsyslog is configured to send logs to a remote log host	Manual
1.0	6.2.3.7 Ensure rsyslog is not configured to receive logs from a remote client	Pass

W	Benchmark Item	Result
	6.2.3.8 Ensure logrotate is configured	Manual
6.2.4 Configure Logfiles		
1.0	6.2.4.1 Ensure access to all logfiles has been configured	Fail
6.3 System Auditing		
6.3.1 Configure auditd Service		
6.3.2 Configure Data Retention		
6.3.3 Configure auditd Rules		
6.3.4 Configure auditd File Access		
7 System Maintenance		
7.1 Configure system file and directory access		
1.0	7.1.1 Ensure access to /etc/passwd is configured	Pass
1.0	7.1.2 Ensure access to /etc/passwd- is configured	Pass
1.0	7.1.3 Ensure access to /etc/group is configured	Pass
1.0	7.1.4 Ensure access to /etc/group- is configured	Pass
1.0	7.1.5 Ensure access to /etc/shadow is configured	Pass
1.0	7.1.6 Ensure access to /etc/shadow- is configured	Pass
1.0	7.1.7 Ensure access to /etc/gshadow is configured	Pass
1.0	7.1.8 Ensure access to /etc/gshadow- is configured	Pass
1.0	7.1.9 Ensure access to /etc/shells is configured	Pass
1.0	7.1.10 Ensure access to /etc/security/opasswd is configured	Pass
1.0	7.1.11 Ensure world writable files and directories are secured	Pass
1.0	7.1.12 Ensure no files or directories without an owner and a group exist	Pass
	7.1.13 Ensure SUID and SGID files are reviewed	Manual
7.2 Local User and Group Settings		
1.0	7.2.1 Ensure accounts in /etc/passwd use shadowed passwords	Pass
1.0	7.2.2 Ensure /etc/shadow password fields are not empty	Pass
1.0	7.2.3 Ensure all groups in /etc/passwd exist in /etc/group	Pass
1.0	7.2.4 Ensure shadow group is empty	Pass
1.0	7.2.5 Ensure no duplicate UIDs exist	Pass
1.0	7.2.6 Ensure no duplicate GIDs exist	Pass
1.0	7.2.7 Ensure no duplicate user names exist	Pass
1.0	7.2.8 Ensure no duplicate group names exist	Pass
1.0	7.2.9 Ensure local interactive user home directories are configured	Fail
1.0	7.2.10 Ensure local interactive user dot files access is configured	Fail



Assessment Details

1 Initial Setup

Items in this section are advised for all systems but may be difficult or require extensive preparation after the initial setup of the system.

1.1 Filesystem

The file system is generally a built-in layer used to handle the data management of the storage.

1.1.1 Configure Filesystem Kernel Modules

Several uncommon filesystem types are supported under Linux. Removing support for unneeded filesystem types reduces the local attack surface of the system. If a filesystem type is not needed it should be disabled. Native Linux file systems are designed to ensure that built-in security controls function as expected. Non-native filesystems can lead to unexpected consequences to both the security and functionality of the system and should be used with caution. Many filesystems are created for niche use cases and are not maintained and supported as the operating systems are updated and patched. Users of non-native filesystems should ensure that there is attention and ongoing support for them, especially in light of frequent operating system changes.

Standard network connectivity and Internet access to cloud storage may make the use of non-standard filesystem formats to directly attach heterogeneous devices much less attractive.

Note : This should not be considered a comprehensive list of filesystems. You may wish to consider additions to those

listed here for your environment. For the current available file system modules on the system see `ls /usr/lib/modules/**/kernel/fs | sort -u`

Start up scripts

Kernel modules loaded directly via `insmod` will ignore what is configured in the relevant `/etc/modprobe.d/*.conf` files. If modules are still being loaded after a reboot whilst having the correctly configured `blacklist` and `install` command, check for `insmod` entries in start up scripts such as `.bashrc`.

Return values

Using `/bin/false` as the command in disabling a particular module serves two purposes; to convey the meaning of the entry to the user and cause a non-zero return value. The latter can be tested for in scripts. Please note that `insmod` will ignore what is configured in the relevant configuration files. The preferred way to load modules is with `modprobe`.

1.1.1.1 Ensure cramfs kernel module is not available

Fail

Description:

The `cramfs` filesystem type is a compressed read-only Linux filesystem embedded in small footprint systems. A `cramfs` image can be used without having to first decompress the image.

Rationale:

Removing support for unneeded filesystem types reduces the local attack surface of the system. If this filesystem type is not needed, disable it.

Remediation:

Run the following to unload and disable the `cramfs` kernel module. This can also be done by running the script included below.

Run the following commands to unload the `cramfs` kernel module:

```
# modprobe -r cramfs 2>/dev/null
# rmmod cramfs 2>/dev/null
```

Perform the following to disable the `cramfs` kernel module:

Create a file ending in `.conf` with `install cramfs /bin/false` in the `/etc/modprobe.d/` directory

Example:

```
# printf '\n%s\n' "install cramfs /bin/false" >> cramfs.conf
```

Create a file ending in `.conf` with `blacklist cramfs` in the `/etc/modprobe.d/` directory

Example:

```
# printf '\n%s\n' "blacklist cramfs" >> cramfs.conf
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53: CM-7 a
- URL: NIST SP 800-53A :: CM-7.1 (ii)
- URL: RHEL 8 STIG Group ID: V-230498
- URL: RHEL 9 STIG Group ID: V-257880

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

1.1.1.2 Ensure freevxfs kernel module is not available

Fail

Description:

The `freevxfs` filesystem type is a free version of the Veritas type filesystem. This is the primary filesystem type for HP-UX operating systems.

Rationale:

Removing support for unneeded filesystem types reduces the local attack surface of the system. If this filesystem type is not needed, disable it.

Remediation:

Run the following to unload and disable the `freevxfs` kernel module. This can also be done by running the script included below.

Run the following commands to unload the `freevxfs` kernel module:

```
# modprobe -r freevxfs 2>/dev/null
# rmmod freevxfs 2>/dev/null
```

Perform the following to disable the `freevxfs` kernel module:

Create a file ending in `.conf` with `install freevxfs /bin/false` in the `/etc/modprobe.d/` directory

Example:

```
# printf '\n%s\n' "install freevxfs /bin/false" >> freevxfs.conf
```

Create a file ending in `.conf` with `blacklist freevxfs` in the `/etc/modprobe.d/` directory

Example:

```
# printf '\n%s\n' "blacklist freevxfs" >> freevxfs.conf
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53: CM-7 a
- URL: NIST SP 800-53A: CM-7.1 (ii)

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

1.1.1.3 Ensure hfs kernel module is not available

Fail

Description:

The `hfs` filesystem type is a hierarchical filesystem that allows you to mount Mac OS filesystems.

Rationale:

Removing support for unneeded filesystem types reduces the local attack surface of the system. If this filesystem type is not needed, disable it.

Remediation:

Run the following to unload and disable the `hfs` kernel module. This can also be done by running the script included below.

Run the following commands to unload the `hfs` kernel module:

```
# modprobe -r hfs 2>/dev/null
# rmmod hfs 2>/dev/null
```

Perform the following to disable the `hfs` kernel module:

Create a file ending in `.conf` with `install hfs /bin/false` in the `/etc/modprobe.d/` directory

Example:

```
# printf '\n%s\n' "install hfs /bin/false" >> hfs.conf
```

Create a file ending in `.conf` with `blacklist hfs` in the `/etc/modprobe.d/` directory

Example:

```
# printf '\n%s\n' "blacklist hfs" >> hfs.conf
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- [URL: NIST SP 800-53: CM-7 a](#)
- [URL: NIST SP 800-53A :: CM-7.1 \(ii\)](#)

CIS Controls V7.0:

- [Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- More](#)
>

CIS Critical Security Controls V8.0:

- [Control 4: Secure Configuration of Enterprise Assets and Software: -- More](#)
>

[Back to Summary](#)

1.1.1.4 Ensure hfsplus kernel module is not available

Fail

Description:

The `hfsplus` filesystem type is a hierarchical filesystem designed to replace `hfs` that allows you to mount Mac OS filesystems.

Rationale:

Removing support for unneeded filesystem types reduces the local attack surface of the system. If this filesystem type is not needed, disable it.

Remediation:

Run the following to unload and disable the `hfsplus` kernel module. This can also be done by running the script included below.

Run the following commands to unload the `hfsplus` kernel module:

```
# modprobe -r hfsplus 2>/dev/null
# rmmod hfsplus 2>/dev/null
```

Perform the following to disable the `hfsplus` kernel module:

Create a file ending in `.conf` with `install hfsplus /bin/false` in the `/etc/modprobe.d/` directory

Example:

```
# printf '\n%s\n' "install hfsplus /bin/false" >> hfsplus.conf
```

Create a file ending in .conf with blacklist hfsplus in the /etc/modprobe.d/ directory

Example:

```
# printf '\n%s\n' "blacklist hfsplus" >> hfsplus.conf
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-7 a
- URL: NIST SP 800-53A :: CM-7.1 (ii)

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

1.1.1.5 Ensure jffs2 kernel module is not available

Fail

Description:

The jffs2 (journaling flash filesystem 2) filesystem type is a log-structured filesystem used in flash memory devices.

Rationale:

Removing support for unneeded filesystem types reduces the local attack surface of the system. If this filesystem type is not needed, disable it.

Remediation:

Run the following to unload and disable the jffs2 kernel module. This can also be done by running the script included below.

Run the following commands to unload the jffs2 kernel module:

```
# modprobe -r jffs2 2>/dev/null
# rmmod jffs2 2>/dev/null
```

Perform the following to disable the jffs2 kernel module:

Create a file ending in .conf with install jffs2 /bin/false in the /etc/modprobe.d/ directory

Example:

```
# printf '\n%s\n' "install jffs2 /bin/false" >> jffs2.conf
```

Create a file ending in .conf with blacklist jffs2 in the /etc/modprobe.d/ directory

Example:

```
# printf '\n%s\n' "blacklist jffs2" >> jffs2.conf
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)**1.1.1.10 Ensure unused filesystems kernel modules are not available**

Manual

Description:

Filesystem kernel modules are pieces of code that can be dynamically loaded into the Linux kernel to extend its filesystem capabilities, or so-called base kernel, of an operating system. Filesystem kernel modules are typically used to add support for new hardware (as device drivers), or for adding system calls.

Rationale:

While loadable filesystem kernel modules are a convenient method of modifying the running kernel, this can be abused by attackers on a compromised system to prevent detection of their processes or files, allowing them to maintain control over the system. Many rootkits make use of loadable filesystem kernel modules in this way.

Removing support for unneeded filesystem types reduces the local attack surface of the system. If this filesystem type is not needed, disable it. The following filesystem kernel modules have known CVE's and should be made unavailable if no dependencies exist:

- `afs` - CVE-2022-37402
- `ceph` - CVE-2022-0670
- `cifs` - CVE-2022-29869
- `exfat` CVE-2022-29973
- `ext` CVE-2022-1184
- `fat` CVE-2022-22043
- `fscache` CVE-2022-3630
- `fuse` CVE-2023-0386
- `gfs2` CVE-2023-3212
- `nfs_common` CVE-2023-6660
- `nfsd` CVE-2022-43945
- `smbfs_common` CVE-2022-2585

Remediation:

- **IF** - the module is available in the running kernel:

- Unload the filesystem kernel module from the kernel
- Create a file ending in `.conf` with install filesystem kernel modules `/bin/false` in the `/etc/modprobe.d/` directory
- Create a file ending in `.conf` with deny list filesystem kernel modules in the `/etc/modprobe.d/` directory

WARNING : unloading, disabling or denylisting filesystem modules that are in use on the system maybe FATAL. It is extremely important to thoroughly review the filesystems returned by the audit before following the remediation procedure.

Example of unloading the `gfs2` kernel module:

```
# modprobe -r gfs2 2>/dev/null
# rmmod gfs2 2>/dev/null
```

Example of fully disabling the `gfs2` kernel module:

```
# printf '%s\n' "blacklist gfs2" "install gfs2 /bin/false" >> /etc/modprobe.d/gfs2.conf
```

Note:

- Disabling a kernel module by modifying the command above for each unused filesystem kernel module
- The example `gfs2` must be updated with the appropriate module name for the command or example script below to run correctly.

Below is an example Script that can be modified to use on various filesystem kernel modules manual remediation process:

Example Script

```
#!/usr/bin/env bash

{
a_output2=(); a_output3=(); l_dl="" # Initialize arrays and clear variables

l_mod_name="gfs2" # set module name

l_mod_type="fs" # set module type

l_mod_path="$(readlink -f /usr/lib/modules/**/kernel/$l_mod_type || readlink -f /lib/modules/**/kernel/$l_mod_type)"

f_module_fix()
{
l_dl="y" # Set to ignore duplicate checks

a_showconfig=() # Create array with modprobe output

while IFS= read -r l_showconfig; do
a_showconfig+=("$l_showconfig")

done < <(modprobe --showconfig | grep -P -- '\b(install|blacklist)\h+' "${l_mod_name//-//_}" '\b')

if lsmod | grep "$l_mod_name" &> /dev/null; then # Check if the module is currently loaded
a_output2+=(" - unloading kernel module: \"$l_mod_name\"")

modprobe -r "$l_mod_name" 2>/dev/null; rmmod "$l_mod_name" 2>/dev/null
fi

if ! grep -Pq -- '\binstall\h+' "${l_mod_name//-//_}" '\b' <<< "${a_showconfig[*]}"; then
a_output2+=(" - setting kernel module: \"$l_mod_name\" to \"$(readlink -f /bin/false)\"")

printf '%s\n' "install $l_mod_name $(readlink -f /bin/false)" >> /etc/modprobe.d/"$l_mod_name".conf
fi

if ! grep -Pq -- '\bblacklist\h+' "${l_mod_name//-//_}" '\b' <<< "${a_showconfig[*]}"; then
a_output2+=(" - denylisting kernel module: \"$l_mod_name\"")

printf '%s\n' "blacklist $l_mod_name" >> /etc/modprobe.d/"$l_mod_name".conf
fi
}

for l_mod_base_directory in $l_mod_path; do # Check if the module exists on the system

if [ -d "$l_mod_base_directory/${l_mod_name//-/\/}" ] && [ -n "$(ls -A "$l_mod_base_directory/${l_mod_name//-/\/}")" ]; then

a_output3+=(" - \"$l_mod_base_directory\"")

[[ "$l_mod_name" =~ overlay ]] && l_mod_name="${l_mod_name::-2}"

[ "$l_dl" != "y" ] && f_module_fix

else

echo -e " - kernel module: \"$l_mod_name\" doesn't exist in \"$l_mod_base_directory\""

fi

done

[ "${#a_output3[@]}" -gt 0 ] && printf '%s\n' "" " -- INFO --" " - module: \"$l_mod_name\" exists in:" "${a_output3[@]}"
```

```
[ "${#a_output2[@]}" -gt 0 ] && printf '%s\n' "" "${a_output2[@]}" || printf '%s\n' "" " - No changes needed"

printf '%s\n' "" " - remediation of kernel module: \"\$1_mod_name\" complete" ""
}
```

Impact:

This list may be quite extensive and covering all edges cases is difficult. Therefore, it's crucial to carefully consider the implications and dependencies before making any changes to the filesystem kernel module configurations.

[Show Rule Result XML](#)

References:

- URL: <https://cve.mitre.org/cgi-bin/cvekey.cgi?keyword=filesystem>

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

1.1.2 Configure Filesystem Partitions

Directories that are used for system-wide functions can be further protected by placing them on separate partitions. This provides protection for resource exhaustion and enables the use of mounting options that are applicable to the directory's intended use. Users' data can be stored on separate partitions and have stricter mount options. A user partition is a filesystem that has been established for use by the users and does not contain software for system operations.

The recommendations in this section are easier to perform during initial system installation. If the system is already installed, it is recommended that a full backup be performed before repartitioning the system.

Note:

-IF- you are repartitioning a system that has already been installed (This may require the system to be in single-user mode):

- Mount the new partition to a temporary mountpoint e.g. `mount /dev/sda2 /mnt`
- Copy data from the original partition to the new partition. e.g. `cp -a /var/tmp/* /mnt`
- Verify that all data is present on the new partition. e.g. `ls -la /mnt`
- Unmount the new partition. e.g. `umount /mnt`
- Remove the data from the original directory that was in the old partition. e.g. `rm -Rf /var/tmp/*` Otherwise it will still consume space in the old partition that will be masked when the new filesystem is mounted.
- Mount the new partition to the desired mountpoint. e.g. `mount /dev/sda2 /var/tmp`
- Update `/etc/fstab` with the new mountpoint. e.g. `/dev/sda2 /var/tmp xfs defaults,rw,nosuid,nodev,noexec,relatime 0 0`

1.1.2.1 Configure /tmp

The `/tmp` directory is a world-writable directory used to store data used by the system and user applications for a short period of time. This data should have no expectation of surviving a reboot, as this directory is intended to be emptied after each reboot.

1.1.2.1.1 Ensure /tmp is a separate partition

Fail

Description:

The `/tmp` directory is a world-writable directory used for temporary storage by all users and some applications.

Note: If an entry for `/tmp` exists in `/etc/fstab` it will take precedence over entries in a systemd unit file. There is an exception to this when a system is diskless and connected to iSCSI, entries in `/etc/fstab` may not take precedence

over a systemd unit file.

Rationale:

Making `/tmp` its own file system allows an administrator to set additional mount options such as the `noexec` option on the mount, making `/tmp` useless for an attacker to install executable code. It would also prevent an attacker from establishing a hard link to a system `setuid` program and wait for it to be updated. Once the program was updated, the hard link would be broken, and the attacker would have his own copy of the program. If the program happened to have a security vulnerability, the attacker could continue to exploit the known flaw.

This can be accomplished by either mounting `tmpfs` to `/tmp` , or creating a separate partition for `/tmp` .

Remediation:

First ensure that systemd is correctly configured to ensure that `/tmp` will be mounted at boot time.

```
# systemctl unmask tmp.mount
```

For specific configuration requirements of the `/tmp` mount for your environment, modify `/etc/fstab` .

Example of using `tmpfs` with specific mount options:

```
tmpfs /tmp tmpfs defaults,rw,nosuid,nodev,noexec,relatime,size=2G 0 0
```

Note: the `size=2G` is an example of setting a specific size for `tmpfs` .

Example of using a volume or disk with specific mount options. The source location of the volume or disk will vary depending on your environment:

```
<device> /tmp <fstype> defaults,nodev,nosuid,noexec 0 0
```

Impact:

By design files saved to `/tmp` should have no expectation of surviving a reboot of the system. `tmpfs` is ram based and all files stored to `tmpfs` will be lost when the system is rebooted.

If files need to be persistent through a reboot, they should be saved to `/var/tmp` not `/tmp` .

Since the `/tmp` directory is intended to be world-writable, there is a risk of resource exhaustion if it is not bound to `tmpfs` or a separate partition.

Running out of `/tmp` space is a problem regardless of what kind of filesystem lies under it, but in a configuration where `/tmp` is not a separate file system it will essentially have the whole disk available, as the default installation only creates a single `/` partition. On the other hand, a RAM-based `/tmp` (as with `tmpfs`) will almost certainly be much smaller, which can lead to applications filling up the filesystem much more easily. Another alternative is to create a dedicated partition for `/tmp` from a separate volume or disk. One of the downsides of a disk-based dedicated partition is that it will be slower than `tmpfs` which is RAM-based.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** <https://www.freedesktop.org/wiki/Software/systemd/APIFileSystems/>
- **URL:** <https://www.freedesktop.org/software/systemd/man/systemd-fstab-generator.html>
- **URL:** <https://www.kernel.org/doc/Documentation/filesystems/tmpfs.txt>
- **URL:** NIST SP 800-53 Rev. 5: CM-7
- **URL:** RHEL 8 STIG Vul ID: V-230295
- **URL:** RHEL 8 Rule ID: SV-30295r627750

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

1.1.2.1.2 Ensure nodev option set on /tmp partition

Pass

Description:

The `nodev` mount option specifies that the filesystem cannot contain special devices.

Rationale:

Since the `/tmp` filesystem is not intended to support devices, set this option to ensure that users cannot create a block or character special devices in `/tmp` .

Remediation:

- **IF** - a separate partition exists for `/tmp` .

Edit the `/etc/fstab` file and add `nodev` to the fourth field (mounting options) for the `/tmp` partition.

Example:

```
<device> /tmp <fstype> defaults,rw,nosuid,nodev,noexec,relatime 0 0
```

Run the following command to remount `/tmp` with the configured options:

```
# mount -o remount /tmp
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** See the `fstab(5)` manual page for more information.
- **URL:** NIST SP 800-53 Rev. 5: CM-7
- **URL:** NIST SP 800-53 Revision 5 :: CM-7 (2)
- **URL:** RHEL 8 STIG Vul ID: V-230511
- **URL:** RHEL 8 STIG Rule ID: SV-230511r854052

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

1.1.2.1.3 Ensure nosuid option set on /tmp partition

Pass

Description:

The `nosuid` mount option specifies that the filesystem cannot contain `setuid` files.

Rationale:

Since the `/tmp` filesystem is only intended for temporary file storage, set this option to ensure that users cannot create `setuid` files in `/tmp` .

Remediation:

- **IF** - a separate partition exists for `/tmp` .

Edit the `/etc/fstab` file and add `nosuid` to the fourth field (mounting options) for the `/tmp` partition.

Example:

```
<device> /tmp <fstype> defaults,rw,nosuid,nodev,noexec,relatime 0 0
```

Run the following command to remount `/tmp` with the configured options:

```
# mount -o remount /tmp
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** See the fstab(5) manual page for more information.
- **URL:** NIST SP 800-53 Rev. 5: AC-3, MP-2
- **URL:** RHEL 8 STIG Vul ID: V-230512
- **URL:** RHEL 8 STIG Rule ID: SV-230512r854053

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

1.1.2.1.4 Ensure noexec option set on /tmp partition

Pass

Description:

The `noexec` mount option specifies that the filesystem cannot contain executable binaries.

Rationale:

Since the `/tmp` filesystem is only intended for temporary file storage, set this option to ensure that users cannot run executable binaries from `/tmp`.

Remediation:

- **IF** - a separate partition exists for `/tmp`.

Edit the `/etc/fstab` file and add `noexec` to the fourth field (mounting options) for the `/tmp` partition.

Example:

```
<device> /tmp <fstype> defaults,rw,nosuid,nodev,noexec,relatime 0 0
```

Run the following command to remount `/tmp` with the configured options:

```
# mount -o remount /tmp
```

Impact:

Setting the `noexec` option on `/tmp` may prevent installation and/or updating of some 3rd party software.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** See the fstab(5) manual page for more information.
- **URL:** NIST SP 800-53 Rev. 5: AC-3, MP-2
- **URL:** NIST SP 800-53 Revision 5 :: CM-7 (2)
- **URL:** RHEL 8 STIG Vul ID: V-230513
- **URL:** RHEL 8 STIG Rule ID: SV-230513r854054

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

1.1.2.2 Configure /dev/shm

The `/dev/shm` directory is a world-writable directory that can function as shared memory that facilitates inter process communication (IPC)

1.1.2.2.1 Ensure /dev/shm is a separate partition Pass

Description:

The `/dev/shm` directory is a world-writable directory that can function as shared memory that facilitates inter process communication (IPC).

Rationale:

Making `/dev/shm` its own file system allows an administrator to set additional mount options such as the `noexec` option on the mount, making `/dev/shm` useless for an attacker to install executable code. It would also prevent an attacker from establishing a hard link to a system `setuid` program and wait for it to be updated. Once the program was updated, the hard link would be broken and the attacker would have his own copy of the program. If the program happened to have a security vulnerability, the attacker could continue to exploit the known flaw.

This can be accomplished by mounting `tmpfs` to `/dev/shm`.

Remediation:

For specific configuration requirements of the `/dev/shm` mount for your environment, modify `/etc/fstab`.

Example:

```
tmpfs /dev/shm tmpfs defaults,rw,nosuid,nodev,noexec,relatime,size=2G 0 0
```

Impact:

Since the `/dev/shm` directory is intended to be world-writable, there is a risk of resource exhaustion if it is not bound to a separate partition.

`/dev/shm` utilizing `tmpfs` can be resized using the `size={size}` parameter in the relevant entry in `/etc/fstab`.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: <https://www.freedesktop.org/wiki/Software/systemd/APIFileSystems/>
- URL: <https://www.freedesktop.org/software/systemd/man/systemd-fstab-generator.html>
- URL: NIST SP 800-53 Rev. 5: CM-7

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

1.1.2.2.2 Ensure nodev option set on /dev/shm partition Pass

Description:

The `nodev` mount option specifies that the filesystem cannot contain special devices.

Rationale:

Since the `/dev/shm` filesystem is not intended to support devices, set this option to ensure that users cannot attempt to create special devices in `/dev/shm` partitions.

Remediation:

- **IF** - a separate partition exists for `/dev/shm` .

Edit the `/etc/fstab` file and add `nodev` to the fourth field (mounting options) for the `/dev/shm` partition. See the `fstab(5)` manual page for more information.

Example:

```
tmpfs /dev/shm tmpfs defaults,rw,nosuid,nodev,noexec,relatime 0 0
```

Run the following command to remount `/dev/shm` with the configured options:

```
# mount -o remount /dev/shm
```

Note: It is recommended to use `tmpfs` as the device/filesystem type as `/dev/shm` is used as shared memory space by applications.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: AC-3, MP-2
- **URL:** NIST SP 800-53 Revision 5 :: CM-7 (2)
- **URL:** RHEL 8 STIG Vul ID: V-230508
- **URL:** RHEL 8 STIG Rule ID: SV-230508r854049

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

1.1.2.2.3 Ensure nosuid option set on /dev/shm partition

Pass

Description:

The `nosuid` mount option specifies that the filesystem cannot contain `setuid` files.

Rationale:

Setting this option on a file system prevents users from introducing privileged programs onto the system and allowing non-root users to execute them.

Remediation:

- **IF** - a separate partition exists for `/dev/shm` .

Edit the `/etc/fstab` file and add `nosuid` to the fourth field (mounting options) for the `/dev/shm` partition. See the `fstab(5)` manual page for more information.

Example:

```
tmpfs /dev/shm tmpfs defaults,rw,nosuid,nodev,noexec,relatime 0 0
```

Run the following command to remount `/dev/shm` with the configured options:

```
# mount -o remount /dev/shm
```

Note: It is recommended to use `tmpfs` as the device/filesystem type as `/dev/shm` is used as shared memory space by applications.

Assessment:
[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: AC-3, MP-2
- **URL:** NIST SP 800-53 Revision 5 :: CM-7 (2)
- **URL:** RHEL 8 STIG Vul ID: V-230509
- **URL:** RHEL 8 STIG Rule ID: SV-230509r854050

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

1.1.2.2.4 Ensure noexec option set on /dev/shm partition

Fail

Description:

The `noexec` mount option specifies that the filesystem cannot contain executable binaries.

Rationale:

Setting this option on a file system prevents users from executing programs from shared memory. This deters users from introducing potentially malicious software on the system.

Remediation:

- **IF** - a separate partition exists for `/dev/shm` .

Edit the `/etc/fstab` file and add `noexec` to the fourth field (mounting options) for the `/dev/shm` partition.

Example:

```
tmpfs /dev/shm tmpfs defaults,rw,nosuid,nodev,noexec,relatime 0 0
```

Run the following command to remount `/dev/shm` with the configured options:

```
# mount -o remount /dev/shm
```

Note: It is recommended to use `tmpfs` as the device/filesystem type as `/dev/shm` is used as shared memory space by applications.

Assessment:
[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** See the `fstab(5)` manual page for more information.
- **URL:** NIST SP 800-53 Rev. 5: AC-3, MP-2
- **URL:** NIST SP 800-53 Revision 5 :: CM-7 (2)
- **URL:** RHEL 8 STIG Vul ID: V-230510
- **URL:** RHEL 8 STIG Rule ID: SV-230510r854051

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

1.1.2.3 Configure /home

Please note that home directories can be mounted anywhere and are not necessarily restricted to /home , nor restricted to a single location, nor is the name restricted in any way.

Finding user home directories can be done by looking in /etc/passwd , looking over the mounted file systems with mount or querying the relevant database with getent .

The following script can be ran to find user's home directories:

```
#!/usr/bin/env bash

{

l_valid_shells="^($(awk -F\| ' $NF != "nologin" {print}' /etc/shells | sed -rn '/^\|/{s/,/\|/g;p}' | paste
-s -d '|' - ))$"

awk -v pat="$l_valid_shells" -F: '($1!~/^(root|halt|sync|shutdown|nfsnobody)$/ && ($3>="$(awk '/
^s*UID_MIN/{print $2}' /etc/login.defs)'" || $3 != 65534) && $(NF) ~ pat) {print $1 " - " $6}' /etc/passwd
}
```

1.1.2.3.2 Ensure nodev option set on /home partition

Pass

Description:

The nodev mount option specifies that the filesystem cannot contain special devices.

Rationale:

Since the /home filesystem is not intended to support devices, set this option to ensure that users cannot create a block or character special devices in /home .

Remediation:

- **IF** - a separate partition exists for /home .

Edit the /etc/fstab file and add nodev to the fourth field (mounting options) for the /home partition.

Example:

```
<device> /home <fstype> defaults,rw,nosuid,nodev,relatime 0 0
```

Run the following command to remount /home with the configured options:

```
# mount -o remount /home
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: See the fstab(5) manual page for more information.
- URL: NIST SP 800-53 Rev. 5: AC-3, MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

1.1.2.3.3 Ensure nosuid option set on /home partition

Pass

Description:

The `nosuid` mount option specifies that the filesystem cannot contain `setuid` files.

Rationale:

Since the `/home` filesystem is only intended for user file storage, set this option to ensure that users cannot create `setuid` files in `/home`.

Remediation:

- **IF** - a separate partition exists for `/home`.

Edit the `/etc/fstab` file and add `nosuid` to the fourth field (mounting options) for the `/home` partition.

Example:

```
<device> /home <fstype> defaults,rw,nosuid,nodev,relatime 0 0
```

Run the following command to remount `/home` with the configured options:

```
# mount -o remount /home
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** See the `fstab(5)` manual page for more information.
- **URL:** NIST SP 800-53 Rev. 5: AC-3, MP-2

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

1.1.2.4 Configure /var

The `/var` directory is used by daemons and other system services to temporarily store dynamic data. Some directories created by these processes may be world-writable.

1.1.2.4.2 Ensure nodev option set on /var partition

Pass

Description:

The `nodev` mount option specifies that the filesystem cannot contain special devices.

Rationale:

Since the `/var` filesystem is not intended to support devices, set this option to ensure that users cannot create a block or character special devices in `/var`.

Remediation:

- **IF** - a separate partition exists for `/var`.

Edit the `/etc/fstab` file and add `nodev` to the fourth field (mounting options) for the `/var` partition.

Example:

```
<device> /var <fstype> defaults,rw,nosuid,nodev,relatime 0 0
```

Run the following command to remount `/var` with the configured options:

```
# mount -o remount /var
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** See the `fstab(5)` manual page for more information.
- **URL:** NIST SP 800-53 Rev. 5: AC-3, MP-2

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

1.1.2.4.3 Ensure nosuid option set on /var partition

Pass

Description:

The `nosuid` mount option specifies that the filesystem cannot contain `setuid` files.

Rationale:

Since the `/var` filesystem is only intended for variable files such as logs, set this option to ensure that users cannot create `setuid` files in `/var`.

Remediation:

- **IF** - a separate partition exists for `/var`.

Edit the `/etc/fstab` file and add `nosuid` to the fourth field (mounting options) for the `/var` partition.

Example:

```
<device> /var <fstype> defaults,rw,nosuid,nodev,relatime 0 0
```

Run the following command to remount `/var` with the configured options:

```
# mount -o remount /var
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** See the `fstab(5)` manual page for more information.
- **URL:** NIST SP 800-53 Rev. 5: AC-3, MP-2

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

1.1.2.5 Configure /var/tmp

The `/var/tmp` directory is a world-writable directory used for temporary storage by all users and some applications. Temporary files residing in `/var/tmp` are to be preserved between reboots.

1.1.2.5.2 Ensure nodev option set on /var/tmp partition

Pass

Description:

The `nodev` mount option specifies that the filesystem cannot contain special devices.

Rationale:

Since the `/var/tmp` filesystem is not intended to support devices, set this option to ensure that users cannot create a block or character special devices in `/var/tmp`.

Remediation:

- **IF** - a separate partition exists for `/var/tmp`.

Edit the `/etc/fstab` file and add `nodev` to the fourth field (mounting options) for the `/var/tmp` partition.

Example:

```
<device> /var/tmp <fstype> defaults,rw,nosuid,nodev,noexec,relatime 0 0
```

Run the following command to remount `/var/tmp` with the configured options:

```
# mount -o remount /var/tmp
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** See the `fstab(5)` manual page for more information.
- **URL:** NIST SP 800-53 Rev. 5: AC-3, MP-2
- **URL:** NIST SP 800-53 Revision 5 :: CM-7 (2)
- **URL:** RHEL 8 STIG Vul ID: V-230520
- **URL:** RHEL 8 STIG Rule ID: SV-230520r854061

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

1.1.2.5.3 Ensure nosuid option set on /var/tmp partition

Pass

Description:

The `nosuid` mount option specifies that the filesystem cannot contain `setuid` files.

Rationale:

Since the `/var/tmp` filesystem is only intended for temporary file storage, set this option to ensure that users cannot create `setuid` files in `/var/tmp`.

Remediation:

- **IF** - a separate partition exists for `/var/tmp`.

Edit the `/etc/fstab` file and add `nosuid` to the fourth field (mounting options) for the `/var/tmp` partition.

Example:

```
<device> /var/tmp <fstype> defaults,rw,nosuid,nodev,noexec,relatime 0 0
```

Run the following command to remount `/var/tmp` with the configured options:

```
# mount -o remount /var/tmp
```

Assessment:
[Show](#) Assessment Evidence

[Show](#) Rule Result XML

- References:**
- **URL:** See the `fstab(5)` manual page for more information.
 - **URL:** NIST SP 800-53 Rev. 5: AC-3, MP-2
 - **URL:** NIST SP 800-53 Revision 5 :: CM-7 (2)
 - **URL:** RHEL 8 STIG Vul ID: V-230521
 - **URL:** RHEL 8 STIG STIG ID: RHEL-08-040133

- CIS Controls V7.0:**
- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

- CIS Critical Security Controls V8.0:**
- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

1.1.2.5.4 Ensure noexec option set on /var/tmp partitionPass

Description:

The `noexec` mount option specifies that the filesystem cannot contain executable binaries.

Rationale:

Since the `/var/tmp` filesystem is only intended for temporary file storage, set this option to ensure that users cannot run executable binaries from `/var/tmp`.

Remediation:

- **IF** - a separate partition exists for `/var/tmp`.

Edit the `/etc/fstab` file and add `noexec` to the fourth field (mounting options) for the `/var/tmp` partition.

Example:

```
<device> /var/tmp <fstype> defaults,rw,nosuid,nodev,noexec,relatime 0 0
```

Run the following command to remount `/var/tmp` with the configured options:

```
# mount -o remount /var/tmp
```

Assessment:
[Show](#) Assessment Evidence

[Show](#) Rule Result XML

- References:**
- **URL:** See the `fstab(5)` manual page for more information.
 - **URL:** NIST SP 800-53 Rev. 5: AC-3, MP-2
 - **URL:** NIST SP 800-53 Revision 5 :: CM-7 (2)
 - **URL:** RHEL 8 STIG Vul ID: V-230522
 - **URL:** RHEL 8 STIG Rule ID: SV-230522r854063

- CIS Controls V7.0:**

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

1.1.2.6 Configure /var/log

The `/var/log` directory is used by system services to store log data.

1.1.2.6.2 Ensure nodev option set on /var/log partition

Pass

Description:

The `nodev` mount option specifies that the filesystem cannot contain special devices.

Rationale:

Since the `/var/log` filesystem is not intended to support devices, set this option to ensure that users cannot create a block or character special devices in `/var/log`.

Remediation:

- **IF** - a separate partition exists for `/var/log`.

Edit the `/etc/fstab` file and add `nodev` to the fourth field (mounting options) for the `/var/log` partition.

Example:

```
<device> /var/log <fstype> defaults,rw,nosuid,nodev,noexec,relatime 0 0
```

Run the following command to remount `/var/log` with the configured options:

```
# mount -o remount /var/log
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** See the `fstab(5)` manual page for more information.
- **URL:** NIST SP 800-53 Rev. 5: AC-3, MP-2
- **URL:** NIST SP 800-53 Revision 5 :: CM-7 (2)
- **URL:** RHEL 8 STIG Vul ID: V-230514
- **URL:** RHEL 8 STIG Rule ID: SV-230514r854055

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

1.1.2.6.3 Ensure nosuid option set on /var/log partition

Pass

Description:

The `nosuid` mount option specifies that the filesystem cannot contain `setuid` files.

Rationale:

Since the `/var/log` filesystem is only intended for log files, set this option to ensure that users cannot create `setuid`

files in `/var/log`.

Remediation:

- **IF** - a separate partition exists for `/var/log`.

Edit the `/etc/fstab` file and add `nosuid` to the fourth field (mounting options) for the `/var/log` partition.

Example:

```
<device> /var/log <fstype> defaults,rw,nosuid,nodev,noexec,relatime 0 0
```

Run the following command to remount `/var/log` with the configured options:

```
# mount -o remount /var/log
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** See the `fstab(5)` manual page for more information.
- **URL:** NIST SP 800-53 Rev. 5: AC-3, MP-2
- **URL:** NIST SP 800-53 Revision 5 :: CM-7 (2)
- **URL:** RHEL 8 STIG Vul ID: V-230515
- **URL:** RHEL 8 STIG Rule ID: SV-230515r854056

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

1.1.2.6.4 Ensure noexec option set on /var/log partition

Pass

Description:

The `noexec` mount option specifies that the filesystem cannot contain executable binaries.

Rationale:

Since the `/var/log` filesystem is only intended for log files, set this option to ensure that users cannot run executable binaries from `/var/log`.

Remediation:

- **IF** - a separate partition exists for `/var/log`.

Edit the `/etc/fstab` file and add `noexec` to the fourth field (mounting options) for the `/var/log` partition.

Example:

```
<device> /var/log <fstype> defaults,rw,nosuid,nodev,noexec,relatime 0 0
```

Run the following command to remount `/var/log` with the configured options:

```
# mount -o remount /var/log
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** See the `fstab(5)` manual page for more information.
- **URL:** NIST SP 800-53 Rev. 5: AC-3, MP-2
- **URL:** NIST SP 800-53 Revision 5 :: CM-7 (2)
- **URL:** RHEL 8 STIG Vul ID: V-230516
- **URL:** RHEL 8 STIG Rule ID: SV-230516r854057

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

1.1.2.7 Configure `/var/log/audit`

The auditing daemon, `auditd`, stores log data in the `/var/log/audit` directory.

1.1.2.7.2 Ensure `nodev` option set on `/var/log/audit` partition

Pass

Description:

The `nodev` mount option specifies that the filesystem cannot contain special devices.

Rationale:

Since the `/var/log/audit` filesystem is not intended to support devices, set this option to ensure that users cannot create a block or character special devices in `/var/log/audit`.

Remediation:

- **IF** - a separate partition exists for `/var/log/audit`.

Edit the `/etc/fstab` file and add `nodev` to the fourth field (mounting options) for the `/var/log/audit` partition.

Example:

```
<device> /var/log/audit <fstype> defaults,rw,nosuid,nodev,noexec,relatime 0 0
```

Run the following command to remount `/var/log/audit` with the configured options:

```
# mount -o remount /var/log/audit
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** See the `fstab(5)` manual page for more information.
- **URL:** NIST SP 800-53 Revision 5 :: CM-7 (2)
- **URL:** STIG ID: RHEL-08-040129 | RULE ID: SV-230517r958804 | CAT II
- **URL:** STIG ID: RHEL-09-231175 | RULE ID: SV-257876r958804 | CAT II
- **URL:** STIG ID: ALMA-09-027190 | RULE ID: SV-269320r1050202 | CAT II

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

1.1.2.7.3 Ensure `nosuid` option set on `/var/log/audit` partition

Pass

Description:

The `nosuid` mount option specifies that the filesystem cannot contain `setuid` files.

Rationale:

Since the `/var/log/audit` filesystem is only intended for variable files such as logs, set this option to ensure that users cannot create `setuid` files in `/var/log/audit`.

Remediation:

- **IF** - a separate partition exists for `/var/log/audit`.

Edit the `/etc/fstab` file and add `nosuid` to the fourth field (mounting options) for the `/var/log/audit` partition.

Example:

```
<device> /var/log/audit <fstype> defaults,rw,nosuid,nodev,noexec,relatime 0 0
```

Run the following command to remount `/var/log/audit` with the configured options:

```
# mount -o remount /var/log/audit
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** See the `fstab(5)` manual page for more information.
- **URL:** NIST SP 800-53 Rev. 5: AC-3, MP-2
- **URL:** NIST SP 800-53 Revision 5 :: CM-7 (2)
- **URL:** STIG ID: RHEL-08-040130 | RULE ID: SV-230518r958804 | CAT II
- **URL:** STIG ID: RHEL-09-231170 | RULE ID: SV-257875r958804 | CAT II

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

1.1.2.7.4 Ensure noexec option set on /var/log/audit partition

Pass

Description:

The `noexec` mount option specifies that the filesystem cannot contain executable binaries.

Rationale:

Since the `/var/log/audit` filesystem is only intended for audit logs, set this option to ensure that users cannot run executable binaries from `/var/log/audit`.

Remediation:

- **IF** - a separate partition exists for `/var/log/audit`.

Edit the `/etc/fstab` file and add `noexec` to the fourth field (mounting options) for the `/var/log/audit` partition.

Example:

```
<device> /var/log/audit <fstype> defaults,rw,nosuid,nodev,noexec,relatime 0 0
```

Run the following command to remount `/var/log/audit` with the configured options:

```
# mount -o remount /var/log/audit
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** See the fstab(5) manual page for more information.
- **URL:** NIST SP 800-53 Rev. 5: AC-3, MP-2
- **URL:** NIST SP 800-53 Revision 5 :: CM-7 (2)
- **URL:** STIG ID: RHEL-08-040131 | RULE ID: SV-230519r958804 | CAT II
- **URL:** STIG ID: RHEL-09-231165 | RULE ID: SV-257874r958804 | CAT II

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

1.2 Package Management

Patch management procedures may vary widely between enterprises. Large enterprises may choose to install a local updates server that can be used in place of their distributions servers, whereas a single deployment of a system may prefer to get updates directly. Updates can be performed automatically or manually, depending on the site's policy for patch management. Organizations may prefer to test patches against their environment on a non-production system before rolling out to production.

Outdated software is vulnerable to cyber criminals and hackers. Software updates help reduce the risk to your organization. The release of software update notes often reveals the patched exploitable entry points to the public. Public knowledge of these exploits can make your organization more vulnerable to malicious actors attempting to gain entry to your system's data.

Software updates often offer new and improved features and speed enhancements.

For the purpose of this benchmark, the requirement is to ensure that a patch management process is defined and maintained, the specifics of which are left to the organization.

1.2.1 Configure Package Repositories

Patch management procedures may vary widely between enterprises. Large enterprises may choose to install a local updates server that can be used in place of their distributions servers, whereas a single deployment of a system may prefer to get updates directly. Updates can be performed automatically or manually, depending on the site's policy for patch management. Organizations may prefer to test patches against their environment on a non-production system before rolling out to production.

Outdated software is vulnerable to cyber criminals and hackers. Software updates help reduce the risk to your organization. The release of software update notes often reveals the patched exploitable entry points to the public. Public knowledge of these exploits can leave your organization more vulnerable to malicious actors attempting to gain access to your system's data.

Note: Creation of an appropriate patch management policy is left to the organization.

1.2.1.1 Ensure GPG keys are configured

Manual

Description:

Most package managers implement GPG key signing to verify package integrity during installation.

Rationale:

It is important to ensure that updates are obtained from a valid source to protect against spoofing that could lead to the inadvertent installation of malware on the system.

Remediation:

Update your package manager GPG keys in accordance with site policy.

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: SI-2
- URL: <https://manpages.debian.org/stretch/apt/sources.list.5.en.html>

CIS Controls V7.0:

- Control 3: Continuous Vulnerability Management: -- [More](#)
 - Control 3: Continuous Vulnerability Management: -- [More](#)
- >

CIS Critical Security Controls V8.0:

- Control 7: Continuous Vulnerability Management: -- [More](#)
 - Control 7: Continuous Vulnerability Management: -- [More](#)
- >

[Back to Summary](#)

1.2.1.2 Ensure package manager repositories are configured

Manual

Description:

Systems need to have package manager repositories configured to ensure they receive the latest patches and updates.

Rationale:

If a system's package repositories are misconfigured important patches may not be identified or a rogue repository could introduce compromised software.

Remediation:

Configure your package manager repositories according to site policy.

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: SI-2

CIS Controls V7.0:

- Control 3: Continuous Vulnerability Management: -- [More](#)
 - Control 3: Continuous Vulnerability Management: -- [More](#)
- >

CIS Critical Security Controls V8.0:

- Control 7: Continuous Vulnerability Management: -- [More](#)
 - Control 7: Continuous Vulnerability Management: -- [More](#)
- >

[Back to Summary](#)

1.2.2 Configure Package Updates

1.2.2.1 Ensure updates, patches, and additional security software are installed

Manual

Description:

Periodically patches are released for included software either due to security flaws or to include additional functionality.

Rationale:

Newer patches may contain security enhancements that would not be available through the latest full update. As a result, it is recommended that the latest software patches be used to take advantage of the latest functionality. As with any software installation, organizations need to determine if a given update meets their requirements and verify the compatibility and supportability of any additional software against the update revision that is selected.

Remediation:

Run the following commands to update all packages following local site policy guidance on applying updates and patches:

Run the following command to update the system with the available patches and updates:

```
# apt update
```

Run one of the following commands to apply the updates and patches:

```
# apt upgrade
```

- OR -

```
# apt dist-upgrade
```

Note: When running the command `apt dist-upgrade` that apt has a "smart" conflict resolution system, and it will attempt to upgrade the most important packages at the expense of less important ones if necessary. So, `dist-upgrade` command may remove some packages.

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: SI-2

CIS Controls V7.0:

- Control 3: Continuous Vulnerability Management: -- [More](#)
- Control 3: Continuous Vulnerability Management: -- [More](#)

>

CIS Critical Security Controls V8.0:

- Control 7: Continuous Vulnerability Management: -- [More](#)
- Control 7: Continuous Vulnerability Management: -- [More](#)

>

[Back to Summary](#)

1.3 Mandatory Access Control

Mandatory Access Control (MAC) provides an additional layer of access restrictions to processes on top of the base Discretionary Access Controls. By restricting how processes can access files and resources on a system the potential impact from vulnerabilities in the processes can be reduced.

Impact: Mandatory Access Control limits the capabilities of applications and daemons on a system, while this can prevent unauthorized access the configuration of MAC can be complex and difficult to implement correctly preventing legitimate access from occurring.

1.3.1 Configure AppArmor

AppArmor provides a Mandatory Access Control (MAC) system that greatly augments the default Discretionary Access Control (DAC) model. Under AppArmor MAC rules are applied by file paths instead of by security contexts as in other MAC systems. As such it does not require support in the filesystem and can be applied to network mounted filesystems for example. AppArmor security policies define what system resources applications can access and what privileges they can do so with. This automatically limits the damage that the software can do to files accessible by the calling user. The user does not need to take any action to gain this benefit. For an action to occur, both the traditional DAC permissions must be satisfied as well as the AppArmor MAC rules. The action will not be allowed if either one of these models does not permit the action. In this way, AppArmor rules can only make a system's permissions more restrictive and secure.

References:

1. AppArmor Documentation: <http://wiki.apparmor.net/index.php/Documentation>
2. Ubuntu AppArmor Documentation: <https://help.ubuntu.com/community/AppArmor>

1.3.1.1 Ensure latest versions of the apparmor packages are installed

Fail

Description:

AppArmor provides Mandatory Access Controls.

Rationale:

Without a Mandatory Access Control system installed only the default Discretionary Access Control system will be available.

Remediation:

Run the following command to install the latest versions of `apparmor` and `apparmor-utils` :

```
# apt install apparmor apparmor-utils
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: AC-3

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

1.3.1.2 Ensure AppArmor is enabled in the bootloader configuration

Fail

Description:

Configure AppArmor to be enabled at boot time and verify that it has not been overwritten by the bootloader boot parameters.

Note: This recommendation is designed around the grub bootloader, if LILO or another bootloader is in use in your environment enact equivalent settings.

Rationale:

AppArmor must be enabled at boot time in your bootloader configuration to ensure that the controls it provides are not overridden.

Remediation:

Edit `/etc/default/grub` and add the `apparmor=1` and `security=apparmor` parameters to the `GRUB_CMDLINE_LINUX=` line

```
GRUB_CMDLINE_LINUX="apparmor=1 security=apparmor"
```

Run the following command to update the `grub2` configuration:

```
# update-grub
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: AC-3

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

1.3.1.3 Ensure all AppArmor Profiles are not disabled

Pass

Description:

AppArmor profiles define what resources applications are able to access.

Rationale:

Security configuration requirements vary from site to site. Some sites may mandate a policy that is stricter than the default policy, which is perfectly acceptable. This item is intended to ensure that any policies that exist on the system are activated.

Remediation:

Run the following command to set all profiles to enforce mode:

```
# aa-enforce /etc/apparmor.d/*
```

- OR -

Run the following command to set all profiles to complain mode:

```
# aa-complain /etc/apparmor.d/*
```

Note: Any unconfined processes may need to have a profile created or activated for them and then be restarted

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AC-3

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

1.4 Configure Bootloader

The recommendations in this section focus on securing the bootloader and settings involved in the boot process directly.

1.4.1 Ensure bootloader password is set

Fail

Description:

Setting the boot loader password will require that anyone rebooting the system must enter a password before being able to set command line boot parameters

Rationale:

Requiring a boot password upon execution of the boot loader will prevent an unauthorized user from entering boot parameters or changing the boot partition. This prevents users from weakening security (e.g. turning off AppArmor at boot time).

Remediation:

Create an encrypted password with `grub-mkpasswd-pbkdf2` :

```
# grub-mkpasswd-pbkdf2 --iteration-count=600000 --salt=64
```

```
Enter password: <password>
```

```
Reenter password: <password>
```

```
PBKDF2 hash of your password is <encrypted-password>
```

Add the following into a custom `/etc/grub.d` configuration file:

```
set superusers="<username>"
password_pbkdf2 <username> <encrypted-password>
```

The superuser/user information and password should not be contained in the `/etc/grub.d/00_header` file as this file could be overwritten in a package update.

If there is a requirement to be able to boot/reboot without entering the password, edit `/etc/grub.d/10_linux` and add `--unrestricted` to the line `CLASS=`

Example:

```
CLASS="--class gnu-linux --class gnu --class os --unrestricted"
```

Run the following command to update the grub2 configuration:

```
# update-grub
```

Impact:

If password protection is enabled, only the designated superuser can edit a GRUB 2 menu item by pressing "e" or access the GRUB 2 command line by pressing "c"

If GRUB 2 is set up to boot automatically to a password-protected menu entry the user has no option to back out of the password prompt to select another menu entry. Holding the SHIFT key will not display the menu in this case. The user must enter the correct username and password. If unable to do so, the configuration files will have to be edited via a LiveCD or other means to fix the problem

You can add `--unrestricted` to the menu entries to allow the system to boot without entering a password. A password will still be required to edit menu items.

More Information: <https://help.ubuntu.com/community/Grub2/Passwords>

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: AC-3
- **URL:** NIST SP 800-53A :: AC-3.1

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

1.4.2 Ensure access to bootloader config is configured

Fail

Description:

The grub configuration file contains information on boot settings and passwords for unlocking boot options.

Rationale:

Setting the permissions to read and write for root only prevents non-root users from seeing the boot parameters or changing them. Non-root users who read the boot parameters may be able to identify weaknesses in security upon boot and be able to exploit them.

Remediation:

Run the following commands to set permissions on your grub configuration:

```
# chown root:root /boot/grub/grub.cfg
# chmod u-x,go-rwx /boot/grub/grub.cfg
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: AC-3

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

1.5 Configure Additional Process Hardening

1.5.1 Ensure address space layout randomization is enabled

Fail

Description:

Address space layout randomization (ASLR) is an exploit mitigation technique which randomly arranges the address space of key data areas of a process.

Rationale:

Randomly placing virtual memory regions will make it difficult to write memory page exploits as the memory placement will be consistently shifting.

Remediation:

Set the following parameter in `/etc/sysctl.conf` or a file in `/etc/sysctl.d/` ending in `.conf`:

- `kernel.randomize_va_space = 2`

Example:

```
# printf "\n%s\n" "kernel.randomize_va_space = 2" >> /etc/sysctl.d/60-kernel_sysctl.conf
```

Run the following command to set the active kernel parameter:

```
# sysctl -w kernel.randomize_va_space=2
```

Note: If these settings appear in a canonically later file, or later in the same file, these settings will be overwritten

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** CCI-000366: The organization implements the security configuration settings
- **URL:** NIST SP 800-53: CM-6
- **URL:** NIST SP 800-53A: CM-6.1 (iv)
- **URL:** NIST SP 800-53: SI-16
- **URL:** STIG ID: UBTU-20-010448 | Rule ID: SV-238369r958928 | CAT II
- **URL:** STIG ID: UBTU-22-213020 | Rule ID: SV-260474r958928 | CAT II

CIS Controls V7.0:

- **Control 8: Malware Defenses:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 10: Malware Defenses:** -- [More](#)
>

[Back to Summary](#)

1.5.2 Ensure ptrace_scope is restricted

Pass

Description:

The `ptrace()` system call provides a means by which one process (the "tracer") may observe and control the execution of another process (the "tracee"), and examine and change the tracee's memory and registers.

The `sysctl` settings (writable only with `CAP_SYS_PTRACE`) are:

- 0 - classic ptrace permissions: a process can `PTRACE_ATTACH` to any other process running under the same uid, as long as it is dumpable (i.e. did not transition uids, start privileged, or have called `prctl(PR_SET_DUMPABLE...)` already). Similarly, `PTRACE_TRACEME` is unchanged.
- 1 - restricted ptrace: a process must have a predefined relationship with the inferior it wants to call `PTRACE_ATTACH` on. By default, this relationship is that of only its descendants when the above classic criteria is also met. To change the relationship, an inferior can call `prctl(PR_SET_PTRACER, debugger, ...)` to declare an allowed debugger PID to call `PTRACE_ATTACH` on the inferior. Using `PTRACE_TRACEME` is unchanged.
- 2 - admin-only attach: only processes with `CAP_SYS_PTRACE` may use ptrace with `PTRACE_ATTACH`, or through children calling `PTRACE_TRACEME`.
- 3 - no attach: no processes may use ptrace with `PTRACE_ATTACH` nor via `PTRACE_TRACEME`. Once set, this `sysctl` value cannot be changed.

Rationale:

If one application is compromised, it would be possible for an attacker to attach to other running processes (e.g. Bash, Firefox, SSH sessions, GPG agent, etc) to extract additional credentials and continue to expand the scope of their attack.

Enabling restricted mode will limit the ability of a compromised process to `PTRACE_ATTACH` on other processes running under the same user. With restricted mode, ptrace will continue to work with root user.

Remediation:

Set the `kernel.yama.ptrace_scope` parameter in `/etc/sysctl.conf` or a file in `/etc/sysctl.d/` ending in `.conf` to a value of 1, 2, or 3:

```
kernel.yama.ptrace_scope = 1
- OR -
kernel.yama.ptrace_scope = 2
- OR -
kernel.yama.ptrace_scope = 3
```

Example:

```
# printf "\n%s\n" "kernel.yama.ptrace_scope = 1" >> /etc/sysctl.d/60-kernel_sysctl.conf
```

Run the following command to set the active kernel parameter:

```
# sysctl -w kernel.yama.ptrace_scope=1
```

Note:

- If a value of 2 or 3 is preferred, or required by local site policy, replace the 1 with the desired value of 2 or 3 in the example above
- If this setting appears in a canonically later file, or later in the same file, the setting will be overwritten

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: <https://www.kernel.org/doc/Documentation/security/Yama.txt>
- URL: <https://github.com/raj3shp/termspy>
- URL: NIST SP 800-53 Rev. 5: CM-6

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

1.5.3 Ensure core dumps are restricted

Fail

Description:

A core dump is the memory of an executable program. It is generally used to determine why a program aborted. It can also be used to glean confidential information from a core file. The system provides the ability to set a soft limit for core dumps, but this can be overridden by the user.

Rationale:

Setting a hard limit on core dumps prevents users from overriding the soft variable. If core dumps are required, consider setting limits for user groups (see `limits.conf(5)`). In addition, setting the `fs.suid_dumpable` variable to 0 will prevent setuid programs from dumping core.

Remediation:

Add the following line to `/etc/security/limits.conf` or a `/etc/security/limits.d/*` file:

```
* hard core 0
```

Set the following parameter in `/etc/sysctl.conf` or a file in `/etc/sysctl.d/` ending in `.conf` :

- `fs.suid_dumpable = 0`

Example:

```
# printf "\n%s" "fs.suid_dumpable = 0" >> /etc/sysctl.d/60-fs_sysctl.conf
```

Run the following command to set the active kernel parameter:

```
# sysctl -w fs.suid_dumpable=0
```

Note: If these settings appear in a canonically later file, or later in the same file, these settings will be overwritten

- IF `-systemd-coredump` is installed:

edit `/etc/systemd/coredump.conf` and add/modify the following lines:

```
Storage=none
ProcessSizeMax=0
```

Run the command:

```
systemctl daemon-reload
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-6

[Back to Summary](#)

1.5.4 Ensure prelink is not installed

Pass

Description:

`prelink` is a program that modifies ELF shared libraries and ELF dynamically linked binaries in such a way that the time needed for the dynamic linker to perform relocations at startup significantly decreases.

Rationale:

The prelinking feature can interfere with the operation of AIDE, because it changes binaries. Prelinking can also increase the vulnerability of the system if a malicious user is able to compromise a common library such as `libc`.

Remediation:

Run the following command to restore binaries to normal:

```
# prelink -ua
```

Uninstall `prelink` using the appropriate package manager or manual installation:

```
# apt purge prelink
```

Assessment:

[Show Assessment Evidence](#)

[Show Rule Result XML](#)

References:

- URL: NIST SP 800-53 Rev. 5: CM-6, CM-1, CM-3

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

1.5.5 Ensure Automatic Error Reporting is not enabled

Fail

Description:

The Apport Error Reporting Service automatically generates crash reports for debugging

Rationale:

Apport collects potentially sensitive data, such as core dumps, stack traces, and log files. They can contain passwords, credit card numbers, serial numbers, and other private material.

Remediation:

Edit `/etc/default/apport` and add or edit the enabled parameter to equal 0 :

```
enabled=0
```

Run the following commands to stop and mask the apport service

```
# systemctl stop apport.service
# systemctl mask apport.service
```

- OR -

Run the following command to remove the apport package:

```
# apt purge apport
```

Assessment:

[Show Assessment Evidence](#)

[Show Rule Result XML](#)

References:

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

1.6 Configure Command Line Warning Banners

Presenting a warning message prior to the normal user login may assist in the prosecution of trespassers on the computer system. Changing some of these login banners also has the side effect of hiding OS version information and other detailed system information from attackers attempting to target specific exploits at a system.

Guidelines published by the US Department of Defense require that warning messages include at least the name of the organization that owns the system, the fact that the system is subject to monitoring and that such monitoring is in compliance with local statutes, and that use of the system implies consent to such monitoring. It is important that the organization's legal counsel review the content of all messages before any system modifications are made, as these warning messages are inherently site-specific. More information (including citations of relevant case law) can be found at <http://www.justice.gov/criminal/cybercrime/>

The `/etc/motd`, `/etc/issue`, and `/etc/issue.net` files govern warning banners for standard command line logins for both local and remote users.

Note: The text provided in the remediation actions for these items is intended as an example only. Please edit to include the specific text for your organization as approved by your legal department.

1.6.1 Ensure `/etc/motd` is configured

Pass

Description:

The contents of the `/etc/motd` file are displayed to users after login and function as a message of the day for authenticated users.

Unix-based systems have typically displayed information about the OS release and patch level upon logging in to the system. This information can be useful to developers who are developing software for a particular OS platform. If `mingetty(8)` supports the following options, they display operating system information: `\m` - machine architecture `\r` - operating system release `\s` - operating system name `\v` - operating system version

Rationale:

Warning messages inform users who are attempting to login to the system of their legal status regarding the system and must include the name of the organization that owns the system and any monitoring policies that are in place. Displaying OS and patch level information in login banners also has the side effect of providing detailed system information to attackers attempting to target specific exploits of a system. Authorized users can easily get this information by running the `"uname -a"` command once they have logged in.

Remediation:

Edit the file found in `/etc/motd.d/*` with the appropriate contents according to your site policy, remove any instances of `\m`, `\r`, `\s`, `\v` or references to the OS platform

- OR -

- IF - the `motd` is not used, this file can be removed.

Run the following command to remove the `motd` file:

```
# rm /etc/motd
```

Run the following script and review and/or update all returned files' contents to:

- Remove all system information (\v , \r ; \m , \s)
- Remove any reference to the operating system
- Ensure contents follow local site policy

```
#!/usr/bin/env bash

{

a_files=()

for l_file in /etc/motd{,.d/*}; do

if grep -Psqi -- "(\v|\r|\m|\s|b$(grep ^ID= /etc/os-release | cut -d= -f2 | sed -e 's/"/'/'g')\b)" "$l_file"; then

echo -e "\n - File: \"$l_file\" includes system information. Edit this file to remove these entries"

else

a_files+=("$l_file")

fi

done

if [ "${#a_files[@]}" -gt 0 ]; then

echo -e "\n- ** Please review the following files and verify their contents follow local site policy\n"

printf '%s\n' "${a_files[@]}"

fi

}
```

Assessment:[Show](#) Assessment Evidence[Show](#) Rule Result XML**References:**

- URL: NIST SP 800-53 Rev. 5: CM-6, CM-1, CM-3

CIS Controls V7.0:

- Not Explicitly Mapped: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Not Explicitly Mapped: -- [More](#)
>

[Back to Summary](#)**1.6.2 Ensure /etc/issue is configured****Fail****Description:**

The contents of the `/etc/issue` file are displayed to users prior to login for local terminals.

Unix-based systems have typically displayed information about the OS release and patch level upon logging in to the system. This information can be useful to developers who are developing software for a particular OS platform. If `mingetty(8)` supports the following options, they display operating system information: `\m` - machine architecture `\r` - operating system release `\s` - operating system name `\v` - operating system version - or the operating system's name

Rationale:

Warning messages inform users who are attempting to login to the system of their legal status regarding the system and must include the name of the organization that owns the system and any monitoring policies that are in place. Displaying OS and patch level information in login banners also has the side effect of providing detailed system

information to attackers attempting to target specific exploits of a system. Authorized users can easily get this information by running the "uname -a" command once they have logged in.

Remediation:

Edit the /etc/issue file with the appropriate contents according to your site policy, remove any instances of \m , \r , \s , \v or references to the OS platform

Example:

```
# echo "Authorized users only. All activity may be monitored and reported." > /etc/issue
```

Assessment:

Show Assessment Evidence

Show Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-6, CM-1, CM-3

CIS Controls V7.0:

- Not Explicitly Mapped: -- More >

CIS Critical Security Controls V8.0:

- Not Explicitly Mapped: -- More >

Back to Summary

1.6.3 Ensure /etc/issue.net is configured

Fail

Description:

The contents of the /etc/issue.net file are displayed to users prior to login for remote connections from configured services.

Unix-based systems have typically displayed information about the OS release and patch level upon logging in to the system. This information can be useful to developers who are developing software for a particular OS platform. If mingetty(8) supports the following options, they display operating system information: \m - machine architecture \r - operating system release \s - operating system name \v - operating system version

Rationale:

Warning messages inform users who are attempting to login to the system of their legal status regarding the system and must include the name of the organization that owns the system and any monitoring policies that are in place. Displaying OS and patch level information in login banners also has the side effect of providing detailed system information to attackers attempting to target specific exploits of a system. Authorized users can easily get this information by running the "uname -a" command once they have logged in.

Remediation:

Edit the /etc/issue.net file with the appropriate contents according to your site policy, remove any instances of \m , \r , \s , \v or references to the OS platform

Example:

```
# echo "Authorized users only. All activity may be monitored and reported." > /etc/issue.net
```

Assessment:

Show Assessment Evidence

Show Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-6, CM-1, CM-3

CIS Controls V7.0:

- Not Explicitly Mapped: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Not Explicitly Mapped: -- [More](#)
>

[Back to Summary](#)**1.6.4 Ensure access to /etc/motd is configured**

Pass

Description:

The contents of the `/etc/motd` file are displayed to users after login and function as a message of the day for authenticated users.

Rationale:

- **IF** - the `/etc/motd` file does not have the correct access configured, it could be modified by unauthorized users with incorrect or misleading information.

Remediation:

Run the following commands to set mode, owner, and group on `/etc/motd`:

```
# chown root:root $(readlink -e /etc/motd)
# chmod u-x,go-wx $(readlink -e /etc/motd)
```

- OR -

Run the following command to remove the `/etc/motd` file:

```
# rm /etc/motd
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AC-3, MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)**1.6.5 Ensure access to /etc/issue is configured**

Pass

Description:

The contents of the `/etc/issue` file are displayed to users prior to login for local terminals.

Rationale:

- **IF** - the `/etc/issue` file does not have the correct access configured, it could be modified by unauthorized users with incorrect or misleading information.

Remediation:

Run the following commands to set mode, owner, and group on `/etc/issue`:

```
# chown root:root $(readlink -e /etc/issue)

# chmod u-x,go-wx $(readlink -e /etc/issue)
```

Assessment:
[Show](#) Assessment Evidence

[Show](#) Rule Result XML

- References:**
- URL: NIST SP 800-53 Rev. 5: AC-3, MP-2

- CIS Controls V7.0:**
- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

- CIS Critical Security Controls V8.0:**
- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

1.6.6 Ensure access to /etc/issue.net is configured

Pass

Description:

The contents of the `/etc/issue.net` file are displayed to users prior to login for remote connections from configured services.

Rationale:

- **IF** - the `/etc/issue.net` file does not have the correct access configured, it could be modified by unauthorized users with incorrect or misleading information.

Remediation:

Run the following commands to set mode, owner, and group on `/etc/issue.net` :

```
# chown root:root $(readlink -e /etc/issue.net)

# chmod u-x,go-wx $(readlink -e /etc/issue.net)
```

Assessment:
[Show](#) Assessment Evidence

[Show](#) Rule Result XML

- References:**
- URL: NIST SP 800-53 Rev. 5: AC-3, MP-2

- CIS Controls V7.0:**
- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

- CIS Critical Security Controls V8.0:**
- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

1.7 Configure GNOME Display Manager

The GNOME Display Manager (GDM) is a program that manages graphical display servers and handles graphical user logins.

This subsection requires user profiles to already exist on the system. A profile is a list of configuration databases.

Sample profile:

```
user-db:user
system-db:local
system-db:site
```

Configuring a single user and multiple system databases allows for layering of preferences. Settings from the user database file take precedence over the settings in the local database file, and the local database file in turn takes precedence over the site database file.

Note:

- IF - GDM is not installed on the system, this section can be skipped
- The Remediation Procedure commands in this section **MUST** be done from a command window on a graphical desktop or an error will be returned.

1.7.2 Ensure GDM login banner is configured

Fail

Description:

GDM is the GNOME Display Manager which handles graphical login for GNOME based systems.

Rationale:

Warning messages inform users who are attempting to login to the system of their legal status regarding the system and must include the name of the organization that owns the system and any monitoring policies that are in place.

Remediation:

- IF - A user profile is already created run the following commands to set and enable the text banner message on the login screen:

```
# gsettings set org.gnome.login-screen banner-message-text 'Authorized uses only. All activity may be monitored and reported'
# gsettings set org.gnome.login-screen banner-message-enable true
```

Note:

- banner-message-text may be set in accordance with local site policy
- gsettings commands in this section **MUST** be done from a command window on a graphical desktop or an error will be returned.
- The system must be restarted after all gsettings configurations have been set in order for CIS-CAT Assessor to appropriately assess.

- OR/IF - A user profile does not exist:

- Create or edit the gdm profile in the /etc/dconf/profile/gdm with the following lines:

```
user-db:user
system-db:gdm
file-db:/usr/share/gdm/greeter-dconf-defaults
```

Note: gdm is the name of a dconf database.

- Create a gdm keyfile for machine-wide settings in /etc/dconf/db/gdm.d/01-banner-message :

```
[org/gnome/login-screen]
banner-message-enable=true
banner-message-text='Type the banner message here.'
```

- Update the system databases

```
# dconf update
```

Note:

- Users must log out and back in again before the system-wide settings take effect.
- There is no character limit for the banner message. gnome-shell autodetects longer stretches of text and enters two column mode.
- The banner message cannot be read from an external file.

Assessment:

Show Assessment Evidence

[Show Rule Result XML](#)

References:

- URL: https://help.gnome.org/admin/system-admin-guide/stable/login-banner.html.en
- URL: NIST SP 800-53 Rev. 5: CM-1,CM-2, CM-6, CM-7, IA-5

[Back to Summary](#)

1.7.3 Ensure GDM disable-user-list option is enabled

Fail

Description:

GDM is the GNOME Display Manager which handles graphical login for GNOME based systems.

The `disable-user-list` option controls if a list of users is displayed on the login screen

Rationale:

Displaying the user list eliminates half of the Userid/Password equation that an unauthorized person would need to log on.

Remediation:

- IF - A user profile exists run the following command to enable the `disable-user-list` :

```
# gsettings set org.gnome.login-screen disable-user-list true
```

Note:

- `gsettings` commands in this section MUST be done from a command window on a graphical desktop or an error will be returned.
- The system must be restarted after all `gsettings` configurations have been set in order for CIS-CAT Assessor to appropriately assess.

- OR/IF - A user profile does not exist:

1. Create or edit the `gdm` profile in `/etc/dconf/profile/gdm` with the following lines:

```
user-db:user
system-db:gdm
file-db:/usr/share/gdm/greeter-dconf-defaults
```

Note: `gdm` is the name of a `dconf` database.

2. Create a `gdm` keyfile for machine-wide settings in `/etc/dconf/db/gdm.d/00-login-screen` :

```
[org/gnome/login-screen]
# Do not show the user list
disable-user-list=true
```

3. Update the system databases:

```
# dconf update
```

Note: When the user profile is created or changed, the user will need to log out and log in again before the changes will be applied.

Assessment:

[Show Assessment Evidence](#)

[Show Rule Result XML](#)

References:

- URL: https://help.gnome.org/admin/system-admin-guide/stable/login-userlist-disable.html.en
- URL: NIST SP 800-53 Rev. 5: CM-1,CM-2, CM-6, CM-7, IA-5

[Back to Summary](#)

1.7.4 Ensure GDM screen locks when the user is idle

Fail

Description:

GNOME Desktop Manager can make the screen lock automatically whenever the user is idle for some amount of time.

Rationale:

Setting a lock-out value reduces the window of opportunity for unauthorized user access to another user's session that has been left unattended.

Remediation:

- IF - A user profile is already created run the following commands to enable screen locks when the user is idle:

```
# gsettings set org.gnome.desktop.screensaver lock-delay 5
# gsettings set org.gnome.desktop.session idle-delay 900
# gsettings set org.gnome.desktop.screensaver lock-enabled true
```

Note:

- gsettings commands in this section MUST be done from a command window on a graphical desktop or an error will be returned.
- The system must be restarted after all gsettings configurations have been set in order for CIS-CAT Assessor to appropriately assess.

- OR/IF- A user profile does not exist:

1. Create or edit the user profile in the /etc/dconf/profile/ and verify it includes the following:

```
user-db:user
system-db:{NAME_OF_DCONF_DATABASE}
```

Note:local is the name of a dconf database used in the examples.

2. Create the directory /etc/dconf/db/local.d/ if it doesn't already exist:
3. Create the key file /etc/dconf/db/local.d/00-screensaver to provide information for the local database:

Example key file:

```
# Specify the dconf path
[org/gnome/desktop/session]

# Number of seconds of inactivity before the screen goes blank
# Set to 0 seconds if you want to deactivate the screensaver.
idle-delay=uint32 180

# Specify the dconf path
[org/gnome/desktop/screensaver]

# Number of seconds after the screen is blank before locking the screen
lock-delay=uint32 0

# Ensure screen locks after inactivity
lock-enabled=true
```

Note: You must include the uint32 along with the integer key values as shown.

4. Run the following command to update the system databases:

```
# dconf update
```

5. Users must log out and back in again before the system-wide settings take effect.

Assessment:
[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** <https://help.gnome.org/admin/system-admin-guide/stable/desktop-lockscreens.html.en>

CIS Controls V7.0:

- **Control 16: Account Monitoring and Control:** -- [More](#) >

CIS Critical Security Controls V8.0:

- **Control 4: Secure Configuration of Enterprise Assets and Software:** -- [More](#) >

[Back to Summary](#)

1.7.5 Ensure GDM screen locks cannot be overridden

Fail

Description:

GNOME Desktop Manager can lock down specific settings by using the lockdown mode in dconf to prevent users from changing specific settings.

To lock down a dconf key or subpath, create a locks subdirectory in the keyfile directory. The files inside this directory contain a list of keys or subpaths to lock. Just as with the keyfiles, you may add any number of files to this directory.

Rationale:

Setting a lock-out value reduces the window of opportunity for unauthorized user access to another user's session that has been left unattended.

Without locking down the system settings, user settings take precedence over the system settings.

Remediation:

1. To prevent the user from overriding these settings, create the file `/etc/dconf/db/local.d/locks/00-screensaver` with the following content:

```
# Lock desktop screensaver settings
/org/gnome/desktop/session/idle-delay
/org/gnome/desktop/screensaver/lock-delay
/org/gnome/desktop/screensaver/lock-enabled
```

2. Update the system databases:

```
# dconf update
```

Note:

- A user profile must exist in order to apply locks. If a user profile does not exist review the remediation steps in the previous recommendation.
- Users must log out and back in again before the system-wide settings take effect.

Assessment:
[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** <https://help.gnome.org/admin/system-admin-guide/stable/desktop-lockscreens.html.en>
- **URL:** <https://help.gnome.org/admin/system-admin-guide/stable/dconf-lockdown.html.en>
- **URL:** NIST SP 800-53 Rev. 5: CM-11

CIS Controls V7.0:

- **Control 16: Account Monitoring and Control:** -- [More](#)

>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- >

[Back to Summary](#)

1.7.8 Ensure GDM autorun-never is enabled

Fail

Description:

The `autorun-never` setting allows the GNOME Desktop Display Manager to disable autorun through GDM.

Rationale:

Malware on removable media may taking advantage of Autorun features when the media is inserted into a system and execute.

Remediation:

- **IF** - A user profile exist run the following command to set `autorun-never` to `true` for GDM users:

```
# gsettings set org.gnome.desktop.media-handling autorun-never true
```

Note:

- `gsettings` commands in this section **MUST** be done from a command window on a graphical desktop or an error will be returned.
- The system must be restarted after all `gsettings` configurations have been set in order for CIS-CAT Assessor to appropriately assess.

- **OR/IF** - A user profile does not exist:

1. create the file `/etc/dconf/db/local.d/locks/00-media-autorun` with the following content:

```
[org/gnome/desktop/media-handling]
autorun-never=true
```

2. Update the systems databases:

```
# dconf update
```

Note: Users must log out and back in again before the system-wide settings take effect.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-1,CM-2, CM-6, CM-7, IA-5

CIS Controls V7.0:

- Control 8: Malware Defenses: -- [More](#)
- >

CIS Critical Security Controls V8.0:

- Control 10: Malware Defenses: -- [More](#)
- >

[Back to Summary](#)

1.7.9 Ensure GDM autorun-never is not overridden

Fail

Description:

The `autorun-never` setting allows the GNOME Desktop Display Manager to disable autorun through GDM.

By using the lockdown mode in `dconf`, you can prevent users from changing specific settings.

To lock down a dconf key or subpath, create a locks subdirectory in the keyfile directory. The files inside this directory contain a list of keys or subpaths to lock. Just as with the keyfiles, you may add any number of files to this directory.

Rationale:

Malware on removable media may taking advantage of Autorun features when the media is inserted into a system and execute.

Remediation:

1. To prevent the user from overriding these settings, create the file `/etc/dconf/db/local.d/locks/00-media-autorun` with the following content:

```
[org/gnome/desktop/media-handling]

autorun-never=true
```

2. Update the systems databases:

```
# dconf update
```

Note:

- A user profile must exist in order to apply locks. If a user profile does not exist review the remediation steps in the previous recommendation.
- Users must log out and back in again before the system-wide settings take effect.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: CM-1,CM-2, CM-6, CM-7, IA-5

CIS Controls V7.0:

- **Control 8: Malware Defenses:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 10: Malware Defenses:** -- [More](#)
>

[Back to Summary](#)

1.7.10 Ensure XDMCP is not enabled

Pass

Description:

X Display Manager Control Protocol (XDMCP) is designed to provide authenticated access to display management services for remote displays

Rationale:

XDMCP is inherently insecure.

- XDMCP is not a ciphered protocol. This may allow an attacker to capture keystrokes entered by a user
- XDMCP is vulnerable to man-in-the-middle attacks. This may allow an attacker to steal the credentials of legitimate users by impersonating the XDMCP server.

Remediation:

Edit all files returned by the audit and remove or comment out the `Enable=true` line in the `[xdmcp]` block:

Example file:

```
# GDM configuration storage

#

# See /usr/share/gdm/gdm.schemas for a list of available options.
```

```
[daemon]

# Uncomment the line below to force the login screen to use Xorg
#WaylandEnable=false

# Enabling automatic login
# AutomaticLoginEnable = true
# AutomaticLogin = user1

# Enabling timed login
# TimedLoginEnable = true
# TimedLogin = user1
# TimedLoginDelay = 10

[security]

[xdmcp]

# Enable=true <- **This line should be removed or commented out**

[chooser]

[debug]

# Uncomment the line below to turn on debugging
# More verbose logs
# Additionally lets the X server dump core if it crashes
# Enable=true
```

Assessment:[Show](#) Assessment Evidence[Show](#) Rule Result XML**References:**

- **URL:** NIST SP 800-53 Rev. 5: SI-4

CIS Controls V7.0:

- **Control 9: Limitation and Control of Network Ports, Protocols, and Services:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 4: Secure Configuration of Enterprise Assets and Software:** -- [More](#)
>

[Back to Summary](#)

2 Services

While applying system updates and patches helps correct known vulnerabilities, one of the best ways to protect the system against as yet unreported vulnerabilities is to disable all services that are not required for normal system operation. This prevents the exploitation of vulnerabilities discovered at a later date. If a service is not enabled, it cannot be exploited. The actions in this section of the document provide guidance on some services which can be safely disabled and under which circumstances, greatly reducing the number of possible threats to the resulting system. Additionally, some services which

should remain enabled but with secure configuration are covered as well as insecure service clients.

2.1 Configure Server Services

This section describes services that are installed on systems that specifically need to run these services. If any of these services are not required, it is recommended that the package be removed.

- **IF** - the package is required for a dependency:

- Ensure the dependent package is approved by local site policy
- Ensure stopping and masking the service and/or socket meets local site policy
- Stop and mask the service and/or socket to reduce the potential attack surface

The following commands can be used to stop and mask the service and socket:

```
# systemctl stop <service_name>.socket <service_name>.service
# systemctl mask <service_name>.socket <service_name>.service
```

Note: This should not be considered a comprehensive list of services not required for normal system operation. You may wish to consider additions to those listed here for your environment

2.1.3 Ensure dhcp server services are not in use

Pass

Description:

The Dynamic Host Configuration Protocol (DHCP) is a service that allows machines to be dynamically assigned IP addresses. There are two versions of the DHCP protocol `DHCPv4` and `DHCPv6`. At startup the server may be started for one or the other via the `-4` or `-6` arguments.

Rationale:

Unless a system is specifically set up to act as a DHCP server, it is recommended that this package be removed to reduce the potential attack surface.

Remediation:

Run the following commands to stop `isc-dhcp-server.service` and `isc-dhcp-server6.service` and remove the `isc-dhcp-server` package:

```
# systemctl stop isc-dhcp-server.service isc-dhcp-server6.service
# apt purge isc-dhcp-server
```

- **OR** -

- **IF** - the `isc-dhcp-server` package is required as a dependency:

Run the following commands to stop and mask `isc-dhcp-server.service` and `isc-dhcp-server6.service`:

```
# systemctl stop isc-dhcp-server.service isc-dhcp-server6.service
# systemctl mask isc-dhcp-server isc-dhcp-server6.service
```

Impact:

There may be packages that are dependent on the `isc-dhcp-server` package. If the `isc-dhcp-server` package is removed, these dependent packages will be removed as well. Before removing the `isc-dhcp-server` package, review any dependent packages to determine if they are required on the system.

- **IF** - a dependent package is required: stop and mask the `isc-dhcp-server.service` and `isc-dhcp-server6.service` leaving the `isc-dhcp-server` package installed.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** More detailed documentation on DHCP is available at <http://www.isc.org/software/dhcp>.
- **URL:** NIST SP 800-53 Rev. 5: CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

2.1.4 Ensure dns server services are not in use

Pass

Description:

The Domain Name System (DNS) is a hierarchical naming system that maps names to IP addresses for computers, services and other resources connected to a network.

Note:bind9 is the package and bind.service is the alias for named.service .

Rationale:

Unless a system is specifically designated to act as a DNS server, it is recommended that the package be deleted to reduce the potential attack surface.

Remediation:

Run the following commands to stop named.service and remove the bind9 package:

```
# systemctl stop named.service
# apt purge bind9
```

- OR -

- IF - the bind9 package is required as a dependency:

Run the following commands to stop and mask bind9.service :

```
# systemctl stop named.service
# systemctl mask named.service
```

Impact:

There may be packages that are dependent on the bind9 package. If the bind9 package is removed, these dependent packages will be removed as well. Before removing the bind9 package, review any dependent packages to determine if they are required on the system.

- IF - a dependent package is required: stop and mask named.service leaving the bind9 package installed.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

2.1.5 Ensure dnsmasq services are not in use

Pass

Description:

dnsmasq is a lightweight tool that provides DNS caching, DNS forwarding and DHCP (Dynamic Host Configuration Protocol) services.

Rationale:

Unless a system is specifically designated to act as a DNS caching, DNS forwarding and/or DHCP server, it is recommended that the package be removed to reduce the potential attack surface.

Remediation:

Run the following commands to stop dnsmasq.service and remove dnsmasq package:

```
# systemctl stop dnsmasq.service
# apt purge dnsmasq
```

- OR -

- IF - the dnsmasq package is required as a dependency:

Run the following commands to stop and mask the dnsmasq.service :

```
# systemctl stop dnsmasq.service
# systemctl mask dnsmasq.service
```

Impact:

There may be packages that are dependent on the dnsmasq package. If the dnsmasq package is removed, these dependent packages will be removed as well. Before removing the dnsmasq package, review any dependent packages to determine if they are required on the system.

- IF - a dependent package is required: stop and mask the dnsmasq.service leaving the dnsmasq package installed.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

2.1.6 Ensure ftp server services are not in use

Pass

Description:

The File Transfer Protocol (FTP) provides networked computers with the ability to transfer files.

Rationale:

FTP does not protect the confidentiality of data or authentication credentials. It is recommended SFTP be used if file transfer is required. Unless there is a need to run the system as a FTP server (for example, to allow anonymous downloads), it is recommended that the package be deleted to reduce the potential attack surface.

Note: Other ftp server packages may exist. They should also be audited, if not required and authorized by local site policy

Remediation:

Run the following commands to stop `vsftpd.service` and remove the `vsftpd` package:

```
# systemctl stop vsftpd.service
# apt purge vsftpd
```

- OR -

- IF - the `vsftpd` package is required as a dependency:

Run the following commands to stop and mask the `vsftpd.service` :

```
# systemctl stop vsftpd.service
# systemctl mask vsftpd.service
```

Note: Other ftp server packages may exist. If not required and authorized by local site policy, they should also be removed. If the package is required for a dependency, the service should be stopped and masked.

Impact:

There may be packages that are dependent on the `vsftpd` package. If the `vsftpd` package is removed, these dependent packages will be removed as well. Before removing the `vsftpd` package, review any dependent packages to determine if they are required on the system.

- IF - a dependent package is required: stop and mask the `vsftpd.service` leaving the `vsftpd` package installed.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

2.1.7 Ensure ldap server services are not in use

Pass

Description:

The Lightweight Directory Access Protocol (LDAP) was introduced as a replacement for NIS/YP. It is a service that provides a method for looking up information from a central database.

Rationale:

If the system will not need to act as an LDAP server, it is recommended that the software be removed to reduce the potential attack surface.

Remediation:

Run the following commands to stop `slapd.service` and remove the `slapd` package:

```
# systemctl stop slapd.service
# apt purge slapd
```

- OR -

- IF - the slapd package is required as a dependency:

Run the following commands to stop and mask slapd.service :

```
# systemctl stop slapd.service
# systemctl mask slapd.service
```

Impact:

There may be packages that are dependent on the slapd package. If the slapd package is removed, these dependent packages will be removed as well. Before removing the slapd package, review any dependent packages to determine if they are required on the system.

- IF - a dependent package is required: stop and mask the slapd.service leaving the slapd package installed.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- [URL](#): For more detailed documentation on OpenLDAP, go to the project homepage at <http://www.openldap.org>.
- [URL](#): NIST SP 800-53 Rev. 5: CM-7

CIS Controls V7.0:

- [Control 9: Limitation and Control of Network Ports, Protocols, and Services](#): -- [More](#)
>

CIS Critical Security Controls V8.0:

- [Control 4: Secure Configuration of Enterprise Assets and Software](#): -- [More](#)
>

[Back to Summary](#)

2.1.8 Ensure message access server services are not in use

Pass

Description:

dovecot-imapd and dovecot-pop3d are an open source IMAP and POP3 server for Linux based systems.

Rationale:

Unless POP3 and/or IMAP servers are to be provided by this system, it is recommended that the package be removed to reduce the potential attack surface.

Note: Several IMAP/POP3 servers exist and can use other service names. These should also be audited and the packages removed if not required.

Remediation:

Run one of the following commands to remove dovecot-imapd and dovecot-pop3d :

Run the following commands to stop dovecot.socket and dovecot.service , and remove the dovecot-imapd and dovecot-pop3d packages:

```
# systemctl stop dovecot.socket dovecot.service
# apt purge dovecot-imapd dovecot-pop3d
```

- OR -

- IF - a package is installed **and** is required for dependencies:

Run the following commands to stop and mask dovecot.socket and dovecot.service :

```
# systemctl stop dovecot.socket dovecot.service
```

```
# systemctl mask dovecot.socket dovecot.service
```

Impact:

There may be packages that are dependent on `dovecot-imapd` and/or `dovecot-pop3d` packages. If `dovecot-imapd` and `dovecot-pop3d` packages are removed, these dependent packages will be removed as well. Before removing `dovecot-imapd` and/or `dovecot-pop3d` packages, review any dependent packages to determine if they are required on the system.

- **IF** - a dependent package is required: stop and mask `dovecot.socket` and `dovecot.service` leaving `dovecot-imapd` and/or `dovecot-pop3d` packages installed.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#) >

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#) >

[Back to Summary](#)

2.1.9 Ensure network file system services are not in use

Pass

Description:

The Network File System (NFS) is one of the first and most widely distributed file systems in the UNIX environment. It provides the ability for systems to mount file systems of other servers through the network.

Rationale:

If the system does not export NFS shares, it is recommended that the `nfs-kernel-server` package be removed to reduce the remote attack surface.

Remediation:

Run the following command to stop `nfs-server.service` and remove `nfs-kernel-server` package:

```
# systemctl stop nfs-server.service
# apt purge nfs-kernel-server
```

- **OR** -

- **IF** - the `nfs-kernel-server` package is required as a dependency:

Run the following commands to stop and mask the `nfs-server.service` :

```
# systemctl stop nfs-server.service
# systemctl mask nfs-server.service
```

Impact:

There may be packages that are dependent on the `nfs-kernel-server` package. If the `nfs-kernel-server` package is removed, these dependent packages will be removed as well. Before removing the `nfs-kernel-server` package, review any dependent packages to determine if they are required on the system.

- **IF** - a dependent package is required: stop and mask the `nfs-server.service` leaving the `nfs-kernel-server` package installed.

Assessment:[Show](#) Assessment Evidence[Show](#) Rule Result XML**References:**

- URL: NIST SP 800-53 Rev. 5: CM-6, CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)**2.1.10 Ensure nis server services are not in use**

Pass

Description:

The Network Information Service (NIS) (formally known as Yellow Pages) is a client-server directory service protocol for distributing system configuration files. The NIS server is a collection of programs that allow for the distribution of configuration files. The NIS client (`ypbind`) was used to bind a machine to an NIS server and receive the distributed configuration files.

Rationale:

`ypserv.service` is inherently an insecure system that has been vulnerable to DOS attacks, buffer overflows and has poor authentication for querying NIS maps. NIS generally has been replaced by such protocols as Lightweight Directory Access Protocol (LDAP). It is recommended that `ypserv.service` be removed and other, more secure services be used

Remediation:

Run the following commands to stop `ypserv.service` and remove `ypserv` package:

```
# systemctl stop ypserv.service
# apt purge ypserv
```

- OR -

- IF - the `ypserv` package is required as a dependency:

Run the following commands to stop and mask `ypserv.service` :

```
# systemctl stop ypserv.service
# systemctl mask ypserv.service
```

Impact:

There may be packages that are dependent on the `ypserv` package. If the `ypserv` package is removed, these dependent packages will be removed as well. Before removing the `ypserv` package, review any dependent packages to determine if they are required on the system.

- IF - a dependent package is required: stop and mask the `ypserv.service` leaving the `ypserv` package installed.

Assessment:[Show](#) Assessment Evidence[Show](#) Rule Result XML**References:**

- URL: NIST SP 800-53 Rev. 5: CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

2.1.12 Ensure rpcbind services are not in use

Pass

Description:

The `rpcbind` utility maps RPC services to the ports on which they listen. RPC processes notify `rpcbind` when they start, registering the ports they are listening on and the RPC program numbers they expect to serve. The client system then contacts `rpcbind` on the server with a particular RPC program number. The `rpcbind.service` redirects the client to the proper port number so it can communicate with the requested service.

Portmapper is an RPC service, which always listens on tcp and udp 111, and is used to map other RPC services (such as `nfs`, `nlockmgr`, `quotad`, `mountd`, etc.) to their corresponding port number on the server. When a remote host makes an RPC call to that server, it first consults with portmap to determine where the RPC server is listening.

Rationale:

A small request (~82 bytes via UDP) sent to the Portmapper generates a large response (7x to 28x amplification), which makes it a suitable tool for DDoS attacks. If `rpcbind` is not required, it is recommended to remove `rpcbind` package to reduce the potential attack surface.

Remediation:

Run the following commands to stop `rpcbind.socket` and `rpcbind.service` , and remove the `rpcbind` package:

```
# systemctl stop rpcbind.socket rpcbind.service
# apt purge rpcbind
```

- OR -

- IF - the `rpcbind` package is required as a dependency:

Run the following commands to stop and mask the `rpcbind.socket` and `rpcbind.service` :

```
# systemctl stop rpcbind.socket rpcbind.service
# systemctl mask rpcbind.socket rpcbind.service
```

Impact:

Many of the `libvirt` packages used by Enterprise Linux virtualization, and the `nfs-utils` package used for The Network File System (NFS), are dependent on the `rpcbind` package. If the `rpcbind` package is removed, these dependent packages will be removed as well. Before removing the `rpcbind` package, review any dependent packages to determine if they are required on the system.

- IF - a dependent package is required: stop and mask the `rpcbind.socket` and `rpcbind.service` leaving the `rpcbind` package installed.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-6, CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)

>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- >

[Back to Summary](#)

2.1.13 Ensure rsync services are not in use

Fail

Description:

The `rsync` service can be used to synchronize files between systems over network links.

Rationale:

`rsync.service` presents a security risk as the `rsync` protocol is unencrypted.

The `rsync` package should be removed to reduce the attack area of the system.

Remediation:

Run the following commands to stop `rsync.service` , and remove the `rsync` package:

```
# systemctl stop rsync.service
# apt purge rsync
```

- OR -

- IF - the `rsync` package is required as a dependency:

Run the following commands to stop and mask `rsync.service` :

```
# systemctl stop rsync.service
# systemctl mask rsync.service
```

Impact:

There may be packages that are dependent on the `rsync` package. If the `rsync` package is removed, these dependent packages will be removed as well. Before removing the `rsync` package, review any dependent packages to determine if they are required on the system.

- IF - a dependent package is required: stop and mask `rsync.service` leaving the `rsync` package installed.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
- >

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- >

[Back to Summary](#)

2.1.14 Ensure samba file server services are not in use

Pass

Description:

The Samba daemon allows system administrators to configure their Linux systems to share file systems and directories with Windows desktops. Samba will advertise the file systems and directories via the Server Message Block (SMB) protocol. Windows desktop users will be able to mount these directories and file systems as letter drives on their systems.

Rationale:

If there is no need to mount directories and file systems to Windows systems, then this service should be deleted to reduce the potential attack surface.

Remediation:

Run the following commands to stop `smbd.service` and remove `samba` package:

```
# systemctl stop smbd.service
# apt purge samba
```

- OR -

- IF - the `samba` package is required as a dependency:

Run the following commands to stop and mask the `smbd.service` :

```
# systemctl stop smbd.service
# systemctl mask smbd.service
```

Impact:

There may be packages that are dependent on the `samba` package. If the `samba` package is removed, these dependent packages will be removed as well. Before removing the `samba` package, review any dependent packages to determine if they are required on the system.

- IF - a dependent package is required: stop and mask the `smbd.service` leaving the `samba` package installed.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-6, CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

2.1.15 Ensure snmp services are not in use

Pass

Description:

Simple Network Management Protocol (SNMP) is a widely used protocol for monitoring the health and welfare of network equipment, computer equipment and devices like UPSs.

Net-SNMP is a suite of applications used to implement SNMPv1 (RFC 1157), SNMPv2 (RFCs 1901-1908), and SNMPv3 (RFCs 3411-3418) using both IPv4 and IPv6.

Support for SNMPv2 classic (a.k.a. "SNMPv2 historic" - RFCs 1441-1452) was dropped with the 4.0 release of the UCD-snmp package.

The Simple Network Management Protocol (SNMP) server is used to listen for SNMP commands from an SNMP management system, execute the commands or collect the information and then send results back to the requesting system.

Rationale:

The SNMP server can communicate using `SNMPv1` , which transmits data in the clear and does not require authentication to execute commands. `SNMPv3` replaces the simple/clear text password sharing used in `SNMPv2` with more securely encoded parameters. If the the SNMP service is not required, the `snmpd` package should be removed to reduce the attack surface of the system.

Note: If SNMP is required:

- The server should be configured for `SNMP v3` only. User Authentication and Message Encryption should be configured.
- If `SNMP v2` is **absolutely** necessary, modify the community strings' values.

Remediation:

Run the following commands to stop `snmpd.service` and remove the `snmpd` package:

```
# systemctl stop snmpd.service
# apt purge snmpd
```

- **OR** - If the package is required for dependencies:

Run the following commands to stop and mask the `snmpd.service` :

```
# systemctl stop snmpd.service
# systemctl mask snmpd.service
```

Impact:

There may be packages that are dependent on the `snmpd` package. If the `snmpd` package is removed, these packages will be removed as well.

Before removing the `snmpd` package, review any dependent packages to determine if they are required on the system. If a dependent package is required, stop and mask the `snmpd.service` leaving the `snmpd` package installed.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#) >

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#) >

[Back to Summary](#)

2.1.16 Ensure tftp server services are not in use

Pass

Description:

Trivial File Transfer Protocol (TFTP) is a simple protocol for exchanging files between two TCP/IP machines. TFTP servers allow connections from a TFTP Client for sending and receiving files.

Rationale:

Unless there is a need to run the system as a TFTP server, it is recommended that the package be removed to reduce the potential attack surface.

TFTP does not have built-in encryption, access control or authentication. This makes it very easy for an attacker to exploit TFTP to gain access to files

Remediation:

Run the following commands to stop `tftpd-hpa.service` , and remove the `tftpd-hpa` package:

```
# systemctl stop tftpd-hpa.service
# apt purge tftpd-hpa
```

- OR -

- IF - the `tftpd-hpa` package is required as a dependency:

Run the following commands to stop and mask `tftpd-hpa.service` :

```
# systemctl stop tftpd-hpa.service
# systemctl mask tftpd-hpa.service
```

Impact:

TFTP is often used to provide files for network booting such as for PXE based installation of servers.

There may be packages that are dependent on the `tftpd-hpa` package. If the `tftpd-hpa` package is removed, these dependent packages will be removed as well. Before removing the `tftpd-hpa` package, review any dependent packages to determine if they are required on the system.

- IF - a dependent package is required: stop and mask `tftpd-hpa.service` leaving the `tftpd-hpa` package installed.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

2.1.17 Ensure web proxy server services are not in use

Pass

Description:

`squid`- OR `-squid-openssl` is a standard proxy server used in many distributions and environments.

Rationale:

Unless a system is specifically set up to act as a proxy server, it is recommended that the `squid` package be removed to reduce the potential attack surface.

Note: Several HTTP proxy servers exist. These should be checked and removed unless required.

Remediation:

Run the following commands to stop `squid.service` and remove the `squid` and `squid-openssl` packages:

```
# systemctl stop squid.service
# apt purge squid squid-openssl
```

- OR - If the `squid` package is required as a dependency:

Run the following commands to stop and mask the `squid.service` :

```
# systemctl stop squid.service
# systemctl mask squid.service
```

Impact:

There may be packages that are dependent on the `squid- OR -squid-openssl` package. If the `squid- OR -squid-openssl` package is removed, these dependent packages will be removed as well. Before removing the `squid` package, review any dependent packages to determine if they are required on the system.

- **IF** - a dependent package is required: stop and mask the `squid.service` leaving the `squid- OR -squid-openssl` package installed.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

2.1.18 Ensure web server services are not in use

Pass

Description:

Web servers provide the ability to host web site content.

Rationale:

Unless there is a local site approved requirement to run a web server service on the system, web server packages should be removed to reduce the potential attack surface.

Remediation:

Run the following commands to stop `httpd.socket` , `httpd.service` , and `nginx.service` , and remove `apache2` and `nginx` packages:

```
# systemctl stop apache2.socket apache2.service nginx.service
# apt purge apache2 nginx
```

- **OR** -

- **IF** - a package is installed **and** is required for dependencies:

Run the following commands to stop and mask `apache2.socket` , `apache2.service` , and `nginx.service` :

```
# systemctl stop apache2.socket apache2.service nginx.service
# systemctl mask apache2.socket apache2.service nginx.service
```

Note: Other web server packages may exist. If not required and authorized by local site policy, they should also be removed. If the package is required for a dependency, the service and socket should be stopped and masked.

Impact:

Removal of web server packages will remove that ability for the server to host web services.

- **IF** - the web server package is required for a dependency, any related service or socket should be stopped and masked.

Note: If the remediation steps to mask a service are followed and that package is not installed on the system, the service and/or socket will still be masked. If the package is installed due to an approved requirement to host a web server, the associated service and/or socket would need to be unmasked before it could be enabled and/or started.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

2.1.19 Ensure xinetd services are not in use

Pass

Description:

The eXtended InterNET Daemon (`xinetd`) is an open source super daemon that replaced the original `inetd` daemon. The `xinetd` daemon listens for well known services and dispatches the appropriate daemon to properly respond to service requests.

Rationale:

If there are no `xinetd` services required, it is recommended that the package be removed to reduce the attack surface are of the system.

Note: If an `xinetd` service or services are required, ensure that any `xinetd` service not required is stopped and masked

Remediation:

Run the following commands to stop `xinetd.service` , and remove the `xinetd` package:

```
# systemctl stop xinetd.service
# apt purge xinetd
```

-OR-

-IF- the `xinetd` package is required as a dependency:

Run the following commands to stop and mask the `xinetd.service` :

```
# systemctl stop xinetd.service
# systemctl mask xinetd.service
```

Impact:

There may be packages that are dependent on the `xinetd` package. If the `xinetd` package is removed, these dependent packages will be removed as well. Before removing the `xinetd` package, review any dependent packages to determine if they are required on the system.

-IF- a dependent package is required: stop and mask `xinetd.service` leaving the `xinetd` package installed.

Assessment:

[Show](#) Assessment Evidence

[Show Rule Result XML](#)

References:

- URL: NIST SP 800-53 Rev. 5: CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

2.1.21 Ensure mail transfer agents are configured for local-only mode

Pass

Description:

Mail Transfer Agents (MTA), such as sendmail and Postfix, are used to listen for incoming mail and transfer the messages to the appropriate user or mail server. If the system is not intended to be a mail server, it is recommended that the MTA be configured to only process local mail.

Rationale:

The software for all Mail Transfer Agents is complex and most have a long history of security issues. While it is important to ensure that the system can process local mail messages, it is not necessary to have the MTA's daemon listening on a port unless the server is intended to be a mail server that receives and processes mail from other systems.

Remediation:

Edit `/etc/postfix/main.cf` and add the following line to the RECEIVING MAIL section. If the line already exists, change it to look like the line below:

```
inet_interfaces = loopback-only
```

Run the following command to restart postfix :

```
# systemctl restart postfix
```

Note:

- This recommendation is designed around the postfix mail server.
- Depending on your environment you may have an alternative MTA installed such as exim4 or sendmail. If this is the case consult the documentation for your installed MTA to configure the recommended state.

Assessment:

[Show Assessment Evidence](#)

[Show Rule Result XML](#)

References:

- URL: NIST SP 800-53 Rev. 5: CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

2.1.22 Ensure only approved services are listening on a network interface

Manual

Description:

A network port is identified by its number, the associated IP address, and the type of the communication protocol such as TCP or UDP.

A listening port is a network port on which an application or process listens on, acting as a communication endpoint.

Each listening port can be open or closed (filtered) using a firewall. In general terms, an open port is a network port that accepts incoming packets from remote locations.

Rationale:

Services listening on the system pose a potential risk as an attack vector. These services should be reviewed, and if not required, the service should be stopped, and the package containing the service should be removed. If required packages have a dependency, the service should be stopped and masked to reduce the attack surface of the system.

Remediation:

Run the following commands to stop the service and remove the package containing the service:

```
# systemctl stop <service_name>.socket <service_name>.service
# apt purge <package_name>
```

- OR - If required packages have a dependency:

Run the following commands to stop and mask the service and socket:

```
# systemctl stop <service_name>.socket <service_name>.service
# systemctl mask <service_name>.socket <service_name>.service
```

Note: replace <service_name> with the appropriate service name.

Impact:

There may be packages that are dependent on the service's package. If the service's package is removed, these dependent packages will be removed as well. Before removing the service's package, review any dependent packages to determine if they are required on the system.

- IF - a dependent package is required: stop and mask the <service_name>.socket and <service_name>.service leaving the service's package installed.

[Show Rule Result XML](#)

References:

- URL: NIST SP 800-53 Rev. 5: CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#) >

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#) >

[Back to Summary](#)

2.2 Configure Client Services

A number of insecure services exist. While disabling the servers prevents a local attack against these services, it is advised to remove their clients unless they are required.

Note: This should not be considered a comprehensive list of insecure service clients. You may wish to consider additions to those listed here for your environment.

2.2.1 Ensure nis client is not installed

Pass

Description:

The Network Information Service (NIS), formerly known as Yellow Pages, is a client-server directory service protocol used to distribute system configuration files. The NIS client was used to bind a machine to an NIS server and receive the distributed configuration files.

Rationale:

The NIS service is inherently an insecure system that has been vulnerable to DOS attacks, buffer overflows and has poor authentication for querying NIS maps. NIS generally has been replaced by such protocols as Lightweight Directory Access Protocol (LDAP). It is recommended that the service be removed.

Remediation:

Uninstall `nis` :

```
# apt purge nis
```

Impact:

Many insecure service clients are used as troubleshooting tools and in testing environments. Uninstalling them can inhibit capability to test and troubleshoot. If they are required it is advisable to remove the clients after use to prevent accidental or intentional misuse.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: CM-7, CM-11

CIS Controls V7.0:

- **Control 2: Inventory and Control of Software Assets:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 4: Secure Configuration of Enterprise Assets and Software:** -- [More](#)
>

[Back to Summary](#)

2.2.2 Ensure rsh client is not installed

Pass

Description:

The `rsh-client` package contains the client commands for the `rsh` services.

Rationale:

These legacy clients contain numerous security exposures and have been replaced with the more secure SSH package. Even if the server is removed, it is best to ensure the clients are also removed to prevent users from inadvertently attempting to use these commands and therefore exposing their credentials. Note that removing the `rsh-client` package removes the clients for `rsh`, `rcp` and `rlogin`.

Remediation:

Uninstall `rsh` :

```
# apt purge rsh-client
```

Impact:

Many insecure service clients are used as troubleshooting tools and in testing environments. Uninstalling them can inhibit capability to test and troubleshoot. If they are required it is advisable to remove the clients after use to prevent accidental or intentional misuse.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)**2.2.3 Ensure talk client is not installed**

Pass

Description:

The `talk` software makes it possible for users to send and receive messages across systems through a terminal session. The `talk` client, which allows initialization of talk sessions, is installed by default.

Rationale:

The software presents a security risk as it uses unencrypted protocols for communication.

Remediation:

Uninstall `talk` :

```
# apt purge talk
```

Impact:

Many insecure service clients are used as troubleshooting tools and in testing environments. Uninstalling them can inhibit capability to test and troubleshoot. If they are required it is advisable to remove the clients after use to prevent accidental or intentional misuse.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)**2.2.4 Ensure telnet client is not installed**

Fail

Description:

The `telnet` package contains the `telnet` client, which allows users to start connections to other systems via the telnet protocol.

Rationale:

The `telnet` protocol is insecure and unencrypted. The use of an unencrypted transmission medium could allow an unauthorized user to steal credentials. The `ssh` package provides an encrypted session and stronger security and is included in most Linux distributions.

Remediation:

Uninstall telnet :

```
# apt purge telnet
```

Impact:

Many insecure service clients are used as troubleshooting tools and in testing environments. Uninstalling them can inhibit capability to test and troubleshoot. If they are required it is advisable to remove the clients after use to prevent accidental or intentional misuse.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-7, CM-11

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

2.2.5 Ensure ldap client is not installed

Pass

Description:

The Lightweight Directory Access Protocol (LDAP) was introduced as a replacement for NIS/YP. It is a service that provides a method for looking up information from a central database.

Rationale:

If the system will not need to act as an LDAP client, it is recommended that the software be removed to reduce the potential attack surface.

Remediation:

Uninstall `ldap-utils` :

```
# apt purge ldap-utils
```

Impact:

Removing the LDAP client will prevent or inhibit using LDAP for authentication in your environment.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

2.2.6 Ensure ftp client is not installed

Fail

Description:

FTP (File Transfer Protocol) is a traditional and widely used standard tool for transferring files between a server and clients over a network, especially where no authentication is necessary (permits anonymous users to connect to a server).

Rationale:

FTP does not protect the confidentiality of data or authentication credentials. It is recommended SFTP be used if file transfer is required. Unless there is a need to run the system as a FTP server (for example, to allow anonymous downloads), it is recommended that the package be removed to reduce the potential attack surface.

Remediation:

Run the following command to uninstall `ftp` :

```
# apt purge ftp
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-7, CM-11

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

2.3 Configure Time Synchronization

It is recommended that systems be configured to synchronize their time using a service such as `systemd-timesyncd` , or `chrony` .

Virtual systems may be configured to receive their time synchronization from their host system.

The host system must be configured to synchronize its time from an authoritative source to be considered compliant with this section.

Any "physical" clock present on a system should be synchronized from an authoritative time source.

Only one time synchronization method should be in use on the system

Notes: Only the section related to the time synchronization method in use on the system should be followed, all other time synchronization recommendations should be skipped

2.3.1 Ensure time synchronization is in use

It is recommended that physical systems and virtual guests lacking direct access to the physical host's clock be configured to synchronize their time using a service such as `systemd-timesyncd` , or `chrony` .

2.3.1.1 Ensure a single time synchronization daemon is in use

Pass

Description:

System time should be synchronized between all systems in an environment. This is typically done by establishing an authoritative time server or set of servers and having all systems synchronize their clocks to them.

Note:

- On virtual systems where host based time synchronization is available consult your virtualization software documentation and verify that host based synchronization is in use and follows local site policy. In this scenario, this section should be skipped
- Only **one** time synchronization method should be in use on the system. Configuring multiple time synchronization methods could lead to unexpected or unreliable results

Rationale:

Time synchronization is important to support time sensitive security mechanisms and ensures log files have consistent time records across the enterprise, which aids in forensic investigations.

Remediation:

On physical systems, and virtual systems where host based time synchronization is not available.

Select **one** of the two time synchronization daemons; `chrony` (1) or `systemd-timesyncd` (2) and following the remediation procedure for the selected daemon.

Note: enabling more than one synchronization daemon could lead to unexpected or unreliable results:

1. `chrony`

Run the following command to install `chrony` :

```
# apt install chrony
```

Run the following commands to stop and mask the `systemd-timesyncd` daemon:

```
# systemctl stop systemd-timesyncd.service
```

```
# systemctl mask systemd-timesyncd.service
```

Note:

- Subsection: **Configure chrony** should be followed
- Subsection: **Configure systemd-timesyncd** should be skipped

2. systemd-timesyncd

Run the following command to remove the chrony package:

```
# apt purge chrony
# apt autoremove chrony
```

Note:

- Subsection: **Configure systemd-timesyncd** should be followed
- Subsection: **Configure chrony** should be skipped

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AU-3, AU-12

CIS Controls V7.0:

- Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 8: Audit Log Management: -- [More](#)
>

[Back to Summary](#)

2.3.2 Configure systemd-timesyncd

`systemd-timesyncd` is a daemon that has been added for synchronizing the system clock across the network. It implements an SNTP client. In contrast to NTP implementations such as `chrony` or the NTP reference server this only implements a client side, and does not bother with the full NTP complexity, focusing only on querying time from one remote server and synchronizing the local clock to it. The daemon runs with minimal privileges, and has been hooked up with `networkd` to only operate when network connectivity is available. The daemon saves the current clock to disk every time a new NTP sync has been acquired, and uses this to possibly correct the system clock early at bootup, in order to accommodate for systems that lack an RTC such as the Raspberry Pi and embedded devices, and make sure that time monotonically progresses on these systems, even if it is not always correct. To make use of this daemon a new system user and group "`systemd-timesync`" needs to be created on installation of `systemd`.

The default configuration is set during compilation, so configuration is only needed when it is necessary to deviate from those defaults. Initially, the main configuration file in `/etc/systemd/` contains commented out entries showing the defaults as a guide to the administrator. Local overrides can be created by editing this file or by creating drop-ins, as described below. Using drop-ins for local configuration is recommended over modifications to the main configuration file.

In addition to the "main" configuration file, drop-in configuration snippets are read from `/usr/lib/systemd/*.conf.d/`, `/usr/local/lib/systemd/*.conf.d/`, and `/etc/systemd/*.conf.d/`. Those drop-ins have higher precedence and override the main configuration file. Files in the `/*.conf.d/` configuration subdirectories are sorted by their filename in lexicographic order, regardless of in which of the subdirectories they reside. When multiple files specify the same option, for options which accept just a single value, the entry in the file sorted last takes precedence, and for options which accept a list of values, entries are collected as they occur in the sorted files.

When packages need to customize the configuration, they can install drop-ins under `/usr/`. Files in `/etc/` are reserved for the local administrator, who may use this logic to override the configuration files installed by vendor packages. Drop-ins have to be used to override package drop-ins, since the main configuration file has lower precedence. It is recommended to prefix all filenames in those subdirectories with a two-digit number and a dash, to simplify the ordering of the files.

To disable a configuration file supplied by the vendor, the recommended way is to place a symlink to `/dev/null` in the configuration directory in `/etc/`, with the same filename as the vendor configuration file.

Note:

- The recommendations in this section only apply if `timesyncd` is in use on the system
- The `systemd-timesyncd` service specifically implements only SNTP.
 - This minimalistic service will set the system clock for large offsets or slowly adjust it for smaller deltas
 - More complex use cases are not covered by `systemd-timesyncd`
- If `chrony` is used, `systemd-timesyncd` should be stopped and masked, and this section skipped
- One, and only one, time synchronization method should be in use on the system

2.3.2.1 Ensure systemd-timesyncd configured with authorized timeserver

Fail

Description:

NTP=

- A space-separated list of NTP server host names or IP addresses. During runtime this list is combined with any per-interface NTP servers acquired from `systemd-networkd.service(8)`. `systemd-timesyncd` will contact all configured system or per-interface servers in turn, until one responds. When the empty string is assigned, the list of NTP servers is reset, and all prior assignments will have no effect. This setting defaults to an empty list.

FallbackNTP=

- A space-separated list of NTP server host names or IP addresses to be used as the fallback NTP servers. Any per-interface NTP servers obtained from `systemd-networkd.service(8)` take precedence over this setting, as do any servers set via `NTP=` above. This setting is hence only relevant if no other NTP server information is known. When the empty string is assigned, the list of NTP servers is reset, and all prior assignments will have no effect. If this option is not given, a compiled-in list of NTP servers is used.

Rationale:

Time synchronization is important to support time sensitive security mechanisms and to ensure log files have consistent time records across the enterprise to aid in forensic investigations

Remediation:

Set NTP and/or FallbackNPT parameters to local site approved authoritative time server(s) in `/etc/systemd/timesyncd.conf` or a file in `/etc/systemd/timesyncd.conf.d/` ending in `.conf` in the `[Time]` section:

Example file:

```
[Time]

NTP=time.nist.gov # Uses the generic name for NIST's time servers

FallbackNTP=time-a-g.nist.gov time-b-g.nist.gov time-c-g.nist.gov # Space separated list of NIST time servers
```

Example script to create systemd drop-in configuration file:

```
#!/usr/bin/env bash

{

a_settings=("NTP=time.nist.gov" "FallbackNTP=time-a-g.nist.gov time-b-g.nist.gov time-c-g.nist.gov")

[ ! -d /etc/systemd/timesyncd.conf.d/ ] && mkdir /etc/systemd/timesyncd.conf.d/

if grep -Psq -- '^h*[Time\]' /etc/systemd/timesyncd.conf.d/60-timesyncd.conf; then

printf '%s\n' " " "${a_settings[@]}" >> /etc/systemd/timesyncd.conf.d/60-timesyncd.conf

else

printf '%s\n' " " "[Time]" "${a_settings[@]}" >> /etc/systemd/timesyncd.conf.d/60-timesyncd.conf

fi

}
```

Note: If this setting appears in a canonically later file, or later in the same file, the setting will be overwritten

Run to following command to update the parameters in the service:

```
# systemctl reload-or-restart systemd-timesyncd
```

Assessment:

[Show](#) Assessment Evidence

[Show Rule Result XML](#)**References:**

- **URL:** <https://www.freedesktop.org/software/systemd/man/timesyncd.conf.html>
- **URL:** <https://tf.nist.gov/tf-cgi/servers.cgi>
- **URL:** NIST SP 800-53 Rev. 5: AU-7, AU-8

CIS Controls V7.0:

- **Control 6: Maintenance, Monitoring and Analysis of Audit Logs:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 8: Audit Log Management:** -- [More](#)
>

[Back to Summary](#)**2.3.2.2 Ensure systemd-timesyncd is enabled and running**

Pass

Description:

systemd-timesyncd is a daemon that has been added for synchronizing the system clock across the network

Rationale:

systemd-timesyncd needs to be enabled and running in order to synchronize the system to a timeserver.

Time synchronization is important to support time sensitive security mechanisms and to ensure log files have consistent time records across the enterprise to aid in forensic investigations

Remediation:

- **IF** `systemd-timesyncd` is in use on the system, run the following commands:

Run the following command to unmask `systemd-timesyncd.service` :

```
# systemctl unmask systemd-timesyncd.service
```

Run the following command to enable and start `systemd-timesyncd.service` :

```
# systemctl --now enable systemd-timesyncd.service
```

- OR -

If another time synchronization service is in use on the system, run the following command to stop and mask `systemd-timesyncd` :

```
# systemctl --now mask systemd-timesyncd.service
```

Assessment:[Show Assessment Evidence](#)[Show Rule Result XML](#)**References:**

- **URL:** NIST SP 800-53 Rev. 5: AU-7, AU-8

CIS Controls V7.0:

- **Control 6: Maintenance, Monitoring and Analysis of Audit Logs:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 8: Audit Log Management:** -- [More](#)
>

[Back to Summary](#)**2.3.3 Configure chrony**

`chrony` is a daemon which implements the Network Time Protocol (NTP) and is designed to synchronize system clocks across a variety of systems and use a source that is highly accurate.

`chrony` can be configured to be a client and/or a server.

More information on `chrony` can be found at: <http://chrony.tuxfamily.org/> .

Note:

- If `systemd-timesyncd` is being used, `chrony` should be removed and this section skipped
- Only one time synchronization method should be in use on the system

2.3.3.1 Ensure `chrony` is configured with authorized timeserver

Pass

Description:

- `server`
 - The `server` directive specifies an NTP server which can be used as a time source. The client-server relationship is strictly hierarchical: a client might synchronize its system time to that of the server, but the server's system time will never be influenced by that of a client.
 - This directive can be used multiple times to specify multiple servers.
 - The directive is immediately followed by either the name of the server, or its IP address.
- `pool`
 - The syntax of this directive is similar to that for the `server` directive, except that it is used to specify a pool of NTP servers rather than a single NTP server. The pool name is expected to resolve to multiple addresses which might change over time.
 - This directive can be used multiple times to specify multiple pools.
 - All options valid in the `server` directive can be used in this directive too.

Rationale:

Time synchronization is important to support time sensitive security mechanisms and to ensure log files have consistent time records across the enterprise to aid in forensic investigations

Remediation:

Edit `/etc/chrony/chrony.conf` or a file ending in `.sources` in `/etc/chrony/sources.d/` and add or edit `server` or `pool` lines as appropriate according to local site policy:

Edit the `Chrony` configuration and add or edit the `server` and/or `pool` lines returned by the Audit Procedure as appropriate according to local site policy

```
<[server|pool]> <[remote-server|remote-pool]>
```

Example script to add a drop-in configuration for the `pool` directive:

```
#!/usr/bin/env bash

{
[ ! -d "/etc/chrony/sources.d/" ] && mkdir /etc/chrony/sources.d/
printf '%s\n' "" "#The maxsources option is unique to the pool directive" \
"pool time.nist.gov iburst maxsources 4" >> /etc/chrony/sources.d/60-sources.sources
chronyc reload sources &>/dev/null
}
```

Example script to add a drop-in configuration for the `server` directive:

```
#!/usr/bin/env bash

{
[ ! -d "/etc/chrony/sources.d/" ] && mkdir /etc/chrony/sources.d/
printf '%s\n' "" "server time-a-g.nist.gov iburst" "server 132.163.97.3 iburst" \
```

```
"server time-d-b.nist.gov iburst" >> /etc/chrony/sources.d/60-sources.sources

chronyc reload sources &>/dev/null

}
```

Run the following command to reload the `chronyd` config:

```
# systemctl reload-or-restart chronyd
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: [chrony.conf\(5\) Manual Page](#)
- URL: <https://tf.nist.gov/tf-cgi/servers.cgi>
- URL: NIST SP 800-53 Rev. 5: AU-3, AU-12

CIS Controls V7.0:

- Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 8: Audit Log Management: -- [More](#)
>

[Back to Summary](#)

2.3.3.2 Ensure chrony is running as user `_chrony`

Pass

Description:

The `chrony` package is installed with a dedicated user account `_chrony` . This account is granted the access required by the `chronyd` service

Rationale:

The `chronyd` service should run with only the required privlidges

Remediation:

Add or edit the user line to `/etc/chrony/chrony.conf` or a file ending in `.conf` in `/etc/chrony/conf.d/` :

```
user _chrony
```

- OR -

If another time synchronization service is in use on the system, run the following command to remove `chrony` from the system:

```
# apt purge chrony

# apt autoremove chrony
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AU-8

CIS Controls V7.0:

- Not Explicitly Mapped: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Not Explicitly Mapped: -- [More](#)
>

[Back to Summary](#)

2.3.3.3 Ensure chrony is enabled and running

Pass

Description:

chrony is a daemon for synchronizing the system clock across the network

Rationale:

chrony needs to be enabled and running in order to synchronize the system to a timeserver.

Time synchronization is important to support time sensitive security mechanisms and to ensure log files have consistent time records across the enterprise to aid in forensic investigations

Remediation:

- **IF** -chrony is in use on the system, run the following commands:

Run the following command to unmask `chrony.service` :

```
# systemctl unmask chrony.service
```

Run the following command to enable and start `chrony.service` :

```
# systemctl --now enable chrony.service
```

- **OR** -

If another time synchronization service is in use on the system, run the following command to remove `chrony` :

```
# apt purge chrony
# apt autoremove chrony
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AU-8

CIS Controls V7.0:

- Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 8: Audit Log Management: -- [More](#)
>

[Back to Summary](#)

2.4 Job Schedulers

A job scheduler is used to execute jobs, commands, or shell scripts, at fixed times, dates, or intervals

2.4.1 Configure cron

`cron` is a time based job scheduler

- **IF** -`cron` is not installed on the system, this sub section can be skipped

Note: Other methods such as `systemd timers` exist for scheduling jobs. If an alternate method is in use, it should be secured in accordance with local site policy

2.4.1.1 Ensure cron daemon is enabled and active

Pass

Description:

The cron daemon is used to execute batch jobs on the system.

Rationale:

While there may not be user jobs that need to be run on the system, the system does have maintenance jobs that may include security monitoring that have to run, and cron is used to execute them.

Remediation:

- IF - cron is installed on the system:

Run the following commands to unmask, enable, and start cron :

```
# systemctl unmask "$(systemctl list-unit-files | awk '$1~/^cron?\.service/{print $1}')"
# systemctl --now enable "$(systemctl list-unit-files | awk '$1~/^cron?\.service/{print $1}')
```

Assessment:

Show Assessment Evidence

Show Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5

CIS Controls V7.0:

- Not Explicitly Mapped: -- More >

CIS Critical Security Controls V8.0:

- Not Explicitly Mapped: -- More >

Back to Summary

2.4.1.2 Ensure access to /etc/crontab is configured

Fail

Description:

The /etc/crontab file is used by cron to control its own jobs. The commands in this item make sure that root is the user and group owner of the file and that only the owner can access the file.

Rationale:

This file contains information on what system jobs are run by cron. Write access to these files could provide unprivileged users with the ability to elevate their privileges. Read access to these files could provide users with the ability to gain insight on system jobs that run on the system and could provide them a way to gain unauthorized privileged access.

Remediation:

- IF - cron is installed on the system:

Run the following commands to set ownership and permissions on /etc/crontab :

```
# chown root:root /etc/crontab
# chmod og-rwx /etc/crontab
```

Assessment:

Show Assessment Evidence

Show Rule Result XML

References:

- [URL: NIST SP 800-53 Rev. 5: AC-3, MP-2](#)

CIS Controls V7.0:

- [Control 14: Controlled Access Based on the Need to Know: -- More](#)
>

CIS Critical Security Controls V8.0:

- [Control 3: Data Protection: -- More](#)
>

[Back to Summary](#)

2.4.1.3 Ensure access to /etc/cron.hourly is configured

Fail

Description:

This directory contains system `cron` jobs that need to run on an hourly basis. The files in this directory cannot be manipulated by the `crontab` command, but are instead edited by system administrators using a text editor. The commands below restrict read/write and search access to user and group root, preventing regular users from accessing this directory.

Rationale:

Granting write access to this directory for non-privileged users could provide them the means for gaining unauthorized elevated privileges. Granting read access to this directory could give an unprivileged user insight in how to gain elevated privileges or circumvent auditing controls.

Remediation:

- **IF** - cron is installed on the system:

Run the following commands to set ownership and permissions on the `/etc/cron.hourly` directory:

```
# chown root:root /etc/cron.hourly/  
# chmod og-rwx /etc/cron.hourly/
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- [URL: NIST SP 800-53 Rev. 5: AC-3. MP-2](#)

CIS Controls V7.0:

- [Control 14: Controlled Access Based on the Need to Know: -- More](#)
>

CIS Critical Security Controls V8.0:

- [Control 3: Data Protection: -- More](#)
>

[Back to Summary](#)

2.4.1.4 Ensure access to /etc/cron.daily is configured

Fail

Description:

The `/etc/cron.daily` directory contains system cron jobs that need to run on a daily basis. The files in this directory cannot be manipulated by the `crontab` command, but are instead edited by system administrators using a text editor. The commands below restrict read/write and search access to user and group root, preventing regular users from accessing this directory.

Rationale:

Granting write access to this directory for non-privileged users could provide them the means for gaining unauthorized elevated privileges. Granting read access to this directory could give an unprivileged user insight in how to gain

elevated privileges or circumvent auditing controls.

Remediation:

- **IF** - cron is installed on the system:

Run the following commands to set ownership and permissions on the `/etc/cron.daily` directory:

```
# chown root:root /etc/cron.daily/
# chmod og-rwx /etc/cron.daily/
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AC-3. MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

2.4.1.5 Ensure access to `/etc/cron.weekly` is configured

Fail

Description:

The `/etc/cron.weekly` directory contains system cron jobs that need to run on a weekly basis. The files in this directory cannot be manipulated by the `crontab` command but are instead edited by system administrators using a text editor. The commands below restrict read/write and search access to user and group root, preventing regular users from accessing this directory.

Rationale:

Granting write access to this directory for non-privileged users could provide them the means for gaining unauthorized elevated privileges. Granting read access to this directory could give an unprivileged user insight in how to gain elevated privileges or circumvent auditing controls.

Remediation:

- **IF** - cron is installed on the system:

Run the following commands to set ownership and permissions on the `/etc/cron.weekly` directory:

```
# chown root:root /etc/cron.weekly/
# chmod og-rwx /etc/cron.weekly/
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AC-3. MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

2.4.1.6 Ensure access to /etc/cron.monthly is configured

Fail

Description:

The `/etc/cron.monthly` directory contains system cron jobs that need to run on a monthly basis. The files in this directory cannot be manipulated by the `crontab` command but are instead edited by system administrators using a text editor. The commands below restrict read/write and search access to user and group root, preventing regular users from accessing this directory.

Rationale:

Granting write access to this directory for non-privileged users could provide them the means for gaining unauthorized elevated privileges. Granting read access to this directory could give an unprivileged user insight in how to gain elevated privileges or circumvent auditing controls.

Remediation:

- **IF** - cron is installed on the system:

Run the following commands to set ownership and permissions on the `/etc/cron.monthly` directory:

```
# chown root:root /etc/cron.monthly/  
# chmod og-rwx /etc/cron.monthly/
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: AC-3. MP-2

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

2.4.1.7 Ensure access to /etc/cron.yearly is configured

Pass

Description:

The `/etc/cron.yearly` directory contains system cron jobs that need to run on an annual basis. The files in this directory cannot be manipulated by the `crontab` command but are instead edited by system administrators using a text editor. The commands below restrict read/write and search access to user and group root, preventing regular users from accessing this directory.

Rationale:

Granting write access to this directory for non-privileged users could provide them the means for gaining unauthorized elevated privileges. Granting read access to this directory could give an unprivileged user insight in how to gain elevated privileges or circumvent auditing controls.

Remediation:

- **IF** - cron is installed on the system:

Run the following commands to set ownership and permissions on the `/etc/cron.yearly` directory:

```
# chown root:root /etc/cron.yearly/

# chmod og-rwx /etc/cron.yearly/
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AC-3. MP-2

[Back to Summary](#)

2.4.1.8 Ensure access to /etc/cron.d is configured

Fail

Description:

The `/etc/cron.d` directory contains system `cron` jobs that need to run in a similar manner to the hourly, daily weekly and monthly jobs from `/etc/crontab` , but require more granular control as to when they run. The files in this directory cannot be manipulated by the `crontab` command but are instead edited by system administrators using a text editor. The commands below restrict read/write and search access to user and group root, preventing regular users from accessing this directory.

Rationale:

Granting write access to this directory for non-privileged users could provide them the means for gaining unauthorized elevated privileges. Granting read access to this directory could give an unprivileged user insight in how to gain elevated privileges or circumvent auditing controls.

Remediation:

- IF - cron is installed on the system:

Run the following commands to set ownership and permissions on the `/etc/cron.d` directory:

```
# chown root:root /etc/cron.d/

# chmod og-rwx /etc/cron.d/
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AC-3. MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

2.4.1.9 Ensure access to crontab is configured

Fail

Description:

`crontab` is the program used to install, deinstall, or list the tables used to drive the cron daemon. Each user can have their own crontab, and though files are created, they are not intended to be edited directly.

If the `/etc/cron.allow` file exists, then you must be listed (one user per line) therein in order to be allowed to use this command. If the `/etc/cron.allow` file does not exist but the `/etc/cron.deny` file does exist, then you must

not be listed in the `/etc/cron.deny` file in order to use this command.

If neither of these files exists, then depending on site-dependent configuration parameters, only the super user will be allowed to use this command, or all users will be able to use this command.

If both files exist then `/etc/cron.allow` takes precedence. Which means that `/etc/cron.deny` is not considered and your user must be listed in `/etc/cron.allow` in order to be able to use the crontab.

Regardless of the existence of any of these files, the root administrative user is always allowed to setup a crontab.

The files `/etc/cron.allow` and `/etc/cron.deny`, if they exist, must be either world-readable, or readable by group `crontab`. If they are not, then cron will deny access to all users until the permissions are fixed.

There is one file for each user's crontab. Users are not allowed to edit the file directly to ensure that only users allowed by the system to run periodic tasks can add them, and only syntactically correct crontabs will be written. This is enforced by having the directory writable only by the `crontab` group and configuring `crontab` command with the `setgid` bit set for that specific group.

Note:

- Even though a given user is not listed in `cron.allow`, cron jobs can still be run as that user
- The files `/etc/cron.allow` and `/etc/cron.deny`, if they exist, only controls administrative access to the `crontab` command for scheduling and modifying cron jobs

Rationale:

On many systems, only the system administrator is authorized to schedule `cron` jobs. Using the `cron.allow` file to control who can run `cron` jobs enforces this policy. It is easier to manage an allow list than a deny list. In a deny list, you could potentially add a user ID to the system and forget to add it to the deny files.

Remediation:

- **IF** - cron is installed on the system:

Run the following script to:

- Create `/etc/cron.allow` if it doesn't exist
- Change owner to user `root`
- Change group owner to group `root`- **OR** - group `crontab` if it exists
- Change mode to `640` or more restrictive

```
#!/usr/bin/env bash

{
[ ! -e "/etc/cron.deny" ] && touch /etc/cron.allow

chmod u-x,g-wx,o-rwx /etc/cron.allow

if grep -Pq -- '^h*crontab\: ' /etc/group; then
chown root:crontab /etc/cron.allow
else
chown root:root /etc/cron.allow
fi
}
```

- **IF** `/etc/cron.deny` exists, run the following script to:

- Change owner to user `root`
- Change group owner to group `root`- **OR** - group `crontab` if it exists
- Change mode to `640` or more restrictive

```
#!/usr/bin/env bash

{
```

```
if [ -e "/etc/cron.deny" ]; then
    chmod u-x,g-wx,o-rwx /etc/cron.deny
    if grep -Pq -- '^h*crontab\: ' /etc/group; then
        chown root:crontab /etc/cron.deny
    else
        chown root:root /etc/cron.deny
    fi
fi
}
```

Note: On systems where cron is configured to use the group `crontab`, if the group `crontab` is not set as the owner of `cron.allow`, then cron will deny access to all users and you will see an error similar to:

```
You (<USERNAME>) are not allowed to use this program (crontab)
See crontab(1) for more information
```

Assessment:

[Show Assessment Evidence](#)

[Show Rule Result XML](#)

References:

- URL: NIST SP 800-53 Rev. 5: AC-3. MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

2.4.2 Configure at

`at` is a command-line utility used to schedule a job for later execution

Note: if `at` is not installed on the system, this section can be skipped

2.4.2.1 Ensure access to at is configured

Pass

Description:

`at` allows fairly complex time specifications, extending the POSIX.2 standard. It accepts times of the form HH:MM to run a job at a specific time of day. (If that time is already past, the next day is assumed.) You may also specify midnight, noon, or teatime (4pm) and you can have a time-of-day suffixed with AM or PM for running in the morning or the evening. You can also say what day the job will be run, by giving a date in the form month-name day with an optional year, or giving a date of the form MMDD[CC]YY, MM/DD/[CC]YY, DD.MM.[CC]YY or [CC]YY-MM-DD. The specification of a date must follow the specification of the time of day. You can also give times like now + count time-units, where the time-units can be minutes, hours, days, or weeks and you can tell `at` to run the job today by suffixing the time with today and to run the job tomorrow by suffixing the time with tomorrow.

The `/etc/at.allow` and `/etc/at.deny` files determine which user can submit commands for later execution via `at` or `batch`. The format of the files is a list of usernames, one on each line. Whitespace is not permitted. If the file `/etc/at.allow` exists, only usernames mentioned in it are allowed to use `at`. If `/etc/at.allow` does not exist, `/etc/at.deny` is checked, every username not mentioned in it is then allowed to use `at`. An empty `/etc/at.deny` means that every user may use `at`. If neither file exists, only the superuser is allowed to use `at`.

Rationale:

On many systems, only the system administrator is authorized to schedule `at` jobs. Using the `at.allow` file to control who can run `at` jobs enforces this policy. It is easier to manage an allow list than a deny list. In a deny list, you could potentially add a user ID to the system and forget to add it to the deny files.

Remediation:

- **IF** - `at` is installed on the system:

Run the following script to:

- `/etc/at.allow`:
 - Create the file if it doesn't exist
 - Change owner or user `root`
 - If group `daemon` exists, change to group `daemon` , else change group to `root`
 - Change mode to `640` or more restrictive
- - **IF** - `/etc/at.deny` exists:
 - Change owner or user `root`
 - If group `daemon` exists, change to group `daemon` , else change group to `root`
 - Change mode to `640` or more restrictive

```
#!/usr/bin/env bash

{
grep -Pq -- '^daemon\b' /etc/group && l_group="daemon" || l_group="root"

[ ! -e "/etc/at.allow" ] && touch /etc/at.allow

chown root:"$l_group" /etc/at.allow

chmod u-x,g-wx,o-rwx /etc/at.allow

[ -e "/etc/at.deny" ] && chown root:"$l_group" /etc/at.deny

[ -e "/etc/at.deny" ] && chmod u-x,g-wx,o-rwx /etc/at.deny
}
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: AC-3. MP-2

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

3 Network

This section provides guidance on for securing the network configuration of the system

3.1 Configure Network Devices

To reduce the attack surface of a system, unused devices should be disabled.

Note: This should not be considered a comprehensive list, you may wish to consider additions to those listed here for your environment.

3.1.1 Ensure IPv6 status is identified

Manual

Description:

Internet Protocol Version 6 (IPv6) is the most recent version of Internet Protocol (IP). It's designed to supply IP addressing and additional security to support the predicted growth of connected devices. IPv6 is based on 128-bit addressing and can support 340 undecillion, which is 340,282,366,920,938,463,463,374,607,431,768,211,456 unique addresses.

Features of IPv6

- Hierarchical addressing and routing infrastructure
- Statefull and Stateless configuration
- Support for quality of service (QoS)
- An ideal protocol for neighboring node interaction

Rationale:

IETF RFC 4038 recommends that applications are built with an assumption of dual stack. It is recommended that IPv6 be enabled and configured in accordance with Benchmark recommendations.

- **IF** - dual stack and IPv6 are not used in your environment, IPv6 may be disabled to reduce the attack surface of the system, and recommendations pertaining to IPv6 can be skipped.

Note: It is recommended that IPv6 be enabled and configured unless this is against local site policy

Remediation:

Enable or disable IPv6 in accordance with system requirements and local site policy

Impact:

IETF RFC 4038 recommends that applications are built with an assumption of dual stack.

When enabled, IPv6 will require additional configuration to reduce risk to the system.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: CM-7

CIS Controls V7.0:

- **Control 9: Limitation and Control of Network Ports, Protocols, and Services:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 4: Secure Configuration of Enterprise Assets and Software:** -- [More](#)
>

[Back to Summary](#)

3.2 Configure Network Kernel Modules

The Linux kernel modules support several network protocols that are not commonly used. If these protocols are not needed, it is recommended that they be disabled in the kernel.

Note: This should not be considered a comprehensive list of uncommon network protocols, you may wish to consider additions to those listed here for your environment.

3.3 Configure Network Kernel Parameters

The following network parameters are intended for use on both host only and router systems. A system acts as a router if it has at least two interfaces and is configured to perform routing functions.

Note:

- sysctl settings are defined through files in `/usr/local/lib`, `/usr/lib/`, `/lib/`, `/run/`, and `/etc/`
- Files are typically placed in the `sysctl.d` directory within the parent directory
- The paths where sysctl preload files usually exist
 - `/run/sysctl.d/*.conf`

- `/etc/sysctl.d/*.conf`
- `/usr/local/lib/sysctl.d/*.conf`
- `/usr/lib/sysctl.d/*.conf`
- `/lib/sysctl.d/*.conf`
- `/etc/sysctl.conf`
- Files must have the " `.conf` " extension
- Vendors settings usually live in `/usr/lib/` or `/usr/local/lib/`
- To override a whole file, create a new file with the same name in `/etc/sysctl.d/` and put new settings there.
- To override only specific settings, add a file with a lexically later name in `/etc/sysctl.d/` and put new settings there.
- The command `/usr/lib/systemd/systemd-sysctl --cat-config` produces output containing The system's loaded kernel parameters and the files they're configured in:
 - Entries listed latter in the file take precedence over the same settings listed earlier in the file
 - Files containing kernel parameters that are over-riden by other files with the same name will not be listed
 - On systems running UncomplicatedFirewall, the kernel parameters may be set or over-written. This will not be visible in the output of the command
- On systems with Uncomplicated Firewall, additional settings may be configured in `/etc/ufw/sysctl.conf`
 - The settings in `/etc/ufw/sysctl.conf` will override settings other settings and **will not** be visible in the output of the `/usr/lib/systemd/systemd-sysctl --cat-config` command
 - This behavior can be changed by updating the `IPT_SYSCTL` parameter in `/etc/default/ufw`

The system's loaded kernel parameters and the files they're configured in can be viewed by running the following command:

```
# /usr/lib/systemd/systemd-sysctl --cat-config
```

3.3.1 Ensure ip forwarding is disabled

Fail

Description:

The `net.ipv4.ip_forward` and `net.ipv6.conf.all.forwarding` flags are used to tell the system whether it can forward packets or not.

Rationale:

Setting `net.ipv4.ip_forward` and `net.ipv6.conf.all.forwarding` to 0 ensures that a system with multiple interfaces (for example, a hard proxy), will never be able to forward packets, and therefore, never serve as a router.

Remediation:

Set the following parameter in `/etc/sysctl.conf` or a file in `/etc/sysctl.d/` ending in `.conf`:

- `net.ipv4.ip_forward = 0`

Example:

```
# printf '%s\n' "net.ipv4.ip_forward = 0" >> /etc/sysctl.d/60-netipv4_sysctl.conf
```

Run the following script to set the active kernel parameters:

```
#!/usr/bin/env bash

{
sysctl -w net.ipv4.ip_forward=0
sysctl -w net.ipv4.route.flush=1
}
```

- IF - IPv6 is enabled on the system:

Set the following parameter in `/etc/sysctl.conf` or a file in `/etc/sysctl.d/` ending in `.conf`:

- `net.ipv6.conf.all.forwarding = 0`

Example:

```
# printf '%s\n' "net.ipv6.conf.all.forwarding = 0" >> /etc/sysctl.d/60-netipv6_sysctl.conf
```

Run the following script to set the active kernel parameters:

```
#!/usr/bin/env bash

{

sysctl -w net.ipv6.conf.all.forwarding=0

sysctl -w net.ipv6.route.flush=1

}
```

Note: If these settings appear in a canonically later file, or later in the same file, these settings will be overwritten

Impact:

IP forwarding is required on systems configured to act as a router. If these parameters are disabled, the system will not be able to perform as a router.

Many Cloud Service Provider (CSP) hosted systems require IP forwarding to be enabled. If the system is running on a CSP platform, this requirement should be reviewed before disabling IP forwarding.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5
- **URL:** NIST SP 800-53A :: CM-6.1 (iv)
- **URL:** RHEL 8 STIG Vul ID: V-230540
- **URL:** RHEL 8 STIG Rule ID: SV-230540r858810

CIS Controls V7.0:

- **Control 9: Limitation and Control of Network Ports, Protocols, and Services:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 4: Secure Configuration of Enterprise Assets and Software:** -- [More](#)
>

[Back to Summary](#)

3.3.2 Ensure packet redirect sending is disabled

Fail

Description:

ICMP Redirects are used to send routing information to other hosts. As a host itself does not act as a router (in a host only configuration), there is no need to send redirects.

Rationale:

An attacker could use a compromised host to send invalid ICMP redirects to other router devices in an attempt to corrupt routing and have users access a system set up by the attacker as opposed to a valid system.

Remediation:

Set the following parameters in /etc/sysctl.conf or a file in /etc/sysctl.d/ ending in .conf :

- net.ipv4.conf.all.send_redirects = 0
- net.ipv4.conf.default.send_redirects = 0

Example:

```
# printf '%s\n' "net.ipv4.conf.all.send_redirects = 0" "net.ipv4.conf.default.send_redirects = 0" >> /
etc/sysctl.d/60-netip4_sysctl.conf
```

Run the following script to set the active kernel parameters:

```
#!/usr/bin/env bash
```

```
{  
  
sysctl -w net.ipv4.conf.all.send_redirects=0  
  
sysctl -w net.ipv4.conf.default.send_redirects=0  
  
sysctl -w net.ipv4.route.flush=1  
  
}
```

Note: If these settings appear in a canonically later file, or later in the same file, these settings will be overwritten

Impact:

IP forwarding is required on systems configured to act as a router. If these parameters are disabled, the system will not be able to perform as a router.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 :: CM-6 b
- **URL:** NIST SP 800-53A :: CM-6.1 (iv)
- **URL:** RHEL 8 STIG GROUP ID: V-230536
- **URL:** RHEL 8 STIG RULE ID: SV-230536r858795
- **URL:** RHEL 8 STIG GROUP ID: V-230543
- **URL:** RHEL 8 STIG RULE ID: SV-230543r858816

CIS Controls V7.0:

- **Control 9: Limitation and Control of Network Ports, Protocols, and Services:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 4: Secure Configuration of Enterprise Assets and Software:** -- [More](#)
>

[Back to Summary](#)

3.3.3 Ensure bogus icmp responses are ignored

Pass

Description:

Setting `net.ipv4.icmp_ignore_bogus_error_responses` to 1 prevents the kernel from logging bogus responses (RFC-1122 non-compliant) from broadcast reframes, keeping file systems from filling up with useless log messages.

Rationale:

Some routers (and some attackers) will send responses that violate RFC-1122 and attempt to fill up a log file system with many useless error messages.

Remediation:

Set the following parameter in `/etc/sysctl.conf` or a file in `/etc/sysctl.d/` ending in `.conf`:

- `net.ipv4.icmp_ignore_bogus_error_responses = 1`

Example:

```
# printf '%s\n' "net.ipv4.icmp_ignore_bogus_error_responses = 1" >> /etc/sysctl.d/60-netipv4_sysctl.conf
```

Run the following script to set the active kernel parameters:

```
#!/usr/bin/env bash  
  
{  
  
sysctl -w net.ipv4.icmp_ignore_bogus_error_responses=1  
  
sysctl -w net.ipv4.route.flush=1
```

```
}
```

Note: If these settings appear in a canonically later file, or later in the same file, these settings will be overwritten

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-1,CM-2, CM-6, CM-7, IA-5

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

3.3.4 Ensure broadcast icmp requests are ignored

Pass

Description:

Setting `net.ipv4.icmp_echo_ignore_broadcasts` to 1 will cause the system to ignore all ICMP echo and timestamp requests to broadcast and multicast addresses.

Rationale:

Accepting ICMP echo and timestamp requests with broadcast or multicast destinations for your network could be used to trick your host into starting (or participating) in a Smurf attack. A Smurf attack relies on an attacker sending large amounts of ICMP broadcast messages with a spoofed source address. All hosts receiving this message and responding would send echo-reply messages back to the spoofed address, which is probably not routable. If many hosts respond to the packets, the amount of traffic on the network could be significantly multiplied.

Remediation:

Set the following parameter in `/etc/sysctl.conf` or a file in `/etc/sysctl.d/` ending in `.conf` :

- `net.ipv4.icmp_echo_ignore_broadcasts = 1`

Example:

```
# printf '%s\n' "net.ipv4.icmp_echo_ignore_broadcasts = 1" >> /etc/sysctl.d/60-netip4_sysctl.conf
```

Run the following script to set the active kernel parameters:

```
#!/usr/bin/env bash

{
sysctl -w net.ipv4.icmp_echo_ignore_broadcasts=1
sysctl -w net.ipv4.route.flush=1
}
```

Note: If these settings appear in a canonically later file, or later in the same file, these settings will be overwritten

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5
- **URL:** NIST SP 800-53A :: CM-6.1 (iv)
- **URL:** RHEL 8 STIG Vul ID: V-230537
- **URL:** RHEL 8 STIG Rule ID: SV-230537r858797

CIS Controls V7.0:

- **Control 9: Limitation and Control of Network Ports, Protocols, and Services:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 4: Secure Configuration of Enterprise Assets and Software:** -- [More](#)
>

[Back to Summary](#)

3.3.5 Ensure icmp redirects are not accepted

Fail

Description:

ICMP redirect messages are packets that convey routing information and tell your host (acting as a router) to send packets via an alternate path. It is a way of allowing an outside routing device to update your system routing tables.

Rationale:

ICMP redirect messages are packets that convey routing information and tell your host (acting as a router) to send packets via an alternate path. It is a way of allowing an outside routing device to update your system routing tables. By setting `net.ipv4.conf.all.accept_redirects`, `net.ipv4.conf.default.accept_redirects`, `net.ipv6.conf.all.accept_redirects`, and `net.ipv6.conf.default.accept_redirects` to 0, the system will not accept any ICMP redirect messages, and therefore, won't allow outsiders to update the system's routing tables.

Remediation:

Set the following parameters in `/etc/sysctl.conf` or a file in `/etc/sysctl.d/` ending in `.conf`:

- `net.ipv4.conf.all.accept_redirects = 0`
- `net.ipv4.conf.default.accept_redirects = 0`

Example:

```
# printf '%s\n' "net.ipv4.conf.all.accept_redirects = 0" "net.ipv4.conf.default.accept_redirects = 0" >>
/etc/sysctl.d/60-netipv4_sysctl.conf
```

Run the following script to set the active kernel parameters:

```
#!/usr/bin/env bash

{
sysctl -w net.ipv4.conf.all.accept_redirects=0
sysctl -w net.ipv4.conf.default.accept_redirects=0
sysctl -w net.ipv4.route.flush=1
}
```

- **IF** - IPv6 is enabled on the system:

Set the following parameters in `/etc/sysctl.conf` or a file in `/etc/sysctl.d/` ending in `.conf`:

- `net.ipv6.conf.all.accept_redirects = 0`
- `net.ipv6.conf.default.accept_redirects = 0`

Example:

```
# printf '%s\n' "net.ipv6.conf.all.accept_redirects = 0" "net.ipv6.conf.default.accept_redirects = 0" >>
/etc/sysctl.d/60-netipv6_sysctl.conf
```

Run the following script to set the active kernel parameters:

```
#!/usr/bin/env bash
```

```
{
sysctl -w net.ipv6.conf.all.accept_redirects=0

sysctl -w net.ipv6.conf.default.accept_redirects=0

sysctl -w net.ipv6.route.flush=1
}
```

Note: If these settings appear in a canonically later file, or later in the same file, these settings will be overwritten

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 :: CM-6 b
- URL: NIST SP 800-53A :: CM-6.1 (iv)
- URL: RHEL 8 STIG GROUP ID: V-230535
- URL: RHEL 8 STIG RULE ID: SV-230535r858793
- URL: RHEL 8 STIG GROUP ID: V-230544
- URL: RHEL 8 STIG RULE ID: SV-230544r858820
- URL: RHEL 8 STIG GROUP ID: V-230550
- URL: RHEL 8 STIG RULE ID: SV-230550r627750
- URL: RHEL 8 STIG GROUP ID: V-230553
- URL: RHEL 8 STIG RULE ID: SV-230553r809324

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

3.3.6 Ensure secure icmp redirects are not accepted

Fail

Description:

Secure ICMP redirects are the same as ICMP redirects, except they come from gateways listed on the default gateway list. It is assumed that these gateways are known to your system, and that they are likely to be secure.

Rationale:

It is still possible for even known gateways to be compromised. Setting `net.ipv4.conf.all.secure_redirects` and `net.ipv4.conf.default.secure_redirects` to 0 protects the system from routing table updates by possibly compromised known gateways.

Remediation:

Set the following parameters in `/etc/sysctl.conf` or a file in `/etc/sysctl.d/` ending in `.conf` :

- `net.ipv4.conf.all.secure_redirects = 0`
- `net.ipv4.conf.default.secure_redirects = 0`

Example:

```
# printf '%s\n' "net.ipv4.conf.all.secure_redirects = 0" "net.ipv4.conf.default.secure_redirects = 0" >>
/etc/sysctl.d/60-netipv4_sysctl.conf
```

Run the following script to set the active kernel parameters:

```
#!/usr/bin/env bash

{
```

```
sysctl -w net.ipv4.conf.all.secure_redirects=0

sysctl -w net.ipv4.conf.default.secure_redirects=0

sysctl -w net.ipv4.route.flush=1

}
```

Note: If these settings appear in a canonically later file, or later in the same file, these settings will be overwritten

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5

CIS Controls V7.0:

- **Control 9: Limitation and Control of Network Ports, Protocols, and Services:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 4: Secure Configuration of Enterprise Assets and Software:** -- [More](#)
>

[Back to Summary](#)

3.3.7 Ensure reverse path filtering is enabled

Fail

Description:

Setting `net.ipv4.conf.all.rp_filter` and `net.ipv4.conf.default.rp_filter` to 1 forces the Linux kernel to utilize reverse path filtering on a received packet to determine if the packet was valid. Essentially, with reverse path filtering, if the return packet does not go out the same interface that the corresponding source packet came from, the packet is dropped (and logged if `log_martians` is set).

Rationale:

Setting `net.ipv4.conf.all.rp_filter` and `net.ipv4.conf.default.rp_filter` to 1 is a good way to deter attackers from sending your system bogus packets that cannot be responded to. One instance where this feature breaks down is if asymmetrical routing is employed. This would occur when using dynamic routing protocols (bgp, ospf, etc) on your system. If you are using asymmetrical routing on your system, you will not be able to enable this feature without breaking the routing.

Remediation:

Set the following parameters in `/etc/sysctl.conf` or a file in `/etc/sysctl.d/` ending in `.conf`:

- `net.ipv4.conf.all.rp_filter = 1`
- `net.ipv4.conf.default.rp_filter = 1`

Example:

```
# printf '%s\n' "net.ipv4.conf.all.rp_filter = 1" "net.ipv4.conf.default.rp_filter = 1" >> /etc/
sysctl.d/60-netip4_sysctl.conf
```

Run the following script to set the active kernel parameters:

```
#!/usr/bin/env bash

{
sysctl -w net.ipv4.conf.all.rp_filter=1

sysctl -w net.ipv4.conf.default.rp_filter=1

sysctl -w net.ipv4.route.flush=1

}
```


Note: If these settings appear in a canonically later file, or later in the same file, these settings will be overwritten

Impact:

If you are using asymmetrical routing on your system, you will not be able to enable this feature without breaking the routing.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: CM-1,CM-2, CM-6, CM-7, IA-5
- **URL:** NIST SP 800-53A :: CM-6.1 (iv)
- **URL:** RHEL 8 STIG Vul ID: V-230549
- **URL:** RHEL 8 STIG Rule ID: SV-230549r858830

CIS Controls V7.0:

- **Control 9: Limitation and Control of Network Ports, Protocols, and Services:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 4: Secure Configuration of Enterprise Assets and Software:** -- [More](#)
>

[Back to Summary](#)

3.3.8 Ensure source routed packets are not accepted

Fail

Description:

In networking, source routing allows a sender to partially or fully specify the route packets take through a network. In contrast, non-source routed packets travel a path determined by routers in the network. In some cases, systems may not be routable or reachable from some locations (e.g. private addresses vs. Internet routable), and so source routed packets would need to be used.

Rationale:

Setting `net.ipv4.conf.all.accept_source_route`, `net.ipv4.conf.default.accept_source_route`, `net.ipv6.conf.all.accept_source_route` and `net.ipv6.conf.default.accept_source_route` to 0 disables the system from accepting source routed packets. Assume this system was capable of routing packets to Internet routable addresses on one interface and private addresses on another interface. Assume that the private addresses were not routable to the Internet routable addresses and vice versa. Under normal routing circumstances, an attacker from the Internet routable addresses could not use the system as a way to reach the private address systems. If, however, source routed packets were allowed, they could be used to gain access to the private address systems as the route could be specified, rather than rely on routing protocols that did not allow this routing.

Remediation:

Set the following parameters in `/etc/sysctl.conf` or a file in `/etc/sysctl.d/` ending in `.conf`:

- `net.ipv4.conf.all.accept_source_route = 0`
- `net.ipv4.conf.default.accept_source_route = 0`

Example:

```
# printf '%s\n' "net.ipv4.conf.all.accept_source_route = 0" "net.ipv4.conf.default.accept_source_route = 0" >> /etc/sysctl.d/60-netipv4_sysctl.conf
```

Run the following script to set the active kernel parameters:

```
#!/usr/bin/env bash

{

sysctl -w net.ipv4.conf.all.accept_source_route=0

sysctl -w net.ipv4.conf.default.accept_source_route=0
```

```
sysctl -w net.ipv4.route.flush=1
}
```

- IF - IPv6 is enabled on the system:

Set the following parameters in /etc/sysctl.conf or a file in /etc/sysctl.d/ ending in .conf :

- net.ipv6.conf.all.accept_source_route = 0
- net.ipv6.conf.default.accept_source_route = 0

Example:

```
# printf '%s\n' "net.ipv6.conf.all.accept_source_route = 0" "net.ipv6.conf.default.accept_source_route = 0" >> /etc/sysctl.d/60-netipv6_sysctl.conf
```

Run the following command to set the active kernel parameters:

```
#!/usr/bin/env bash

{

sysctl -w net.ipv6.conf.all.accept_source_route=0

sysctl -w net.ipv6.conf.default.accept_source_route=0

sysctl -w net.ipv6.route.flush=1

}
```

Note: If these settings appear in a canonically later file, or later in the same file, these settings will be overwritten

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 :: CM-6 b
- URL: NIST SP 800-53A :: CM-6.1 (iv)
- URL: REHL 8 STIG GROUP ID: V-230538
- URL: REHL 8 STIG RULE ID: SV-230538r858801
- URL: REHL 8 STIG GROUP ID: V-230539
- URL: REHL 8 STIG RULE ID: SV-230539r861085
- URL: REHL 8 STIG GROUP ID: V-230541
- URL: REHL 8 STIG RULE ID: SV-230541r858812
- URL: REHL 8 STIG GROUP ID: V-230542
- URL: REHL 8 STIG RULE ID: SV-230542r858814

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

3.3.9 Ensure suspicious packets are logged

Fail

Description:

When enabled, this feature logs packets with un-routable source addresses to the kernel log.

Rationale:

Setting net.ipv4.conf.all.log_martians and net.ipv4.conf.default.log_martians to 1 enables this feature. Logging these packets allows an administrator to investigate the possibility that an attacker is sending spoofed packets to their system.

Remediation:

Set the following parameters in `/etc/sysctl.conf` or a file in `/etc/sysctl.d/` ending in `.conf` :

- `net.ipv4.conf.all.log_martians = 1`
- `net.ipv4.conf.default.log_martians = 1`

Example:

```
# printf '%s\n' "net.ipv4.conf.all.log_martians = 1" "net.ipv4.conf.default.log_martians = 1" >> /etc/
sysctl.d/60-netipv4_sysctl.conf
```

Run the following script to set the active kernel parameters:

```
#!/usr/bin/env bash

{
sysctl -w net.ipv4.conf.all.log_martians=1
sysctl -w net.ipv4.conf.default.log_martians=1
sysctl -w net.ipv4.route.flush=1
}
```

Note: If these settings appear in a canonically later file, or later in the same file, these settings will be overwritten

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: AU-3

CIS Controls V7.0:

- **Control 6: Maintenance, Monitoring and Analysis of Audit Logs:** -- [More](#)
- **Control 6: Maintenance, Monitoring and Analysis of Audit Logs:** -- [More](#)

>

CIS Critical Security Controls V8.0:

- **Control 8: Audit Log Management:** -- [More](#)

>

[Back to Summary](#)

3.3.10 Ensure tcp syn cookies is enabled

Fail

Description:

When `tcp_syncookies` is set, the kernel will handle TCP SYN packets normally until the half-open connection queue is full, at which time, the SYN cookie functionality kicks in. SYN cookies work by not using the SYN queue at all. Instead, the kernel simply replies to the SYN with a SYN/ACK, but will include a specially crafted TCP sequence number that encodes the source and destination IP address and port number and the time the packet was sent. A legitimate connection would send the ACK packet of the three way handshake with the specially crafted sequence number. This allows the system to verify that it has received a valid response to a SYN cookie and allow the connection, even though there is no corresponding SYN in the queue.

Rationale:

Attackers use SYN flood attacks to perform a denial of service attacked on a system by sending many SYN packets without completing the three way handshake. This will quickly use up slots in the kernel's half-open connection queue and prevent legitimate connections from succeeding. Setting `net.ipv4.tcp_syncookies` to 1 enables SYN cookies, allowing the system to keep accepting valid connections, even if under a denial of service attack.

Remediation:

Set the following parameter in `/etc/sysctl.conf` or a file in `/etc/sysctl.d/` ending in `.conf` :

- `net.ipv4.tcp_syncookies = 1`

Example:

```
# printf '%s\n' "net.ipv4.tcp_syncookies = 1" >> /etc/sysctl.d/60-netipv4_sysctl.conf
```

Run the following script to set the active kernel parameters:

```
#!/usr/bin/env bash

{
sysctl -w net.ipv4.tcp_syncookies=1
sysctl -w net.ipv4.route.flush=1
}
```

Note: If these settings appear in a canonically later file, or later in the same file, these settings will be overwritten

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- [URL](#): NIST SP 800-53 Rev. 5: CM-1,CM-2, CM-6, CM-7, IA-5
- [URL](#): STIG ID: UBTU-20-010412 | Rule ID: SV-238333r958528 | CAT II
- [URL](#): STIG ID: UBTU-22-253010 | Rule ID: SV-260522r958528 | CAT II

CIS Controls V7.0:

- [Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- More](#)
>

CIS Critical Security Controls V8.0:

- [Control 4: Secure Configuration of Enterprise Assets and Software: -- More](#)
>

[Back to Summary](#)

3.3.11 Ensure ipv6 router advertisements are not accepted

Fail

Description:

Routers periodically multicast Router Advertisement messages to announce their availability and convey information to neighboring nodes that enable them to be automatically configured on the network.

`net.ipv6.conf.all.accept_ra` and `net.ipv6.conf.default.accept_ra` determine the systems ability to accept these advertisements

Rationale:

It is recommended that systems do not accept router advertisements as they could be tricked into routing traffic to compromised machines. Setting hard routes within the system (usually a single default route to a trusted router) protects the system from bad routes. Setting `net.ipv6.conf.all.accept_ra` and `net.ipv6.conf.default.accept_ra` to 0 disables the system's ability to accept IPv6 router advertisements.

Remediation:

- **IF** - IPv6 is enabled on the system:

Set the following parameters in `/etc/sysctl.conf` or a file in `/etc/sysctl.d/` ending in `.conf` :

- `net.ipv6.conf.all.accept_ra = 0`
- `net.ipv6.conf.default.accept_ra = 0`

Example:

```
# printf '%s\n' "net.ipv6.conf.all.accept_ra = 0" "net.ipv6.conf.default.accept_ra = 0" >> /etc/sysctl.d/60-netipv6_sysctl.conf
```

Run the following script to set the active kernel parameters:

```
#!/usr/bin/env bash

{

sysctl -w net.ipv6.conf.all.accept_ra=0

sysctl -w net.ipv6.conf.default.accept_ra=0

sysctl -w net.ipv6.route.flush=1

}
```

Note: If these settings appear in a canonically later file, or later in the same file, these settings will be overwritten

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 :: CM-6 b
- **URL:** NIST SP 800-53A :: CM-6.1 (iv)
- **URL:** RHEL 8 STIG Vul ID: V-230541
- **URL:** RHEL 8 STIG Rule ID: SV-230541r858812
- **URL:** RHEL 8 STIG Vul ID: V-230542
- **URL:** RHEL 8 STIG Rule ID: SV-230542r858814

CIS Controls V7.0:

- **Control 9: Limitation and Control of Network Ports, Protocols, and Services:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 4: Secure Configuration of Enterprise Assets and Software:** -- [More](#)
>

[Back to Summary](#)

4 Host Based Firewall

A firewall is a set of rules. When a data packet moves into or out of a protected network space, its contents (in particular, information about its origin, target, and the protocol it plans to use) are tested against the firewall rules to see if it should be allowed through

To provide a Host Based Firewall, the Linux kernel includes support for:

- **Netfilter** - A set of hooks inside the Linux kernel that allows kernel modules to register callback functions with the network stack. A registered callback function is then called back for every packet that traverses the respective hook within the network stack. Includes the `ip_tables`, `ip6_tables`, `arp_tables`, and `ebtables` kernel modules. These modules are some of the significant parts of the Netfilter hook system.
- **nftables** - A subsystem of the Linux kernel providing filtering and classification of network packets/datagrams/frames. nftables is supposed to replace certain parts of Netfilter, while keeping and reusing most of it. nftables utilizes the building blocks of the Netfilter infrastructure, such as the existing hooks into the networking stack, connection tracking system, userspace queueing component, and logging subsystem. **Is available in Linux kernels 3.13 and newer** .

In order to configure firewall rules for Netfilter or nftables, a firewall utility needs to be installed. Guidance has been included for the following firewall utilities:

- **UncomplicatedFirewall (ufw)** - Provides firewall features by acting as a front-end for the Linux kernel's netfilter framework via the iptables backend. ufw supports both IPv4 and IPv6 networks
- **nftables** - Includes the nft utility for configuration of the nftables subsystem of the Linux kernel
- **iptables** - Includes the iptables, ip6tables, arptables and ebtables utilities for configuration Netfilter and the ip_tables, ip6_tables, arp_tables, and ebtables kernel modules.

Notes:

- Only **one** method should be used to configure a firewall on the system. Use of more than one method could produce unexpected results
- This section is intended only to ensure the resulting firewall rules are in place, not how they are configured

4.1 Configure a single firewall utility

Only one method should be used to configure a firewall on the system. Use of more than one method could produce unexpected results.

This section ensures that only one firewall is in use on the system and provides guidance to the subsequent subsection that should be followed for a single firewall utility configuration.

4.1.1 Ensure a single firewall configuration utility is in use

Pass

Description:

In Linux security, employing a single, effective firewall configuration utility ensures that only legitimate traffic gets processed, reducing the system's exposure to potential threats. The choice between `ufw`, `nftables`, and `iptables` depends on organizational needs.

Note: `iptables` is being phased out, and support for `iptables` will be reduced over time. It is recommended to transition towards either `nftables` or `ufw` as the default firewall management tool.

Rationale:

Proper configuration of a single firewall utility minimizes cyber threats and protects services and data, while avoiding vulnerabilities like open ports or exposed services. Standardizing on a single tool simplifies management, reduces errors, and fortifies security across Linux systems.

Remediation:

Remediating to a single firewall configuration is a complex process and involves several steps. The following provides the basic steps to follow for a single firewall configuration:

- Determine which firewall utility best fits organizational needs
- Follow the recommendations in the subsequent subsection for the single firewall to be used

Note: Review the firewall subsection overview for the selected firewall to be used, it contains a script to simplify this process.

- Return to this recommendation to ensure a single firewall configuration utility is in use

Impact:

The use of more than one firewall utility may produce unexpected results.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: <https://wiki.debian.org/DebianFirewall>
- URL: <https://wiki.ubuntu.com/UncomplicatedFirewall>
- URL: <https://assets.ubuntu.com/v1/544d9904-ubuntu-server-guide-2024-01-22.pdf>
- URL: <https://www.debian.org/doc/manuals/debian-reference/debian-reference.en.pdf>

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

4.2 Configure UncomplicatedFirewall

If `nftables` or `iptables` are being used in your environment, please follow the guidance in their respective section and pass-over the guidance in this section.

Uncomplicated Firewall (UFW) is a program for managing a netfilter firewall designed to be easy to use.

- Uses a command-line interface consisting of a small number of simple commands
- Uses `iptables` for configuration
- Rules are processed until first matching rule. The first matching rule will be applied.

Notes:

- Configuration of a live system's firewall directly over a remote connection will often result in being locked out
- Rules should be ordered so that ALLOW rules come before DENY rules.

4.2.1 Ensure ufw is installed

Pass

Description:

The Uncomplicated Firewall (ufw) is a frontend for iptables and is particularly well-suited for host-based firewalls. ufw provides a framework for managing netfilter, as well as a command-line interface for manipulating the firewall

Rationale:

A firewall utility is required to configure the Linux kernel's netfilter framework via the iptables or nftables back-end.

The Linux kernel's netfilter framework host-based firewall can protect against threats originating from within a corporate network to include malicious mobile code and poorly configured software on a host.

Note: Only one firewall utility should be installed and configured. UFW is dependent on the iptables package

Remediation:

Run the following command to install Uncomplicated Firewall (UFW):

```
# apt install ufw
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: SC-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

4.2.2 Ensure nftables is not in use with ufw

Pass

Description:

nftables is a subsystem of the Linux kernel providing filtering and classification of network packets/datagrams/frames and is the successor to iptables.

Rationale:

Running both ufw and nftables may lead to conflict.

Remediation:

Run the following command to remove nftables :

```
# apt purge nftables
```

- OR -

Run the following commands to stop and mask nftables.service :

```
# systemctl stop nftables.service  
# systemctl mask nftables.service
```

Assessment:[Show](#) Assessment Evidence[Show](#) Rule Result XML**References:**

- URL: NIST SP 800-53 Rev. 5: CA-9, CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)**4.2.3 Ensure iptables-persistent is not installed with ufw**

Pass

Description:

The `iptables-persistent` is a boot-time loader for netfilter rules, iptables plugin

Rationale:

Running both `ufw` and the services included in the `iptables-persistent` package may lead to conflict

Remediation:

Run the following command to remove the `iptables-persistent` package:

```
# apt purge iptables-persistent
```

Assessment:[Show](#) Assessment Evidence[Show](#) Rule Result XML**References:**

- URL: NIST SP 800-53 Rev. 5: SC-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)**4.2.4 Ensure ufw service is enabled**

Fail

Description:

UncomplicatedFirewall (ufw) is a frontend for iptables. ufw provides a framework for managing netfilter, as well as a command-line and available graphical user interface for manipulating the firewall.

Note:

- When running `ufw enable` or starting `ufw` via its initscript, `ufw` will flush its chains. This is required so `ufw` can maintain a consistent state, but it may drop existing connections (eg `ssh`). `ufw` does support adding rules before enabling the firewall.
- Run the following command before running `ufw enable`.

```
# ufw allow proto tcp from any to any port 22
```


- The rules will still be flushed, but the ssh port will be open after enabling the firewall. Please note that once ufw is 'enabled', ufw will not flush the chains when adding or removing rules (but will when modifying a rule or changing the default policy)
- By default, ufw will prompt when enabling the firewall while running under ssh. This can be disabled by using `ufw --force enable`

Rationale:

The ufw service must be enabled and running in order for ufw to protect the system

Remediation:

Run the following command to unmask the ufw daemon:

```
# systemctl unmask ufw.service
```

Run the following command to enable and start the ufw daemon:

```
# systemctl --now enable ufw.service
```

```
active
```

Run the following command to enable ufw:

```
# ufw enable
```

Impact:

Changing firewall settings while connected over network can result in being locked out of the system.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** <http://manpages.ubuntu.com/manpages/precise/en/man8/ufw.8.html>
- **URL:** NIST SP 800-53 Rev. 5: SC-7
- **URL:** STIG ID: UBTU-20-010434 | Rule ID: SV-238355r958672 | CAT II
- **URL:** STIG ID: UBTU-22-251015 | Rule ID: SV-260515r958672 | CAT II

CIS Controls V7.0:

- **Control 9: Limitation and Control of Network Ports, Protocols, and Services:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 4: Secure Configuration of Enterprise Assets and Software:** -- [More](#)
- **Control 4: Secure Configuration of Enterprise Assets and Software:** -- [More](#)
>

[Back to Summary](#)

4.2.5 Ensure ufw loopback traffic is configured

Fail

Description:

Configure the loopback interface to accept traffic. Configure all other interfaces to deny traffic to the loopback network (127.0.0.0/8 for IPv4 and ::1/128 for IPv6).

Rationale:

Loopback traffic is generated between processes on machine and is typically critical to operation of the system. The loopback interface is the only place that loopback network (127.0.0.0/8 for IPv4 and ::1/128 for IPv6) traffic should be seen, all other interfaces should ignore traffic on this network as an anti-spoofing measure.

Remediation:

Run the following commands to configure the loopback interface to accept traffic:

```
# ufw allow in on lo
# ufw allow out on lo
```

Run the following commands to configure all other interfaces to deny traffic to the loopback network:

```
# ufw deny in from 127.0.0.0/8
# ufw deny in from ::1
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: SC-7
- **URL:** <https://manpages.ubuntu.com/manpages/jammy/en/man8/ufw-framework.8.html>

CIS Controls V7.0:

- **Control 9: Limitation and Control of Network Ports, Protocols, and Services:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 4: Secure Configuration of Enterprise Assets and Software:** -- [More](#)
- **Control 4: Secure Configuration of Enterprise Assets and Software:** -- [More](#)
>

[Back to Summary](#)

4.2.6 Ensure ufw outbound connections are configured

Manual

Description:

Configure the firewall rules for new outbound connections.

Note:

- Changing firewall settings while connected over network can result in being locked out of the system.
- Unlike iptables, when a new outbound rule is added, ufw automatically takes care of associated established connections, so no rules for the latter kind are required.

Rationale:

If rules are not in place for new outbound connections all packets will be dropped by the default policy preventing network usage.

Remediation:

Configure ufw in accordance with site policy. The following commands will implement a policy to allow all outbound connections on all interfaces:

```
# ufw allow out on all
```

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: SC-7

CIS Controls V7.0:

- **Control 9: Limitation and Control of Network Ports, Protocols, and Services:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 4: Secure Configuration of Enterprise Assets and Software:** -- [More](#)
- **Control 4: Secure Configuration of Enterprise Assets and Software:** -- [More](#)
>

[Back to Summary](#)

4.2.7 Ensure ufw firewall rules exist for all open ports

Fail

Description:

Services and ports can be accepted or explicitly rejected.

Note:

- Changing firewall settings while connected over network can result in being locked out of the system
- The remediation command opens up the port to traffic from all sources. Consult ufw documentation and set any restrictions in compliance with site policy

Rationale:

To reduce the attack surface of a system, all services and ports should be blocked unless required.

- Any ports that have been opened on non-loopback addresses need firewall rules to govern traffic.
- Without a firewall rule configured for open ports, the default firewall policy will drop all packets to these ports.
- Required ports should have a firewall rule created to allow approved connections in accordance with local site policy.
- Unapproved ports should have an explicit deny rule created.

Remediation:

For each port identified in the audit which does not have a firewall rule, evaluate the service listening on the port and add a rule for accepting or denying inbound connections in accordance with local site policy:

Examples:

```
# ufw allow in <port>/<tcp or udp protocol>

# ufw deny in <port>/<tcp or udp protocol>
```

Note: Examples create rules for from any, to any. More specific rules should be centered when allowing inbound traffic e.g only traffic from this network.

Example to allow traffic on port 443 using the tcp protocol from the 192.168.1.0 network:

```
ufw allow from 192.168.1.0/24 to any proto tcp port 443
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: SC-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

4.2.8 Ensure ufw default deny firewall policy

Fail

Description:

A default deny policy on connections ensures that any unconfigured network usage will be rejected.

Note: Any port or protocol without a explicit allow before the default deny will be blocked

Rationale:

With a default accept policy the firewall will accept any packet that is not configured to be denied. It is easier to allow list acceptable usage than to deny list unacceptable usage.

Remediation:

Run the following commands to implement a default *deny* policy:

```
# ufw default deny incoming
# ufw default deny outgoing
# ufw default deny routed
```

Impact:

Any port and protocol not explicitly allowed will be blocked. The following rules should be considered before applying the default deny.

```
ufw allow out http
ufw allow out https
ufw allow out ntp # Network Time Protocol
ufw allow out to any port 53 # DNS
ufw allow out to any port 853 # DNS over TLS
ufw logging on
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: SC-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

4.3 Configure nftables

If Uncomplicated Firewall (UFW) or iptables are being used in your environment, please follow the guidance in their respective section and pass-over the guidance in this section.

`nftables` is a subsystem of the Linux kernel providing filtering and classification of network packets/datagrams/frames and is the successor to `iptables`. The biggest change with the successor `nftables` is its simplicity. With `iptables`, we have to configure every single rule and use the syntax which can be compared with normal commands. With `nftables`, the simpler syntax, much like BPF (Berkely Packet Filter) means shorter lines and less repetition. Support for `nftables` should also be compiled into the kernel, together with the related `nftables` modules. Please ensure that your kernel supports `nf_tables` before choosing this option.

Notes:

- This section broadly assumes starting with an empty `nftables` firewall ruleset (established by flushing the rules with `nft flush ruleset`).
- Remediation steps included only affect the live system, you will also need to configure your default firewall configuration to apply on boot.
- Configuration of a live systems firewall directly over a remote connection will often result in being locked out. It is advised to have a known good firewall configuration set to run on boot and to configure an entire firewall structure in a script that is then run and tested before saving to boot.

The following script will implement the firewall rules of this section and open port 22(ssh) from anywhere. Opening the ports for port 22(ssh) needs to be updated in accordance with local site policy. **Allow port 22(ssh) needs to be updated to only allow systems requiring ssh connectivity to connect, as per site policy .**

Save the script below as `/etc/nftables.rules`

```
#!/sbin/nft -f
```

```
# This nftables.rules config should be saved as /etc/nftables.rules

# flush nftables ruleset

flush ruleset

# Load nftables ruleset

# nftables config with inet table named filter

table inet filter {

# Base chain for input hook named input (Filters inbound network packets)

chain input {

type filter hook input priority 0; policy drop;


# Ensure loopback traffic is configured

iif "lo" accept

ip saddr 127.0.0.0/8 counter packets 0 bytes 0 drop

# If IPv6 is enabled on the system ensure IPv6 loopback traffic is configured

ip6 saddr ::1 counter packets 0 bytes 0 drop


# Ensure established connections are configured

ip protocol tcp ct state established accept

ip protocol udp ct state established accept


# Accept port 22 (SSH) traffic from anywhere

tcp dport ssh accept

}


# Base chain for hook forward named forward (Filters forwarded network packets)

chain forward {

type filter hook forward priority 0; policy drop;

}


# Base chain for hook output named output (Filters outbound network packets)

chain output {

type filter hook output priority 0; policy drop;

# Ensure outbound and established connections are configured

ip protocol tcp ct state established,related,new accept

ip protocol udp ct state established,related,new accept

}

}
```

Run the following command to load the file into nftables

```
# nft -f /etc/nftables.rules
```

All changes in the nftables subsections are temporary.

To make these changes permanent:

Run the following command to create the nftables.rules file

```
nft list ruleset > /etc/nftables.rules
```

Add the following line to /etc/nftables.conf

```
include "/etc/nftables.rules"
```

4.3.1 Ensure nftables is installed

Fail

Description:

nftables provides a new in-kernel packet classification framework that is based on a network-specific Virtual Machine (VM) and a new nft userspace command line tool. nftables reuses the existing Netfilter subsystems such as the existing hook infrastructure, the connection tracking system, NAT, userspace queuing and logging subsystem.

Notes:

- nftables is available in Linux kernel 3.13 and newer
- Only one firewall utility should be installed and configured
- Changing firewall settings while connected over the network can result in being locked out of the system

Rationale:

nftables is a subsystem of the Linux kernel that can protect against threats originating from within a corporate network to include malicious mobile code and poorly configured software on a host.

Remediation:

Run the following command to install nftables :

```
# apt install nftables
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CA-9

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

4.3.2 Ensure ufw is uninstalled or disabled with nftables

Fail

Description:

Uncomplicated Firewall (UFW) is a program for managing a netfilter firewall designed to be easy to use.

Rationale:

Running both the nftables service and ufw may lead to conflict and unexpected results.

Remediation:

Run **one** of the following to either remove ufw or disable ufw and mask ufw.service :

Run the following command to remove ufw :

```
# apt purge ufw
```

-OR-

Run the following commands to disable `ufw` and mask `ufw.service` :

```
# ufw disable
# systemctl stop ufw.service
# systemctl mask ufw.service
```

Note: `ufw disable` needs to be run before `systemctl mask ufw.service` in order to correctly disable UFW

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: SC-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

4.3.3 Ensure iptables are flushed with nftables

Manual

Description:

nftables is a replacement for iptables, ip6tables, ebtables and arptables

Rationale:

It is possible to mix iptables and nftables. However, this increases complexity and also the chance to introduce errors. For simplicity flush out all iptables rules, and ensure it is not loaded

Remediation:

Run the following commands to flush iptables:

For iptables:

```
# iptables -F
```

For ip6tables:

```
# ip6tables -F
```

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CA-9, SC-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

4.3.4 Ensure a nftables table exists

Fail

Description:

Tables hold chains. Each table only has one address family and only applies to packets of this family. Tables can have one of five families.

Rationale:

nftables doesn't have any default tables. Without a table being built, nftables will not filter network traffic.

Remediation:

Run the following command to create a table in nftables

```
# nft create table inet <table name>
```

Example:

```
# nft create table inet filter
```

Impact:

Adding rules to a running nftables can cause loss of connectivity to the system

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CA-9, SC-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

4.3.5 Ensure nftables base chains exist

Fail

Description:

Chains are containers for rules. They exist in two kinds, base chains and regular chains. A base chain is an entry point for packets from the networking stack, a regular chain may be used as jump target and is used for better rule organization.

Rationale:

If a base chain doesn't exist with a hook for input, forward, and delete, packets that would flow through those chains will not be touched by nftables.

Remediation:

Run the following command to create the base chains:

```
# nft create chain inet <table name> <base chain name> { type filter hook <(input|forward|output)>
priority 0 \; }
```

Example:

```
# nft create chain inet filter input { type filter hook input priority 0 \; }

# nft create chain inet filter forward { type filter hook forward priority 0 \; }

# nft create chain inet filter output { type filter hook output priority 0 \; }
```

Impact:

If configuring nftables over ssh, creating a base chain with a policy of drop will cause loss of connectivity.

Ensure that a rule allowing ssh has been added to the base chain prior to setting the base chain's policy to drop

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

4.3.6 Ensure nftables loopback traffic is configured

Fail

Description:

Configure the loopback interface to accept traffic. Configure all other interfaces to deny traffic to the loopback network

Rationale:

Loopback traffic is generated between processes on machine and is typically critical to the operation of the system. The loopback interface is the only place that loopback network traffic should be seen, all other interfaces should ignore traffic on this network as an anti-spoofing measure.

Remediation:

Run the following commands to implement the loopback rules:

```
# nft add rule inet filter input iif lo accept
# nft add rule inet filter input ip saddr 127.0.0.0/8 counter drop
```

- IF - IPv6 is enabled on the system:

Run the following command to implement the IPv6 loopback rule:

```
# nft add rule inet filter input ip6 saddr ::1 counter drop
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CA-9, SC-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

4.3.7 Ensure nftables outbound and established connections are configured

Manual

Description:

Configure the firewall rules for new outbound, and established connections

Rationale:

If rules are not in place for new outbound, and established connections all packets will be dropped by the default policy preventing network usage.

Remediation:

Configure nftables in accordance with site policy. The following commands will implement a policy to allow all outbound connections and all established connections:

```
# nft add rule inet filter input ip protocol tcp ct state established accept

# nft add rule inet filter input ip protocol udp ct state established accept

# nft add rule inet filter output ip protocol tcp ct state new,related,established accept
```

```
# nft add rule inet filter output ip protocol udp ct state new,related,established accept
```

[Show Rule Result XML](#)

References:

- URL: NIST SP 800-53 Rev. 5: CA-9, SC-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

4.3.8 Ensure nftables default deny firewall policy

Fail

Description:

Base chain policy is the default verdict that will be applied to packets reaching the end of the chain.

Rationale:

There are two policies: accept (Default) and drop. If the policy is set to `accept` , the firewall will accept any packet that is not configured to be denied and the packet will continue transversing the network stack.

It is easier to allow list acceptable usage than to deny list unacceptable usage.

Note:

- Allow port 22(ssh) needs to be updated to only allow systems requiring ssh connectivity to connect, as per site policy.
- Changing firewall settings while connected over network can result in being locked out of the system.

Remediation:

Run the following command for the base chains with the input, forward, and output hooks to implement a default DROP policy:

```
# nft chain <table family> <table name> <chain name> { policy drop \; }
```

Example:

```
# nft chain inet filter input { policy drop \; }

# nft chain inet filter forward { policy drop \; }

# nft chain inet filter output { policy drop \; }
```

Impact:

If configuring nftables over ssh, creating a base chain with a policy of drop will cause loss of connectivity.

Ensure that a rule allowing ssh has been added to the base chain prior to setting the base chain's policy to drop

Assessment:

[Show Assessment Evidence](#)

[Show Rule Result XML](#)

References:

- URL: Manual Page nft
- URL: NIST SP 800-53 Rev. 5: CA-9, SC-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)**4.3.9 Ensure nftables service is enabled****Fail****Description:**

The nftables service allows for the loading of nftables rulesets during boot, or starting on the nftables service

Rationale:

The nftables service restores the nftables rules from the rules files referenced in the `/etc/nftables.conf` file during boot or the starting of the nftables service

Remediation:

Run the following command to enable the nftables service:

```
# systemctl enable nftables
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CA-9, SC-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)**4.3.10 Ensure nftables rules are permanent****Fail****Description:**

nftables is a subsystem of the Linux kernel providing filtering and classification of network packets/datagrams/frames.

The nftables service reads the `/etc/nftables.conf` file for a nftables file or files to include in the nftables ruleset.

A nftables ruleset containing the input, forward, and output base chains allow network traffic to be filtered.

Note: Saving the script and following the instruction in the Configure nftables section overview will implement the rules in the configure nftable section, open port 22(ssh) from anywhere, and applies nftables ruleset on boot.

Rationale:

Changes made to nftables ruleset only affect the live system, you will also need to configure the nftables ruleset to apply on boot

Remediation:

Edit the `/etc/nftables.conf` file and un-comment or add a line with `include <Absolute path to nftables rules file>` for each nftables file you want included in the nftables ruleset on boot

Example:

```
# vi /etc/nftables.conf
```

Add the line:

```
include "/etc/nftables.rules"
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CA-9, SC-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

4.4 Configure iptables

If Uncomplicated Firewall (UFW) or nftables are being used in your environment, please follow the guidance in their respective section and pass-over the guidance in this section.

IPtables is an application that allows a system administrator to configure the IPv4 and IPv6 tables, chains and rules provided by the Linux kernel firewall. While several methods of configuration exist this section is intended only to ensure the resulting IPtables rules are in place, not how they are configured. If IPv6 is in use in your environment, similar settings should be applied to the IP6tables as well.

Note:

- Configuration of a live system's firewall directly over a remote connection will often result in being locked out.
- iptables is being phased out, and support for iptables will be reduced over time. It is recommended to transition towards either nftables or ufw as the default firewall management tool.

4.4.1 Configure iptables software

This section provides guidance for installing, enabling, removing, and disabling software packages necessary for using IPTables as the method for configuring and maintaining a Host Based Firewall on the system.

Note: Using more than one method to configure and maintain a Host Based Firewall can cause unexpected results. If Uncomplicated Firewall (UFW) or NFTables are being used for configuration and maintenance, this section should be skipped and the guidance in their respective section followed.

4.4.1.1 Ensure iptables packages are installed

Fail

Description:

iptables is a utility program that allows a system administrator to configure the tables provided by the Linux kernel firewall, implemented as different Netfilter modules, and the chains and rules it stores. Different kernel modules and programs are used for different protocols; iptables applies to IPv4, ip6tables to IPv6, arptables to ARP, and ebtables to Ethernet frames.

Rationale:

A method of configuring and maintaining firewall rules is necessary to configure a Host Based Firewall.

Remediation:

Run the following command to install iptables and iptables-persistent :

```
# apt install iptables iptables-persistent
```

Assessment:
[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CA-9, SC-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

4.4.1.2 Ensure nftables is not in use with iptables

Pass

Description:

nftables is a subsystem of the Linux kernel providing filtering and classification of network packets/datagrams/frames and is the successor to iptables.

Rationale:

Running both iptables and nftables may lead to conflict.

Remediation:

Run the following command to remove nftables :

```
# apt purge nftables
```

- OR -

Run the following commands to stop and mask nftables.service :

```
# systemctl stop nftables.service
# systemctl mask nftables.service
```

Assessment:
[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CA-9, CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

4.4.1.3 Ensure ufw is not in use with iptables

Fail

Description:

Uncomplicated Firewall (UFW) is a program for managing a netfilter firewall designed to be easy to use.

- Uses a command-line interface consisting of a small number of simple commands
- Uses iptables for configuration

Rationale:

Running `iptables.persistent` with `ufw` enabled may lead to conflict and unexpected results.

Remediation:

Run the following command to remove `ufw` :

```
# apt purge ufw
```

- OR -

Run the following commands to disable `ufw`, and stop and mask `ufw.service` :

```
# ufw disable
# systemctl stop ufw.service
# systemctl mask ufw.service
```

Note: `ufw disable` needs to be run before `systemctl mask ufw.service` in order to correctly disable UFW

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CA-9, CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

4.4.2 Configure IPv4 iptables

`iptables` is used to set up, maintain, and inspect the tables of IP packet filter rules in the Linux kernel. Several different tables may be defined. Each table contains a number of built-in chains and may also contain user-defined chains.

Each chain is a list of rules which can match a set of packets. Each rule specifies what to do with a packet that matches. This is called a 'target', which may be a jump to a user-defined chain in the same table.

Note: This section broadly assumes starting with an empty `iptables` firewall ruleset (established by flushing the rules with `iptables -F`). Remediation steps included only affect the live system, you will also need to configure your default firewall configuration to apply on boot. Configuration of a live systems firewall directly over a remote connection will often result in being locked out. It is advised to have a known good firewall configuration set to run on boot and to configure an entire firewall structure in a script that is then run and tested before saving to boot. The following script will implement the firewall rules of this section and open port 22(ssh) from anywhere:

```
#!/bin/bash

# Flush IPtables rules

iptables -F

# Ensure default deny firewall policy
```

```
iptables -P INPUT DROP

iptables -P OUTPUT DROP

iptables -P FORWARD DROP


# Ensure loopback traffic is configured

iptables -A INPUT -i lo -j ACCEPT

iptables -A OUTPUT -o lo -j ACCEPT

iptables -A INPUT -s 127.0.0.0/8 -j DROP


# Ensure outbound and established connections are configured

iptables -A OUTPUT -p tcp -m state --state NEW,ESTABLISHED -j ACCEPT

iptables -A OUTPUT -p udp -m state --state NEW,ESTABLISHED -j ACCEPT

iptables -A INPUT -p tcp -m state --state ESTABLISHED -j ACCEPT

iptables -A INPUT -p udp -m state --state ESTABLISHED -j ACCEPT


# Open inbound ssh(tcp port 22) connections

iptables -A INPUT -p tcp --dport 22 -m state --state NEW -j ACCEPT
```

4.4.2.1 Ensure iptables default deny firewall policy

Fail

Description:

A default deny all policy on connections ensures that any unconfigured network usage will be rejected.

Notes:

- Changing firewall settings while connected over network can result in being locked out of the system
- Remediation will only affect the active system firewall, be sure to configure the default policy in your firewall management to apply on boot as well

Rationale:

With a default accept policy the firewall will accept any packet that is not configured to be denied. It is easier to allow list acceptable usage than to deny list unacceptable usage.

Remediation:

Run the following commands to implement a default DROP policy:

```
# iptables -P INPUT DROP

# iptables -P OUTPUT DROP

# iptables -P FORWARD DROP
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CA-9, SC-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#) >

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

4.4.2.2 Ensure iptables loopback traffic is configured

Fail

Description:

Configure the loopback interface to accept traffic. Configure all other interfaces to deny traffic to the loopback network (127.0.0.0/8).

Note:

- Changing firewall settings while connected over network can result in being locked out of the system
- Remediation will only affect the active system firewall, be sure to configure the default policy in your firewall management to apply on boot as well

Rationale:

Loopback traffic is generated between processes on machine and is typically critical to the operation of the system. The loopback interface is the only place that loopback network (127.0.0.0/8) traffic should be seen, all other interfaces should ignore traffic on this network as an anti-spoofing measure.

Remediation:

Run the following commands to implement the loopback rules:

```
# iptables -A INPUT -i lo -j ACCEPT
# iptables -A OUTPUT -o lo -j ACCEPT
# iptables -A INPUT -s 127.0.0.0/8 -j DROP
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CA-9, SC-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

4.4.2.3 Ensure iptables outbound and established connections are configured

Manual

Description:

Configure the firewall rules for new outbound, and established connections.

Note:

- Changing firewall settings while connected over network can result in being locked out of the system
- Remediation will only affect the active system firewall, be sure to configure the default policy in your firewall management to apply on boot as well

Rationale:

If rules are not in place for new outbound, and established connections all packets will be dropped by the default policy preventing network usage.

Remediation:

Configure iptables in accordance with site policy. The following commands will implement a policy to allow all outbound connections and all established connections:

```
# iptables -A OUTPUT -p tcp -m state --state NEW,ESTABLISHED -j ACCEPT
# iptables -A OUTPUT -p udp -m state --state NEW,ESTABLISHED -j ACCEPT
# iptables -A INPUT -p tcp -m state --state ESTABLISHED -j ACCEPT
# iptables -A INPUT -p udp -m state --state ESTABLISHED -j ACCEPT
```

[Show Rule Result XML](#)

References:

- [URL: NIST SP 800-53 Rev. 5: CA-9, SC-7](#)

CIS Controls V7.0:

- [Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- More](#)
>

CIS Critical Security Controls V8.0:

- [Control 4: Secure Configuration of Enterprise Assets and Software: -- More](#)
- [Control 4: Secure Configuration of Enterprise Assets and Software: -- More](#)
>

[Back to Summary](#)

4.4.2.4 Ensure iptables firewall rules exist for all open ports

Fail

Description:

Any ports that have been opened on non-loopback addresses need firewall rules to govern traffic.

Notes:

- Changing firewall settings while connected over network can result in being locked out of the system
- Remediation will only affect the active system firewall, be sure to configure the default policy in your firewall management to apply on boot as well
- The remediation command opens up the port to traffic from all sources. Consult iptables documentation and set any restrictions in compliance with site policy

Rationale:

Without a firewall rule configured for open ports default firewall policy will drop all packets to these ports.

Remediation:

For each port identified in the audit which does not have a firewall rule establish a proper rule for accepting inbound connections:

```
# iptables -A INPUT -p <protocol> --dport <port> -m state --state NEW -j ACCEPT
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CA-9, SC-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#) >

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#) >

[Back to Summary](#)

4.4.3 Configure IPv6 ip6tables

Ip6tables is used to set up, maintain, and inspect the tables of IPv6 packet filter rules in the Linux kernel. Several different tables may be defined. Each table contains a number of built-in chains and may also contain user-defined chains. Each chain is a list of rules which can match a set of packets. Each rule specifies what to do with a packet that matches. This is called a 'target', which may be a jump to a user-defined chain in the same table.

If IPv6 is enabled on the system, the ip6tables should be configured.

Note: This section broadly assumes starting with an empty ip6tables firewall ruleset (established by flushing the rules with ip6tables -F). Remediation steps included only affect the live system, you will also need to configure your default firewall configuration to apply on boot. Configuration of a live systems firewall directly over a remote connection will often result in being locked out. It is advised to have a known good firewall configuration set to run on boot and to configure an entire firewall structure in a script that is then run and tested before saving to boot.

The following script will implement the firewall rules of this section and open port 22(ssh) from anywhere:

```
#!/bin/bash

# Flush ip6tables rules
ip6tables -F

# Ensure default deny firewall policy
```

```
ip6tables -P INPUT DROP

ip6tables -P OUTPUT DROP

ip6tables -P FORWARD DROP


# Ensure loopback traffic is configured

ip6tables -A INPUT -i lo -j ACCEPT

ip6tables -A OUTPUT -o lo -j ACCEPT

ip6tables -A INPUT -s ::1 -j DROP


# Ensure outbound and established connections are configured

ip6tables -A OUTPUT -p tcp -m state --state NEW,ESTABLISHED -j ACCEPT

ip6tables -A OUTPUT -p udp -m state --state NEW,ESTABLISHED -j ACCEPT

ip6tables -A INPUT -p tcp -m state --state ESTABLISHED -j ACCEPT

ip6tables -A INPUT -p udp -m state --state ESTABLISHED -j ACCEPT


# Open inbound ssh(tcp port 22) connections

ip6tables -A INPUT -p tcp --dport 22 -m state --state NEW -j ACCEPT
```

4.4.3.1 Ensure ip6tables default deny firewall policy

Fail

Description:

A default deny all policy on connections ensures that any unconfigured network usage will be rejected.

Note:

- Changing firewall settings while connected over network can result in being locked out of the system
- Remediation will only affect the active system firewall, be sure to configure the default policy in your firewall management to apply on boot as well

Rationale:

With a default accept policy the firewall will accept any packet that is not configured to be denied. It is easier to allow list acceptable usage than to deny list unacceptable usage.

Remediation:

- IF - IPv6 is enabled on your system:

Run the following commands to implement a default DROP policy:

```
# ip6tables -P INPUT DROP

# ip6tables -P OUTPUT DROP

# ip6tables -P FORWARD DROP
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CA-9, SC-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
 - Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- >

[Back to Summary](#)

4.4.3.2 Ensure iptables loopback traffic is configured

Fail

Description:

Configure the loopback interface to accept traffic. Configure all other interfaces to deny traffic to the loopback network (::1).

Note:

- Changing firewall settings while connected over network can result in being locked out of the system
- Remediation will only affect the active system firewall, be sure to configure the default policy in your firewall management to apply on boot as well

Rationale:

Loopback traffic is generated between processes on machine and is typically critical to operation of the system. The loopback interface is the only place that loopback network (::1) traffic should be seen, all other interfaces should ignore traffic on this network as an anti-spoofing measure.

Remediation:

Run the following commands to implement the loopback rules:

```
# iptables -A INPUT -i lo -j ACCEPT
# iptables -A OUTPUT -o lo -j ACCEPT
# iptables -A INPUT -s ::1 -j DROP
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CA-9, SC-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
- >

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
 - Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- >

[Back to Summary](#)

4.4.3.3 Ensure iptables outbound and established connections are configured

Manual

Description:

Configure the firewall rules for new outbound, and established IPv6 connections.

Note:

- Changing firewall settings while connected over network can result in being locked out of the system
- Remediation will only affect the active system firewall, be sure to configure the default policy in your firewall management to apply on boot as well

Rationale:

If rules are not in place for new outbound, and established connections all packets will be dropped by the default policy preventing network usage.

Remediation:

Configure iptables in accordance with site policy. The following commands will implement a policy to allow all outbound connections and all established connections:

```
# iptables -A OUTPUT -p tcp -m state --state NEW,ESTABLISHED -j ACCEPT
# iptables -A OUTPUT -p udp -m state --state NEW,ESTABLISHED -j ACCEPT
# iptables -A INPUT -p tcp -m state --state ESTABLISHED -j ACCEPT
# iptables -A INPUT -p udp -m state --state ESTABLISHED -j ACCEPT
```

Show Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CA-9, SC-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- More >

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- More
- Control 4: Secure Configuration of Enterprise Assets and Software: -- More >

[Back to Summary](#)

4.4.3.4 Ensure iptables firewall rules exist for all open ports

Fail

Description:

Any ports that have been opened on non-loopback addresses need firewall rules to govern traffic.

Notes:

- Changing firewall settings while connected over network can result in being locked out of the system
- Remediation will only affect the active system firewall, be sure to configure the default policy in your firewall management to apply on boot as well
- The remediation command opens up the port to traffic from all sources. Consult iptables documentation and set any restrictions in compliance with site policy

Rationale:

Without a firewall rule configured for open ports default firewall policy will drop all packets to these ports.

Remediation:

For each port identified in the audit which does not have a firewall rule establish a proper rule for accepting inbound connections:

```
# iptables -A INPUT -p <protocol> --dport <port> -m state --state NEW -j ACCEPT
```

Assessment:

Show Assessment Evidence

Show Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CA-9, SC-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- More >

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- More
- Control 4: Secure Configuration of Enterprise Assets and Software: -- More >

5 Access Control

5.1 Configure SSH Server

Secure Shell (SSH) is a secure, encrypted replacement for common login services such as `telnet`, `ftp`, `rlogin`, `rsh`, and `rcp`. It is strongly recommended that sites abandon older clear-text login protocols and use SSH to prevent session hijacking and sniffing of sensitive data off the network.

The recommendations in this section only apply if the SSH daemon is installed on the system, **if remote access is not required the SSH daemon can be removed and this section skipped**.

`sshd_config`:

- The openSSH daemon configuration directives, `Include` and `Match`, may cause the audits in this section's recommendations to report incorrectly. It is recommended that these options only be used if they're needed and fully understood. If these options are configured in accordance with local site policy, they should be accounted for when following the recommendations in this section.
- The default `Include` location is the `/etc/ssh/sshd_config.d` directory. This default has been accounted for in this section. If a file has an additional `Include` that isn't this default location, the files should be reviewed to verify that the recommended setting is not being over-ridden.
- The audits of the running configuration in this section are run in the context of the root user, the local host name, and the local host's IP address. If a `Match` block exists that matches one of these criteria, the output of the audit will be from the match block. The respective matched criteria should be replaced with a non-matching substitution.
- `Include`:
 - Include the specified configuration file(s).
 - Multiple pathnames may be specified and each pathname may contain glob(7) wildcards that will be expanded and processed in lexical order.
 - Files without absolute paths are assumed to be in `/etc/ssh/`.
 - An `Include` directive may appear inside a `Match` block to perform conditional inclusion.
- `Match`:
 - Introduces a conditional block. If all of the criteria on the `Match` line are satisfied, the keywords on the following lines override those set in the global section of the config file, until either another `Match` line or the end of the file. If a keyword appears in multiple `Match` blocks that are satisfied, only the first instance of the keyword is applied.
 - The arguments to `Match` are one or more criteria-pattern pairs or the single token `All` which matches all criteria. The available criteria are `User`, `Group`, `Host`, `LocalAddress`, `LocalPort`, and `Address`.
 - The match patterns may consist of single entries or comma-separated lists and may use the wildcard and negation operators described in the PATTERNS section of `ssh_config(5)`.
 - The patterns in an `Address` criteria may additionally contain addresses to match in CIDR address/masklen format, such as `192.0.2.0/24` or `2001:db8::/32`. Note that the mask length provided must be consistent with the address - it is an error to specify a mask length that is too long for the address or one with bits set in this host portion of the address. For example, `192.0.2.0/33` and `192.0.2.0/8`, respectively.
 - Only a subset of keywords may be used on the lines following a `Match` keyword. Available keywords are available in the `ssh_config` man page.
- Once all configuration changes have been made to `/etc/ssh/sshd_config` or any included configuration files, the `sshd` configuration must be reloaded

Command to re-load the SSH daemon configuration:

```
# systemctl reload-or-restart sshd
```

`sshd` command:

- `-T` - Extended test mode. Check the validity of the configuration file, output the effective configuration to stdout and then exit. Optionally, `Match` rules may be applied by specifying the connection parameters using one or more `-C` options.
- `-C` - `connection_spec`. Specify the connection parameters to use for the `-T` extended test mode. If provided, any `Match` directives in the configuration file that would apply are applied before the configuration is written to standard output. The connection parameters are supplied as `keyword=value` pairs and may be supplied in any order, either with multiple `-C` options or as a comma-separated list. The keywords are `addr`, `user`, `host`, `laddr`, `lport`, and `rdomain` and correspond to source address, user, resolved source host name, local address, local port number and routing domain respectively.

5.1.1 Ensure access to `/etc/ssh/sshd_config` is configured

Pass

Description:

The file `/etc/ssh/sshd_config`, and files ending in `.conf` in the `/etc/ssh/sshd_config.d` directory, contain configuration specifications for `sshd`.

Rationale:

configuration specifications for `sshd` need to be protected from unauthorized changes by non-privileged users.

Remediation:

Run the following script to set ownership and permissions on /etc/ssh/sshd_config and files ending in .conf in the /etc/ssh/sshd_config.d directory:

```
#!/usr/bin/env bash

{
chmod u-x,og-rwx /etc/ssh/sshd_config
chown root:root /etc/ssh/sshd_config
while IFS= read -r -d $'\0' l_file; do
if [ -e "$l_file" ]; then
chmod u-x,og-rwx "$l_file"
chown root:root "$l_file"
fi
done < <(find /etc/ssh/sshd_config.d -type f -print0 2>/dev/null)
}
```

- IF - other locations are listed in an Include statement, *.conf files in these locations access should also be modified.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AC-3. MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

5.1.2 Ensure access to SSH private host key files is configured

Pass

Description:

An SSH private key is one of two files used in SSH public key authentication. In this authentication method, the possession of the private key is proof of identity. Only a private key that corresponds to a public key will be able to authenticate successfully. The private keys need to be stored and handled carefully, and no copies of the private key should be distributed.

Rationale:

If an unauthorized user obtains the private SSH host key file, the host could be impersonated

Remediation:

Run the following script to set mode, ownership, and group on the private SSH host key files:

```
#!/usr/bin/env bash

{
```



```
a_output=(); a_output2=(); l_ssh_group_name="$(awk -F: '($1 ~ /^(ssh_keys|_?ssh)$/) {print $1}' /etc/group)"

f_file_access_fix()

{
while IFS= read -r l_file_mode l_file_owner l_file_group; do
a_out2=()

[ "$l_file_group" = "$l_ssh_group_name" ] && l_pmask="0137" || l_pmask="0177"

l_maxperm="$( printf '%o' $(( 0777 & ~$l_pmask )) )"

if [ $(( $l_file_mode & $l_pmask )) -gt 0 ]; then
a_out2+=(" Mode: \"$l_file_mode\" should be mode: \"$l_maxperm\" or more restrictive" \
" updating to mode: \:$l_maxperm\"")

if [ "l_file_group" = "$l_ssh_group_name" ]; then
chmod u-x,g-wx,o-rwx "$l_file"
else
chmod u-x,go-rwx "$l_file"
fi
fi

if [ "$l_file_owner" != "root" ]; then
a_out2+=(" Owned by: \"$l_file_owner\" should be owned by \"root\"" \
" Changing ownership to \"root\"")

chown root "$l_file"
fi

if [[ ! "$l_file_group" =~ ($l_ssh_group_name|root) ]]; then
[ -n "$l_ssh_group_name" ] && l_new_group="$l_ssh_group_name" || l_new_group="root"

a_out2+=(" Owned by group \"$l_file_group\" should be group owned by: \"$l_ssh_group_name\" or \"root\"" \
" Changing group ownership to \"$l_new_group\"")

chgrp "$l_new_group" "$l_file"
fi

if [ "${#a_out2[@]}" -gt 0 ]; then
a_output2+=(" - File: \"$l_file\" "${a_out2[@]}")
else
a_output+=(" - File: \"$l_file\" \
"Correct: mode: \"$l_file_mode\", owner: \"$l_file_owner\", and group owner: \"$l_file_group\" configured")
fi

done <<(stat -Lc '%a:%U:%G' "$l_file")
}

while IFS= read -r -d $'\0' l_file; do
if ssh-keygen -lf &>/dev/null "$l_file"; then
file "$l_file" | grep -Piq -- '\bopenssh\b+([^\n\r]+\b)?private\b+key\b' && f_file_access_fix
fi

done <<(find -L /etc/ssh -xdev -type f -print0 2>/dev/null)

if [ "${#a_output2[@]}" -le 0 ]; then
```

```
printf '%s\n' " " - No access changes required" ""
else
printf '%s\n' " " - Remediation results:" "${a_output2[@]}" ""
fi
}
```

Assessment:[Show Assessment Evidence](#)[Show Rule Result XML](#)**References:**

- **URL:** NIST SP 800-53 Rev. 5: AC-3. MP-2

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)**5.1.3 Ensure access to SSH public host key files is configured**

Pass

Description:

An SSH public key is one of two files used in SSH public key authentication. In this authentication method, a public key is a key that can be used for verifying digital signatures generated using a corresponding private key. Only a public key that corresponds to a private key will be able to authenticate successfully.

Rationale:

If a public host key file is modified by an unauthorized user, the SSH service may be compromised.

Remediation:

Run the following script to set mode, ownership, and group on the public SSH host key files:

```
#!/usr/bin/env bash

{
a_output=(); a_output2=()

l_pmask="0133"; l_maxperm="$( printf '%o' $(( 0777 & ~$l_pmask )) )"
f_file_access_fix()
{
while IFS=: read -r l_file_mode l_file_owner l_file_group; do
a_out2=()

[ $(( $l_file_mode & $l_pmask )) -gt 0 ] && \
a_out2+=(" Mode: \"$l_file_mode\" should be mode: \"$l_maxperm\" or more restrictive" \
" updating to mode: \"$l_maxperm\"") && chmod u-x,go-wx "$l_file"

[ "$l_file_owner" != "root" ] && \
a_out2+=(" Owned by: \"$l_file_owner\" should be owned by \"root\" \" \" \
" Changing ownership to \"root\"") && chown root "$l_file"
```

```
[ "$l_file_group" != "root" ] && \
a_out2+=(" Owned by group \"$l_file_group\" should be group owned by: \"root\" \" \" \
" Changing group ownership to \"root\"") && chgrp root "$l_file"
if [ "${#a_out2[@]}" -gt "0" ]; then
a_output2+=(" - File: \"$l_file\" \" "${a_out2[@]}")
else
a_output+=(" - File: \"$l_file\" \" \" \
" Correct: mode: \"$l_file_mode\", owner: \"$l_file_owner\", and group owner: \"$l_file_group\" configured")
fi
done < <(stat -Lc '%a:%U:%G' "$l_file")
}
while IFS= read -r -d $'\0' l_file; do
if ssh-keygen -lf &>/dev/null "$l_file"; then
file "$l_file" | grep -Piq -- '\bopenssh\b+([^\n\r]+\b)?public\b+key\b' && f_file_access_fix
fi
done < <(find -L /etc/ssh -xdev -type f -print0 2>/dev/null)
if [ "${#a_output2[@]}" -le "0" ]; then
printf '%s\n' " " - No access changes required" "
else
printf '%s\n' " - Remediation results:" "${a_output2[@]}" "
fi
}
```

Assessment:[Show](#) Assessment Evidence[Show](#) Rule Result XML**References:**

- URL: NIST SP 800-53 Rev. 5: AC-3. MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)**5.1.4 Ensure sshd access is configured**

Pass

Description:

There are several options available to limit which users and group can access the system via SSH. It is recommended that at least one of the following options be leveraged:

- AllowUsers :
 - The AllowUsers variable gives the system administrator the option of allowing specific users to ssh into the system. The list consists of space separated user names. Numeric user IDs are not recognized with this variable. If a system administrator wants to restrict user access further by only allowing the allowed users to log in from a particular host, the entry can be specified in the form of user@host.
- AllowGroups :

- The `AllowGroups` variable gives the system administrator the option of allowing specific groups of users to `ssh` into the system. The list consists of space separated group names. Numeric group IDs are not recognized with this variable.
- `DenyUsers` :
 - The `DenyUsers` variable gives the system administrator the option of denying specific users to `ssh` into the system. The list consists of space separated user names. Numeric user IDs are not recognized with this variable. If a system administrator wants to restrict user access further by specifically denying a user's access from a particular host, the entry can be specified in the form of `user@host`.
- `DenyGroups` :
 - The `DenyGroups` variable gives the system administrator the option of denying specific groups of users to `ssh` into the system. The list consists of space separated group names. Numeric group IDs are not recognized with this variable.

Rationale:

Restricting which users can remotely access the system via SSH will help ensure that only authorized users access the system.

Remediation:

Edit the `/etc/ssh/sshd_config` file to set one or more of the parameters above any `Include` and `Match` set statements as follows:

```
AllowUsers <userlist>

- AND/OR -

AllowGroups <grouplist>
```

Note:

- First occurrence of a option takes precedence, `Match` set statements withstanding. If `Include` locations are enabled, used, and order of precedence is understood in your environment, the entry may be created in a `.conf` file in a `Include` directory.
- **Be advised** that these options are "ANDed" together. If both `AllowUsers` and `AllowGroups` are set, connections will be limited to the list of users that are also a member of an allowed group. It is recommended that only one be set for clarity and ease of administration.
- It is easier to manage an allow list than a deny list. In a deny list, you could potentially add a user or group and forget to add it to the deny list.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** SSHD_CONFIG(5)
- **URL:** NIST SP 800-53 Rev. 5: AC-3. MP-2
- **URL:** SSHD(8)
- **URL:** <https://documentation.suse.com/en-us/sles/15-SP6/html/SLES-all/cha-ssh.html>

CIS Controls V7.0:

- **Control 4: Controlled Use of Administrative Privileges:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

5.1.5 Ensure sshd Banner is configured

Pass

Description:

The `Banner` parameter specifies a file whose contents must be sent to the remote user before authentication is permitted. By default, no banner is displayed.

Rationale:

Banners are used to warn connecting users of the particular site's policy regarding connection. Presenting a warning message prior to the normal user login may assist the prosecution of trespassers on the computer system.

Remediation:

Edit the `/etc/ssh/sshd_config` file to set the `Banner` parameter above any `Include` and `Match` entries as follows:

```
Banner /etc/issue.net
```

Note: First occurrence of a option takes precedence, `Match` set statements withstanding. If `Include` locations are enabled, used, and order of precedence is understood in your environment, the entry may be created in a file in `Include` location.

Edit the file being called by the `Banner` argument with the appropriate contents according to your site policy, remove any instances of `\m`, `\r`, `\s`, `\v` or references to the `OS platform`

Example:

```
# printf '%s\n' "Authorized users only. All activity may be monitored and reported." > "$(ssh -T | awk '$1 == "banner" {print $2}')
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5

CIS Controls V7.0:

- **Not Explicitly Mapped:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Not Explicitly Mapped:** -- [More](#)
>

[Back to Summary](#)

5.1.6 Ensure sshd Ciphers are configured

Pass

Description:

This variable limits the ciphers that SSH can use during communication.

Notes:

- Some organizations may have stricter requirements for approved ciphers.
- Ensure that ciphers used are in compliance with site policy.
- The only "strong" ciphers currently FIPS 140 compliant are:
 - aes256-gcm@openssh.com
 - aes128-gcm@openssh.com
 - aes256-ctr
 - aes192-ctr
 - aes128-ctr

Rationale:

Weak ciphers that are used for authentication to the cryptographic module cannot be relied upon to provide confidentiality or integrity, and system data may be compromised.

- The Triple DES ciphers, as used in SSH, have a birthday bound of approximately four billion blocks, which makes it easier for remote attackers to obtain clear text data via a birthday attack against a long-duration encrypted session, aka a "Sweet32" attack.
- Error handling in the SSH protocol; Client and Server, when using a block cipher algorithm in Cipher Block Chaining (CBC) mode, makes it easier for remote attackers to recover certain plain text data from an arbitrary block of cipher text in an SSH session via unknown vectors.

Remediation:

Edit the `/etc/ssh/sshd_config` file and add/modify the `Ciphers` line to contain a comma separated list of the site unapproved (weak) Ciphers preceded with a `-` above any `Include` entries:

Example:

```
Ciphers -3des-cbc,aes128-cbc,aes192-cbc,aes256-cbc,chacha20-poly1305@openssh.com
```

- **IF** -CVE-2023-48795 has been addressed, and it meets local site policy, `chacha20-poly1305@openssh.com` may be removed from the list of excluded ciphers.

Note: First occurrence of an option takes precedence. If Include locations are enabled, used, and order of precedence is understood in your environment, the entry may be created in a file in Include location.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** <https://nvd.nist.gov/vuln/detail/CVE-2023-48795>
- **URL:** <https://nvd.nist.gov/vuln/detail/CVE-2019-1543>
- **URL:** <https://nvd.nist.gov/vuln/detail/CVE-2016-2183>
- **URL:** <https://nvd.nist.gov/vuln/detail/CVE-2008-5161>
- **URL:** <https://www.openssh.com/txt/cbc.adv>
- **URL:** <https://www.openssh.com/txt/cbc.adv>
- **URL:** SSHD_CONFIG(5)
- **URL:** NIST SP 800-53 Rev. 5: SC-8

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

5.1.7 Ensure `sshd ClientAliveInterval` and `ClientAliveCountMax` are configured

Pass

Description:

Note: To clarify, the two settings described below are only meant for idle connections from a protocol perspective and are not meant to check if the user is active or not. An idle user does not mean an idle connection. SSH does not and never had, intentionally, the capability to drop idle users. In SSH versions before 8.2p1 there was a bug that caused these values to behave in such a manner that they were abused to disconnect idle users. This bug has been resolved in 8.2p1 and thus it can no longer be abused disconnect idle users.

The two options `ClientAliveInterval` and `ClientAliveCountMax` control the timeout of SSH sessions. Taken directly from `man 5 sshd_config`:

- `ClientAliveInterval` Sets a timeout interval in seconds after which if no data has been received from the client, `sshd(8)` will send a message through the encrypted channel to request a response from the client. The default is 0, indicating that these messages will not be sent to the client.
- `ClientAliveCountMax` Sets the number of client alive messages which may be sent without `sshd(8)` receiving any messages back from the client. If this threshold is reached while client alive messages are being sent, `sshd` will disconnect the client, terminating the session. It is important to note that the use of client alive messages is very different from `TCPKeepAlive`. The client alive messages are sent through the encrypted channel and therefore will not be spoofable. The `TCP keepalive` option enabled by `TCPKeepAlive` is spoofable. The client alive mechanism is valuable when the client or server depend on knowing when a connection has become unresponsive. The default value is 3. If `ClientAliveInterval` is set to 15, and `ClientAliveCountMax` is left at the default, unresponsive SSH clients will be disconnected after approximately 45 seconds. Setting a zero `ClientAliveCountMax` disables connection termination.

Rationale:

In order to prevent resource exhaustion, appropriate values should be set for both `ClientAliveInterval` and `ClientAliveCountMax`. Specifically, looking at the source code, `ClientAliveCountMax` must be greater than zero in order to utilize the ability of SSH to drop idle connections. If connections are allowed to stay open indefinitely, this can potentially be used as a DDOS attack or simple resource exhaustion could occur over unreliable networks.

The example set here is a 45 second timeout. Consult your site policy for network timeouts and apply as appropriate.

Remediation:

Edit the `/etc/ssh/sshd_config` file to set the `ClientAliveInterval` and `ClientAliveCountMax` parameters

above any Include and Match entries according to site policy.

Example:

```
ClientAliveInterval 15
ClientAliveCountMax 3
```

Note: First occurrence of a option takes precedence, Match set statements withstanding. If Include locations are enabled, used, and order of precedence is understood in your environment, the entry may be created in a file in Include location.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: SSHD_CONFIG(5)
- URL: SSHD(8)
- URL: NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5

CIS Controls V7.0:

- Not Explicitly Mapped: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Not Explicitly Mapped: -- [More](#)
>

[Back to Summary](#)

5.1.8 Ensure sshd DisableForwarding is enabled

Pass

Description:

The DisableForwarding parameter disables all forwarding features, including X11, ssh-agent(1), TCP and StreamLocal. This option overrides all other forwarding-related options and may simplify restricted configurations.

- X11Forwarding provides the ability to tunnel X11 traffic through the connection to enable remote graphic connections.
- ssh-agent is a program to hold private keys used for public key authentication. Through use of environment variables the agent can be located and automatically used for authentication when logging in to other machines using ssh.
- SSH port forwarding is a mechanism in SSH for tunneling application ports from the client to the server, or servers to clients. It can be used for adding encryption to legacy applications, going through firewalls, and some system administrators and IT professionals use it for opening backdoors into the internal network from their home machines.

Rationale:

Disable X11 forwarding unless there is an operational requirement to use X11 applications directly. There is a small risk that the remote X11 servers of users who are logged in via SSH with X11 forwarding could be compromised by other users on the X11 server. Note that even if X11 forwarding is disabled, users can always install their own forwarders.

anyone with root privilege on the the intermediate server can make free use of ssh-agent to authenticate them to other servers

Leaving port forwarding enabled can expose the organization to security risks and backdoors. SSH connections are protected with strong encryption. This makes their contents invisible to most deployed network monitoring and traffic filtering solutions. This invisibility carries considerable risk potential if it is used for malicious purposes such as data exfiltration. Cybercriminals or malware could exploit SSH to hide their unauthorized communications, or to exfiltrate stolen data from the target network.

Remediation:

Edit the /etc/ssh/sshd_config file to set the DisableForwarding parameter to yes above any Include entry as follows:

```
DisableForwarding yes
```

Note: First occurrence of a option takes precedence. If Include locations are enabled, used, and order of precedence is understood in your environment, the entry may be created in a file in Include location.

Impact:

SSH tunnels are widely used in many corporate environments. In some environments the applications themselves may have very limited native support for security. By utilizing tunneling, compliance with SOX, HIPAA, PCI-DSS, and other standards can be achieved without having to modify the applications.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: sshd_config(5)
- URL: SSHD(8)
- URL: NIST SP 800-53 Rev. 5: CM-7

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

5.1.9 Ensure sshd GSSAPIAuthentication is disabled

Pass

Description:

The `GSSAPIAuthentication` parameter specifies whether user authentication based on GSSAPI is allowed

Rationale:

Allowing GSSAPI authentication through SSH exposes the system's GSSAPI to remote hosts, and should be disabled to reduce the attack surface of the system

Remediation:

Edit the `/etc/ssh/sshd_config` file to set the `GSSAPIAuthentication` parameter to `no` above any `Include` and `Match` entries as follows:

```
GSSAPIAuthentication no
```

Note: First occurrence of an option takes precedence, `Match` set statements withstanding. If `Include` locations are enabled, used, and order of precedence is understood in your environment, the entry may be created in a file in `Include` location.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: SSHD_CONFIG(5)
- URL: SSHD(8)
- URL: NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5

CIS Controls V7.0:

- Control 4: Controlled Use of Administrative Privileges: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 5: Account Management: -- [More](#)
>

[Back to Summary](#)

5.1.10 Ensure sshd HostbasedAuthentication is disabled

Pass

Description:

The HostbasedAuthentication parameter specifies if authentication is allowed through trusted hosts via the user of .rhosts , Or /etc/hosts.equiv , along with successful public key client host authentication.

Rationale:

Even though the .rhosts files are ineffective if support is disabled in /etc/pam.conf , disabling the ability to use .rhosts files in SSH provides an additional layer of protection.

Remediation:

Edit the /etc/ssh/sshd_config file to set the HostbasedAuthentication parameter to no above any Include and Match entries as follows:

```
HostbasedAuthentication no
```

Note: First occurrence of a option takes precedence, Match set statements withstanding. If Include locations are enabled, used, and order of precedence is understood in your environment, the entry may be created in a file in Include location.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: SSHD_CONFIG(5)
- URL: SSHD(8)
- URL: NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5

CIS Controls V7.0:

- Control 4: Controlled Use of Administrative Privileges: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 5: Account Management: -- [More](#)
>

[Back to Summary](#)

5.1.11 Ensure sshd IgnoreRhosts is enabled

Pass

Description:

The IgnoreRhosts parameter specifies that .rhosts and .shosts files will not be used in RhostsRSAAuthentication Or HostbasedAuthentication .

Rationale:

Setting this parameter forces users to enter a password when authenticating with SSH.

Remediation:

Edit the /etc/ssh/sshd_config file to set the IgnoreRhosts parameter to yes above any Include entry as follows:

```
IgnoreRhosts yes
```

Note: First occurrence of a option takes precedence. If Include locations are enabled, used, and order of precedence is understood in your environment, the entry may be created in a file in Include location.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: SSHD_CONFIG(5)
- URL: SSHD(8)
- URL: NIST SP 800-53 Rev. 5: CM-1,CM-2, CM-6, CM-7, IA-5

CIS Controls V7.0:

- Control 4: Controlled Use of Administrative Privileges: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 5: Account Management: -- [More](#)
>

[Back to Summary](#)

5.1.12 Ensure sshd KexAlgorithms is configured

Pass

Description:

Key exchange is any method in cryptography by which cryptographic keys are exchanged between two parties, allowing use of a cryptographic algorithm. If the sender and receiver wish to exchange encrypted messages, each must be equipped to encrypt messages to be sent and decrypt messages received

Notes:

- Kex algorithms have a higher preference the earlier they appear in the list
- Some organizations may have stricter requirements for approved Key exchange algorithms
- Ensure that Key exchange algorithms used are in compliance with site policy
- The only Key Exchange Algorithms currently FIPS 140 approved are:
 - ecdh-sha2-nistp256
 - ecdh-sha2-nistp384
 - ecdh-sha2-nistp521
 - diffie-hellman-group-exchange-sha256
 - diffie-hellman-group16-sha512
 - diffie-hellman-group18-sha512
 - diffie-hellman-group14-sha256

Rationale:

Key exchange methods that are considered weak should be removed. A key exchange method may be weak because too few bits are used, or the hashing algorithm is considered too weak. Using weak algorithms could expose connections to man-in-the-middle attacks

Remediation:

Edit the `/etc/ssh/sshd_config` file and add/modify the `KexAlgorithms` line to contain a comma separated list of the site unapproved (weak) `KexAlgorithms` preceded with a `-` above any `Include` entries:

Example:

```
KexAlgorithms -diffie-hellman-group1-sha1,diffie-hellman-group14-sha1,diffie-hellman-group-exchange-sha1
```

Note: First occurrence of an option takes precedence. If `Include` locations are enabled, used, and order of precedence is understood in your environment, the entry may be created in a file in `Include` location.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: <https://ubuntu.com/server/docs/openssh-crypto-configuration>
- URL: NIST SP 800-53 Rev. 5: SC-8
- URL: SSHD(8)
- URL: SSHD_CONFIG(5)

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

5.1.13 Ensure sshd LoginGraceTime is configured

Pass

Description:

The LoginGraceTime parameter specifies the time allowed for successful authentication to the SSH server. The longer the Grace period is the more open unauthenticated connections can exist. Like other session controls in this session the Grace Period should be limited to appropriate organizational limits to ensure the service is available for needed access.

Rationale:

Setting the LoginGraceTime parameter to a low number will minimize the risk of successful brute force attacks to the SSH server. It will also limit the number of concurrent unauthenticated connections While the recommended setting is 60 seconds (1 Minute), set the number based on site policy.

Remediation:

Edit the /etc/ssh/sshd_config file to set the LoginGraceTime parameter to 60 seconds or less above any Include entry as follows:

```
LoginGraceTime 60
```

Note: First occurrence of a option takes precedence. If Include locations are enabled, used, and order of precedence is understood in your environment, the entry may be created in a file in Include location.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: SSHD_CONFIG(5)
- URL: NIST SP 800-53 Rev. 5: CM-6
- URL: SSHD(8)

CIS Controls V7.0:

- Not Explicitly Mapped: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Not Explicitly Mapped: -- [More](#)
>

[Back to Summary](#)

5.1.14 Ensure sshd LogLevel is configured

Pass

Description:

SSH provides several logging levels with varying amounts of verbosity. The DEBUG options are specifically not recommended other than strictly for debugging SSH communications. These levels provide so much data that it is difficult to identify important security information, and may violate the privacy of users.

Rationale:

The INFO level is the basic level that only records login activity of SSH users. In many situations, such as Incident Response, it is important to determine when a particular user was active on a system. The logout record can eliminate those users who disconnected, which helps narrow the field.

The VERBOSE level specifies that login and logout activity as well as the key fingerprint for any SSH key used for login will be logged. This information is important for SSH key management, especially in legacy environments.

Remediation:

Edit the `/etc/ssh/sshd_config` file to set the `LogLevel` parameter to `VERBOSE` or `INFO` above any `Include` and `Match` entries as follows:

```
LogLevel VERBOSE

- OR -

LogLevel INFO
```

Note: First occurrence of an option takes precedence, `Match` set statements withstanding. If `Include` locations are enabled, used, and order of precedence is understood in your environment, the entry may be created in a file in `Include` location.

Assessment:
[Show](#) Assessment Evidence

[Show](#) Rule Result XML

- References:**
- **URL:** https://www.ssh.com/ssh/sshd_config/
 - **URL:** NIST SP 800-53 Rev. 5: AU-3, AU-12, SI-5

- CIS Controls V7.0:**
- **Control 6: Maintenance, Monitoring and Analysis of Audit Logs:** -- [More](#)
 - **Control 6: Maintenance, Monitoring and Analysis of Audit Logs:** -- [More](#)

- CIS Critical Security Controls V8.0:**
- **Control 8: Audit Log Management:** -- [More](#)

[Back to Summary](#)

5.1.15 Ensure sshd MACs are configured

Pass

Description:

This variable limits the types of MAC algorithms that SSH can use during communication.

- Notes:**
- Some organizations may have stricter requirements for approved MACs.
 - Ensure that MACs used are in compliance with site policy.
 - The only "strong" MACs currently FIPS 140 approved are:
 - HMAC-SHA1
 - HMAC-SHA2-256
 - HMAC-SHA2-384
 - HMAC-SHA2-512

Rationale:

MD5 and 96-bit MAC algorithms are considered weak and have been shown to increase exploitability in SSH downgrade attacks. Weak algorithms continue to have a great deal of attention as a weak spot that can be exploited with expanded computing power. An attacker that breaks the algorithm could take advantage of a MiTM position to decrypt the SSH tunnel and capture credentials and information.

Remediation:

Edit the `/etc/ssh/sshd_config` file and add/modify the `MACs` line to contain a comma separated list of the site unapproved (weak) MACs preceded with a `-` above any `Include` entries:

Example:

```
MACs -hmac-md5,hmac-md5-96,hmac-ripemd160,hmac-sha1-96,umac-64@openssh.com,hmac-md5-etm@openssh.com,hmac-md5-96-etm@openssh.com,hmac-ripemd160-etm@openssh.com,hmac-sha1-96-etm@openssh.com,umac-64-etm@openssh.com
```

- IF -CVE-2023-48795 has not been reviewed and addressed, the following `etm` MACs should be added to the exclude list: [hmac-sha1-etm@openssh.com](#) , [hmac-sha2-256-etm@openssh.com](#) , [hmac-sha2-512-etm@openssh.com](#)

Note: First occurrence of an option takes precedence. If Include locations are enabled, used, and order of precedence is understood in your environment, the entry may be created in a file in Include location.

Assessment:
[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** <https://nvd.nist.gov/vuln/detail/CVE-2023-48795>
- **URL:** More information on SSH downgrade attacks can be found here: <http://www.mitls.org/pages/attacks/SLOTH>
- **URL:** [SSHD_CONFIG\(5\)](#)
- **URL:** [NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5](#)

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
- **Control 16: Account Monitoring and Control:** -- [More](#)

>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)

>

[Back to Summary](#)

5.1.16 Ensure sshd MaxAuthTries is configured

Pass

Description:

The `MaxAuthTries` parameter specifies the maximum number of authentication attempts permitted per connection. When the login failure count reaches half the number, error messages will be written to the `syslog` file detailing the login failure.

Rationale:

Setting the `MaxAuthTries` parameter to a low number will minimize the risk of successful brute force attacks to the SSH server. While the recommended setting is 4, set the number based on site policy.

Remediation:

Edit the `/etc/ssh/sshd_config` file to set the `MaxAuthTries` parameter to 4 or less above any `Include` and `Match` entries as follows:

```
MaxAuthTries 4
```

Note: First occurrence of an option takes precedence, `Match` set statements withstanding. If `Include` locations are enabled, used, and order of precedence is understood in your environment, the entry may be created in a file in `Include` location.

Assessment:
[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** [SSHD_CONFIG\(5\)](#)
- **URL:** [NIST SP 800-53 Rev. 5: AU-3](#)

CIS Controls V7.0:

- **Control 16: Account Monitoring and Control:** -- [More](#)

>

CIS Critical Security Controls V8.0:

- **Control 8: Audit Log Management:** -- [More](#)

>

[Back to Summary](#)

5.1.17 Ensure sshd MaxSessions is configured

Pass

Description:

The `MaxSessions` parameter specifies the maximum number of open sessions permitted from a given connection.

Rationale:

To protect a system from denial of service due to a large number of concurrent sessions, use the rate limiting function of `MaxSessions` to protect availability of `sshd` logins and prevent overwhelming the daemon.

Remediation:

Edit the `/etc/ssh/sshd_config` file to set the `MaxSessions` parameter to 10 or less above any `Include` and `Match` entries as follows:

```
MaxSessions 10
```

Note: First occurrence of an option takes precedence, `Match` set statements withstanding. If `Include` locations are enabled, used, and order of precedence is understood in your environment, the entry may be created in a file in `Include` location.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** SSHD_CONFIG(5)
- **URL:** NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5

CIS Controls V7.0:

- **Not Explicitly Mapped:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Not Explicitly Mapped:** -- [More](#)
>

[Back to Summary](#)

5.1.18 Ensure sshd MaxStartups is configured

Pass

Description:

The `MaxStartups` parameter specifies the maximum number of concurrent unauthenticated connections to the `SSH` daemon.

Rationale:

To protect a system from denial of service due to a large number of pending authentication connection attempts, use the rate limiting function of `MaxStartups` to protect availability of `sshd` logins and prevent overwhelming the daemon.

Remediation:

Edit the `/etc/ssh/sshd_config` file to set the `MaxStartups` parameter to `10:30:60` or more restrictive above any `Include` entries as follows:

```
MaxStartups 10:30:60
```

Note: First occurrence of a option takes precedence. If `Include` locations are enabled, used, and order of precedence is understood in your environment, the entry may be created in a file in `Include` location.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: SSHD_CONFIG(5)
- URL: NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5

CIS Controls V7.0:

- Not Explicitly Mapped: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Not Explicitly Mapped: -- [More](#)
>

[Back to Summary](#)

5.1.19 Ensure sshd PermitEmptyPasswords is disabled

Pass

Description:

The `PermitEmptyPasswords` parameter specifies if the SSH server allows login to accounts with empty password strings.

Rationale:

Disallowing remote shell access to accounts that have an empty password reduces the probability of unauthorized access to the system.

Remediation:

Edit `/etc/ssh/sshd_config` and set the `PermitEmptyPasswords` parameter to `no` above any `Include` and `Match` entries as follows:

```
PermitEmptyPasswords no
```

Note: First occurrence of an option takes precedence, `Match` set statements withstanding. If `Include` locations are enabled, used, and order of precedence is understood in your environment, the entry may be created in a file in `Include` location.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: SSHD_CONFIG(5)
- URL: NIST SP 800-53 Rev. 5: CM-1,CM-2, CM-6, CM-7, IA-5

CIS Controls V7.0:

- Control 4: Controlled Use of Administrative Privileges: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 5: Account Management: -- [More](#)
>

[Back to Summary](#)

5.1.20 Ensure sshd PermitRootLogin is disabled

Pass

Description:

The `PermitRootLogin` parameter specifies if the root user can log in using SSH. The default is `prohibit-password`.

Rationale:

Disallowing `root` logins over SSH requires system admins to authenticate using their own individual account, then escalating to `root`. This limits opportunity for non-repudiation and provides a clear audit trail in the event of a security incident.

Remediation:

Edit the `/etc/ssh/sshd_config` file to set the `PermitRootLogin` parameter to `no` above any `Include` and `Match` entries as follows:

```
PermitRootLogin no
```

Note: First occurrence of an option takes precedence, `Match` set statements withstanding. If `Include` locations are enabled, used, and order of precedence is understood in your environment, the entry may be created in a file in `Include` location.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: SSHD_CONFIG(5)
- URL: NIST SP 800-53 Rev. 5:AC-6

CIS Controls V7.0:

- Control 4: Controlled Use of Administrative Privileges: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 5: Account Management: -- [More](#)
>

[Back to Summary](#)

5.1.21 Ensure sshd PermitUserEnvironment is disabled

Pass

Description:

The `PermitUserEnvironment` option allows users to present environment options to the SSH daemon.

Rationale:

Permitting users the ability to set environment variables through the SSH daemon could potentially allow users to bypass security controls (e.g. setting an execution path that has SSH executing trojan'd programs)

Remediation:

Edit the `/etc/ssh/sshd_config` file to set the `PermitUserEnvironment` parameter to `no` above any `Include` entries as follows:

```
PermitUserEnvironment no
```

Note: First occurrence of an option takes precedence. If `Include` locations are enabled, used, and order of precedence is understood in your environment, the entry may be created in a file in `Include` location.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: SSHD_CONFIG(5)
- URL: NIST SP 800-53 Rev. 5: CM-1,CM-2, CM-6, CM-7, IA-5
- URL: SSHD(8)

CIS Controls V7.0:

- Not Explicitly Mapped: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Not Explicitly Mapped: -- [More](#)
>

5.1.22 Ensure sshd UsePAM is enabled

Pass

Description:

The `UsePAM` directive enables the Pluggable Authentication Module (PAM) interface. If set to `yes` this will enable PAM authentication using `ChallengeResponseAuthentication` and `PasswordAuthentication` directives in addition to PAM account and session module processing for all authentication types.

Rationale:

When `usePAM` is set to `yes` , PAM runs through account and session types properly. This is important if you want to restrict access to services based off of IP, time or other factors of the account. Additionally, you can make sure users inherit certain environment variables on login or disallow access to the server

Remediation:

Edit the `/etc/ssh/sshd_config` file to set the `UsePAM` parameter to `yes` above any `Include` entries as follows:

```
UsePAM yes
```

Note: First occurrence of an option takes precedence. If `Include` locations are enabled, used, and order of precedence is understood in your environment, the entry may be created in a file in `Include` location.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: SSHD_CONFIG(5)
- URL: NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5
- URL: SSHD(8)

CIS Controls V7.0:

- Control 4: Controlled Use of Administrative Privileges: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 5: Account Management: -- [More](#)
>

5.2 Configure privilege escalation

There are various tools which allows a permitted user to execute a command as the superuser or another user, as specified by the security policy.

sudo

[sudo documentation](#)

The invoking user's real (not effective) user ID is used to determine the user name with which to query the security policy.

`sudo` supports a plug-in architecture for security policies and input/output logging. Third parties can develop and distribute their own policy and I/O logging plug-ins to work seamlessly with the `sudo` front end. The default security policy is `sudoers` , which is configured via the file `/etc/sudoers` and any entries in `/etc/sudoers.d` .

pkexec

[pkexec documentation](#)

`pkexec` allows an authorized user to execute *PROGRAM* as another user. If *username* is not specified, then the program will be executed as the administrative super user, `root` .

5.2.1 Ensure sudo is installed

Pass

Description:

`sudo` allows a permitted user to execute a command as the superuser or another user, as specified by the security policy. The invoking user's real (not effective) user ID is used to determine the user name with which to query the security policy.

Rationale:

`sudo` supports a plug-in architecture for security policies and input/output logging. Third parties can develop and distribute their own policy and I/O logging plug-ins to work seamlessly with the `sudo` front end. The default security policy is `sudoers`, which is configured via the file `/etc/sudoers` and any entries in `/etc/sudoers.d`.

The security policy determines what privileges, if any, a user has to run `sudo`. The policy may require that users authenticate themselves with a password or another authentication mechanism. If authentication is required, `sudo` will exit if the user's password is not entered within a configurable time limit. This limit is policy-specific.

Remediation:

First determine if LDAP functionality is required. If so, then install `sudo-ldap`, else install `sudo`.

Example:

```
# apt install sudo
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: SUDO(8)
- URL: NIST SP 800-53 Rev. 5: AC-6

CIS Controls V7.0:

- Control 4: Controlled Use of Administrative Privileges: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 5: Account Management: -- [More](#)
>

[Back to Summary](#)

5.2.2 Ensure sudo commands use pty

Fail

Description:

`sudo` can be configured to run only from a pseudo terminal (`pseudo-pty`).

Rationale:

Attackers can run a malicious program using `sudo` which would fork a background process that remains even when the main program has finished executing.

Remediation:

Edit the file `/etc/sudoers` with `visudo` or a file in `/etc/sudoers.d/` with `visudo -f <PATH TO FILE>` and add the following line:

```
Defaults use_pty
```

Edit the file `/etc/sudoers` with `visudo` and any files in `/etc/sudoers.d/` with `visudo -f <PATH TO FILE>` and remove any occurrence of `!use_pty`

Note:

- `sudo` will read each file in `/etc/sudoers.d`, skipping file names that end in `~` or contain a `.` character to avoid causing problems with package manager or editor temporary/backup files.
- Files are parsed in sorted lexical order. That is, `/etc/sudoers.d/01_first` will be parsed before `/etc/`

sudoers.d/10_second.

- Be aware that because the sorting is lexical, not numeric, /etc/sudoers.d/1_whoops would be loaded after /etc/sudoers.d/10_second.
- Using a consistent number of leading zeroes in the file names can be used to avoid such problems.

Impact:

WARNING: Editing the `sudo` configuration incorrectly can cause `sudo` to stop functioning. Always use `visudo` to modify `sudo` configuration files.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: SUDO(8)
- URL: VISUDO(8)
- URL: sudoers(5)
- URL: NIST SP 800-53 Rev. 5: AC-6

CIS Controls V7.0:

- Control 4: Controlled Use of Administrative Privileges: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 5: Account Management: -- [More](#)
>

[Back to Summary](#)

5.2.3 Ensure sudo log file exists

Fail

Description:

sudo can use a custom log file

Rationale:

A sudo log file simplifies auditing of sudo commands

Remediation:

Edit the file `/etc/sudoers` or a file in `/etc/sudoers.d/` with `visudo` or `visudo -f <PATH TO FILE>` and add the following line:

Example:

```
Defaults logfile="/var/log/sudo.log"
```

Note:

- sudo will read each file in `/etc/sudoers.d`, skipping file names that end in `~` or contain a `.` character to avoid causing problems with package manager or editor temporary/backup files.
- Files are parsed in sorted lexical order. That is, `/etc/sudoers.d/01_first` will be parsed before `/etc/sudoers.d/10_second`.
- Be aware that because the sorting is lexical, not numeric, `/etc/sudoers.d/1_whoops` would be loaded after `/etc/sudoers.d/10_second`.
- Using a consistent number of leading zeroes in the file names can be used to avoid such problems.

Impact:

WARNING: Editing the `sudo` configuration incorrectly can cause `sudo` to stop functioning. Always use `visudo` to modify `sudo` configuration files.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** SUDO(8)
- **URL:** VISUDO(8)
- **URL:** sudoers(5)
- **URL:** NIST SP 800-53 Rev. 5: AU-3, AU-12

CIS Controls V7.0:

- **Control 6: Maintenance, Monitoring and Analysis of Audit Logs:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 8: Audit Log Management:** -- [More](#)
>

[Back to Summary](#)**5.2.5 Ensure re-authentication for privilege escalation is not disabled globally**

Pass

Description:

The operating system must be configured so that users must re-authenticate for privilege escalation.

Rationale:

Without re-authentication, users may access resources or perform tasks for which they do not have authorization.

When operating systems provide the capability to escalate a functional capability, it is critical the user re-authenticate.

Remediation:

Configure the operating system to require users to reauthenticate for privilege escalation.

Based on the outcome of the audit procedure, use `visudo -f <PATH TO FILE>` to edit the relevant sudoers file.

Remove any occurrences of `!authenticate` tags in the file(s).

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: AC-6

CIS Controls V7.0:

- **Control 4: Controlled Use of Administrative Privileges:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 5: Account Management:** -- [More](#)
>

[Back to Summary](#)**5.2.6 Ensure sudo authentication timeout is configured**

Pass

Description:

`sudo` caches used credentials for a default of 5 minutes. This is for ease of use when there are multiple administrative tasks to perform. The timeout can be modified to suit local security policies.

Rationale:

A timeout value reduces the window of opportunity for unauthorized privileged `sudo` access.

Remediation:

- **IF** - the currently configured timeout is a negative number (disabled), greater than 15 minutes, or doesn't follow local

site polity:

Run the `visudo` command and edit or add the following line:

```
Defaults timestamp_timeout=<N>
```

Example:

```
Defaults timestamp_timeout=15
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** <https://www.sudo.ws/man/1.9.0/sudoers.man.html>
- **URL:** NIST SP 800-53 Rev. 5: AC-6

CIS Controls V7.0:

- **Control 4: Controlled Use of Administrative Privileges:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 5: Account Management:** -- [More](#)
>

[Back to Summary](#)

5.2.7 Ensure access to the su command is restricted

Fail

Description:

The `su` command allows a user to run a command or shell as another user. The program has been superseded by `sudo`, which allows for more granular control over privileged access. Normally, the `su` command can be executed by any user. By uncommenting the `pam_wheel.so` statement in `/etc/pam.d/su`, the `su` command will only allow users in a specific groups to execute `su`. This group should be empty to reinforce the use of `sudo` for privileged access.

Rationale:

Restricting the use of `su`, and using `sudo` in its place, provides system administrators better control of the escalation of user privileges to execute privileged commands. The `sudo` utility also provides a better logging and audit mechanism, as it can log each command executed via `sudo`, whereas `su` can only record that a user executed the `su` program.

Remediation:

Create an empty group that will be specified for use of the `su` command. The group should be named according to site policy.

Example:

```
# groupadd sugroup
```

Add the following line to the `/etc/pam.d/su` file, specifying the empty group:

```
auth required pam_wheel.so use_uid group=sugroup
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: AC-3. MP-2

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

5.3 Pluggable Authentication Modules

Pluggable Authentication Modules (PAM) is a service that implements modular authentication modules on Linux systems. PAM is implemented as a set of shared objects that are loaded and executed when a program needs to authenticate a user. Files for PAM are typically located in the `/etc/pam.d` directory. PAM must be carefully configured to secure system authentication. While this section covers some of PAM, please consult other PAM resources to fully understand the configuration capabilities.

Note: This section includes guidance that requires the additional modules and the latest version of software packages as covered in the "Configure PAM software packages" section

5.3.1 Configure PAM software packages

Updated versions of PAM include additional functionality

5.3.1.1 Ensure latest version of pam is installed

Pass

Description:

Linux Pluggable Authentication Modules (PAM) is a suite of libraries that allow a Linux system administrator to configure methods to authenticate users. It provides a flexible and centralized way to switch authentication methods for secured applications by using configuration files instead of changing application code

The `libpam-runtime` provides the runtime support for the PAM library

Rationale:

Older versions of the `libpam-runtime` package may not include the latest security and feature patches and updates.

Note: This Benchmark includes Recommendations that depend on newer `libpam-runtime` features. These Recommendation were written and tested against version `1.3.1-5ubuntu4.7`. Latest available version of the package should be used

Remediation:

Run the following command to install the latest version of `libpam-runtime` :

```
# apt install libpam-runtime
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** PAM(7)

CIS Controls V7.0:

- **Not Explicitly Mapped:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Not Explicitly Mapped:** -- [More](#)
>

[Back to Summary](#)

5.3.1.2 Ensure latest version of libpam-modules is installed

Pass

Description:

`libpam-modules` is a package containing a set of Pluggable Authentication Modules (PAM) which allows system administrators to configure different authentication methods for user logins, providing flexibility in how users can access applications by using various authentication modules to include password checks and other security mechanisms.

Rationale:

Older versions of the `libpam-modules` package may not include the latest security and feature patches and updates.

Note: This Benchmark includes Recommendations that depend on newer `libpam-modules` features. These recommendations were written and tested against version `1.3.1-5ubuntu4.7`. The latest available version should be used.

Remediation:

Run the following command to install the latest version of `libpam-modules`:

```
# apt install libpam-modules
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

CIS Controls V7.0:

- Not Explicitly Mapped: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Not Explicitly Mapped: -- [More](#)
>

[Back to Summary](#)

5.3.1.3 Ensure latest version of libpam-pwquality is installed

Fail

Description:

`libpwquality` provides common functions for password quality checking and scoring them based on their apparent randomness. The library also provides a function for generating random passwords with good pronounceability.

This module can be plugged into the password stack of a given service to provide some plug-in strength-checking for passwords. The code was originally based on `pam_cracklib` module and the module is backwards compatible with its options.

Rationale:

Strong passwords reduce the risk of systems being hacked through brute force methods.

Older versions of the `libpam-pwquality` package may not include the latest security and feature patches and updates.

Recommendations were written and tested against version `1.3.1-5ubuntu4.7`. The latest available version should be used.

Remediation:

Run the following command to install the latest version of `libpam-pwquality`:

```
# apt install libpam-pwquality
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** <https://packages.debian.org/buster/libpam-pwquality>

CIS Controls V7.0:

- **Control 4: Controlled Use of Administrative Privileges:** -- [More](#) >

CIS Critical Security Controls V8.0:

- **Control 5: Account Management:** -- [More](#) >

[Back to Summary](#)

5.3.2 Configure pam-auth-update profiles

`pam-auth-update` is a utility that permits configuring the central authentication policy for the system using pre-defined profiles as supplied by PAM module packages.

Profiles - Shipped in the `/usr/share/pam-configs/` directory specify the modules, with options, to enable; the preferred ordering with respect to other profiles; and whether a profile should be enabled by default. Packages providing PAM modules register their profiles at install time by calling `pam-auth-update --package`.

Selection of profiles is done using the standard `debconf` interface. The profile selection question will be asked at medium priority when packages are added or removed, so no user interaction is required by default. Users may invoke `pam-auth-update` directly to change their authentication configuration.

The `pam-auth-update` script makes every effort to respect local changes to `/etc/pam.d/common-*`. Local modifications to the list of module options will be preserved, and additions of modules within the managed portion of the stack will cause `pam-auth-update` to treat the config files as locally modified and not make further changes to the config files unless given the `--force` option.

If the user specifies that `pam-auth-update` should override local configuration changes, the locally-modified files will be saved in `/etc/pam.d/` with a suffix of `.pam-old`.

5.3.2.1 Ensure pam_unix module is enabled

Pass

Description:

`pam_unix` is the standard Unix authentication module. It uses standard calls from the system's libraries to retrieve and set account information as well as authentication. Usually this is obtained from the `/etc/passwd` and if shadow is enabled, the `/etc/shadow` file as well.

The account component performs the task of establishing the status of the user's account and password based on the following shadow elements: `expire`, `last_change`, `max_change`, `min_change`, `warn_change`. In the case of the latter, it may offer advice to the user on changing their password or, through the `PAM_AUTHTOKEN_REQD` return, delay giving service to the user until they have established a new password. The entries listed above are documented in the shadow(5) manual page. Should the user's record not contain one or more of these entries, the corresponding shadow check is not performed.

The authentication component performs the task of checking the users credentials (password). The default action of this module is to not permit the user access to a service if their official password is blank.

Rationale:

The system should only provide access after performing authentication of a user.

Remediation:

Run the following command to enable the `pam_unix` module:

```
# pam-auth-update --enable unix
```

Note: If a site specific custom profile is being used in your environment to configure PAM that includes the

152 of 22705/08/2025, 20:18

configuration for the `pam_faillock` module, enable that module instead

Assessment:
[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: IA-5(1)

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

5.3.2.2 Ensure `pam_faillock` module is enabled

Fail

Description:

The `pam_faillock.so` module maintains a list of failed authentication attempts per user during a specified interval and locks the account in case there were more than the configured number of consecutive failed authentications (this is defined by the `deny` parameter in the faillock configuration). It stores the failure records into per-user files in the tally directory.

Rationale:

Locking out user IDs after n unsuccessful consecutive login attempts mitigates brute force password attacks against your systems.

Remediation:

Create two `pam-auth-update` profiles in `/usr/share/pam-configs/`:

1. Create the `faillock` profile in `/usr/share/pam-configs/` with the following lines:

```
Name: Enable pam_faillock to deny access
Default: yes
Priority: 0
Auth-Type: Primary
Auth:
[default=die] pam_faillock.so authfail
```

Example Script:

```
#!/usr/bin/env bash

{
arr=('Name: Enable pam_faillock to deny access' 'Default: yes' 'Priority: 0' 'Auth-Type: Primary'
'Auth:' ' [default=die] pam_faillock.so authfail')

printf '%s\n' "${arr[@]}" > /usr/share/pam-configs/faillock
}
```

2. Create the `faillock_notify` profile in `/usr/share/pam-configs/` with the following lines:

```
Name: Notify of failed login attempts and reset count upon success
Default: yes
```

```
Priority: 1024

Auth-Type: Primary

Auth:

requisite pam_faillock.so preauth

Account-Type: Primary

Account:

required pam_faillock.so
```

Example Script:

```
#!/usr/bin/env bash

{

arr=('Name: Notify of failed login attempts and reset count upon success' 'Default: yes' 'Priority:
1024' 'Auth-Type: Primary' 'Auth:' ' ' requisite pam_faillock.so preauth' 'Account-Type: Primary'
'Account:' ' ' required pam_faillock.so')

printf '%s\n' "${arr[@]}" > /usr/share/pam-configs/faillock_notify

}
```

Run the following command to update the common-auth and common-account PAM files with the new profiles:

```
# pam-auth-update --enable <profile_filename>
```

Example:

```
# pam-auth-update --enable faillock

# pam-auth-update --enable faillock_notify
```

Note:

- The name used for the file must be used in the `pam-auth-update --enable` command
- The `Name :` line should be easily recognizable and understood
- The `Priority :` Line is important as it effects the order of the lines in the `/etc/pam.d/` files
- If a site specific custom profile is being used in your environment to configure PAM that includes the configuration for the `pam_faillock` module, enable that module instead

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

CIS Controls V7.0:

- [Control 16: Account Monitoring and Control: -- More](#)
>

CIS Critical Security Controls V8.0:

- [Control 6: Access Control Management: -- More](#)
>

[Back to Summary](#)

5.3.2.3 Ensure pam_pwquality module is enabled

Fail

Description:

The `pam_pwquality.so` module performs password quality checking. This module can be plugged into the password stack of a given service to provide strength-checking for passwords. The code was originally based on `pam_cracklib` module and the module is backwards compatible with its options.

The action of this module is to prompt the user for a password and check its strength against a system dictionary and

a set of rules for identifying poor choices.

The first action is to prompt for a single password, check its strength and then, if it is considered strong, prompt for the password a second time (to verify that it was typed correctly on the first occasion). All being well, the password is passed on to subsequent modules to be installed as the new authentication token.

Rationale:

Use of a unique, complex passwords helps to increase the time and resources required to compromise the password.

Remediation:

Run the following script to verify the `pam_pwquality.so` line exists in a `pam-auth-update` profile:

```
# grep -P -- '\bpam_pwquality\.so\b' /usr/share/pam-configs/*
```

Output should be similar to:

```
/usr/share/pam-configs/pwquality: requisite pam_pwquality.so retry=3
/usr/share/pam-configs/pwquality: requisite pam_pwquality.so retry=3
```

- **IF** - similar output is returned:

Run the following command to update `/etc/pam.d/common-password` with the returned profile:

```
# pam-auth-update --enable {PROFILE_NAME}
```

Example:

```
# pam-auth-update pwquality
```

- **IF** - similar output is **NOT** returned:

Create a `pam-auth-update` profile in `/usr/share/pam-configs/` with the following lines:

```
Name: Pwquality password strength checking
Default: yes
Priority: 1024
Conflicts: cracklib
Password-Type: Primary
Password:
requisite pam_pwquality.so retry=3
```

Example:

```
#!/usr/bin/env bash

{
arr=('Name: Pwquality password strength checking' 'Default: yes' 'Priority: 1024' 'Conflicts: cracklib'
'Password-Type: Primary' 'Password:' 'requisite pam_pwquality.so retry=3')

printf '%s\n' "${arr[@]}" > /usr/share/pam-configs/pwquality
}
```

Run the following command to update `/etc/pam.d/common-password` with the `pwquality` profile:

```
# pam-auth-update --enable pwquality
```

Note:

- The name used for the file must be used in the `pam-auth-update --enable` command
- The `Name :` line should be easily recognizable and understood
- The `Priority :` Line is important as it effects the order of the lines in the `/etc/pam.d/` files
- If a site specific custom profile is being used in your environment to configure PAM that includes the configuration for the `pam_pwquality` module, enable that module instead

Assessment:[Show](#) Assessment Evidence[Show](#) Rule Result XML**References:****CIS Controls V7.0:**

- Control 4: Controlled Use of Administrative Privileges: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 5: Account Management: -- [More](#)
>

[Back to Summary](#)**5.3.2.4 Ensure pam_pwhistory module is enabled****Fail****Description:**

The `pam_pwhistory.so` module saves the last passwords for each user in order to force password change history and keep the user from alternating between the same password too frequently.

This module does not work together with `kerberos`. In general, it does not make much sense to use this module in conjunction with `NIS` or `LDAP`, since the old passwords are stored on the local machine and are not available on another machine for password history checking.

Rationale:

Use of a unique, complex passwords helps to increase the time and resources required to compromise the password.

Remediation:

Run the following script to verify the `pam_pwhistory.so` line exists in a `pam-auth-update` profile:

```
# grep -P -- '\bpam_pwhistory\.so\b' /usr/share/pam-configs/*
```

Output should be similar to:

```
/usr/share/pam-configs/pwhistory: requisite pam_pwhistory.so remember=24 enforce_for_root use_authtok
```

- **IF** - similar output is returned:

Run the following command to update `/etc/pam.d/common-password` with the returned profile:

```
# pam-auth-update --enable {PROFILE_NAME}
```

Example:

```
# pam-auth-update pwhistory
```

- **IF** - similar output is **NOT** returned:

Create a `pwhistory` profile in `/usr/share/pam-configs/` with the following lines:

```
Name: pwhistory password history checking
Default: yes
Priority: 1024
Password-Type: Primary
Password: requisite pam_pwhistory.so remember=24 enforce_for_root use_authtok
```

Example Script:

```
#!/usr/bin/env bash

{
```

```
arr=('Name: pwhistory password history checking' 'Default: yes' 'Priority: 1024' 'Password-Type:
Primary' 'Password:' ' requisite pam_pwhistory.so remember=24 enforce_for_root use_authtok')

printf '%s\n' "${arr[@]}" > /usr/share/pam-configs/pwhistory

}
```

Run the following command to update `/etc/pam.d/common-password` with the `pwhistory` profile:

```
# pam-auth-update --enable pwhistory
```

Note:

- The name used for the file must be used in the `pam-auth-update --enable` command
- The `Name:` line should be easily recognizable and understood
- The `Priority:` Line is important as it effects the order of the lines in the `/etc/pam.d/` files
- If a site specific custom profile is being used in your environment to configure PAM that includes the configuration for the `pam_pwhistory` module, enable that module instead

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

CIS Controls V7.0:

- **Control 4: Controlled Use of Administrative Privileges:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 5: Account Management:** -- [More](#)
>

[Back to Summary](#)

5.3.3 Configure PAM Arguments

Pluggable Authentication Modules (PAM) uses arguments to pass information to a pluggable module during authentication for a particular module type. These arguments allow the PAM configuration files for particular programs to use a common PAM module but in different ways.

Invalid arguments are ignored and do not otherwise affect the success or failure of the PAM module. When an invalid argument is passed, an error is usually written to `/var/log/messages` file. However, since the reporting method is controlled by the PAM module, the module must be written correctly to log the error to this file.

Note: If custom PAM files are being used, for this section's remediation, the corresponding files in `/etc/pam.d/` would need to be edited directly, and the `pam-auth-update --enable <EDITED_PROFILE_NAME>` command skipped

5.3.3.1 Configure pam_faillock module

`pam_faillock.so` provides a way to configure the default settings for locking the user after multiple failed authentication attempts.

Options:

- `<dir=/path/to/tally-directory>` - The directory where the user files with the failure records are kept. The default is `/var/run/faillock`. Note: These files will disappear after reboot on systems configured with directory `/var/run/faillock` mounted on virtual memory.
- `audit` - Will log the user name into the system log if the user is not found.
- `silent` - Don't print informative messages to the user. Please note that when this option is not used there will be difference in the authentication behavior for users which exist on the system and non-existing users.
- `no_log_info` - Don't log informative messages via syslog(3).
- `local_users_only` - Only track failed user authentications attempts for local users in `/etc/passwd` and ignore centralized (AD, IdM, LDAP, etc.) users. The `faillock(8)` command will also no longer track user failed authentication attempts. Enabling this option will prevent a double-lockout scenario where a user is locked out locally and in the centralized mechanism.
- `nodelay` - Don't enforce a delay after authentication failures.
- `deny=<n>` - Deny access if the number of consecutive authentication failures for this user during the recent interval exceeds . The default is 3.
- `fail_interval=n` - The length of the interval during which the consecutive authentication failures must happen for the user account

- lock out is `n` seconds. The default is 900 (15 minutes).
- `unlock_time=n` - The access will be re-enabled after `n` seconds after the lock out. The value 0 has the same meaning as value never - the access will not be re-enabled without resetting the faillock entries by the `faillock(8)` command. The default is 600 (10 minutes). Note that the default directory that `pam_faillock` uses is usually cleared on system boot so the access will be also re-enabled after system reboot. If that is undesirable a different tally directory must be set with the `dir` option. Also note that it is usually undesirable to permanently lock out users as they can become easily a target of denial of service attack unless the usernames are random and kept secret to potential attackers.
 - `even_deny_root` - Root account can become locked as well as regular accounts.
 - `root_unlock_time=n` - This option implies `even_deny_root` option. Allow access after `n` seconds to root account after the account is locked. In case the option is not specified the value is the same as of the `unlock_time` option.
 - `admin_group=name` - If a group name is specified with this option, members of the group will be handled by this module the same as the root account (the options `even_deny_root` and `root_unlock_time` will apply to them. By default the option is not set.

5.3.3.1.1 Ensure password failed attempts lockout is configured

Fail

Description:

The `deny=<n>` option will deny access if the number of consecutive authentication failures for this user during the recent interval exceeds .

Rationale:

Locking out user IDs after *n* unsuccessful consecutive login attempts mitigates brute force password attacks against your systems.

Remediation:

Create or edit the following line in `/etc/security/faillock.conf` setting the `deny` option to 5 or less:

```
deny = 5
```

Run the following command:

```
# grep -Pl -- '\bpam_faillock\.so\h+([^\n\r]+\h+)?deny\b' /usr/share/pam-configs/*
```

Edit any returned files and remove the `deny=<N>` arguments from the `pam_faillock.so` line(s):

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

CIS Controls V7.0:

- Control 16: Account Monitoring and Control: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 6: Access Control Management: -- [More](#)
>

[Back to Summary](#)

5.3.3.1.2 Ensure password unlock time is configured

Fail

Description:

`unlock_time=<n>` - The access will be re-enabled after seconds after the lock out. The value 0 has the same meaning as value never - the access will not be re-enabled without resetting the faillock entries by the `faillock(8)` command.

Note:

- The default directory that `pam_faillock` uses is usually cleared on system boot so the access will be also re-enabled after system reboot. If that is undesirable a different tally directory must be set with the `dir` option.
- It is usually undesirable to permanently lock out users as they can become easily a target of denial of service attack unless the usernames are random and kept secret to potential attackers.
- The maximum configurable value for `unlock_time` is 604800

Rationale:

Locking out user IDs after n unsuccessful consecutive login attempts mitigates brute force password attacks against your systems.

Remediation:

Set password unlock time to conform to site policy. `unlock_time` should be 0 (never), or 900 seconds or greater.

Edit `/etc/security/faillock.conf` and update or add the following line:

```
unlock_time = 900
```

Run the following command: remove the `unlock_time` argument from the `pam_faillock.so` module in the PAM files:

```
# grep -Pl -- '\bpam_faillock\.so\h+([\^#\n\r]+\h+)?unlock_time\b' /usr/share/pam-configs/*
```

Edit any returned files and remove the `unlock_time=<N>` argument from the `pam_faillock.so` line(s):

Impact:

Use of `unlock_time=0` may allow an attacker to cause denial of service to legitimate users. This will also require a systems administrator with elevated privileges to unlock the account.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

CIS Controls V7.0:

- Control 16: Account Monitoring and Control: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 6: Access Control Management: -- [More](#)
>

[Back to Summary](#)

5.3.3.2 Configure pam_pwquality module

The `pam_pwquality.so` module checks the strength of passwords. It performs checks such as making sure a password is not a dictionary word, it is a certain length, contains a mix of characters (e.g. alphabet, numeric, other) and more.

These checks are configurable by either:

- use of the module arguments
- modifying the `/etc/security/pwquality.conf` configuration file
- creating a `.conf` file in the `/etc/security/pwquality.conf.d/` directory.

Note: The module arguments override the settings in the `/etc/security/pwquality.conf` configuration file. Settings in the `/etc/security/pwquality.conf` configuration file override settings in a `.conf` file in the `/etc/security/pwquality.conf.d/` directory.

The possible options in the file are:

- `difok` - Number of characters in the new password that must not be present in the old password. (default 1). The special value of 0 disables all checks of similarity of the new password with the old password except the new password being exactly the same as the old one.
- `minlen` - Minimum acceptable size for the new password (plus one if credits are not disabled which is the default). (See `pam_pwquality(8)`.) Cannot be set to lower value than 6. (default 8)
- `dcredit` - The maximum credit for having digits in the new password. If less than 0 it is the minimum number of digits in the new password. (default 0)
- `ucredit` - The maximum credit for having uppercase characters in the new password. If less than 0 it is the minimum number of uppercase characters in the new password. (default 0)
- `lcredit` - The maximum credit for having lowercase characters in the new password. If less than 0 it is the minimum number of lowercase characters in the new password. (default 0)
- `ocredit` - The maximum credit for having other characters in the new password. If less than 0 it is the minimum number of other characters in the new password. (default 0)
- `minclass` - The minimum number of required classes of characters for the new password (digits, uppercase, lowercase, others). (default

- 0)
- **maxrepeat** - The maximum number of allowed same consecutive characters in the new password. The check is disabled if the value is 0. (default 0)
- **maxsequence** - The maximum length of monotonic character sequences in the new password. Examples of such sequence are '12345' or 'fedcb'. Note that most such passwords will not pass the simplicity check unless the sequence is only a minor part of the password. The check is disabled if the value is 0. (default 0)
- **maxclassrepeat** - The maximum number of allowed consecutive characters of the same class in the new password. The check is disabled if the value is 0. (default 0)
- **gecoscheck** - If nonzero, check whether the words longer than 3 characters from the GECOS field of the user's passwd(5) entry are contained in the new password. The check is disabled if the value is 0. (default 0)
- **dictcheck** - If nonzero, check whether the password (with possible modifications) matches a word in a dictionary. Currently the dictionary check is performed using the cracklib library. (default 1)
- **usercheck=<N>** - If nonzero, check whether the password (with possible modifications) contains the user name in some form. It is not performed for user names shorter than 3 characters. (default 1)
- **usersubstr=<N>** - If greater than 3 (due to the minimum length in usercheck), check whether the password contains a substring of at least N length in some form. (default 0)
- **enforcing=<N>** - If nonzero, reject the password if it fails the checks, otherwise only print the warning. This setting applies only to the pam_pwquality module and possibly other applications that explicitly change their behavior based on it. It does not affect pwmake(1) and pwscore(1). (default 1)
- **badwords** - Space separated list of words that must not be contained in the password. These are additional words to the cracklib dictionary check. This setting can be also used by applications to emulate the gecost check for user accounts that are not created yet.
- **dictpath** - Path to the cracklib dictionaries. Default is to use the cracklib default.
- **retry=<N>** - Prompt user at most N times before returning with error. The default is 1.
- **enforce_for_root** - The module will return error on failed check even if the user changing the password is root. This option is off by default which means that just the message about the failed check is printed but root can change the password anyway. Note that root is not asked for an old password so the checks that compare the old and new password are not performed.
- **local_users_only** - The module will not test the password quality for users that are not present in the /etc/passwd file. The module still asks for the password so the following modules in the stack can use the use_authok option. This option is off by default.

5.3.3.2.1 Ensure password number of changed characters is configured

Fail

Description:

The `pwqualitydifok` option sets the number of characters in a password that must not be present in the old password.

Rationale:

Use of a complex password helps to increase the time and resources required to compromise the password. Password complexity, or strength, is a measure of the effectiveness of a password in resisting attempts at guessing and brute-force attacks.

Password complexity is one factor of several that determines how long it takes to crack a password. The more complex the password, the greater the number of possible combinations that need to be tested before the password is compromised.

Remediation:

Create or modify a file ending in `.conf` in the `/etc/security/pwquality.conf.d/` directory or the file `/etc/security/pwquality.conf` and add or modify the following line to set `difok` to 2 or more. Ensure setting conforms to local site policy:

Example:

```
#!/usr/bin/env bash

{
sed -ri 's/^\s*difok\s*=/# &/' /etc/security/pwquality.conf

[ ! -d /etc/security/pwquality.conf.d/ ] && mkdir /etc/security/pwquality.conf.d/

printf '\n%s' "difok = 2" > /etc/security/pwquality.conf.d/50-pwdifok.conf
}
```

Run the following command:

```
# grep -Pl -- '\bpam_pwquality\.so\h+([^\n\r]+\h+)?difok\b' /usr/share/pam-configs/*
```

Edit any returned files and remove the `difok` argument from the `pam_pwquality.so` line(s):

Assessment:
[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: IA-5

CIS Controls V7.0:

- Control 4: Controlled Use of Administrative Privileges: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 5: Account Management: -- [More](#)
>

[Back to Summary](#)

5.3.3.2.2 Ensure minimum password length is configured

Fail

Description:

The minimum password length setting determines the lowers number of characters that make up a password for a user account. There are many different theories about how to determine the best password length for an organization, but perhaps "passphrase" is a better term than "password".

The `minlen` option sets the minimum acceptable size for the new password (plus one if credits are not disabled which is the default). Cannot be set to lower value than 6.

Rationale:

Strong passwords help protect systems from password attacks. Types of password attacks include dictionary attacks, which attempt to use common words and phrases, and brute force attacks, which try every possible combination of characters. Also attackers may try to obtain the account database so they can use tools to discover the accounts and passwords.

Remediation:

Create or modify a file ending in `.conf` in the `/etc/security/pwquality.conf.d/` directory or the file `/etc/security/pwquality.conf` and add or modify the following line to set password length of 14 or more characters. Ensure that password length conforms to local site policy:

Example:

```
#!/usr/bin/env bash

{
sed -ri 's/^\s*minlen\s*=/# &/' /etc/security/pwquality.conf
[ ! -d /etc/security/pwquality.conf.d/ ] && mkdir /etc/security/pwquality.conf.d/
printf '\n%s' "minlen = 14" > /etc/security/pwquality.conf.d/50-pwlength.conf
}
```

Run the following command:

```
# grep -Pl -- '\bpam_pwquality\.so\h+([^\n\r]+\h+)?minlen\b' /usr/share/pam-configs/*
```

Edit any returned files and remove the `minlen` argument from the `pam_pwquality.so` line(s):

Impact:

In general, it is true that longer passwords are better (harder to crack), but it is also true that forced password length requirements can cause user behavior that is predictable and undesirable. For example, requiring users to have a minimum 16-character password may cause them to choose repeating patterns like fourfourfourfour or passwordpassword that meet the requirement but aren't hard to guess. Additionally, length requirements increase the chances that users will adopt other insecure practices, like writing them down, re-using them or storing them

unencrypted in their documents.

Having a reasonable minimum length with no maximum character limit increases the resulting average password length used (and therefore the strength).6

Assessment:

Show Assessment Evidence

Show Rule Result XML

References:

- URL: pam_pwquality(8)
- URL: CIS Password Policy Guide
- URL: NIST SP 800-53 Rev. 5: 1A-5(1)

CIS Controls V7.0:

- Control 4: Controlled Use of Administrative Privileges: -- More >

CIS Critical Security Controls V8.0:

- Control 5: Account Management: -- More >

Back to Summary

5.3.3.2.3 Ensure password complexity is configured

Manual

Description:

Password complexity can be set through:

- minclass - The minimum number of classes of characters required in a new password. (digits, uppercase, lowercase, others). e.g. minclass = 4 requires digits, uppercase, lower case, and special characters.
- dcredit - The maximum credit for having digits in the new password. If less than 0 it is the minimum number of digits in the new password. e.g. dcredit = -1 requires at least one digit
- ucredit - The maximum credit for having uppercase characters in the new password. If less than 0 it is the minimum number of uppercase characters in the new password. e.g. ucredit = -1 requires at least one uppercase character
- ocredit - The maximum credit for having other characters in the new password. If less than 0 it is the minimum number of other characters in the new password. e.g. ocredit = -1 requires at least one special character
- lcredit - The maximum credit for having lowercase characters in the new password. If less than 0 it is the minimum number of lowercase characters in the new password. e.g. lcredit = -1 requires at least one lowercase character

Rationale:

Strong passwords protect systems from being hacked through brute force methods.

Requiring at least one non-alphabetic character increases the search space beyond pure dictionary words, which makes the resulting password harder to crack.

Forcing users to choose an excessively complex password, e.g. some combination of upper-case, lower-case, numbers, and special characters, has a negative impact. It places an extra burden on users and many will use predictable patterns (for example, a capital letter in the first position, followed by lowercase letters, then one or two numbers, and a "special character" at the end). Attackers know this, so dictionary attacks will often contain these common patterns and use the most common substitutions like, \$ for s, @ for a, 1 for l, 0 for o.

Remediation:

Run the following command:

```
# grep -Pl -- '\bpam_pwquality\.so\h+([^\n\r]+\h+)?(minclass|[duloclcredit)\b' /usr/share/pam-configs/*
```

Edit any returned files and remove the minclass , dcredit , ucredit , lcredit , and ocredit arguments from the pam_pwquality.so line(s)

Create or modify a file ending in .conf in the /etc/security/pwquality.conf.d/ directory or the file /etc/security/pwquality.conf and add or modify the following line(s) to set complexity according to local site policy:

- minclass = _N_
- dcredit = _N_ # Value should be either 0 or a number proceeded by a minus (-) symbol

- ucredit = -1 # Value should be either 0 or a number proceeded by a minus (-) symbol
- ocredit = -1 # Value should be either 0 or a number proceeded by a minus (-) symbol
- lcredit = -1 # Value should be either 0 or a number proceeded by a minus (-) symbol

Example 1 - Set minclass = 3:

```
#!/usr/bin/env bash

{

sed -ri 's/^\s*minclass\s*=/# &/' /etc/security/pwquality.conf

sed -ri 's/^\s*[dulo]credit\s*=/# &/' /etc/security/pwquality.conf

[ ! -d /etc/security/pwquality.conf.d/ ] && mkdir /etc/security/pwquality.conf.d/

printf '\n%s' "minclass = 3" > /etc/security/pwquality.conf.d/50-pwcomplexity.conf

}
```

Example 2 - set dcredit = -1, ucredit = -1, and lcredit = -1:

```
#!/usr/bin/env bash

{

sed -ri 's/^\s*minclass\s*=/# &/' /etc/security/pwquality.conf

sed -ri 's/^\s*[dulo]credit\s*=/# &/' /etc/security/pwquality.conf

[ ! -d /etc/security/pwquality.conf.d/ ] && mkdir /etc/security/pwquality.conf.d/

printf '%s\n' "dcredit = -1" "ucredit = -1" "lcredit = -1" > /etc/security/pwquality.conf.d/50-
pwcomplexity.conf

}
```

Impact:

Passwords that are too complex in nature make it harder for users to remember, leading to bad practices. In addition, composition requirements provide no defense against common attack types such as social engineering or insecure storage of passwords

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: pam_pwquality(8)
- URL: PWQUALITY.CONF(5)
- URL: <https://www.cisecurity.org/insights/white-papers/cis-password-policy-guide>
- URL: NIST SP 800-53 Rev. 5: 1A-5

CIS Controls V7.0:

- Control 4: Controlled Use of Administrative Privileges: -- [More](#)

>

CIS Critical Security Controls V8.0:

- Control 5: Account Management: -- [More](#)

>

[Back to Summary](#)

5.3.3.2.4 Ensure password same consecutive characters is configured

Fail

Description:

The pwqualitymaxrepeat option sets the maximum number of allowed same consecutive characters in a new password.

Rationale:

Use of a complex password helps to increase the time and resources required to compromise the password. Password complexity, or strength, is a measure of the effectiveness of a password in resisting attempts at guessing and brute-force attacks.

Password complexity is one factor of several that determines how long it takes to crack a password. The more complex the password, the greater the number of possible combinations that need to be tested before the password is compromised.

Remediation:

Create or modify a file ending in `.conf` in the `/etc/security/pwquality.conf.d/` directory or the file `/etc/security/pwquality.conf` and add or modify the following line to set `maxrepeat` to 3 or less and not 0 . Ensure setting conforms to local site policy:

Example:

```
#!/usr/bin/env bash

{
sed -ri 's/^\s*maxrepeat\s*=/# &/' /etc/security/pwquality.conf

[ ! -d /etc/security/pwquality.conf.d/ ] && mkdir /etc/security/pwquality.conf.d/

printf '\n%s' "maxrepeat = 3" > /etc/security/pwquality.conf.d/50-pwrepeat.conf
}
```

Run the following command:

```
# grep -Pl -- '\bpam_pwquality\.so\h+([^\n\r]+\h+)?maxrepeat\b' /usr/share/pam-configs/*
```

Edit any returned files and remove the `maxrepeat` argument from the `pam_pwquality.so` line(s):

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: IA-5

CIS Controls V7.0:

- Control 4: Controlled Use of Administrative Privileges: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 5: Account Management: -- [More](#)
>

[Back to Summary](#)

5.3.3.2.5 Ensure password maximum sequential characters is configured

Fail

Description:

The `pwqualitymaxsequence` option sets the maximum length of monotonic character sequences in the new password. Examples of such sequence are 12345 or fedcb . The check is disabled if the value is 0 .

Note: Most such passwords will not pass the simplicity check unless the sequence is only a minor part of the password.

Rationale:

Use of a complex password helps to increase the time and resources required to compromise the password. Password complexity, or strength, is a measure of the effectiveness of a password in resisting attempts at guessing and brute-force attacks.

Password complexity is one factor of several that determines how long it takes to crack a password. The more complex the password, the greater the number of possible combinations that need to be tested before the password is compromised.

Remediation:

Create or modify a file ending in `.conf` in the `/etc/security/pwquality.conf.d/` directory or the file `/etc/security/pwquality.conf` and add or modify the following line to set `maxsequence` to 3 or less and not 0 .
Ensure setting conforms to local site policy:

Example:

```
#!/usr/bin/env bash

{
sed -ri 's/^\s*maxsequence\s*=/# &/' /etc/security/pwquality.conf

[ ! -d /etc/security/pwquality.conf.d/ ] && mkdir /etc/security/pwquality.conf.d/

printf '\n%s' "maxsequence = 3" > /etc/security/pwquality.conf.d/50-pwmaxsequence.conf
}
```

Run the following command:

```
# grep -Pl -- '\bpam_pwquality\.so\h+([^\n\r]+\h+)?maxsequence\b' /usr/share/pam-configs/*
```

Edit any returned files and remove the `maxsequence` argument from the `pam_pwquality.so` line(s):

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: IA-5

CIS Controls V7.0:

- Control 4: Controlled Use of Administrative Privileges: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 5: Account Management: -- [More](#)
>

[Back to Summary](#)

5.3.3.2.6 Ensure password dictionary check is enabled

Pass

Description:

The `pwqualitydictcheck` option sets whether to check for the words from the `cracklib` dictionary.

Rationale:

If the operating system allows the user to select passwords based on dictionary words, this increases the chances of password compromise by increasing the opportunity for successful guesses, and brute-force attacks.

Remediation:

Edit any file ending in `.conf` in the `/etc/security/pwquality.conf.d/` directory and/or the file `/etc/security/pwquality.conf` and comment out or remove any instance of `dictcheck = 0`:

Example:

```
# sed -ri 's/^\s*dictcheck\s*=/# &/' /etc/security/pwquality.conf /etc/security/pwquality.conf.d/*.conf
```

Run the following command:

grep -Pl -- '\bpam_pwquality\.so\h+([^\n\r]+\h+)?dictcheck\b' /usr/share/pam-configs/*

Edit any returned files and remove the dictcheck argument from the pam_pwquality.so line(s)

Assessment:

Show Assessment Evidence

Show Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: IA-5

CIS Controls V7.0:

- Control 4: Controlled Use of Administrative Privileges: -- More >

CIS Critical Security Controls V8.0:

- Control 5: Account Management: -- More >

Back to Summary

5.3.3.2.7 Ensure password quality checking is enforced

Pass

Description:

The pam_pwquality module can be configured to either reject a password if it fails the checks, or only print a warning.

This is configured by setting the enforcing=<N> argument. If nonzero, a password will be rejected if it fails the checks, otherwise only a warning message will be provided.

This setting applies only to the pam_pwquality module and possibly other applications that explicitly change their behavior based on it. It does not affect pwmake(1) and pwscore(1).

Rationale:

Strong passwords help protect systems from password attacks. Types of password attacks include dictionary attacks, which attempt to use common words and phrases, and brute force attacks, which try every possible combination of characters. Also attackers may try to obtain the account database so they can use tools to discover the accounts and passwords.

Remediation:

Run the following command:

grep -Pl -- '\bpam_pwquality\.so\h+([^\n\r]+\h+)?enforcing=0\b' /usr/share/pam-configs/*

Edit any returned files and remove the enforcing=0 argument from the pam_pwquality.so line(s)

Edit /etc/security/pwquality.conf and all files ending in .conf in the /etc/security/pwquality.conf.d/ directory and remove or comment out any line containing the enforcing = 0 argument:

Example:

sed -ri 's/^\s*enforcing\s*=\s*0/# &/' /etc/security/pwquality.conf /etc/security/pwquality.conf.d/*.conf

Assessment:

Show Assessment Evidence

Show Rule Result XML

References:

- URL: pam_pwquality(8)
- URL: PWQUALITY.CONF(5)
- URL: NIST SP 800-53 Rev. 5: IA-5

CIS Controls V7.0:

- **Control 4: Controlled Use of Administrative Privileges:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 5: Account Management:** -- [More](#)
>

[Back to Summary](#)

5.3.3.2.8 Ensure password quality is enforced for the root user

Fail

Description:

If the `pwqualityenforce_for_root` option is enabled, the module will return error on failed check even if the user changing the password is root.

This option is off by default which means that just the message about the failed check is printed but root can change the password anyway.

Note: The root is not asked for an old password so the checks that compare the old and new password are not performed.

Rationale:

Use of a complex password helps to increase the time and resources required to compromise the password. Password complexity, or strength, is a measure of the effectiveness of a password in resisting attempts at guessing and brute-force attacks.

Password complexity is one factor of several that determines how long it takes to crack a password. The more complex the password, the greater the number of possible combinations that need to be tested before the password is compromised.

Remediation:

Edit or add the following line in a `*.conf` file in `/etc/security/pwquality.conf.d` or in `/etc/security/pwquality.conf`:

Example:

```
#!/usr/bin/env bash

{
[ ! -d /etc/security/pwquality.conf.d/ ] && mkdir /etc/security/pwquality.conf.d/
printf '\n%s\n' "enforce_for_root" > /etc/security/pwquality.conf.d/50-pwroot.conf
}
```

Assessment:

[Show](#) Assessment Evidence[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: IA-5

CIS Controls V7.0:

- **Control 4: Controlled Use of Administrative Privileges:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 5: Account Management:** -- [More](#)
>

[Back to Summary](#)

5.3.3.3 Configure pam_pwhistory module

pam_pwhistory - PAM module to remember last passwords

pam_history.so module - This module saves the last passwords for each user in order to force password change history and keep the user from alternating between the same password too frequently.

This module does not work together with kerberos. In general, it does not make much sense to use this module in conjunction with NIS or LDAP, since the old passwords are stored on the local machine and are not available on another machine for password history checking.

Options:

- debug - Turns on debugging via syslog(3).
- use_authtok - When password changing enforce the module to use the new password provided by a previously stacked password module (this is used in the example of the stacking of the pam_passwdqc module documented below).
- enforce_for_root - If this option is set, the check is enforced for root, too.
- remember=<N> - The last <N> passwords for each user are saved. The default is 10. Value of 0 makes the module to keep the existing contents of the opasswd file unchanged.
- retry=<N> - Prompt user at most <N> times before returning with error. The default is 1.
- authtok_type=<STRING> - See pam_get_authtok(3) for more details.

Examples:

An example password section would be:

```
##PAM-1.0
```

```
password required pam_pwhistory.so
```

```
password required pam_unix.so use_authtok
```

In combination with pam_passwdqc:

```
##PAM-1.0
```

```
password required pam_passwdqc.so config=/etc/passwdqc.conf
```

```
password required pam_pwhistory.so use_authtok
```

```
password required pam_unix.so use_authtok
```

5.3.3.3.1 Ensure password history remember is configured

Fail

Description:

The /etc/security/opasswd file stores the users' old passwords and can be checked to ensure that users are not recycling recent passwords. The number of passwords remembered is set via the remember argument value in set for the pam_pwhistory module.

- remember=<N> - <N> is the number of old passwords to remember

Rationale:

Requiring users not to reuse their passwords make it less likely that an attacker will be able to guess the password or use a compromised password.

Note: These change only apply to accounts configured on the local system.

Remediation:

Run the following command:

```
# awk '/Password-Type:/{ f = 1;next } /-Type:/{ f = 0 } f {if (/pam_pwhistory\.so/) print FILENAME}' /usr/share/pam-configs/*
```

Edit any returned files and edit or add the remember= argument, with a value of 24 or more, that meets local site policy to the pam_pwhistory line in the Password section:

Example File:


```
Name: pwhistory password history checking
Default: yes
Priority: 1024
Password-Type: Primary
Password:

requisite pam_pwhistory.so remember=24 enforce_for_root use_authtok # <- **ensure line includes
remember=<N>**
```

Run the following command to update the files in the `/etc/pam.d/` directory:

```
# pam-auth-update --enable <MODIFIED_PROFILE_NAME>
```

Example:

```
# pam-auth-update --enable pwhistory
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: IA-5(1)

CIS Controls V7.0:

- **Control 4: Controlled Use of Administrative Privileges:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 5: Account Management:** -- [More](#)
>

[Back to Summary](#)

5.3.3.3.2 Ensure password history is enforced for the root user

Fail

Description:

If the `pwhistoryenforce_for_root` option is enabled, the module will enforce password history for the root user as well

Rationale:

Requiring users not to reuse their passwords make it less likely that an attacker will be able to guess the password or use a compromised password

Note: These change only apply to accounts configured on the local system.

Remediation:

Run the following command:

```
# awk '/Password-Type:{ f = 1;next } /-Type:{ f = 0 } f {if (/pam_pwhistory\.so/) print FILENAME}' /usr/share/pam-configs/*
```

Edit any returned files and add the `enforce_for_root` argument to the `pam_pwhistory` line in the `Password` section:

Example File:

```
Name: pwhistory password history checking
Default: yes
Priority: 1024
Password-Type: Primary
```

```
Password:

requisite pam_pwhistory.so remember=24 enforce_for_root use_authtok # <- **ensure line includes
enforce_for_root**
```

Run the following command to update the files in the `/etc/pam.d/` directory:

```
# pam-auth-update --enable <MODIFIED_PROFILE_NAME>
```

Example:

```
# pam-auth-update --enable pwhistory
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: IA-5

CIS Controls V7.0:

- **Control 4: Controlled Use of Administrative Privileges:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 5: Account Management:** -- [More](#)
>

[Back to Summary](#)

5.3.3.3.3 Ensure pam_pwhistory includes use_authtok

Fail

Description:

`use_authtok` - When password changing enforce the module to set the new password to the one provided by a previously stacked password module

Rationale:

`use_authtok` allows multiple pam modules to confirm a new password before it is accepted.

Remediation:

Run the following command:

```
# awk '/Password-Type:/{ f = 1;next } /-Type:/{ f = 0 } f {if (/pam_pwhistory\.so/) print FILENAME}' /
usr/share/pam-configs/*
```

Edit any returned files and add the `use_authtok` argument to the `pam_pwhistory` line in the `Password` section:

Example File:

```
Name: pwhistory password history checking

Default: yes

Priority: 1024

Password-Type: Primary

Password:

requisite pam_pwhistory.so remember=24 enforce_for_root use_authtok # <- **ensure line includes
use_authtok**
```

Run the following command to update the files in the `/etc/pam.d/` directory:

```
# pam-auth-update --enable <MODIFIED_PROFILE_NAME>
```

Example:

```
# pam-auth-update --enable pwhistory
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: IA-5

CIS Controls V7.0:

- Control 16: Account Monitoring and Control: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

5.3.3.4 Configure pam_unix module

The `pam_unix.so` module is the standard Unix authentication module. It uses standard calls from the system's libraries to retrieve and set account information as well as authentication. Usually this is obtained from the `/etc/passwd` and the `/etc/shadow` file as well if shadow is enabled.

5.3.3.4.1 Ensure pam_unix does not include nullok

Fail

Description:

The `nullok` argument overrides the default action of `pam_unix.so` to not permit the user access to a service if their official password is blank.

Rationale:

Using a strong password is essential to helping protect personal and sensitive information from unauthorized access

Remediation:

Run the following command:

```
# grep -PH -- '^\h*([^\n\r]+\h+)?pam_unix\.so\h+([^\n\r]+\h+)?nullok\b' /usr/share/pam-configs/*
```

Edit any files returned and remove the `nullok` argument for the `pam_unix` lines

Example File:

```
Name: Unix authentication
Default: yes
Priority: 256
Auth-Type: Primary
Auth:
[success=end default=ignore] pam_unix.so try_first_pass # <- **ensure line does not include nullok
nullok**
Auth-Initial:
[success=end default=ignore] pam_unix.so # <- **ensure line does not include nullok nullok**
Account-Type: Primary
Account:
[success=end new_authtok_reqd=done default=ignore] pam_unix.so
Account-Initial:
```

```
[success=end new_authtok_reqd=done default=ignore] pam_unix.so

Session-Type: Additional

Session:

required pam_unix.so

Session-Initial:

required pam_unix.so

Password-Type: Primary

Password:

[success=end default=ignore] pam_unix.so obscure use_authtok try_first_pass yescrypt

Password-Initial:

[success=end default=ignore] pam_unix.so obscure yescrypt
```

Run the following command to update the files in the `/etc/pam.d/` directory:

```
# pam-auth-update --enable <EDITED_PROFILE_NAME>
```

Example:

```
# pam-auth-update --enable unix
```

Note: If custom files are being used, the corresponding files in `/etc/pam.d/` would need to be edited directly, and the `pam-auth-update --enable <EDITED_PROFILE_NAME>` command skipped

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

CIS Controls V7.0:

- Control 4: Controlled Use of Administrative Privileges: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 5: Account Management: -- [More](#)
>

[Back to Summary](#)

5.3.3.4.2 Ensure pam_unix does not include remember

Pass

Description:

The `remember=n` argument saves the last `n` passwords for each user in `/etc/security/opasswd` in order to force password change history and keep the user from alternating between the same password too frequently. The MD5 password hash algorithm is used for storing the old passwords. Instead of this option the `pam_pwhistory` module should be used. The `pam_pwhistory` module saves the last `n` passwords for each user in `/etc/security/opasswd` using the password hash algorithm set on the `pam_unix` module. This allows for the `yescrypt` or `sha512` hash algorithm to be used.

Rationale:

The `remember=n` argument should be removed to ensure a strong password hashing algorithm is being used. A stronger hash provides additional protection to the system by increasing the level of effort needed for an attacker to successfully determine local user's old passwords stored in `/etc/security/opasswd`.

Remediation:

Run the following command:

```
# grep -PH -- '^h*([^\n\r]+\h+)?pam_unix\.so\h+([^\n\r]+\h+)?remember\b' /usr/share/pam-configs/*
```

Edit any files returned and remove the `remember=<N>` argument for the `pam_unix` lines

Example output:

```
[success=end default=ignore] pam_unix.so obscure use_authtok try_first_pass yescrypt remember=5 # **<-
remove remember=<N>**

[success=end default=ignore] pam_unix.so obscure yescrypt remember=5 # **<- remove remember=<N>**
```

Run the following command to update the files in the `/etc/pam.d/` directory:

```
# pam-auth-update --enable <EDITED_PROFILE_NAME>
```

Example:

```
# pam-auth-update --enable unix
```

Note: If custom files are being used, the corresponding files in `/etc/pam.d/` would need to be edited directly, and the `pam-auth-update --enable <EDITED_PROFILE_NAME>` command skipped

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

CIS Controls V7.0:

- Control 4: Controlled Use of Administrative Privileges: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 5: Account Management: -- [More](#)
>

[Back to Summary](#)

5.3.3.4.3 Ensure pam_unix includes a strong password hashing algorithm

Pass

Description:

A cryptographic hash function converts an arbitrary-length input into a fixed length output. Password hashing performs a one-way transformation of a password, turning the password into another string, called the hashed password.

The `pam_unix` module can be configured to use one of the following hashing algorithms for user's passwords:

- `md5` - When a user changes their password next, encrypt it with the MD5 algorithm.
- `bigcrypt` - When a user changes their password next, encrypt it with the DEC C2 algorithm.
- `sha256` - When a user changes their password next, encrypt it with the SHA256 algorithm. The SHA256 algorithm must be supported by the `crypt(3)` function.
- `sha512` - When a user changes their password next, encrypt it with the SHA512 algorithm. The SHA512 algorithm must be supported by the `crypt(3)` function.
- `blowfish` - When a user changes their password next, encrypt it with the blowfish algorithm. The blowfish algorithm must be supported by the `crypt(3)` function.

Rationale:

The SHA-512 algorithm provides a stronger hash than other algorithms used for password hash generation. A stronger hash provides additional protection to the system by increasing the level of effort needed for an attacker to successfully determine local user passwords.

Note: These changes only apply to the local system.

Remediation:

Run the following command:

```
# awk '/Password-Type:{ f = 1;next } /-Type:{ f = 0 } f {if (/pam_unix\.so/) print FILENAME}' /usr/
share/pam-configs/*
```

Edit any returned files and edit or add a the strong sha512 hashing algorithm to the pam_unix lines in the Password section:

Example File:

```
Name: Unix authentication

Default: yes

Priority: 256

Auth-Type: Primary # <- Start of "Auth" section

Auth:

[success=end default=ignore] pam_unix.so try_first_pass

Auth-Initial:

[success=end default=ignore] pam_unix.so

Account-Type: Primary # <- Start of "Account" section

Account:

[success=end new_authtok_reqd=done default=ignore] pam_unix.so

Account-Initial:

[success=end new_authtok_reqd=done default=ignore] pam_unix.so

Session-Type: Additional # <- Start of "Session" section

Session:

required pam_unix.so

Session-Initial:

required pam_unix.so

Password-Type: Primary # <- Start of "Password" section

Password:

[success=end default=ignore] pam_unix.so obscure use_authtok try_first_pass sha512 # <- **ensure hashing
algorithm is sha512**

Password-Initial:

[success=end default=ignore] pam_unix.so obscure sha512 # <- **ensure hashing algorithm is sha512**
```

Note: yescrypt will be acceptable - IF - it becomes supported

Run the following command to update the files in the /etc/pam.d/ directory:

```
# pam-auth-update --enable <MODIFIED_PROFILE_NAME>
```

Example:

```
# pam-auth-update --enable unix
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- [URL: NIST SP 800-53 Rev. 5: IA-5](#)

CIS Controls V7.0:

- [Control 16: Account Monitoring and Control: -- More](#)
>

CIS Critical Security Controls V8.0:

- [Control 3: Data Protection: -- More](#)

>

[Back to Summary](#)

5.3.3.4.4 Ensure pam_unix includes use_authtok

Fail

Description:

use_authtok - When password changing enforce the module to set the new password to the one provided by a previously stacked password module

Rationale:

use_authtok allows multiple pam modules to confirm a new password before it is accepted.

Remediation:

Run the following command:

```
# awk '/Password-Type:/{ f = 1;next } /-Type:/{ f = 0 } f {if (/pam_unix\.so/) print FILENAME}' /usr/share/pam-configs/*
```

Edit any returned files add use_authtok to the pam_unix line in the Password section under Password: subsection:

Note: The if the file's Password section includes a Password-Initial: subsection, use_authtok should not be added to the pam_unix line in the Password-Initial: subsection

Example File:

```
Name: Unix authentication

Default: yes

Priority: 256

Auth-Type: Primary # <- Start of "Auth" section

Auth:

[success=end default=ignore] pam_unix.so try_first_pass

Auth-Initial:

[success=end default=ignore] pam_unix.so

Account-Type: Primary # <- Start of "Account" section

Account:

[success=end new_authtok_reqd=done default=ignore] pam_unix.so

Account-Initial:

[success=end new_authtok_reqd=done default=ignore] pam_unix.so

Session-Type: Additional # <- Start of "Session" section

Session:

required pam_unix.so

Session-Initial:

required pam_unix.so

Password-Type: Primary # <- Start of "Password" section

Password:

[success=end default=ignore] pam_unix.so obscure use_authtok try_first_pass yescrypt # <- **ensure line includes use_authtok**

Password-Initial:

[success=end default=ignore] pam_unix.so obscure yescrypt # <- **Password-Initial: subsection does not include use_authtok
```

Run the following command to update the files in the `/etc/pam.d/` directory:

```
# pam-auth-update --enable <MODIFIED_PROFILE_NAME>
```

Example:

```
# pam-auth-update --enable unix
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: IA-5

CIS Controls V7.0:

- Control 16: Account Monitoring and Control: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

5.4 User Accounts and Environment

This section provides guidance on setting up secure defaults for system and user accounts and their environment.

5.4.1 Configure shadow password suite parameters

While a majority of the password control parameters have been moved to PAM, some parameters are still available through the shadow password suite. Any changes made to `/etc/login.defs` will only be applied if the `usermod` command is used. If user IDs are added a different way, use the `chage` command to effect changes to individual user IDs.

5.4.1.1 Ensure password expiration is configured

Fail

Description:

The `PASS_MAX_DAYS` parameter in `/etc/login.defs` allows an administrator to force passwords to expire once they reach a defined age.

`PASS_MAX_DAYS<N>` - The maximum number of days a password may be used. If the password is older than this, a password change will be forced. If not specified, -1 will be assumed (which disables the restriction).

Rationale:

The window of opportunity for an attacker to leverage compromised credentials or successfully compromise credentials via an online brute force attack is limited by the age of the password. Therefore, reducing the maximum age of a password also reduces an attacker's window of opportunity.

We recommend a yearly password change. This is primarily because for all their good intentions users will share credentials across accounts. Therefore, even if a breach is publicly identified, the user may not see this notification, or forget they have an account on that site. This could leave a shared credential vulnerable indefinitely. Having an organizational policy of a 1-year (annual) password expiration is a reasonable compromise to mitigate this with minimal user burden.

Remediation:

Set the `PASS_MAX_DAYS` parameter to conform to site policy in `/etc/login.defs`:

```
PASS_MAX_DAYS 365
```

Modify user parameters for all users with a password set to match:

```
# chage --maxdays 365 <user>
```

Edit `/etc/login.defs` and set `PASS_MAX_DAYS` to a value greater than 0 that follows local site policy:

Example:

```
PASS_MAX_DAYS 365
```

Run the following command to modify user parameters for all users with a password set to a maximum age no greater than 365 or less than 1 that follows local site policy:

```
# chage --maxdays <N> <user>
```

Example:

```
# awk -F: '($2~/^\$.+\/) {if($5 > 365 || $5 < 1)system ("chage --maxdays 365 " $1)}' /etc/shadow
```

Warning: If a password has been set at system install or kickstart, the last change date field is not set, In this case, setting PASS_MAX_DAYS will immediately expire the password. One possible solution is to populate the last change date field through a command like: `chage -d "$(date +%Y-%m-%d)" root`

Impact:

The password expiration must be greater than the minimum days between password changes or users will be unable to change their password.

Excessive password expiration requirements do more harm than good, because these requirements make users select predictable passwords, composed of sequential words and numbers that are closely related to each other. In these cases, the next password can be predicted based on the previous one (incrementing a number used in the password for example). Also, password expiration requirements offer no containment benefits because attackers will often use credentials as soon as they compromise them. Instead, immediate password changes should be based on key events including, but not limited to:

- Indication of compromise
- Change of user roles
- When a user leaves the organization.

Not only does changing passwords every few weeks or months frustrate the user, but it's also been suggested that it does more harm than good, because it could lead to bad practices by the user such as adding a character to the end of their existing password.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** CIS Password Policy Guide
- **URL:** NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5

CIS Controls V7.0:

- **Control 4: Controlled Use of Administrative Privileges:** -- [More](#) >

CIS Critical Security Controls V8.0:

- **Control 5: Account Management:** -- [More](#) >

[Back to Summary](#)

5.4.1.3 Ensure password expiration warning days is configured

Pass

Description:

The PASS_WARN_AGE parameter in /etc/login.defs allows an administrator to notify users that their password will expire in a defined number of days.

PASS_WARN_AGE<N> - The number of days warning given before a password expires. A zero means warning is given only upon the day of expiration, a negative value means no warning is given. If not specified, no warning will be provided.

Rationale:

Providing an advance warning that a password will be expiring gives users time to think of a secure password. Users

caught unaware may choose a simple password or write it down where it may be discovered.

Remediation:

Edit `/etc/login.defs` and set `PASS_WARN_AGE` to a value of 7 or more that follows local site policy:

Example:

```
PASS_WARN_AGE 7
```

Run the following command to modify user parameters for all users with a password set to a minimum warning to 7 or more days that follows local site policy:

```
# chage --warndays <N> <user>
```

Example:

```
# awk -F: '($2~/^\$.+\$/ ) {if($6 < 7)system ("chage --warndays 7 " $1)}' /etc/shadow
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

CIS Controls V7.0:

- Control 4: Controlled Use of Administrative Privileges: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 5: Account Management: -- [More](#)
>

[Back to Summary](#)

5.4.1.4 Ensure strong password hashing algorithm is configured

Pass

Description:

A cryptographic hash function converts an arbitrary-length input into a fixed length output. Password hashing performs a one-way transformation of a password, turning the password into another string, called the hashed password.

`ENCRYPT_METHOD` (string) - This defines the system default encryption algorithm for encrypting passwords (if no algorithm are specified on the command line). It can take one of these values:

- MD5 - MD5-based algorithm will be used for encrypting password
- SHA256 - SHA256-based algorithm will be used for encrypting password
- SHA512 - SHA512-based algorithm will be used for encrypting password
- BCrypt - BCrypt-based algorithm will be used for encrypting password
- Yescrypt - Yescrypt-based algorithm will be used for encrypting password
- DES - DES-based algorithm will be used for encrypting password (default)

Note:

- This parameter overrides the deprecated `MD5_CRYPT_ENAB` variable.
- This parameter will only affect the generation of group passwords.
- The generation of user passwords is done by PAM and subject to the PAM configuration.
- It is recommended to set this variable consistently with the PAM configuration.

Rationale:

The `SHA-512` and `yescrypt` algorithms provide a stronger hash than other algorithms used by Linux for password hash generation. A stronger hash provides additional protection to the system by increasing the level of effort needed for an attacker to successfully determine local group passwords.

Remediation:

Edit `/etc/login.defs` and set the `ENCRYPT_METHOD` to `SHA512` or `YESCRIPT` :

```
ENCRYPT_METHOD <HASHING_ALGORITHM>
```

Example:

```
ENCRYPT_METHOD YESCRYPT
```

Note:

- This only effects local groups' passwords created after updating the file to use sha512 or yescrypt .
- If it is determined that the password algorithm being used is not sha512 or yescrypt , once it is changed, it is recommended that all group passwords be updated to use the stronger hashing algorithm.
- It is recommended that the chosen hashing algorithm is consistent across /etc/login.defs and the PAM configuration

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: IA-5

CIS Controls V7.0:

- Control 16: Account Monitoring and Control: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

5.4.1.5 Ensure inactive password lock is configured

Fail

Description:

User accounts that have been inactive for over a given period of time can be automatically disabled.

INACTIVE - Defines the number of days after the password exceeded its maximum age where the user is expected to replace this password.

The value is stored in the shadow password file. An input of 0 will disable an expired password with no delay. An input of -1 will blank the respective field in the shadow password file.

Rationale:

Inactive accounts pose a threat to system security since the users are not logging in to notice failed login attempts or other anomalies.

Remediation:

Run the following command to set the default password inactivity period to 45 days or less that meets local site policy:

```
# useradd -D -f <N>
```

Example:

```
# useradd -D -f 45
```

Run the following command to modify user parameters for all users with a password set to a inactive age of 45 days or less that follows local site policy:

```
# chage --inactive <N> <user>
```

Example:

```
# awk -F: '($2~/^\$.+\$/){if($7 > 45 || $7 < 0)system ("chage --inactive 45 " $1)}' /etc/shadow
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML**References:**

- **URL:** CIS Password Policy Guide

CIS Controls V7.0:

- **Control 4: Controlled Use of Administrative Privileges:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 5: Account Management:** -- [More](#)
>

[Back to Summary](#)**5.4.1.6 Ensure all users last password change date is in the past**

Pass

Description:

All users should have a password change date in the past.

Rationale:

If a user's recorded password change date is in the future, then they could bypass any set password expiration.

Remediation:

Investigate any users with a password change date in the future and correct them. Locking the account, expiring the password, or resetting the password manually may be appropriate.

Assessment:[Show](#) Assessment Evidence[Show](#) Rule Result XML**References:****CIS Controls V7.0:**

- **Control 4: Controlled Use of Administrative Privileges:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 5: Account Management:** -- [More](#)
>

[Back to Summary](#)**5.4.2 Configure root and system accounts and environment****5.4.2.1 Ensure root is the only UID 0 account**

Pass

Description:

Any account with UID 0 has superuser privileges on the system.

Rationale:

This access must be limited to only the default `root` account and only from the system console. Administrative access must be through an unprivileged account using an approved mechanism as noted in the Recommendation "Ensure access to the `su` command is restricted".

Remediation:

Run the following command to change the `root` account UID to 0 :

```
# usermod -u 0 root
```

Modify any users other than `root` with UID 0 and assign them a new UID.

Assessment:[Show](#) Assessment Evidence[Show](#) Rule Result XML**References:**

- URL: NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)**5.4.2.2 Ensure root is the only GID 0 account**

Pass

Description:

The `usermod` command can be used to specify which group the `root` account belongs to. This affects permissions of files that are created by the `root` account.

Rationale:

Using GID 0 for the `root` account helps prevent `root` -owned files from accidentally becoming accessible to non-privileged users.

Remediation:

Run the following command to set the `root` user's GID to 0 :

```
# usermod -g 0 root
```

Run the following command to set the `root` group's GID to 0 :

```
# groupmod -g 0 root
```

Remove any users other than the `root` user with GID 0 or assign them a new GID if appropriate.

Assessment:[Show](#) Assessment Evidence[Show](#) Rule Result XML**References:**

- URL: NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)**5.4.2.3 Ensure group root is the only GID 0 group**

Pass

Description:

The `groupmod` command can be used to specify which group the `root` group belongs to. This affects permissions of files that are group owned by the `root` group.

Rationale:

Using GID 0 for the `root` group helps prevent `root` group owned files from accidentally becoming accessible to non-privileged users.

Remediation:

Run the following command to set the `root` group's GID to 0 :

```
# groupmod -g 0 root
```

Remove any groups other than the `root` group with GID 0 or assign them a new GID if appropriate.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

5.4.2.4 Ensure root account access is controlled

Pass

Description:

There are a number of methods to access the root account directly. Without a password set any user would be able to gain access and thus control over the entire system.

Rationale:

Access to `root` should be secured at all times.

Remediation:

Run the following command to set a password for the `root` user:

```
# passwd root
```

- OR -

Run the following command to lock the `root` user account:

```
# usermod -L root
```

Impact:

If there are any automated processes that relies on access to the root account without authentication, they will fail after remediation.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:**CIS Controls V7.0:**

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

Control 3: Data Protection: -- More

>

Back to Summary

5.4.2.5 Ensure root path integrity

Pass

Description:

The `root` user can execute any command on the system and could be fooled into executing programs unintentionally if the `PATH` is not set correctly.

Rationale:

Including the current working directory (`.`) or other writable directory in `root` 's executable path makes it likely that an attacker can gain superuser access by forcing an administrator operating as `root` to execute a Trojan horse program.

Remediation:

Correct or justify any:

- Locations that are not directories
- Empty directories (`:`)
- Trailing (`:`)
- Current working directory (`.`)
- Non `root` owned directories
- Directories that less restrictive than mode `0755`

Assessment:

Show Assessment Evidence

Show Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5

CIS Controls V7.0:

Not Explicitly Mapped: -- More

>

CIS Critical Security Controls V8.0:

Not Explicitly Mapped: -- More

>

Back to Summary

5.4.2.6 Ensure root user umask is configured

Fail

Description:

The user file-creation mode mask (`umask`) is used to determine the file permission for newly created directories and files. In Linux, the default permissions for any newly created directory is `0777` (`rwXrwXrwX`), and for any newly created file it is `0666` (`rw-rw-rw-`). The `umask` modifies the default Linux permissions by restricting (masking) these permissions. The `umask` is not simply subtracted, but is processed bitwise. Bits set in the `umask` are cleared in the resulting file mode.

`umask` can be set with either `Octal` or `Symbolic` values:

- `Octal` (Numeric) Value - Represented by either three or four digits. ie `umask 0027` or `umask 027` . If a four digit `umask` is used, the first digit is ignored. The remaining three digits effect the resulting permissions for user, group, and world/other respectively.
- `Symbolic` Value - Represented by a comma separated list for User `u` , group `g` , and world/other `o` . The permissions listed are not masked by `umask` . ie a `umask` set by `umask u=rwx,g=rx,o=` is the `Symbolic` equivalent of the `Octal` `umask 027` . This `umask` would set a newly created directory with file mode `drwxr-x---` and a newly created file with file mode `rw-r-----` .

root user Shell Configuration Files:

- /root/.bash_profile - Is executed to configure the root users' shell before the initial command prompt. **Is only read by login shells.**
- /root/.bashrc - Is executed for interactive shells. **only read by a shell that's both interactive and non-login**

umask is set by order of precedence. If umask is set in multiple locations, this order of precedence will determine the system's default umask .

Order of precedence:

1. /root/.bash_profile
2. /root/.bashrc
3. The system default umask

Rationale:

Setting a secure value for umask ensures that users make a conscious choice about their file permissions. A permissive umask value could result in directories or files with excessive permissions that can be read and/or written to by unauthorized users.

Remediation:

Edit /root/.bash_profile and /root/.bashrc and remove, comment out, or update any line with umask to be 0027 or more restrictive.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AC-3, MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

5.4.2.7 Ensure system accounts do not have a valid login shell

Pass

Description:

There are a number of accounts provided with most distributions that are used to manage applications and are not intended to provide an interactive shell. Furthermore, a user may add special accounts that are not intended to provide an interactive shell.

Rationale:

It is important to make sure that accounts that are not being used by regular users are prevented from being used to provide an interactive shell. By default, most distributions set the password field for these accounts to an invalid string, but it is also recommended that the shell field in the password file be set to the nologin shell. This prevents the account from potentially being used to run any commands.

Remediation:

Run the following command to set the shell for any service accounts returned by the audit to nologin :

```
# usermod -s $(command -v nologin) <user>
```

Example script:

```
#!/usr/bin/env bash
```



```
{
l_valid_shells="^${( awk -F\ / '$NF != "nologin" {print}' /etc/shells | sed -rn '/^\\/{s,/,\|\\|/,g;p}' |
paste -s -d '|' - )}$"

awk -v pat="$l_valid_shells" -F: '($1!~/^(root|halt|sync|shutdown|nfsnobody)$/ && ($3<"$(awk '/
^\\s*UID_MIN/{print $2}' /etc/login.defs)" || $3 == 65534) && $(NF) ~ pat) {system ("usermod -s
"$$(command -v nologin)" " $1)}' /etc/passwd
}
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AC-2(5), AC-3, AC-11, MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

5.4.2.8 Ensure accounts without a valid login shell are locked

Pass

Description:

There are a number of accounts provided with most distributions that are used to manage applications and are not intended to provide an interactive shell. Furthermore, a user may add special accounts that are not intended to provide an interactive shell.

Rationale:

It is important to make sure that accounts that are not being used by regular users are prevented from being used to provide an interactive shell. By default, most distributions set the password field for these accounts to an invalid string, but it is also recommended that the shell field in the password file be set to the `nologin` shell. This prevents the account from potentially being used to run any commands.

Remediation:

Run the following command to lock any non-root accounts without a valid login shell returned by the audit:

```
# usermod -L <user>
```

Example script: :

```
#!/usr/bin/env bash

{
l_valid_shells="^${(awk -F\ / '$NF != "nologin" {print}' /etc/shells | sed -rn '/^\\/{s,/,\|\\|/,g;p}' |
paste -s -d '|' - )}$"

while IFS= read -r l_user; do

passwd -S "$l_user" | awk '$2 !~ /^L/ {system ("usermod -L " $1)}'

done <<{(awk -v pat="$l_valid_shells" -F: '($1 != "root" && $(NF) != pat) {print $1}' /etc/passwd)
}
```

Assessment:

[Show](#) Assessment Evidence

[Show Rule Result XML](#)

References:

- URL: NIST SP 800-53 Rev. 5: AC-2(5), AC-3, AC-11, MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#) >

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#) >

[Back to Summary](#)

5.4.3 Configure user default environment

5.4.3.2 Ensure default user shell timeout is configured

Fail

Description:

TMOUT is an environmental setting that determines the timeout of a shell in seconds.

- TMOUT=*n* - Sets the shell timeout to *n* seconds. A setting of TMOUT=0 disables timeout.
- readonly TMOUT- Sets the TMOUT environmental variable as readonly, preventing unwanted modification during run-time.
- export TMOUT - exports the TMOUT variable

System Wide Shell Configuration Files:

- /etc/profile - used to set system wide environmental variables on users shells. The variables are sometimes the same ones that are in the .bash_profile , however this file is used to set an initial PATH or PS1 for all shell users of the system. **is only executed for interactive login shells, or shells executed with the --login parameter.**
- /etc/profile.d - /etc/profile will execute the scripts within /etc/profile.d/*.sh . It is recommended to place your configuration in a shell script within /etc/profile.d to set your own system wide environmental variables.
- /etc/bash.bashrc - System wide version of .bashrc . In Fedora derived distributions, /etc/bashrc also invokes /etc/profile.d/*.sh if *non-login* shell, but redirects output to /dev/null if *non-interactive*. **Is only executed for interactive shells or if BASH_ENV is set to /etc/bash.bashrc .**

Rationale:

Setting a timeout value reduces the window of opportunity for unauthorized user access to another user's shell session that has been left unattended. It also ends the inactive session and releases the resources associated with that session.

Remediation:

Review /etc/bashrc , /etc/profile , and all files ending in *.sh in the /etc/profile.d/ directory and remove or edit all TMOUT=_n_ entries to follow local site policy. TMOUT should not exceed 900 or be equal to 0 .

Configure TMOUT in **one** of the following files:

- A file in the /etc/profile.d/ directory ending in .sh
- /etc/profile
- /etc/bashrc

Example command to set TMOUT to 900 seconds in a file in /etc/profile.d/ :

```
# printf '%s\n' "# Set TMOUT to 900 seconds" "typeset -xr TMOUT=900" > /etc/profile.d/50-tmout.sh
```

TMOUT configuration examples:

```
typeset -xr TMOUT=900
```

Deprecated methods:

- As multiple lines:

```
TMOUT=900
```

```
readonly TMOUT
```

```
export TMOUT
```

- As a single line:

```
readonly TMOUT=900 ; export TMOUT
```

Assessment:

[Show Assessment Evidence](#)

[Show Rule Result XML](#)

References:

CIS Controls V7.0:

- Control 16: Account Monitoring and Control: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

5.4.3.3 Ensure default user umask is configured

Fail

Description:

The user file-creation mode mask (`umask`) is used to determine the file permission for newly created directories and files. In Linux, the default permissions for any newly created directory is 0777 (`rw-rw-rw-`), and for any newly created file it is 0666 (`rw-rw-rw-`). The `umask` modifies the default Linux permissions by restricting (masking) these permissions. The `umask` is not simply subtracted, but is processed bitwise. Bits set in the `umask` are cleared in the resulting file mode.

`umask` can be set with either `Octal` or `Symbolic` values:

- `Octal` (Numeric) Value - Represented by either three or four digits. ie `umask 0027` or `umask 027` . If a four digit `umask` is used, the first digit is ignored. The remaining three digits effect the resulting permissions for user, group, and world/other respectively.
- `Symbolic` Value - Represented by a comma separated list for User `u` , group `g` , and world/other `o` . The permissions listed are not masked by `umask` . ie a `umask` set by `umask u=rwx,g=rx,o=` is the `Symbolic` equivalent of the `Octal` `umask 027` . This `umask` would set a newly created directory with file mode `drwxr-x---` and a newly created file with file mode `rw-r-----` .

The default `umask` can be set to use the `pam_umask` module or in a `System Wide Shell Configuration File` . The user creating the directories or files has the discretion of changing the permissions via the `chmod` command, or choosing a different default `umask` by adding the `umask` command into a `User Shell Configuration File` , (`.bash_profile` or `.bashrc`), in their home directory.

Setting the default umask:

- `pam_umask` module:
 - will set the `umask` according to the system default in `/etc/login.defs` and user settings, solving the problem of different `umask` settings with different shells, display managers, remote sessions etc.
 - `umask=<mask>` value in the `/etc/login.defs` file is interpreted as `Octal`
 - Setting `USERGROUPS_ENAB` to `yes` in `/etc/login.defs` (default):
 - will enable setting of the `umask` group bits to be the same as owner bits. (examples: `022 -> 002`, `077 -> 007`) for non-root users, if the `uid` is the same as `gid` , and `username` is the same as the `<primary group name>`
 - `userdel` will remove the user's group if it contains no more members, and `useradd` will create by default a group with the name of the user
- `System Wide Shell Configuration File` :
 - `/etc/profile` - used to set system wide environmental variables on users shells. The variables are sometimes the same ones that are in the `.bash_profile` , however this file is used to set an initial `PATH` or `PS1` for all shell users of the system. **is only executed for interactive login shells, or shells executed with the `--login` parameter.**
 - `/etc/profile.d` - `/etc/profile` will execute the scripts within `/etc/profile.d/*.sh` . It is recommended to place your configuration in a shell script within `/etc/profile.d` to set your own system wide environmental variables.
 - `/etc/bashrc` - System wide version of `.bashrc` . In Fedora derived distributions, `/etc/bashrc` also invokes `/etc/profile.d/*.sh` if *non-login* shell, but redirects output to `/dev/null` if *non-interactive*. **Is only executed for interactive shells**

or if BASH_ENV is set to /etc/bashrc .

User Shell Configuration Files:

- ~/.bash_profile - Is executed to configure your shell before the initial command prompt. **Is only read by login shells.**
- ~/.bashrc - Is executed for interactive shells. **only read by a shell that's both interactive and non-login**

umask is set by order of precedence. If umask is set in multiple locations, this order of precedence will determine the system's default umask .

Order of precedence:

1. A file in /etc/profile.d/ ending in .sh - This will override any other system-wide umask setting
2. In the file /etc/profile
3. On the pam_umask.so module in /etc/pam.d/postlogin
4. In the file /etc/login.defs
5. In the file /etc/default/login

Rationale:

Setting a secure default value for umask ensures that users make a conscious choice about their file permissions. A permissive umask value could result in directories or files with excessive permissions that can be read and/or written to by unauthorized users.

Remediation:

Run the following script and perform the instructions in the output to set the default umask to 027 or more restrictive:

```
#!/usr/bin/env bash

{
l_output="" l_output2="" l_out=""
file_umask_chk ()
{
if grep -Psiq -- '^\\h*umask\\h+(0?[0-7][2-7]7|u(=[rxw]{0,3}),g(=[rx]{0,2}),o=(\\h*#.*)?\\$' "$l_file"; then
l_out="$l_out\\n - umask is set correctly in \"$l_file\\n\"
elif grep -Psiq -- '^\\h*umask\\h+(((0-7)[0-7][01][0-7]\\b|[0-7][0-7][0-7][0-6]\\b)|([0-7][01][0-7]\\b|[0-7][0-7][0-6]\\b)|(u(=[rxw]{1,3}),)?(((g(=[rx]?[rx]?w[rx]?[rx]?\\b)(,o=[rxw]{1,3})?)|((g(=[rxw]{1,3}),)?o=[rxw]{1,3}\\b))))' "$l_file"; then
l_output2="$l_output2\\n - \"$l_file\\n\"
fi
}

while IFS= read -r -d $'\0' l_file; do
file_umask_chk
done < <(find /etc/profile.d/ -type f -name '*.sh' -print0)

[ -n "$l_out" ] && l_output="$l_out"
l_file="/etc/profile" && file_umask_chk
l_file="/etc/bashrc" && file_umask_chk
l_file="/etc/bash.bashrc" && file_umask_chk
l_file="/etc/pam.d/postlogin"

if grep -Psiq '^\\h*session\\h+([^#\\n\\r]+\\h+pam_umask\\.so\\h+([^#\\n\\r]+\\h+)?umask=(((0-7)[0-7][01][0-7]\\b|[0-7][0-7][0-7][0-6]\\b)|([0-7][01][0-7]\\b)))' "$l_file"; then
l_output2="$l_output2\\n - \"$l_file\\n\"
fi
```

```
l_file="/etc/login.defs" && file_umask_chk

l_file="/etc/default/login" && file_umask_chk

if [ -z "$l_output2" ]; then

echo -e " - No files contain a UMASK that is not restrictive enough\n No UMASK updates required to
existing files"

else

echo -e "\n - UMASK is not restrictive enough in the following file(s):$l_output2\n\n- Remediation
Procedure:\n - Update these files and comment out the UMASK line\n or update umask to be \"0027\" or more
restrictive"

fi

if [ -n "$l_output" ]; then

echo -e "$l_output"

else

echo -e " - Configure UMASK in a file in the \"/etc/profile.d/" directory ending in \".sh\"\n\n Example
Command (Hash to represent being run at a root prompt):\n\n# printf '%s\\n' \"umask 027\" > /etc/
profile.d/50-systemwide_umask.sh\n"

fi

}
```

Notes:

- This method only applies to bash and shell. If other shells are supported on the system, it is recommended that their configuration files also are checked
- If the `pam_umask .so` module is going to be used to set `umask`, ensure that it's not being overridden by another setting. Refer to the `PAM_UMASK(8)` man page for more information

Assessment:[Show](#) Assessment Evidence[Show](#) Rule Result XML**References:**

- **URL:** NIST SP 800-53 Rev. 5: AC-3, MP-2

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

6 Logging and Auditing

The items in this section describe how to configure logging, log monitoring, and auditing, using tools included in most distributions.

It is recommended that `rsyslog` be used for logging (with `logwatch` providing summarization) and `auditd` be used for auditing (with `aureport` providing summarization) to automatically monitor logs for intrusion attempts and other suspicious system behavior.

In addition to the local log files created by the steps in this section, it is also recommended that sites collect copies of their system logs on a secure, centralized log server via an encrypted connection. Not only does centralized logging help sites correlate events that may be occurring on multiple systems, but having a second copy of the system log information may be critical after a system compromise where the attacker has modified the local log files on the affected system(s). If a log correlation system is deployed, configure it to process the logs described in this section.

Because it is often necessary to correlate log information from many different systems (particularly after a security incident) it is recommended that the time be synchronized among systems and devices connected to the local network. The standard Internet protocol for time synchronization is the Network Time Protocol (NTP), which is supported by most

network-ready devices. Reference < <http://chrony.tuxfamily.org/> > manual page for more information on configuring chrony.

It is important that all logs described in this section be monitored on a regular basis and correlated to determine trends. A seemingly innocuous entry in one log could be more significant when compared to an entry in another log.

Note on log file permissions: There really isn't a "one size fits all" solution to the permissions on log files. Many sites utilize group permissions so that administrators who are in a defined security group, such as "wheel" do not have to elevate privileges to root in order to read log files. Also, if a third party log aggregation tool is used, it may need to have group permissions to read the log files, which is preferable to having it run setuid to root. Therefore, there are two remediation and audit steps for log file permissions. One is for systems that do not have a secured group method implemented that only permits root to read the log files (`root:root 600`). The other is for sites that do have such a setup and are designated as `root:securegrp 640` where `securegrp` is the defined security group (in some cases `wheel`).

6.1 Configure Integrity Checking

AIDE is a file integrity checking tool, similar in nature to Tripwire. While it cannot prevent intrusions, it can detect unauthorized changes to configuration files by alerting when the files are changed. When setting up AIDE, decide internally what the site policy will be concerning integrity checking. Review the AIDE quick start guide and AIDE documentation before proceeding.

6.1.1 Ensure AIDE is installed

Fail

Description:

AIDE takes a snapshot of filesystem state including modification times, permissions, and file hashes which can then be used to compare against the current state of the filesystem to detect modifications to the system.

Rationale:

By monitoring the filesystem state compromised files can be detected to prevent or limit the exposure of accidental or malicious misconfigurations or modified binaries.

Remediation:

Install AIDE using the appropriate package manager or manual installation:

```
# apt install aide aide-common
```

Configure AIDE as appropriate for your environment. Consult the AIDE documentation for options.

Run the following commands to initialize AIDE:

```
# aideinit
# mv /var/lib/aide/aide.db.new /var/lib/aide/aide.db
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: AU-2
- **URL:** STIG ID: UBTU-20-010450 | Rule ID: SV-238371r958944 | CAT II
- **URL:** STIG ID: UBTU-22-651010 | Rule ID: SV-260582r958944 | CAT II

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

Back to Summary

6.1.2 Ensure filesystem integrity is regularly checked

Fail

Description:

Periodic checking of the filesystem integrity is needed to detect changes to the filesystem.

Rationale:

Periodic file checking allows the system administrator to determine on a regular basis if critical files have been changed in an unauthorized fashion.

Remediation:

Run the following command to unmask `dailyaidecheck.timer` and `dailyaidecheck.service` :

```
# systemctl unmask dailyaidecheck.timer dailyaidecheck.service
```

Run the following command to enable and start `dailyaidecheck.timer` :

```
# systemctl --now enable dailyaidecheck.timer
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** <https://github.com/konstruktoid/hardening/blob/master/config/aidecheck.service>
- **URL:** <https://github.com/konstruktoid/hardening/blob/master/config/aidecheck.timer>
- **URL:** NIST SP 800-53 Rev. 5: AU-2

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

6.1.3 Ensure cryptographic mechanisms are used to protect the integrity of audit tools

Fail

Description:

Audit tools include, but are not limited to, vendor-provided and open source audit tools needed to successfully view and manipulate audit information system activity and records. Audit tools include custom queries and report generators.

Rationale:

Protecting the integrity of the tools used for auditing purposes is a critical step toward ensuring the integrity of audit information. Audit information includes all information (e.g., audit records, audit settings, and audit reports) needed to successfully audit information system activity.

Attackers may replace the audit tools or inject code into the existing tools with the purpose of providing the capability to hide or erase system activity from the audit logs.

Audit tools should be cryptographically signed in order to provide the capability to identify when the audit tools have been modified, manipulated, or replaced. An example is a checksum hash of the file or files.

Remediation:

Add or update the following selection lines for to a file ending in `.conf` in the `/etc/aide/aide.conf.d/` or to `/etc/aide/aide.conf` to protect the integrity of the audit tools:

```
# Audit Tools

/sbin/auditctl p+i+n+u+g+s+b+acl+xattr+sha512

/sbin/auditd p+i+n+u+g+s+b+acl+xattr+sha512

/sbin/ausearch p+i+n+u+g+s+b+acl+xattr+sha512

/sbin/aureport p+i+n+u+g+s+b+acl+xattr+sha512
```

```
/sbin/autrace p+i+n+u+g+s+b+acl+xattrs+sha512
/sbin/augenrules p+i+n+u+g+s+b+acl+xattrs+sha512
```

Assessment:[Show Assessment Evidence](#)[Show Rule Result XML](#)**References:**

- **URL:** NIST SP 800-53 Rev. 5: AU-3

[Back to Summary](#)

6.2 System Logging

Logging services should be configured to prevent information leaks and to aggregate logs on a remote server so that they can be reviewed in the event of a system compromise. A centralized log server provides a single point of entry for further analysis, monitoring and filtering.

Security principals for logging

- Ensure transport layer security is implemented between the client and the log server.
- Ensure that logs are rotated as per the environment requirements.
- Ensure all locally generated logs have the appropriate permissions.
- Ensure all security logs are sent to a remote log server.
- Ensure the required events are logged.

What is covered

This section will cover the minimum best practices for the usage of **either** `rsyslog` **OR** `journald`. The recommendations are written such that each is wholly independent of each other and **only one is implemented**.

- If your organization makes use of an enterprise wide logging system completely outside of `rsyslog` or `journald`, then the following recommendations do not directly apply. However, the principals of the recommendations should be followed regardless of what solution is implemented. If the enterprise solution incorporates either of these tools, careful consideration should be given to the following recommendations to determine exactly what applies.
- Should your organization make use of both `rsyslog` and `journald`, take care how the recommendations may or may not apply to you.

What is not covered

- Enterprise logging systems not utilizing `rsyslog` or `journald`. As logging is very situational and dependent on the local environment, not everything can be covered here.
- Transport layer security should be applied to all remote logging functionality. Both `rsyslog` and `journald` supports secure transport and should be configured as such.
- The log server. There are a multitude of reasons for a centralized log server (and keeping a short period logging on the local system), but the log server is out of scope for these recommendations.

6.2.1 Configure systemd-journald service

`systemd-journald` is a system service that collects and stores logging data. It creates and maintains structured, indexed journals based on logging information that is received from a variety of sources:

- Kernel log messages, via `kmsg`
- Simple system log messages, via the `libc syslog` call
- Structured system log messages via the native Journal API
- Standard output and standard error of service units
- Audit records, originating from the kernel audit subsystem

The daemon will implicitly collect numerous metadata fields for each log messages in a secure and unfakeable way. See `systemd.journal-fields` man page for more information about the collected metadata.

The journal service stores log data either persistently below `/var/log/journal` or in a volatile way below `/run/log/journal/`. By default, log data is stored persistently if `/var/log/journal/` exists during boot, with an implicit fallback to volatile storage. Use `Storage=` in `journald.conf` to configure where log data is placed, independently of the existence of `/var/log/journal/`.

On systems where `/var/log/journal/` does not exist but where persistent logging is desired, and the default `journald.conf` is used, it is sufficient to create the directory and ensure it has the correct access modes and ownership.

Note:systemd-journald.service must be configured appropriately for either journald- **OR** -rsyslog to operate effectively.

6.2.1.1 Ensure journald service is enabled and active

Pass

Description:

Ensure that the systemd-journald service is enabled to allow capturing of logging events.

Rationale:

If the systemd-journald service is not enabled to start on boot, the system will not capture logging events.

Remediation:

Run the following commands to unmask and start systemd-journald.service

```
# systemctl unmask systemd-journald.service
# systemctl start systemd-journald.service
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AU-2, AU-7 AU-12

CIS Controls V7.0:

- Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
- Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
- >

CIS Critical Security Controls V8.0:

- Control 8: Audit Log Management: -- [More](#)
- >

[Back to Summary](#)

6.2.1.2 Ensure journald log file access is configured

Manual

Description:

Journald will create logfiles that do not already exist on the system. This setting controls what permissions will be applied to these newly created files.

Rationale:

It is important to ensure that log files have the correct permissions to ensure that sensitive data is archived and protected.

Remediation:

If the default configuration is not appropriate for the site specific requirements, copy /usr/lib/tmpfiles.d/systemd.conf to /etc/tmpfiles.d/systemd.conf and modify as required. Recommended mode for logfiles is 0640 or more restrictive.

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AC-3, AU-2, AU-12, MP-2, SI-5

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
- >

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

6.2.1.3 Ensure journald log file rotation is configured

Manual

Description:

Journald includes the capability of rotating log files regularly to avoid filling up the system with logs or making the logs unmanageably large. The file `/etc/systemd/journald.conf` is the configuration file used to specify how logs generated by Journald should be rotated.

Rationale:

By keeping the log files smaller and more manageable, a system administrator can easily archive these files to another system and spend less time looking through inordinately large log files.

Remediation:

Edit `/etc/systemd/journald.conf` or a file ending in `.conf` the `/etc/systemd/journald.conf.d/` directory. Set the following parameters in the `[Journal]` section to ensure logs are rotated according to site policy. The settings should be carefully understood as there are specific edge cases and prioritization of parameters.

Example Configuration :

```
[Journal]

SystemMaxUse=1G

SystemKeepFree=500M

RuntimeMaxUse=200M

RuntimeKeepFree=50M

MaxFileSec=1month
```

Example script to create systemd drop-in configuration file:

```
{

a_settings=("SystemMaxUse=1G" "SystemKeepFree=500M" "RuntimeMaxUse=200M" "RuntimeKeepFree=50M"
"MaxFileSec=1month")

[ ! -d /etc/systemd/journald.conf.d/ ] && mkdir /etc/systemd/journald.conf.d/

if grep -Psq -- '^h*\[Journal\]' /etc/systemd/journald.conf.d/60-journald.conf; then
printf '%s\n' "" "${a_settings[@]}" >> /etc/systemd/journald.conf.d/60-journald.conf
else
printf '%s\n' "" "[Journal]" "${a_settings[@]}" >> /etc/systemd/journald.conf.d/60-journald.conf
fi
}
```

Note:

- If these settings appear in a canonically later file, or later in the same file, the setting will be overwritten
- Logfile size and configuration to move logfiles to a remote log server should be accounted for when configuring these settings

Run to following command to update the parameters in the service:

```
# systemctl reload-or-restart systemd-journald
```

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AU-2, AU-7, AU-12

CIS Controls V7.0:

- Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
 - Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
- >

CIS Critical Security Controls V8.0:

- Control 8: Audit Log Management: -- [More](#)
- >

[Back to Summary](#)

6.2.2 Configure journald

Included in the systemd suite is a journaling service called systemd-journald.service for the collection and storage of logging data. It creates and maintains structured, indexed journals based on logging information that is received from a variety of sources such as:

Classic RFC3164 BSD syslog via the /dev/log socket
STDOUT/STDERR of programs via StandardOutput=journal + StandardError=journal in service files (both of which are default settings)
Kernel log messages via the /dev/kmsg device
node Audit records via the kernel's audit subsystem
Structured log messages via journald's native protocol
Any changes made to the systemd-journald configuration will require a re-start of systemd-journald

Note:

- IF `rsyslog` will be used for remote logging on the system this subsection can be skipped

6.2.2.1 Configure systemd-journal-remote

The `systemd-journal-remote` package includes `systemd-journal-upload`.

`systemd-journal-upload` will upload journal entries to the URL specified with `--url=`. This program reads journal entries from one or more journal files, similarly to `journalctl`.

`systemd-journal-upload` transfers the raw content of journal file and uses HTTP as a transport protocol.

`systemd-journal-upload.service` is a system service that uses `systemd-journal-upload` to upload journal entries to a server. It uses the configuration in `journal-upload.conf`.

Note:

- IF `rsyslog` is in use this subsection can be skipped.
- `systemd-journal-remote` package is part of the universe component, this may impact support and update frequency which should be considered when assessing organizational risk.

6.2.2.1.1 Ensure systemd-journal-remote is installed

Pass

Description:

Journald `systemd-journal-remote` supports the ability to send log events it gathers to a remote log host or to receive messages from remote hosts, thus enabling centralized log management.

Rationale:

Storing log data on a remote host protects log integrity from local attacks. If an attacker gains root access on the local system, they could tamper with or remove log data that is stored on the local system.

Note: This recommendation **only applies if** `journald` **is the chosen method for client side logging** . Do not apply this recommendation if `rsyslog` is used.

Remediation:

Run the following command to install `systemd-journal-remote` :

```
# apt install systemd-journal-remote
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AU-2, AU-7 AU-12

CIS Controls V7.0:

- Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
 - Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
- >

CIS Critical Security Controls V8.0:

- Control 8: Audit Log Management: -- [More](#)
- >

[Back to Summary](#)

6.2.2.1.2 Ensure systemd-journal-upload authentication is configured

Manual

Description:

Journald `systemd-journal-upload` supports the ability to send log events it gathers to a remote log host.

Rationale:

Storing log data on a remote host protects log integrity from local attacks. If an attacker gains root access on the local system, they could tamper with or remove log data that is stored on the local system.

Note: This recommendation **only applies if** `journald` **is the chosen method for client side logging** . Do not apply this recommendation if `rsyslog` is used.

Remediation:

Edit the `/etc/systemd/journal-upload.conf` file or a file in `/etc/systemd/journal-upload.conf.d` ending in `.conf` and ensure the following lines are set in the `[Upload]` section per your environment:

Example settings:

```
[Upload]
URL=192.168.50.42
ServerKeyFile=/etc/ssl/private/journal-upload.pem
ServerCertificateFile=/etc/ssl/certs/journal-upload.pem
TrustedCertificateFile=/etc/ssl/ca/trusted.pem
```

Run the following command to update the parameters in the service:

```
# systemctl reload-or-restart systemd-journal-upload
```

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AU-2, AU-12

CIS Controls V7.0:

- Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
 - Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
- >

CIS Critical Security Controls V8.0:

- Control 8: Audit Log Management: -- [More](#)
- >

[Back to Summary](#)

6.2.2.1.3 Ensure systemd-journal-upload is enabled and active

Pass

Description:

Journald `systemd-journal-upload` supports the ability to send log events it gathers to a remote log host.

Rationale:

Storing log data on a remote host protects log integrity from local attacks. If an attacker gains root access on the local system, they could tamper with or remove log data that is stored on the local system.

Note: This recommendation **only applies if `journald` is the chosen method for client side logging** . Do not apply this recommendation if `rsyslog` is used.

Remediation:

Run the following commands to unmask, enable and start `systemd-journal-upload` :

```
# systemctl unmask systemd-journal-upload.service
# systemctl --now enable systemd-journal-upload.service
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AU-2, AU-12

CIS Controls V7.0:

- Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
 - Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
- >

CIS Critical Security Controls V8.0:

- Control 8: Audit Log Management: -- [More](#)
- >

[Back to Summary](#)

6.2.2.1.4 Ensure `systemd-journal-remote` service is not in use

Pass

Description:

Journald `systemd-journal-remote` supports the ability to receive messages from remote hosts, thus acting as a log server. Clients should not receive data from other hosts.

Note:

- The same package, `systemd-journal-remote` , is used for both sending logs to remote hosts and receiving incoming logs.
- With regards to receiving logs, there are two services; `systemd-journal-remote.socket` and `systemd-journal-remote.service` .

Rationale:

If a client is configured to also receive data, thus turning it into a server, the client system is acting outside it's operational boundary.

Note: This recommendation **only applies if `journald` is the chosen method for client side logging** . Do not apply this recommendation if `rsyslog` is used.

Remediation:

Run the following commands to stop and mask `systemd-journal-remote.socket` and `systemd-journal-remote.service`:

```
# systemctl stop systemd-journal-remote.socket systemd-journal-remote.service
# systemctl mask systemd-journal-remote.socket systemd-journal-remote.service
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AU-2, AU-7 AU-12

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
>

[Back to Summary](#)

6.2.2.2 Ensure journald ForwardToSyslog is disabled

Pass

Description:

Data from `journald` should be kept in the confines of the service and not forwarded to other services.

Rationale:

- IF `-journald` is the method for capturing logs, all logs of the system should be handled by `journald` and not forwarded to other logging mechanisms.

Note: This recommendation **only applies if** `journald` is the chosen method for client side logging . Do not apply this recommendation if `rsyslog` is used.

Remediation:

- IF `-rsyslog` is the preferred method for capturing logs, this section and Recommendation should be skipped and the "Configure rsyslog" section followed.

- IF `-journald` is the preferred method for capturing logs:

Set the following parameter in the `[Journal]` section in `/etc/systemd/journald.conf` or a file in `/etc/systemd/journald.conf.d/` ending in `.conf` :

```
ForwardToSyslog=no
```

Note: If this setting appears in a canonically later file, or later in the same file, the setting will be overwritten

Run to following command to update the parameters in the service:

```
# systemctl reload-or-restart systemd-journald
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AU-2, AU-6, AU-7, AU-12

CIS Controls V7.0:

- Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 8: Audit Log Management: -- [More](#)
>

[Back to Summary](#)

6.2.2.3 Ensure journald Compress is configured

Pass

Description:

The journald system includes the capability of compressing overly large files to avoid filling up the system with logs or making the logs unmanageably large.

Rationale:

Uncompressed large files may unexpectedly fill a filesystem leading to resource unavailability. Compressing logs prior to write can prevent sudden, unexpected filesystem impacts.

Note: This recommendation **only applies if** `journald` **is the chosen method for client side logging** . Do not apply this recommendation if `rsyslog` is used.

Remediation:

- **IF** `rsyslog` is the preferred method for capturing logs, this section and Recommendation should be skipped and the "Configure rsyslog" section followed.
- **IF** `journald` is the preferred method for capturing logs:

Set the following parameter in the `[Journal]` section in `/etc/systemd/journald.conf` or a file in `/etc/systemd/journald.conf.d/` ending in `.conf` :

```
Compress=yes
```

Note: If this setting appears in a canonically later file, or later in the same file, the setting will be overwritten

Run to following command to update the parameters in the service:

```
# systemctl reload-or-restart systemd-journald
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: AU-4

CIS Controls V7.0:

- Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
- Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
- Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)

CIS Critical Security Controls V8.0:

- Control 8: Audit Log Management: -- [More](#)
- Control 8: Audit Log Management: -- [More](#)

[Back to Summary](#)

6.2.2.4 Ensure journald Storage is configured

Pass

Description:

Data from journald may be stored in volatile memory or persisted locally on the server. Logs in memory will be lost upon a system reboot. By persisting logs to local disk on the server they are protected from loss due to a reboot.

Rationale:

Writing log data to disk will provide the ability to forensically reconstruct events which may have impacted the operations or security of a system even after a system crash or reboot.

Note: This recommendation **only applies if** `journald` **is the chosen method for client side logging** . Do not apply this recommendation if `rsyslog` is used.

Remediation:

- IF `-rsyslog` is the preferred method for capturing logs, this section and Recommendation should be skipped and the "Configure rsyslog" section followed.

- IF `-journal` is the preferred method for capturing logs:

Set the following parameter in the `[Journal]` section in `/etc/systemd/journald.conf` or a file in `/etc/systemd/journald.conf.d/` ending in `.conf` :

```
Storage=persistent
```

Note: If this setting appears in a canonically later file, or later in the same file, the setting will be overwritten

Run to following command to update the parameters in the service:

```
# systemctl reload-or-restart systemd-journald
```

Assessment:
[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AU-3, AU-12

CIS Controls V7.0:

- Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
 - Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
- >

CIS Critical Security Controls V8.0:

- Control 8: Audit Log Management: -- [More](#)
- >

[Back to Summary](#)

6.2.3 Configure rsyslog

The `rsyslog` software package may be used instead of the default `journald` logging mechanism.

Rsyslog has evolved over several decades. For this reason it supports three different configuration formats ("languages"):

- `basic` - previously known as the `sysklogd` format, this is the format best used to express basic things, such as where the statement fits on a single line.
 - It stems back to the original `syslog.conf` format, in use now for several decades.
 - The most common use case is matching on facility/severity and writing matching messages to a log file.
- `advanced` - previously known as the `RainerScript` format, this format was first available in `rsyslog v6` and is the current, best and most precise format for non-trivial use cases where more than one line is needed.
 - Prior to `v7`, there was a performance impact when using this format that encouraged use of the basic format for best results. Current versions of `rsyslog` do not suffer from this (historical) performance impact.
 - This new style format is specifically targeted towards more advanced use cases like forwarding to remote hosts that might be partially offline.
- `obsolete legacy` - previously known simply as the `legacy` format, this format is exactly what its name implies: it is obsolete and should not be used when writing new configurations. It was created in the early days (up to `rsyslog version 5`) where we expected that `rsyslog` would extend `sysklogd` just mildly. Consequently, it was primarily aimed at small additions to the original `sysklogd` format.
 - Practice has shown that it was notoriously hard to use for more advanced use cases, and thus we replaced it with the advanced format.
 - In essence, everything that needs to be written on a single line that starts with a dollar sign is legacy format. Users of this format are encouraged to migrate to the basic or advanced formats.

Note: This section only applies if `rsyslog` is the chosen method for client side logging. Do not apply this section if `journald` is used.

6.2.3.1 Ensure rsyslog service is enabled and active

Pass

Description:

Once the `rsyslog` package is installed, ensure that the service is enabled.

Rationale:

If the `rsyslog` service is not enabled to start on boot, the system will not capture logging events.

Remediation:

- IF `-rsyslog` is being used for logging on the system:

Run the following commands to unmask, enable, and start `rsyslog.service` :

```
# systemctl unmask rsyslog.service
# systemctl enable rsyslog.service
# systemctl start rsyslog.service
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AU-2, AU-3, AU-12

CIS Controls V7.0:

- Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
- Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)

>

CIS Critical Security Controls V8.0:

- Control 8: Audit Log Management: -- [More](#)

>

[Back to Summary](#)

6.2.3.2 Ensure rsyslog is installed

Pass

Description:

The `rsyslog` software is recommended in environments where `journald` does not meet operation requirements.

Rationale:

The security enhancements of `rsyslog` such as connection-oriented (i.e. TCP) transmission of logs, the option to log to database formats, and the encryption of log data en route to a central logging server) justify installing and configuring the package.

Remediation:

Run the following command to install `rsyslog` :

```
# apt install rsyslog
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AU-2, AU-3, AU-12
- URL: Ubuntu 20.04 STIG Vuln ID: V-238353 Rule ID: SV-238353r991562 STIG ID: UBTU-20-010432
- URL: Ubuntu 22.04 STIG Vuln ID: V-260588 Rule ID: SV-260588r991562 STIG ID: UBTU-22-652010

CIS Controls V7.0:

- Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
- Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)

>

CIS Critical Security Controls V8.0:

- Control 8: Audit Log Management: -- [More](#)

>

6.2.3.3 Ensure journald is configured to send logs to rsyslog

Fail

Description:

Data from `systemd-journald` may be stored in volatile memory or persisted locally on the server. Utilities exist to accept remote export of `systemd-journald` logs, however, use of the `rsyslog` service provides a consistent means of log collection and export.

Rationale:

- IF `rsyslog` is the preferred method for capturing logs, all logs of the system should be sent to it for further processing.

Note: This recommendation **only applies if `rsyslog` is the chosen method for client side logging** . Do not apply this recommendation if `systemd-journald` is used.

Remediation:

- IF `Journald` is the preferred method for capturing logs, this section and Recommendation should be skipped and the "Configure Journald" section followed.

- IF `rsyslog` is the preferred method for capturing logs:

Set the following parameter in the `[Journal]` section in `/etc/systemd/journald.conf` or a file in `/etc/systemd/journald.conf.d/` ending in `.conf` :

```
ForwardToSyslog=yes
```

Note: If this setting appears in a canonically later file, or later in the same file, the setting will be overwritten

Run to following command to update the parameters in the service:

Restart `systemd-journald.service` :

```
# systemctl reload-or-restart systemd-journald.service
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AC-3, AU-2, AU-4, AU-12, MP-2
- URL: SYSTEMD-JOURNALD.SERVICE(8)
- URL: JOURNALD.CONF(5)

CIS Controls V7.0:

- Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
 - Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
 - Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
- >

CIS Critical Security Controls V8.0:

- Control 8: Audit Log Management: -- [More](#)
 - Control 8: Audit Log Management: -- [More](#)
- >

6.2.3.4 Ensure rsyslog log file creation mode is configured

Pass

Description:

The `$FileCreateMode` parameter allows to specify the creation mode with which `rsyslogd` creates new files.

Rationale:

It is important to ensure that log files have the correct permissions to ensure that sensitive data is archived and protected.

Remediation:

Edit either `/etc/rsyslog.conf` or a dedicated `.conf` file in `/etc/rsyslog.d/` and set `$FileCreateMode` to `0640` or more restrictive:

```
$FileCreateMode 0640
```

Reload the service:

```
# systemctl reload-or-restart rsyslog
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** RSYSLOG.CONF(5)
- **URL:** NIST SP 800-53 Rev. 5: AC-3, AC-6, MP-2
- **URL:** https://www.rsyslog.com/doc/configuration/action/rsconf1_filecreatemode.html

CIS Controls V7.0:

- **Control 5: Secure Configuration for Hardware and Software on Mobile Devices, Laptops, Workstations and Servers:** -- [More](#)
 - **Control 6: Maintenance, Monitoring and Analysis of Audit Logs:** -- [More](#)
 - **Control 6: Maintenance, Monitoring and Analysis of Audit Logs:** -- [More](#)
- >

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
 - **Control 8: Audit Log Management:** -- [More](#)
- >

[Back to Summary](#)

6.2.3.5 Ensure rsyslog logging is configured

Manual

Description:

The `rsyslog` and configuration files specifies rules for logging and which files are to be used to log certain classes of messages.

Rationale:

A great deal of important security-related information is sent via `rsyslog` (e.g., successful and failed su attempts, failed login attempts, root login attempts, etc.).

Remediation:

Edit the following lines in the configuration file(s) returned by the audit as appropriate for your environment.

Note: The below configuration is shown for example purposes only. Due care should be given to how the organization wishes to store log data.

```
*.emerg :omusrmsg:*

auth,authpriv.* /var/log/secure

mail.* -/var/log/mail

mail.info -/var/log/mail.info

mail.warning -/var/log/mail.warn

mail.err /var/log/mail.err

cron.* /var/log/cron

*.=warning;*.=err -/var/log/warn

*.crit /var/log/warn
```

```
*.*;mail.none;news.none -/var/log/messages
local0,local1.* -/var/log/localmessages
local2,local3.* -/var/log/localmessages
local4,local5.* -/var/log/localmessages
local6,local7.* -/var/log/localmessages
```

Run the following command to reload the `rsyslogd` configuration:

```
# systemctl reload-or-restart rsyslog
```

[Show Rule Result XML](#)

References:

- **URL:** See the `rsyslog.conf(5)` man page for more information.
- **URL:** NIST SP 800-53 Rev. 5: AU-2, AU-7, AU-12

CIS Controls V7.0:

- **Control 6: Maintenance, Monitoring and Analysis of Audit Logs:** -- [More](#)
 - **Control 6: Maintenance, Monitoring and Analysis of Audit Logs:** -- [More](#)
- >

CIS Critical Security Controls V8.0:

- **Control 8: Audit Log Management:** -- [More](#)
- >

[Back to Summary](#)

6.2.3.6 Ensure rsyslog is configured to send logs to a remote log host

Manual

Description:

`rsyslog` supports the ability to send log events it gathers to a remote log host or to receive messages from remote hosts, thus enabling centralized log management.

Rationale:

Storing log data on a remote host protects log integrity from local attacks. If an attacker gains root access on the local system, they could tamper with or remove log data that is stored on the local system.

Remediation:

Edit the `rsyslog` configuration and add the following line (where `loghost.example.com` is the name of your central log host). The `target` directive may either be a fully qualified domain name or an IP address.

```
*.* action(type="omfwd" target="192.168.2.100" port="514" protocol="tcp"
action.resumeRetryCount="100"
queue.type="LinkedList" queue.size="1000")
```

Run the following command to reload `rsyslog.service`:

```
# systemctl reload-or-restart rsyslog.service
```

[Show Rule Result XML](#)

References:

- **URL:** See the `rsyslog.conf(5)` man page for more information.
- **URL:** NIST SP 800-53 Rev. 5: AU-6
- **URL:** <https://www.rsyslog.com/doc/>

CIS Controls V7.0:

- **Control 6: Maintenance, Monitoring and Analysis of Audit Logs:** -- [More](#)
- **Control 6: Maintenance, Monitoring and Analysis of Audit Logs:** -- [More](#)

>

CIS Critical Security Controls V8.0:

- Control 8: Audit Log Management: -- [More](#)
- >

[Back to Summary](#)**6.2.3.7 Ensure rsyslog is not configured to receive logs from a remote client**

Pass

Description:

`rsyslog` supports the ability to receive messages from remote hosts, thus acting as a log server. Clients should not receive data from other hosts.

Rationale:

If a client is configured to also receive data, thus turning it into a server, the client system is acting outside its operational boundary.

Remediation:

Unless the system's primary function is to serve as a logfile server , modify the files returned by the Audit Procedure and remove the specific lines highlighted by the audit. Verify none of the following entries are present in the `rsyslog` configuration.

advanced format

```
# module(load="imtcp")
# input(type="imtcp" port="514")
# module(load="imudp")
# input(type="imudp" port="514")
```

Note:`obsolete legacy` - previously known as the legacy format. This format is obsolete and should not be used when writing new configurations. It was aimed at small additions to the original `sysklogd` format and has been replaced due to its limitations.

Reload the service:

```
# systemctl reload-or-restart rsyslog
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AU-2, AU-7, AU-12, CM-6
- URL: <https://www.rsyslog.com/doc/index.html>
- URL: https://www.rsyslog.com/doc/configuration/conf_formats.html

CIS Controls V7.0:

- Control 9: Limitation and Control of Network Ports, Protocols, and Services: -- [More](#)
- >

CIS Critical Security Controls V8.0:

- Control 4: Secure Configuration of Enterprise Assets and Software: -- [More](#)
- >

[Back to Summary](#)**6.2.3.8 Ensure logrotate is configured**

Manual

Description:

The system includes the capability of rotating log files regularly to avoid filling up the system with logs or making the

logs unmanageably large. The file `/etc/logrotate.d/rsyslog` is the configuration file used to rotate log files created by `rsyslog`.

Rationale:

By keeping the log files smaller and more manageable, a system administrator can easily archive these files to another system and spend less time looking through inordinately large log files.

Remediation:

Edit `/etc/logrotate.conf`, or the appropriate configuration file provided by the script in the Audit Procedure, as necessary to ensure logs are rotated according to site policy.

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AU-8

CIS Controls V7.0:

- Control 6: Maintenance, Monitoring and Analysis of Audit Logs: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 8: Audit Log Management: -- [More](#)
>

[Back to Summary](#)

6.2.4 Configure Logfiles

6.2.4.1 Ensure access to all logfiles has been configured

Fail

Description:

Log files contain information from many services on the the local system, or in the event of a centralized log server, others systems logs as well.

In general log files are found in `/var/log/`, although application can be configured to store logs elsewhere. Should your application store logs in another, ensure to run the same test on that location.

Rationale:

It is important that log files have the correct permissions to ensure that sensitive data is protected and that only the appropriate users / groups have access to them.

Remediation:

Run the following script to update permissions and ownership on files in `/var/log`.

Although the script is not destructive, ensure that the output of the audit procedure is captured in the event that the remediation causes issues.

```
#!/usr/bin/env bash

{
a_output2=()
f_file_test_fix()
{
a_out2=()
maxperm="$( printf '%o' $(( 0777 & ~$perm_mask)) )"
if [ $(( $1_mode & $perm_mask )) -gt 0 ]; then
```

```
a_out2+=(" o Mode: \"${l_mode}\" should be \"${maxperm}\" or more restrictive" " x Removing excess
permissions")

chmod "${l_rperms}" "${l_fname}"

fi

if [[ ! "${l_user}" =~ ${l_auser} ]]; then

a_out2+=(" o Owned by: \"${l_user}\" and should be owned by \"${l_auser} or \" " x Changing ownership
to: \"${l_fix_account}\")

chown "${l_fix_account}" "${l_fname}"

fi

if [[ ! "${l_group}" =~ ${l_agroup} ]]; then

a_out2+=(" o Group owned by: \"${l_group}\" and should be group owned by \"${l_agroup} or \" " x
Changing group ownership to: \"${l_fix_account}\")

chgrp "${l_fix_account}" "${l_fname}"

fi

[ "${a_out2[@]}" -gt 0 ] && a_output2+(" - File: \"${l_fname}\" is: \"${a_out2[@]}")

}

l_fix_account='root'

while IFS= read -r -d $'\0' l_file; do

while IFS=: read -r l_fname l_mode l_user l_group; do

if grep -Pq -- '\/(apt)\h*$' <<< "$(dirname "${l_fname}")"; then

perm_mask='0133' l_rperms="u-x,go-wx" l_auser="root" l_agroup="(root|adm)"; f_file_test_fix

else

case "$(basename "${l_fname}")" in

lastlog | lastlog.* | wtmp | wtmp.* | wtmp-* | btmp | btmp.* | btmp-* | README)

perm_mask='0113' l_rperms="ug-x,o-wx" l_auser="root" l_agroup="(root|utmp) "

f_file_test_fix ;;

cloud-init.log* | localmessages* | waagent.log*)

perm_mask='0133' l_rperms="u-x,go-wx" l_auser="(root|syslog) " l_agroup="(root|adm) "

file_test_fix ;;

secure | auth.log | syslog | messages)

perm_mask='0137' l_rperms="u-x,g-wx,o-rwx" l_auser="(root|syslog) " l_agroup="(root|adm) "

f_file_test_fix ;;

SSSD | sssd)

perm_mask='0117' l_rperms="ug-x,o-rwx" l_auser="(root|SSSD) " l_agroup="(root|SSSD) "

f_file_test_fix ;;

gdm | gdm3)

perm_mask='0117' l_rperms="ug-x,o-rwx" l_auser="root" l_agroup="(root|gdm|gdm3) "

f_file_test_fix ;;

*.journal | *.journal~)

perm_mask='0137' l_rperms="u-x,g-wx,o-rwx" l_auser="root" l_agroup="(root|systemd-journal) "

f_file_test_fix ;;

*)

perm_mask='0137' l_rperms="u-x,g-wx,o-rwx" l_auser="(root|syslog) " l_agroup="(root|adm) "
```

```
if [ "$l_user" = "root" ] || ! grep -Pq -- "^\h*$(awk -F: ' $1=="'"$l_user"'" {print $7}' /etc/passwd)\b" /etc/shells; then

! grep -Pq -- "$l_auser" <<< "$l_user" && l_auser="(root|syslog|$l_user)"

! grep -Pq -- "$l_agroup" <<< "$l_group" && l_agroup="(root|adm|$l_group)"

fi

f_file_test_fix ;;

esac

fi

done < <(stat -Lc '%n:%#a:%U:%G' "$l_file")

done < <(find -L /var/log -type f \( -perm /0137 -o ! -user root -o ! -group root \) -print0)

if [ "${#a_output2[@]}" -le 0 ]; then # If all files passed, then we report no changes
a_output+="( - All files in \"/var/log/" have appropriate permissions and ownership)"

printf '\n%s' "- All files in \"/var/log/" have appropriate permissions and ownership" " o No changes
required" ""

else

printf '\n%s' "${a_output2[@]}" ""

fi

}
```

Note: You may also need to change the configuration for your logging software or services for any logs that had incorrect permissions.

If there are services that log to other locations, ensure that those log files have the appropriate permissions.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AC-3, MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

6.3 System Auditing

The Linux Auditing System operates on a set of rules that collects certain types of system activity to facilitate incident investigation, detect unauthorized access or modification of data. By default events will be logged to `/var/log/audit/audit.log`, which can be configured in `/etc/audit/auditd.conf`.

The following types of audit rules can be specified:

- Control rules: Configuration of the auditing system.
- File system rules: Allow the auditing of access to a particular file or a directory. Also known as file watches.
- System call rules: Allow logging of system calls that any specified program makes.

Audit rules can be set:

- On the command line using the `auditctl` utility. These rules are not persistent across reboots.
- In `/etc/audit/audit.rules`. These rules have to be merged and loaded before they are active.

Notes:

- For 64 bit systems that have `arch` as a rule parameter, you will need two rules: one for 64 bit and one for 32 bit systems calls.
- If the auditing system is configured to be locked (`-e 2`), a system reboot will be required in order to load any changes.
- Key names are optional on the rules and will not be used as a compliance auditing. The usage of key names is highly recommended as it facilitates organization and searching, as such, all remediation steps will have key names supplied.
- It is best practice to store the rules, in number prepended files, in `/etc/audit/rules.d/`. Rules must end in a `.rules` suffix. This then requires the use of `augenrules` to merge all the rules into `/etc/audit/audit.rules` based on their alphabetical (lexical) sort order. All benchmark recommendations follow this best practice for remediation, specifically using the prefix of 50 which is centre weighed if all rule sets make use of the number prepending naming convention.
- Your system may have been customized to change the default `UID_MIN`. All samples output uses `1000`, but this value will not be used in compliance auditing. To confirm the `UID_MIN` for your system, run the following command: `awk '/^\s*UID_MIN/{print $2}' /etc/login.defs`

Normalization

The Audit system normalizes some entries, so when you look at the sample output keep in mind that:

- With regards to users whose login UID is not set, the values `-1 / unset / 4294967295` are equivalent and normalized to `-1`.
- When comparing field types and both sides of the comparison is valid fields types, such as `euid!=uid`, then the auditing system may normalize such that the output is `uid!=euid`.
- Some parts of the rule may be rearranged whilst others are dependant on previous syntax. For example, the following two statements are the same:

```
-a always,exit -F arch=b64 -S execve -C uid!=euid -F auid!=-1 -F key=user_emulation
```

and

```
-a always,exit -F arch=b64 -C euid!=uid -F auid!=unset -S execve -k user_emulation
```

Capacity planning

The recommendations in this section implement auditing policies that not only produces large quantities of logged data, but may also negatively impact system performance. Capacity planning is critical in order not to adversely impact production environments.

- Disk space. If a significantly large set of events are captured, additional on system or off system storage may need to be allocated. If the logs are not sent to a remote log server, ensure that log rotation is implemented else the disk will fill up and the system will halt. Even when logs are sent to a log server, ensure sufficient disk space to allow caching of logs in the case of temporary network outages.
- Disk IO. It is not just the amount of data collected that should be considered, but the rate at which logs are generated.
- CPU overhead. System call rules might incur considerable CPU overhead. Test the systems open/close syscalls per second with and without the rules to gauge the impact of the rules.

6.3.1 Configure auditd Service

The capturing of system events provides system administrators with information to allow them to determine if unauthorized access to their system is occurring.

6.3.2 Configure Data Retention

When auditing, it is important to carefully configure the storage requirements for audit logs. By default, `auditd` will max out the log files at 5MB and retain only 4 copies of them. Older versions will be deleted. It is possible on a system that the 20 MBs of audit logs may fill up the system causing loss of audit data. While the recommendations here provide guidance, check your site policy for audit storage requirements.

6.3.3 Configure auditd Rules

The Audit system operates on a set of rules that define what is to be captured in the log files.

The following types of Audit rules can be specified:

- Control rules: Allow the Audit system's behavior and some of its configuration to be modified.
- File system rules: Allow the auditing of access to a particular file or a directory. (Also known as file watches)
- System call rules: Allow logging of system calls that any specified program makes.

Audit rules can be set:

- on the command line using the `auditctl` utility. Note that these rules are not persistent across reboots.
- in a file ending in `.rules` in the `/etc/audit/rules.d/` directory.

Note: The Linux Benchmarks are written and tested against x86_64 processor architecture. If you are running a different processor type, please review and update the audit rules for the processor architecture of the system

6.3.4 Configure auditd File Access

Without the capability to restrict which roles and individuals can select which events are audited, unauthorized personnel may be able to prevent the auditing of critical events.

7 System Maintenance

Recommendations in this section are intended as maintenance and are intended to be checked on a frequent basis to ensure system stability. Many recommendations do not have quick remediations and require investigation into the cause and best fix available and may indicate an attempted breach of system security.

7.1 Configure system file and directory access

This section provides guidance on securing aspects of system files and directories.

7.1.1 Ensure access to /etc/passwd is configured

Pass

Description:

The `/etc/passwd` file contains user account information that is used by many system utilities and therefore must be readable for these utilities to operate.

Rationale:

It is critical to ensure that the `/etc/passwd` file is protected from unauthorized write access. Although it is protected by default, the file permissions could be changed either inadvertently or through malicious actions.

Remediation:

Run the following commands to remove excess permissions, set owner, and set group on `/etc/passwd` :

```
# chmod u-x,go-wx /etc/passwd
# chown root:root /etc/passwd
```

Assessment:

Show Assessment Evidence

Show Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AC-3, MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

Back to Summary

7.1.2 Ensure access to /etc/passwd- is configured

Pass

Description:

The `/etc/passwd-` file contains backup user account information.

Rationale:

It is critical to ensure that the `/etc/passwd-` file is protected from unauthorized access. Although it is protected by default, the file permissions could be changed either inadvertently or through malicious actions.

Remediation:

Run the following commands to remove excess permissions, set owner, and set group on `/etc/passwd-` :

```
# chmod u-x,go-wx /etc/passwd-  
# chown root:root /etc/passwd-
```

Assessment:
[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AC-3, MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

7.1.3 Ensure access to /etc/group is configured

Pass

Description:

The /etc/group file contains a list of all the valid groups defined in the system. The command below allows read/write access for root and read access for everyone else.

Rationale:

The /etc/group file needs to be protected from unauthorized changes by non-privileged users, but needs to be readable as this information is used with many non-privileged programs.

Remediation:

Run the following commands to remove excess permissions, set owner, and set group :

```
# chmod u-x,go-wx /etc/group  
# chown root:root /etc/group
```

Assessment:
[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AC-3, MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

7.1.4 Ensure access to /etc/group- is configured

Pass

Description:

The /etc/group- file contains a backup list of all the valid groups defined in the system.

Rationale:

It is critical to ensure that the `/etc/group-` file is protected from unauthorized access. Although it is protected by default, the file permissions could be changed either inadvertently or through malicious actions.

Remediation:

Run the following commands to remove excess permissions, set owner, and set group on `/etc/group-` :

```
# chmod u-x,go-wx /etc/group-  
# chown root:root /etc/group-
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AC-3. MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

7.1.5 Ensure access to `/etc/shadow` is configured

Pass

Description:

The `/etc/shadow` file is used to store the information about user accounts that is critical to the security of those accounts, such as the hashed password and other security information.

Rationale:

If attackers can gain read access to the `/etc/shadow` file, they can easily run a password cracking program against the hashed password to break it. Other security information that is stored in the `/etc/shadow` file (such as expiration) could also be useful to subvert the user accounts.

Remediation:

Run **one** of the following commands to set ownership of `/etc/shadow` to `root` and group to either `root` or `shadow` :

```
# chown root:shadow /etc/shadow  
  
-OR-  
  
# chown root:root /etc/shadow
```

Run the following command to remove excess permissions form `/etc/shadow` :

```
# chmod u-x,g-wx,o-rwx /etc/shadow
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AC-3. MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
- >

[Back to Summary](#)

7.1.6 Ensure access to /etc/shadow- is configured

Pass

Description:

The `/etc/shadow-` file is used to store backup information about user accounts that is critical to the security of those accounts, such as the hashed password and other security information.

Rationale:

It is critical to ensure that the `/etc/shadow-` file is protected from unauthorized access. Although it is protected by default, the file permissions could be changed either inadvertently or through malicious actions.

Remediation:

Run **one** of the following commands to set ownership of `/etc/shadow-` to `root` and group to either `root` or `shadow` :

```
# chown root:shadow /etc/shadow-  
  
-OR-  
  
# chown root:root /etc/shadow-
```

Run the following command to remove excess permissions form `/etc/shadow-` :

```
# chmod u-x,g-wx,o-rwx /etc/shadow-
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AC-3. MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
- >

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
- >

[Back to Summary](#)

7.1.7 Ensure access to /etc/gshadow is configured

Pass

Description:

The `/etc/gshadow` file is used to store the information about groups that is critical to the security of those accounts, such as the hashed password and other security information.

Rationale:

If attackers can gain read access to the `/etc/gshadow` file, they can easily run a password cracking program against the hashed password to break it. Other security information that is stored in the `/etc/gshadow` file (such as group administrators) could also be useful to subvert the group.

Remediation:

Run **one** of the following commands to set ownership of `/etc/gshadow` to `root` and group to either `root` or `shadow` :

```
# chown root:shadow /etc/gshadow
```

-OR-

```
# chown root:root /etc/gshadow
```

Run the following command to remove excess permissions form /etc/gshadow :

```
# chmod u-x,g-wx,o-rwx /etc/gshadow
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AC-3. MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

7.1.8 Ensure access to /etc/gshadow- is configured

Pass

Description:

The /etc/gshadow- file is used to store backup information about groups that is critical to the security of those accounts, such as the hashed password and other security information.

Rationale:

It is critical to ensure that the /etc/gshadow- file is protected from unauthorized access. Although it is protected by default, the file permissions could be changed either inadvertently or through malicious actions.

Remediation:

Run **one** of the following commands to set ownership of /etc/gshadow- to root and group to either root or shadow :

```
# chown root:shadow /etc/gshadow-
```

-OR-

```
# chown root:root /etc/gshadow-
```

Run the following command to remove excess permissions form /etc/gshadow- :

```
# chmod u-x,g-wx,o-rwx /etc/gshadow-
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AC-3. MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

7.1.9 Ensure access to /etc/shells is configured

Pass

Description:

/etc/shells is a text file which contains the full pathnames of valid login shells. This file is consulted by chsh and available to be queried by other programs.

Rationale:

It is critical to ensure that the /etc/shells file is protected from unauthorized access. Although it is protected by default, the file permissions could be changed either inadvertently or through malicious actions.

Remediation:

Run the following commands to remove excess permissions, set owner, and set group on /etc/shells :

```
# chmod u-x,go-wx /etc/shells
# chown root:root /etc/shells
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AC-3, MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

7.1.10 Ensure access to /etc/security/opasswd is configured

Pass

Description:

/etc/security/opasswd and it's backup /etc/security/opasswd.old hold user's previous passwords if pam_unix or pam_pwhistory is in use on the system

Rationale:

It is critical to ensure that /etc/security/opasswd is protected from unauthorized access. Although it is protected by default, the file permissions could be changed either inadvertently or through malicious actions.

Remediation:

Run the following commands to remove excess permissions, set owner, and set group on /etc/security/opasswd and /etc/security/opasswd.old is they exist:

```
# [ -e "/etc/security/opasswd" ] && chmod u-x,go-rwx /etc/security/opasswd
# [ -e "/etc/security/opasswd" ] && chown root:root /etc/security/opasswd
# [ -e "/etc/security/opasswd.old" ] && chmod u-x,go-rwx /etc/security/opasswd.old
# [ -e "/etc/security/opasswd.old" ] && chown root:root /etc/security/opasswd.old
```

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AC-3, MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

7.1.11 Ensure world writable files and directories are secured

Pass

Description:

World writable files are the least secure. Data in world-writable files can be modified and compromised by any user on the system. World writable files may also indicate an incorrectly written script or program that could potentially be the cause of a larger compromise to the system's integrity. See the `chmod(2)` man page for more information.

Setting the sticky bit on world writable directories prevents users from deleting or renaming files in that directory that are not owned by them.

Rationale:

Data in world-writable files can be modified and compromised by any user on the system. World writable files may also indicate an incorrectly written script or program that could potentially be the cause of a larger compromise to the system's integrity.

This feature prevents the ability to delete or rename files in world writable directories (such as `/tmp`) that are owned by another user.

Remediation:

- World Writable Files:
 - It is recommended that write access is removed from `other` with the command (`chmod o-w <filename>`), but always consult relevant vendor documentation to avoid breaking any application dependencies on a given file.
- World Writable Directories:
 - Set the sticky bit on all world writable directories with the command (`chmod a+t <directory_name>`)

Run the following script to:

- Remove other write permission from any world writable files
- Add the sticky bit to all world writable directories

```
#!/usr/bin/env bash

{
l_smask='01000'

a_file=(); a_dir=() # Initialize arrays

a_path=( ! -path "/run/user/*" -a ! -path "/proc/*" -a ! -path "*/containerd/*" -a ! -path "*/kubelet/pods/*" -a ! -path "*/kubelet/plugins/*" -a ! -path "/sys/*" -a ! -path "/snap/*")

while IFS= read -r l_mount; do
while IFS= read -r -d $'\0' l_file; do

if [ -e "$l_file" ]; then

l_mode="$(stat -Lc '%#a' "$l_file")"

if [ -f "$l_file" ]; then # Remove excess permissions from WW files
```



```
echo -e " - File: \"$1_file\" is mode: \"$1_mode\"\\n - removing write permission on \"$1_file\" from \\n\"other\\n\"

chmod o-w \"$1_file"

fi

if [ -d "$1_file" ]; then # Add sticky bit

if [ ! $(( $1_mode & $1_smask )) -gt 0 ]; then

echo -e " - Directory: \"$1_file\" is mode: \"$1_mode\" and doesn't have the sticky bit set\\n - Adding the sticky bit"

chmod a+t "$1_file"

fi

fi

fi

done < <(find "$1_mount" -xdev \( "${a_path[@]}" \) \( -type f -o -type d \) -perm -0002 -print0 2> /dev/null)

done < <(findmnt -Dkerno fstype,target | awk '($1 !~ /^s*(nfs|proc|smb|vfat|iso9660|efivarfs|selinuxfs)/ && $2 !~ /^(\\/run\\/user\\/|\\/tmp|\\/var\\/tmp)/){print $2}')

}
```

Assessment:
[Show](#) Assessment Evidence

[Show](#) Rule Result XML

- References:**
- **URL:** NIST SP 800-53 Rev. 5: AC-3. MP-2
 - **URL:** STIG ID: UBTU-20-010411 | Rule ID: SV-238332r958524 | CAT ||
 - **URL:** STIG ID: UBTU-22-232145 | Rule ID: SV-260513r958524 | CAT ||

- CIS Controls V7.0:**
- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

- CIS Critical Security Controls V8.0:**
- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

7.1.12 Ensure no files or directories without an owner and a group exist

Pass

Description:

Administrators may delete users or groups from the system and neglect to remove all files and/or directories owned by those users or groups.

Rationale:

A new user or group who is assigned a deleted user's user ID or group ID may then end up "owning" a deleted user or group's files, and thus have more access on the system than was intended.

Remediation:

Remove or set ownership and group ownership of these files and/or directories to an active user on the system as appropriate.

Assessment:
[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: AC-3, MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

7.1.13 Ensure SUID and SGID files are reviewed

Manual

Description:

The owner of a file can set the file's permissions to run with the owner's or group's permissions, even if the user running the program is not the owner or a member of the group. The most common reason for a SUID or SGID program is to enable users to perform functions (such as changing their password) that require root privileges.

Rationale:

There are valid reasons for SUID and SGID programs, but it is important to identify and review such programs to ensure they are legitimate. Review the files returned by the action in the audit section and check to see if system binaries have a different checksum than what from the package. This is an indication that the binary may have been replaced.

Remediation:

Ensure that no rogue SUID or SGID programs have been introduced into the system. Review the files returned by the action in the Audit section and confirm the integrity of these binaries.

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5, AC-3, MP-2

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

7.2 Local User and Group Settings

This section provides guidance on securing aspects of the local users and groups.

Note: The recommendations in this section check local users and groups. Any users or groups from other sources such as LDAP will not be audited. In a domain environment similar checks should be performed against domain users and groups.

7.2.1 Ensure accounts in /etc/passwd use shadowed passwords

Pass

Description:

Local accounts can uses shadowed passwords. With shadowed passwords, The passwords are saved in shadow password file, `/etc/shadow` , encrypted by a salted one-way hash. Accounts with a shadowed password have an `x` in the second field in `/etc/passwd` .

Rationale:

The `/etc/passwd` file also contains information like user ID's and group ID's that are used by many system programs. Therefore, the `/etc/passwd` file must remain world readable. In spite of encoding the password with a

randomly-generated one-way hash function, an attacker could still break the system if they got access to the `/etc/passwd` file. This can be mitigated by using shadowed passwords, thus moving the passwords in the `/etc/passwd` file to `/etc/shadow`. The `/etc/shadow` file is set so only root will be able to read and write. This helps mitigate the risk of an attacker gaining access to the encoded passwords with which to perform a dictionary attack.

Note:

- All accounts must have passwords or be locked to prevent the account from being used by an unauthorized user.
- A user account with an empty second field in `/etc/passwd` allows the account to be logged into by providing only the username.

Remediation:

Run the following command to set accounts to use shadowed passwords and migrate passwords in `/etc/passwd` to `/etc/shadow`:

```
# pwconv
```

Investigate to determine if the account is logged in and what it is being used for, to determine if it needs to be forced off.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: IA-5
- **URL:** PWCONV(8)

CIS Controls V7.0:

- **Control 16: Account Monitoring and Control:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)

7.2.2 Ensure `/etc/shadow` password fields are not empty

Pass

Description:

An account with an empty password field means that anybody may log in as that user without providing a password.

Rationale:

All accounts must have passwords or be locked to prevent the account from being used by an unauthorized user.

Remediation:

If any accounts in the `/etc/shadow` file do not have a password, run the following command to lock the account until it can be determined why it does not have a password:

```
# passwd -l <username>
```

Also, check to see if the account is logged in and investigate what it is being used for to determine if it needs to be forced off.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- **URL:** NIST SP 800-53 Rev. 5: IA-5
- **URL:** NIST SP 800-53 Revision 5 :: CM-6 b
- **URL:** NIST SP 800-53A :: CM-6.1 (iv)

CIS Controls V7.0:

- Control 4: Controlled Use of Administrative Privileges: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 5: Account Management: -- [More](#)
>

[Back to Summary](#)**7.2.3 Ensure all groups in /etc/passwd exist in /etc/group**

Pass

Description:

Over time, system administration errors and changes can lead to groups being defined in `/etc/passwd` but not in `/etc/group`.

Rationale:

Groups defined in the `/etc/passwd` file but not in the `/etc/group` file pose a threat to system security since group permissions are not properly managed.

Remediation:

Analyze the output of the Audit step above and perform the appropriate action to correct any discrepancies found.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
- Control 14: Security Awareness and Skills Training: -- [More](#)
>

[Back to Summary](#)**7.2.4 Ensure shadow group is empty**

Pass

Description:

The shadow group allows system programs which require access the ability to read the `/etc/shadow` file. No users should be assigned to the shadow group.

Rationale:

Any users assigned to the shadow group would be granted read access to the `/etc/shadow` file. If attackers can gain read access to the `/etc/shadow` file, they can easily run a password cracking program against the hashed passwords to break them. Other security information that is stored in the `/etc/shadow` file (such as expiration) could also be useful to subvert additional user accounts.

Remediation:

Run the following command to remove all users from the shadow group

```
# sed -ri 's/(^shadow:[^:]*:[^:]*:)([^\:]+$)/\1/' /etc/group
```

Change the primary group of any users with shadow as their primary group.

```
# usermod -g <primary group> <user>
```

Assessment:

[Show](#) Assessment Evidence

[Show Rule Result XML](#)**References:**

- **URL:** NIST SP 800-53 Rev. 5: IA-5

CIS Controls V7.0:

- **Control 14: Controlled Access Based on the Need to Know:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Control 3: Data Protection:** -- [More](#)
>

[Back to Summary](#)**7.2.5 Ensure no duplicate UIDs exist**

Pass

Description:

Although the `useradd` program will not let you create a duplicate User ID (UID), it is possible for an administrator to manually edit the `/etc/passwd` file and change the UID field.

Rationale:

Users must be assigned unique UIDs for accountability and to ensure appropriate access protections.

Satisfies: SRG-OS-000104-GPOS-00051, SRG-OS-000121-GPOS-00062, SRG-OS-000042-GPOS-00020

Remediation:

Based on the results of the audit script, establish unique UIDs and review all files owned by the shared UIDs to determine which UID they are supposed to belong to.

Assessment:

[Show](#) Assessment Evidence

[Show Rule Result XML](#)**References:**

- **URL:** NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5
- **URL:** NIST SP 800-53A :: IA-2.1

CIS Controls V7.0:

- **Not Explicitly Mapped:** -- [More](#)
>

CIS Critical Security Controls V8.0:

- **Not Explicitly Mapped:** -- [More](#)
>

[Back to Summary](#)**7.2.6 Ensure no duplicate GIDs exist**

Pass

Description:

Although the `groupadd` program will not let you create a duplicate Group ID (GID), it is possible for an administrator to manually edit the `/etc/group` file and change the GID field.

Rationale:

User groups must be assigned unique GIDs for accountability and to ensure appropriate access protections.

Remediation:

Based on the results of the audit script, establish unique GIDs and review all files owned by the shared GID to determine which group they are supposed to belong to.

Assessment:[Show](#) Assessment Evidence[Show](#) Rule Result XML**References:**

- URL: NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5

CIS Controls V7.0:

- Not Explicitly Mapped: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Not Explicitly Mapped: -- [More](#)
>

[Back to Summary](#)**7.2.7 Ensure no duplicate user names exist**

Pass

Description:

Although the `useradd` program will not let you create a duplicate user name, it is possible for an administrator to manually edit the `/etc/passwd` file and change the user name.

Rationale:

If a user is assigned a duplicate user name, it will create and have access to files with the first UID for that username in `/etc/passwd`. For example, if "test4" has a UID of 1000 and a subsequent "test4" entry has a UID of 2000, logging in as "test4" will use UID 1000. Effectively, the UID is shared, which is a security problem.

Remediation:

Based on the results of the audit script, establish unique user names for the users. File ownerships will automatically reflect the change as long as the users have unique UIDs.

Assessment:[Show](#) Assessment Evidence[Show](#) Rule Result XML**References:**

- URL: NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5

CIS Controls V7.0:

- Not Explicitly Mapped: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Not Explicitly Mapped: -- [More](#)
>

[Back to Summary](#)**7.2.8 Ensure no duplicate group names exist**

Pass

Description:

Although the `groupadd` program will not let you create a duplicate group name, it is possible for an administrator to manually edit the `/etc/group` file and change the group name.

Rationale:

If a group is assigned a duplicate group name, it will create and have access to files with the first GID for that group in `/etc/group`. Effectively, the GID is shared, which is a security problem.

Remediation:

Based on the results of the audit script, establish unique names for the user groups. File group ownerships will automatically reflect the change as long as the groups have unique GIDs.

Assessment:

[Show](#) Assessment Evidence

[Show](#) Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5

CIS Controls V7.0:

- Not Explicitly Mapped: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Not Explicitly Mapped: -- [More](#)
>

[Back to Summary](#)

7.2.9 Ensure local interactive user home directories are configured

Fail

Description:

The user home directory is space defined for the particular user to set local environment variables and to store personal files. While the system administrator can establish secure permissions for users' home directories, the users can easily override these. Users can be defined in `/etc/passwd` without a home directory or with a home directory that does not actually exist.

Rationale:

Since the user is accountable for files stored in the user home directory, the user must be the owner of the directory. Group or world-writable user home directories may enable malicious users to steal or modify other users' data or to gain another user's system privileges. If the user's home directory does not exist or is unassigned, the user will be placed in `/` and will not be able to write any files or have local environment variables set.

Remediation:

If a local interactive users' home directory is undefined and/or doesn't exist, follow local site policy and perform one of the following:

- Lock the user account
- Remove the user from the system
- create a directory for the user. If undefined, edit `/etc/passwd` and add the absolute path to the directory to the last field of the user.

Run the following script to:

- Remove excessive permissions from local interactive users home directories
- Update the home directory's owner

```
#!/usr/bin/env bash

{

a_output=() a_output2=() a_exists2=() a_mode2=() a_owner2=()

l_valid_shells="^( $( awk -F\ / '$NF != "nologin" {print}' /etc/shells | sed -rn '/^\/{s,/,\|\\/,g;p}' |
paste -s -d '|' - ) )$"

l_mask='0027'; l_max="$( printf '%o' $(( 0777 & ~$l_mask)) )"

l_users="$(awk -v pat="$l_valid_shells" -F: '$(NF) ~ pat { print $1 " " $(NF-1) }' /etc/passwd | wc -l)"

[ "$l_users" -gt 10000 ] && printf '%s\n' "" " ** INFO ** \

" $l_users Local interactive users found on the system" " This may be a long running process" "
***** "
```

```
while IFS=" " read -r l_user l_home; do

if [ -d "$l_home" ]; then

while IFS=: read -r l_own l_mode; do

if [ "$l_user" != "$l_own" ]; then

a_owner2+=(" - User: \"$l_user\" Home \"$l_home\" is owned by: \"$l_own\" \" \" \

" changing owner to: \"$l_user\"") && chown "$l_user" "$l_home"

fi

if [ $(( $l_mode & $l_mask )) -gt 0 ]; then

a_mode2+=(" - User: \"$l_user\" Home \"$l_home\" is mode: \"$l_mode\" \" \" \

" changing to mode: \"$l_max\" or more restrictive")

chmod g-w,o-rwx "$l_home"

fi

done <<< "$(stat -Lc '%U:%#a' "$l_home")"

else

a_exists2+=(" - User: \"$l_user\" Home Directory: \"$l_home\" Doesn't exist")

fi

done <<< "$(awk -v pat="$l_valid_shells" -F: '{(NF) ~ pat { print $1 " " $(NF-1) }' /etc/passwd)"

[ "${#a_exists2[@]}" -gt 0 ] && a_output2+=("${a_exists2[@]}")

[ "${#a_mode2[@]}" -gt 0 ] && a_output2+=("${a_mode2[@]}")

[ "${#a_owner2[@]}" -gt 0 ] && a_output2+=("${a_owner2[@]}")

if [ "${#a_output2[@]}" -gt 0 ]; then

printf '%s\n' " " "${a_output2[@]}"

else

printf '%s\n' " " "- No changes required"

fi

}
```

Assessment:
[Show](#) Assessment Evidence

[Show](#) Rule Result XML

- References:**
- URL: NIST SP 800-53 Revision 5 :: CM-6 b
 - URL: NIST SP 800-53A :: CM-6.1 (iv)

- CIS Controls V7.0:**
- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

- CIS Critical Security Controls V8.0:**
- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

7.2.10 Ensure local interactive user dot files access is configured

Fail

Description:

While the system administrator can establish secure permissions for users' "dot" files, the users can easily override

these.

- `.forward` file specifies an email address to forward the user's mail to.
- `.rhost` file provides the "remote authentication" database for the `rcp`, `rlogin`, and `rsh` commands and the `rcmd()` function. These files bypass the standard password-based user authentication mechanism. They specify remote hosts and users that are considered trusted (i.e. are allowed to access the local system without supplying a password)
- `.netrc` file contains data for logging into a remote host or passing authentication to an API.
- `.bash_history` file keeps track of the user's commands.

Rationale:

User configuration files with excessive or incorrect access may enable malicious users to steal or modify other users' data or to gain another user's system privileges.

Remediation:

Making global modifications to users' files without alerting the user community can result in unexpected outages and unhappy users. Therefore, it is recommended that a monitoring policy be established to report user dot file permissions and determine the action to be taken in accordance with site policy.

The following script will:

- remove excessive permissions on dot files within interactive users' home directories
- change ownership of dot files within interactive users' home directories to the user
- change group ownership of dot files within interactive users' home directories to the user's primary group
- list `.forward` and `.rhost` files to be investigated and manually deleted

```
#!/usr/bin/env bash

{
a_output2=(); a_output3=()

l_maxsize="1000" # Maximum number of local interactive users before warning (Default 1,000)

l_valid_shells="^( awk -F\ ' $NF != "nologin" {print}' /etc/shells | sed -rn '/^\/{s,/,\,\,\,\,g;p}' |
paste -s -d '|' - )$"

a_user_and_home=() # Create array with local users and their home directories

while read -r l_local_user l_local_user_home; do # Populate array with users and user home location

[[ -n "$l_local_user" && -n "$l_local_user_home" ]] && a_user_and_home+=("$l_local_user:
$l_local_user_home")

done <<< "$(awk -v pat="$l_valid_shells" -F: '{(NF) ~ pat { print $1 " " $(NF-1) }' /etc/passwd)"

l_asize="${#a_user_and_home[@]}" # Here if we want to look at number of users before proceeding

[ "${#a_user_and_home[@]}" -gt "$l_maxsize" ] && printf '%s\n' "" " ** INFO **" \

" - \"$l_asize\" Local interactive users found on the system" \

" - This may be a long running check" ""

file_access_fix()

{

a_access_out=()

l_max="$( printf '%o' $(( 0777 & ~$l_mask )) )"

if [ $(( $l_mode & $l_mask )) -gt 0 ]; then

printf '%s\n' "" " - File: \"$l_hdfilename\" is mode: \"$l_mode\" and should be mode: \"$l_max\" or more
restrictive" \

" Updating file: \"$l_hdfilename\" to be mode: \"$l_max\" or more restrictive"

chmod "$l_change" "$l_hdfilename"

fi

if [[ ! "$l_owner" =~ ($l_user) ]]; then

printf '%s\n' "" " - File: \"$l_hdfilename\" owned by: \"$l_owner\" and should be owned by \"$l_user\" or
```

```

} \ " \
" Updating file: \"${l_hdfilename}\" to be owned by \"${l_user} // or \"${l_group}\"
chown \"${l_user}\" \"${l_hdfilename}\"
fi
if [ [ ! \"${l_gowner}\" =~ (${l_group}) ]]; then
printf '%s\\n' \"\" \" - File: \"${l_hdfilename}\" group owned by: \"${l_gowner}\" and should be group owned by
\"${l_group} // or \"${l_group}\" \
" Updating file: \"${l_hdfilename}\" to be group owned by \"${l_group} // or \"${l_group}\"
chgrp \"${l_group}\" \"${l_hdfilename}\"
fi
}
while IFS=: read -r l_user l_home; do
a_dot_file=(); a_netrc=(); a_netrc_warn=(); a_bhout=(); a_hdirout=()
if [ -d \"${l_home}\" ]; then
l_group=\"$(id -gn \"${l_user}\" | xargs)\"; l_group=\"${l_group} // \"${l_group}\"
while IFS= read -r -d $'\\0' l_hdfilename; do
while read -r l_mode l_owner l_gowner; do
case \"$(basename \"${l_hdfilename}\")\" in
.forward | .rhost )
a_dot_file+=(\" - File: \"${l_hdfilename}\" exists\" \" Please review and manually delete this file\") ;;
.netrc )
l_mask='0177'; l_change=\"u-x,go-rwx\"; file_access_fix
a_netrc_warn+=(\" - File: \"${l_hdfilename}\" exists\") ;;
.bash_history )
l_mask='0177'; l_change=\"u-x,go-rwx\"; file_access_fix ;;
* )
l_mask='0133'; l_change=\"u-x,go-wx\"; file_access_fix ;;
esac
done << (stat -lc '%#a %U %G' \"${l_hdfilename}\")
done << (find \"${l_home}\" -xdev -type f -name '.*' -print0)
fi
[ \"${#a_dot_file[@]}\" -gt 0 ] && a_output2+=(\" - User: \"${l_user}\" Home Directory: \"${l_home}\" \"
\"${a_dot_file[@]}\")
[ \"${#a_netrc_warn[@]}\" -gt 0 ] && a_output3+=(\" - User: \"${l_user}\" Home Directory: \"${l_home}\" \"
\"${a_netrc_warn[@]}\")
done <<< \"$(printf '%s\\n' \"${a_user_and_home[@]}\")\"
[ \"${#a_output3[@]}\" -gt 0 ] && printf '%s\\n' \"\" \" ** WARNING **\" \"${a_output3[@]}\" \"
[ \"${#a_output2[@]}\" -gt 0 ] && printf '%s\\n' \"\" \"${a_output2[@]}\"
}

```

Assessment:

Show Assessment Evidence

Show Rule Result XML

References:

- URL: NIST SP 800-53 Rev. 5: CM-1, CM-2, CM-6, CM-7, IA-5

CIS Controls V7.0:

- Control 14: Controlled Access Based on the Need to Know: -- [More](#)
>

CIS Critical Security Controls V8.0:

- Control 3: Data Protection: -- [More](#)
>

[Back to Summary](#)

